

FluCoMa



Ted Moore

tedmooremusic.com

ted@tedmooremusic.com

FluCoMa: Fluid Corpus Manipulation



- “enable techno-fluent musicians to use machine listening and machine learning in their creative practices”
- Integrating Machine Listening and Machine Learning in...
- Max, SuperCollider & Pure Data
- Learning Resources (learn.flucoma.org)
- Discourse Community (discourse.flucoma.org)



FluCoMa: Fluid Corpus Manipulation



Slice Audio

- onset slice
- transient slice
- novelty slice
- amplitude slice
- amplitude gate

Decompose Audio

- extract transients
- harmonic/percussive separation
- model as sine waves
- non-negative matrix factorisation

Analyze Audio

- pitch
- loudness
- mel-bands
- mel-frequency cepstral coefficients
- spectral centroid
- spectral flatness
- chromagram

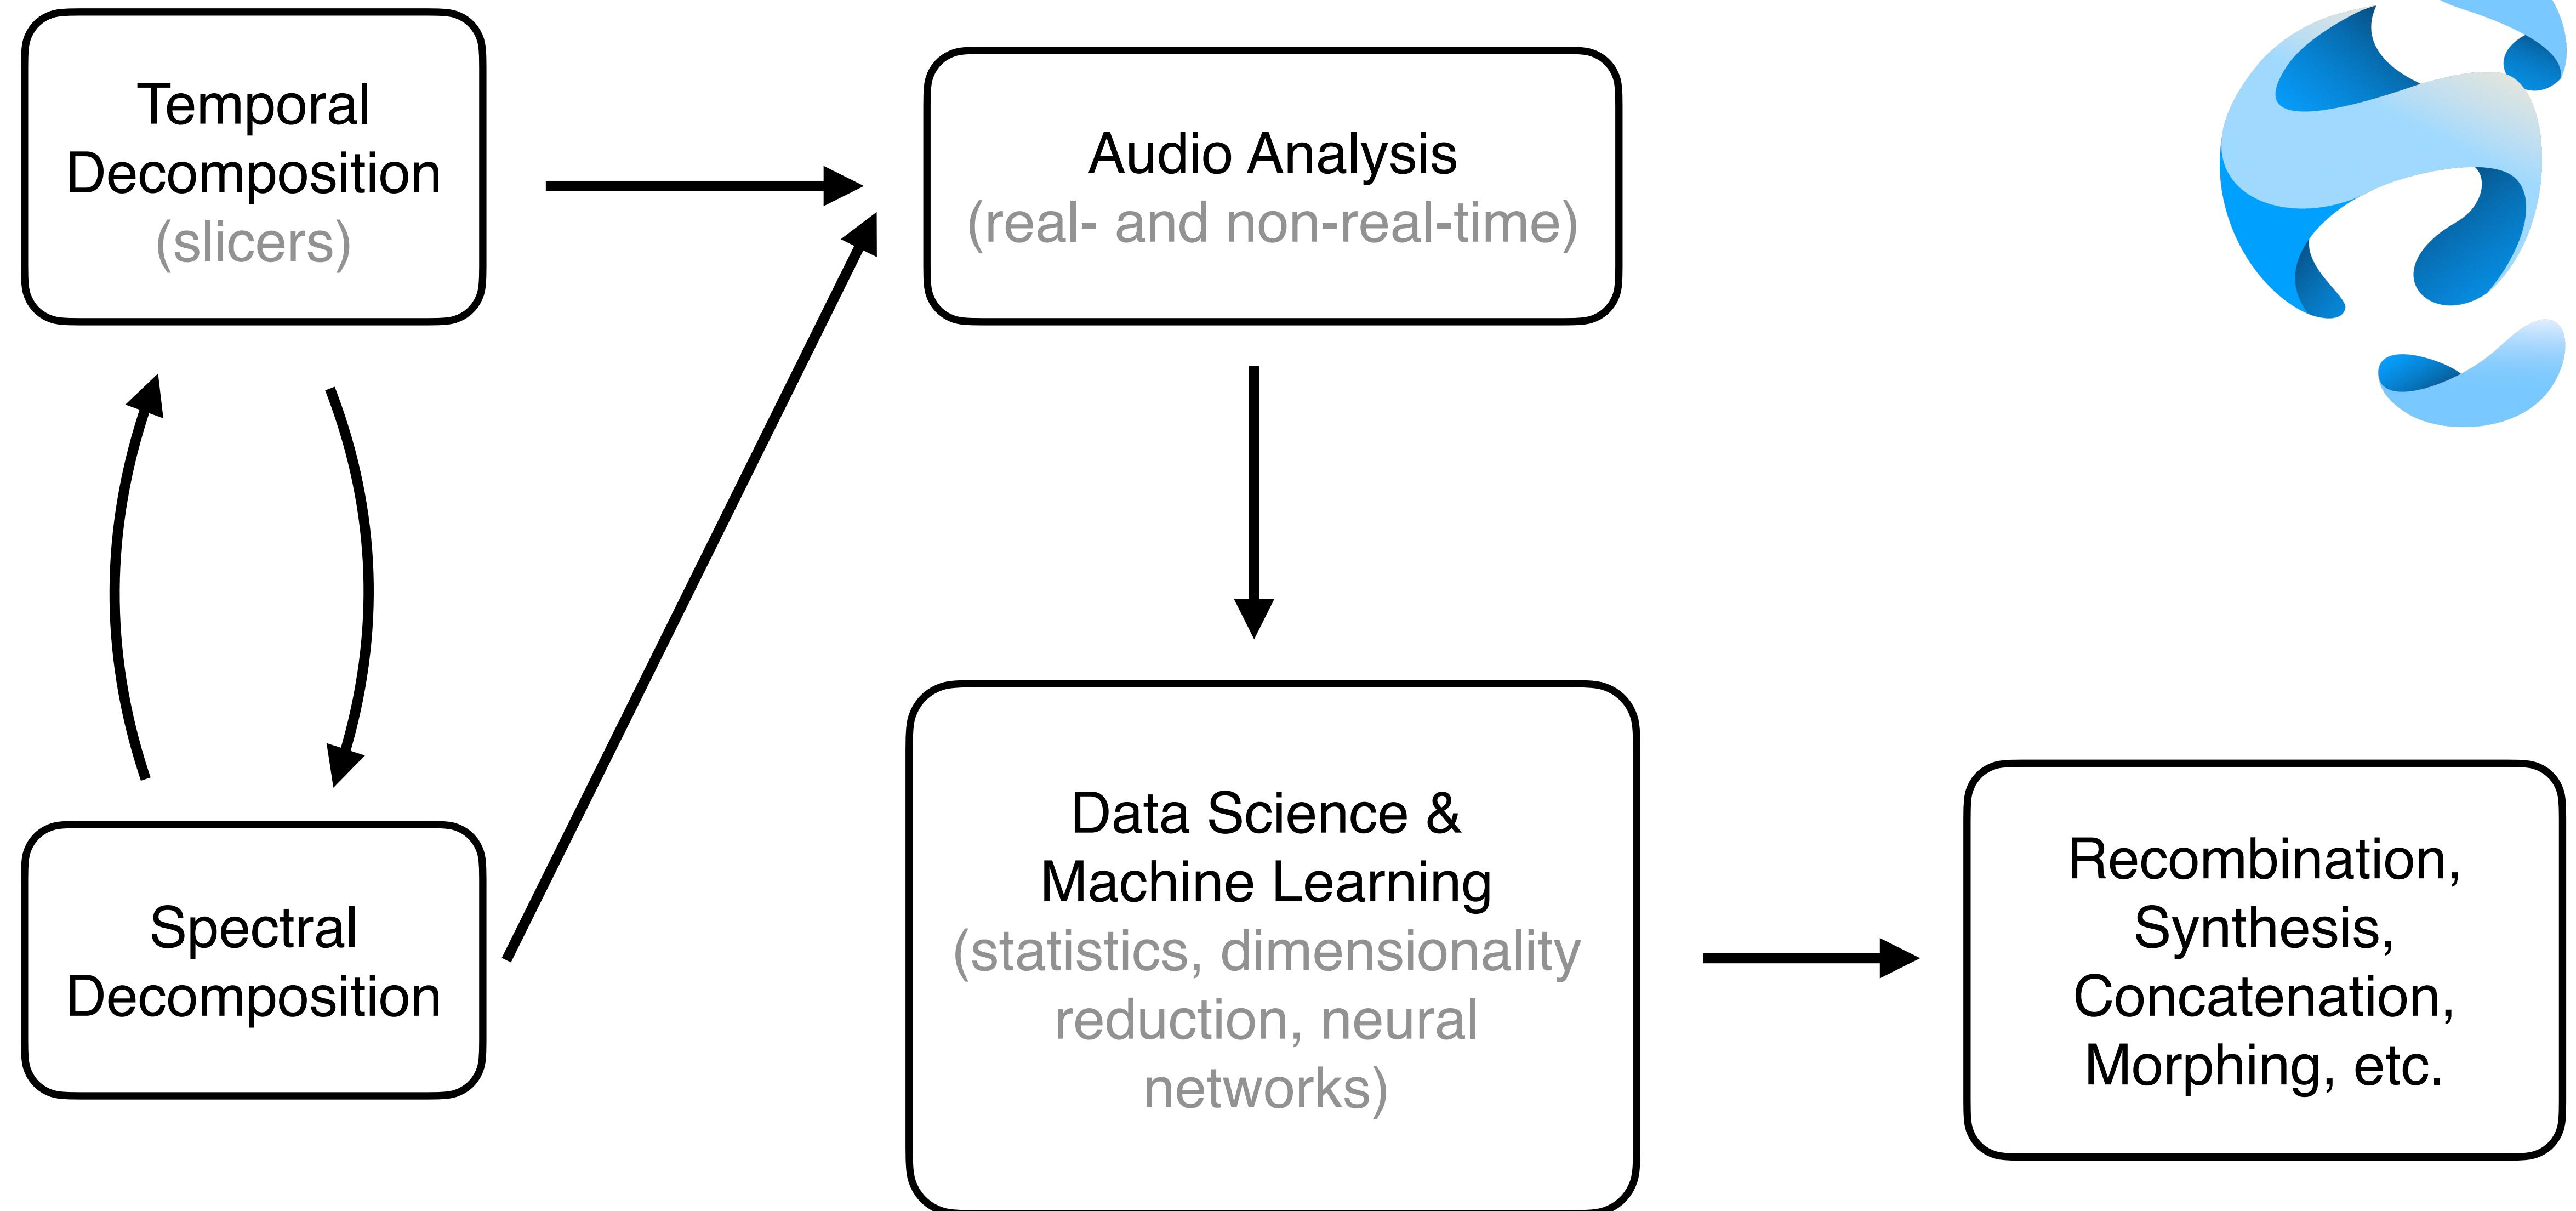
Transform Audio

- audio transport
- non-negative matrix factorisation filters & morphing

Analyze Data

- datasets
- labelsets
- statistical analyses
- normalization
- standardization
- robust scaler
- principal component analysis
- MDS
- KDTree
- K Nearest Neighbours
- neural networks
- SQL-type query
- KMeans
- UMAP
- grid

... and more



Supervised Learning

Classification

- fluid.mlpclassifier~
- fluid.knnclassifier~

Regression

- fluid.mlpregressor~
- fluid.knnregressor~

Preprocessing

- fluid.normalize~
- fluid.standardize~
- fluid.robustscale~

Unsupervised Learning

Dimensionality Reduction

- fluid.pca~
- fluid.umap~
- fluid.mds~
- autoencoder
(fluid.mlpregressor~)

Clustering

- fluid.kmeans~
- fluid.skmeans~

Linear Assignment

- fluid.grid~

Also...
fluid.bufnmf~

scikit-learn

Machine Learning in Python

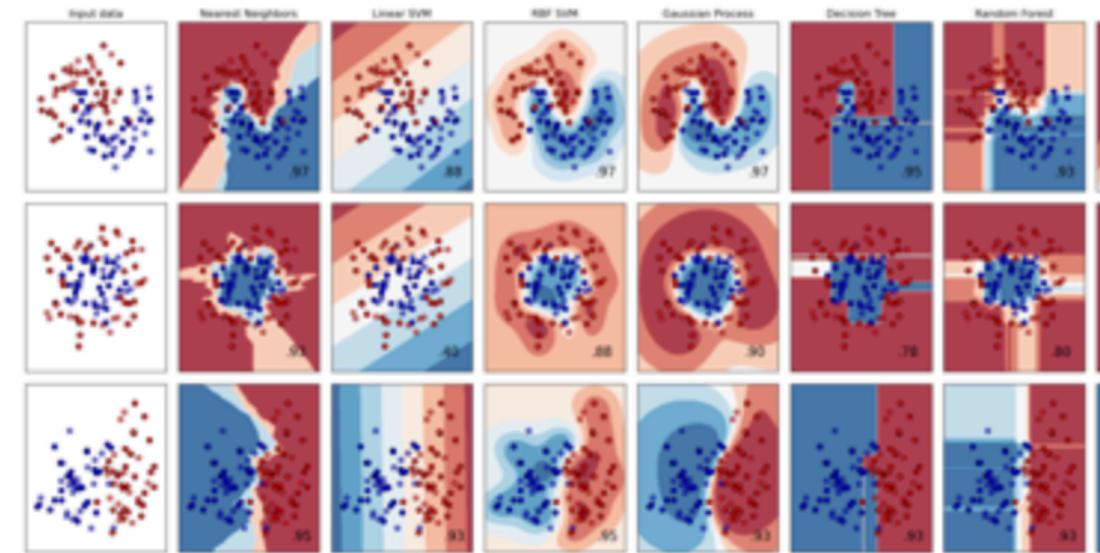
[Getting Started](#)[Release Highlights for 1.2](#)[GitHub](#)

- Simple and efficient tools for predictive data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying which category an object belongs to.

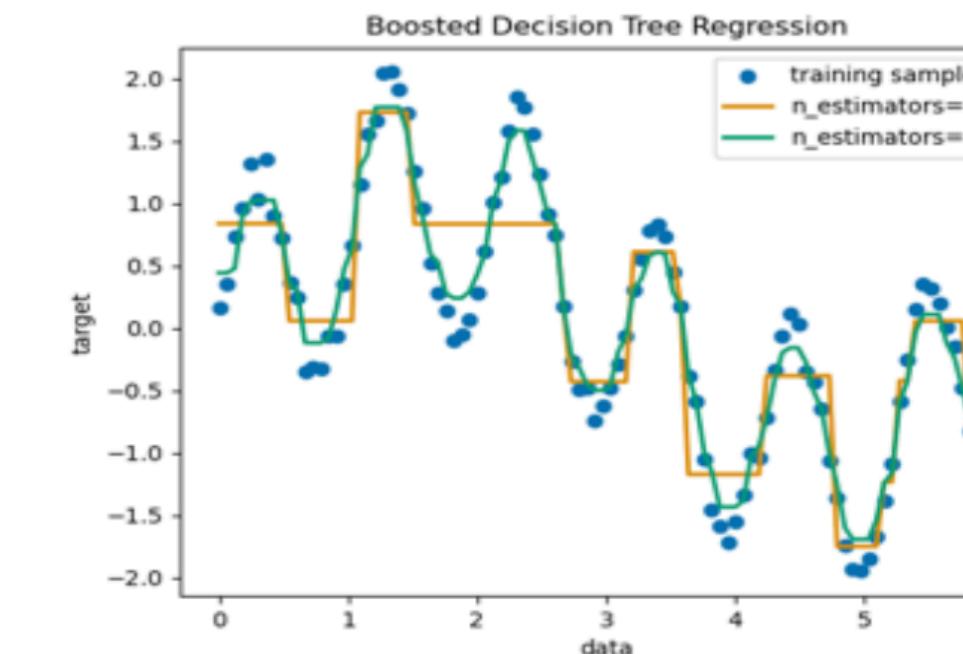
Applications: Spam detection, image recognition.
Algorithms: SVM, nearest neighbors, random forest, and more...

[Examples](#)

Regression

Predicting a continuous-valued attribute associated with an object.

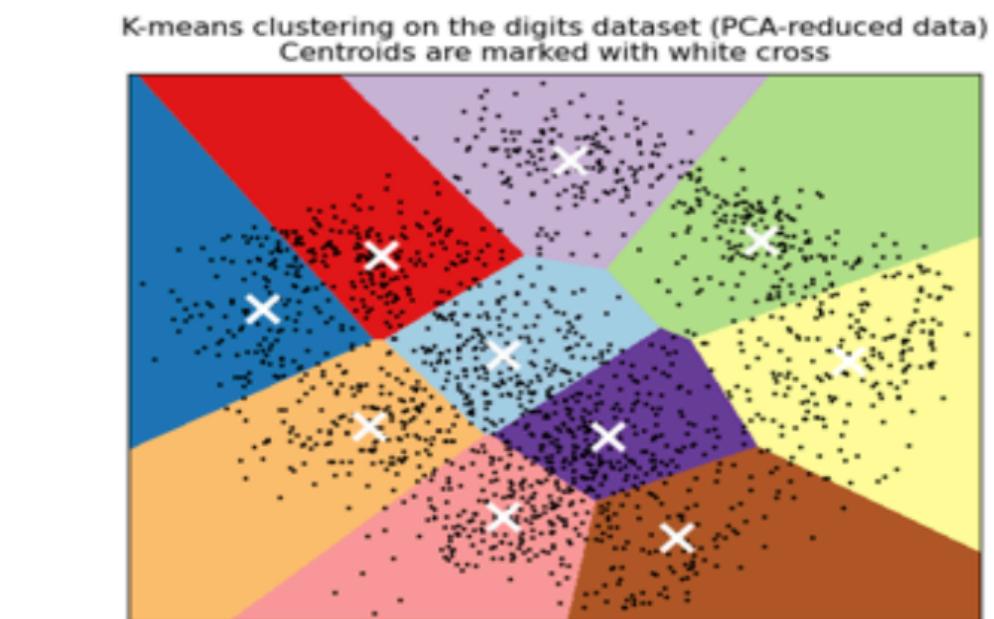
Applications: Drug response, Stock prices.
Algorithms: SVR, nearest neighbors, random forest, and more...

[Examples](#)

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes
Algorithms: k-Means, spectral clustering, mean-shift, and more...

[Examples](#)

Dimensionality reduction

Reducing the number of random variables to consider.

[stable/auto_examples/release_highlights/plot_release_highlights_1_2_0.html](#)

Model selection

Comparing, validating and choosing parameters and models.

...

Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text

Supervised Learning

Learning patterns, associations, or relationships from data that is pre-labeled

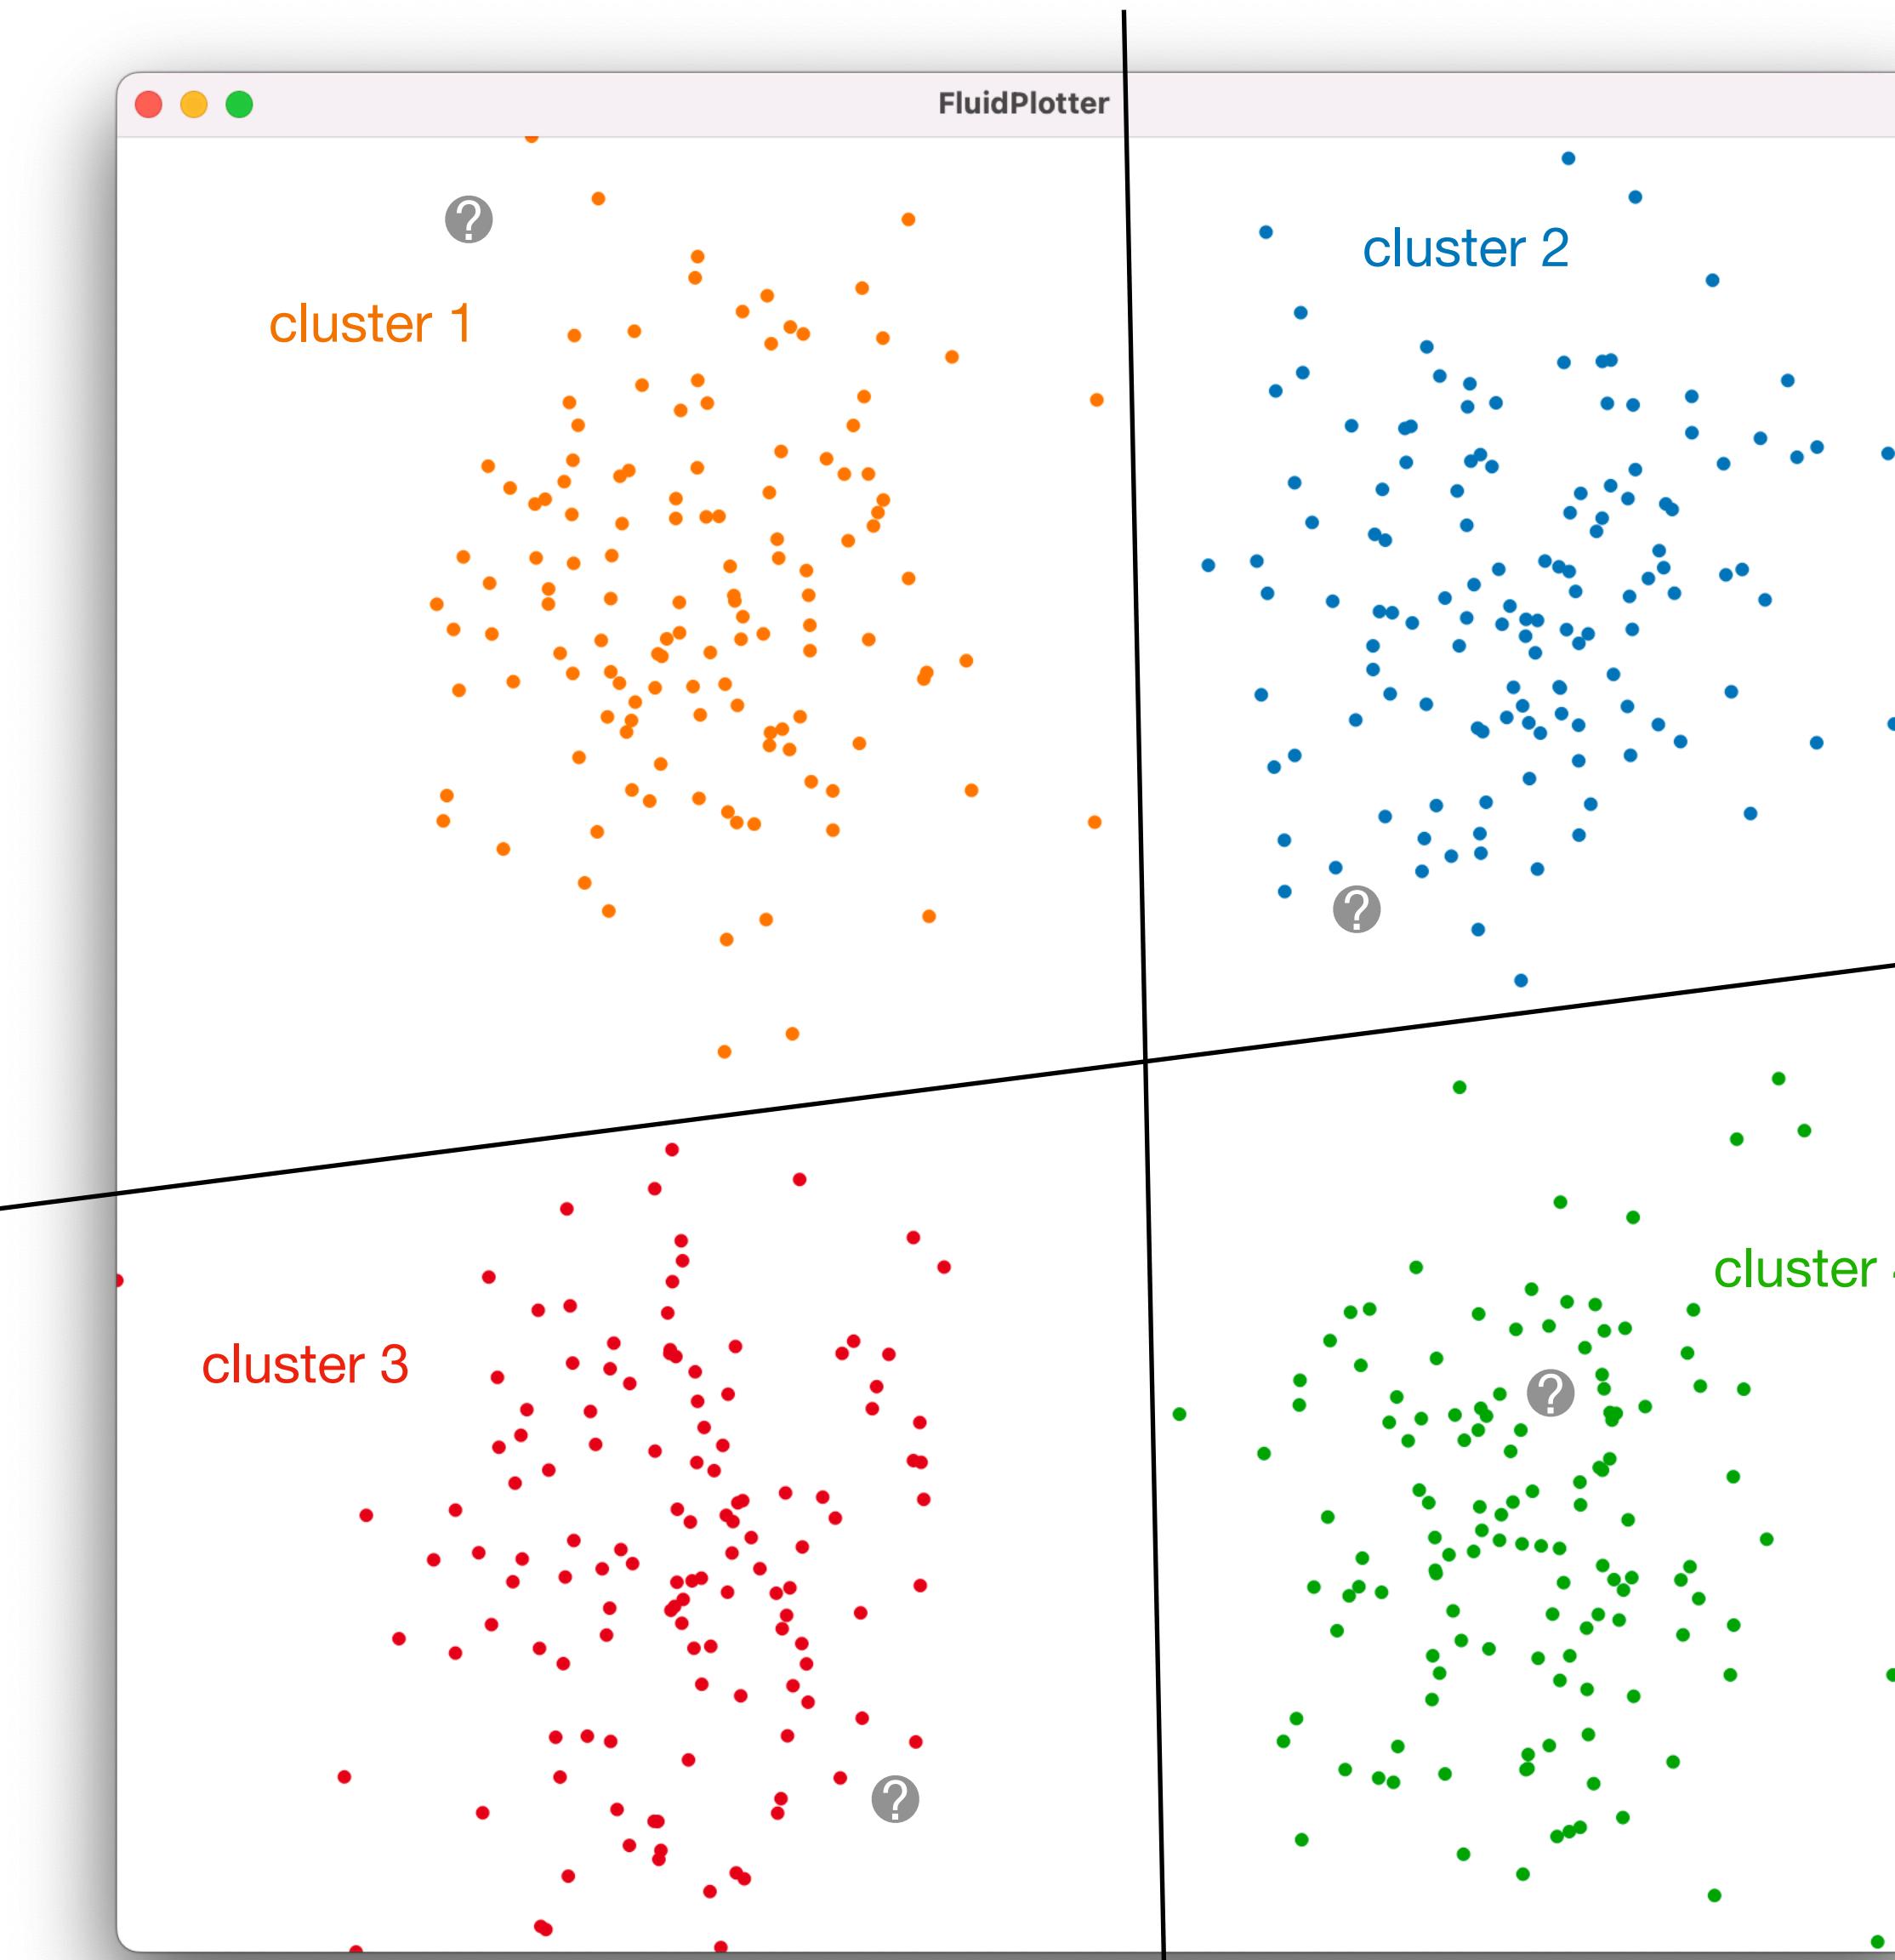


Classification
Regression

Neural Networks
KNN

Unsupervised Learning

Learning/finding patterns and relationships in data that is not labeled



Dimensionality Reduction

Clustering

Feature Learning



Fluid Corpus Manipulation

More help / discussion at [our forum](#)

Check out the [learn platform](#)

Tutorials

Objects

Tasks

Inspiration

Examples

We've worked toward making a set of concise tutorials that cover a very wide range of objects and workflows found in the FluCoMa toolkit. We think that if you at least watch these three tutorials, you'll be familiar with most of the objects and be able to start building your own patches and systems out of the toolkit. Select a tutorial to explore below.

Select a tutorial...

Building a 2D Corpus Explorer

Classifying Sounds with a Neural Network

Controlling a Synth with a Neural Network

Be aware, some of the patches you can open from within the package will differ slightly from the tutorial video's final result.

Imagine you have a large corpus of sounds that you've collected from a studio session, some outside sound walks or experimenting with a synthesiser on a rainy afternoon. This tutorial teaches you how to build a '2D Corpus Explorer', a patch that will enable you to interrogate and listen to those sounds in a structured manner. The end result is similar to CataRT and AudioStellar in that small segments of your corpus sounds are analysed and mapped to a two-dimensional space that can be explored using your mouse. It will cover topics such as segmentation, analysis, data processing and visualisation.

[YouTube Video](#)

[Open Patch](#)



Fluid Corpus Manipulation

More help / discussion at [our forum](#)

Check out the [learn platform](#)

Tutorials

Objects

Tasks

Inspiration

Examples

Select a category...

Slice Audio	Transform Audio
Analyse Audio	Analyse Data
Decompose Audio	Helpers and Utilites

If you have some corpus data, you'll probably want to analyse, manipulate and transform it. A variety of objects fall under this category and help you to make different manoeuvres around your data.

Containers

[fluid.dataset~](#)

[fluid.labelset~](#)

Statistical Analysis

[fluid.bufstats~](#)

[fluid.stats](#)

Scaling and Preprocessing

[fluid.normalize~](#)

[fluid.standardize~](#)

[fluid.robustscale~](#)

Searching and Querying

[fluid.kdtree~](#)

[fluid.datasetquery~](#)

Unsupervised Machine Learning

[fluid.pca~](#)

[fluid.kmeans~](#)

[fluid.skmeans~](#)

[fluid.mds~](#)

[fluid.umap~](#)

[fluid.grid~](#)

Supervised Machine Learning

[fluid.knnregressor~](#)

[fluid.knnclassifier~](#)

[fluid.mlpregressor~](#)

[fluid.mlpclassifier~](#)



Fluid Corpus Manipulation

More help / discussion at [our forum](#)

Check out the [learn platform](#)

Tutorials

Objects

Tasks

Inspiration

Examples

A major part of the FluCoMa project involved commissioning composers to create works using the toolkit. These works are recorded and have also been investigated in detail by Jacob Hart (one of the team members) in "explore" articles. These long-form articles draw out how objects were used, and he has even created patches for you to plunder for your own creative endeavours. This is a great place to start if you're searching for musical inspiration with the FluCoMa Toolkit.

Select a composer(s)...

Lauren Sarah Hayes

Leafcutter John

Olivier Pasquet

Rodrigo Constanzo

Alex Harker

Sam Pluta

Richard Devine

Hans Tutschku

Alice Eldridge & Chris Kiefer

Lauren Sarah Hayes is a Scottish improviser and sound artist. Her music is a mix of experimental pop/live electronics/techno/noise/free improvisation and has been described as 'voracious' and 'exhilarating'. She is a sculptress of sound, manipulating, remixing, and bending voice, drum machines, analogue synths and self-built software live and physically. She is excited by what can happen in the vulnerable relationships between sound, space, and audience. Her shows are highly physical, making the performance of live electronic music more engaging for audiences. Over the last decade she has developed and honed a deliberately challenging and unpredictable performance system that explores the relationships between bodies, sound, environments, and technology.

YouTube Video

Learn Article

Fluid Corpus Manipulation Toolkit

The FluCoMa toolkit to analyse, transform and learn from sounds

Description

The Fluid Corpus Manipulation toolkit provides an open-ended, loosely coupled set of objects for experimental learning.

Almost all objects for audio analysis or transformation have audio-rate and buffer-based versions, a

Many useful examples can be found in the help files as well as in the example folder, which is here:

```
| File.realpath(FluidFilesPath("../..../Examples")).openOS;
```

Contents

- Slice Audio
- Analyse Audio
- Decompose Audio
- Transform Audio
- Analyse Data
- Helpers

Slice Audio

on signals	on buffers	digest
FluidAmpGate	FluidBufAmpGate	Events from amplitude envelope
FluidAmpSlice	FluidBufAmpSlice	Onsets from amplitude envelope
FluidOnsetSlice	FluidBufOnsetSlice	Spectral onset detector
FluidTransientSlice	FluidBufTransientSlice	Transient model onset detector
FluidNoveltySlice	FluidBufNoveltySlice	Novelty based onset detection on a choice of des

Analyse Audio

on signals	on buffers	digest
FluidPitch	FluidBufPitch	Choice of pitch descriptors
FluidLoudness	FluidBufLoudness	Loudness Descriptor

Plan for the Week...



MLPRegressor

control *many* synthesizer
parameters from a smaller
control space



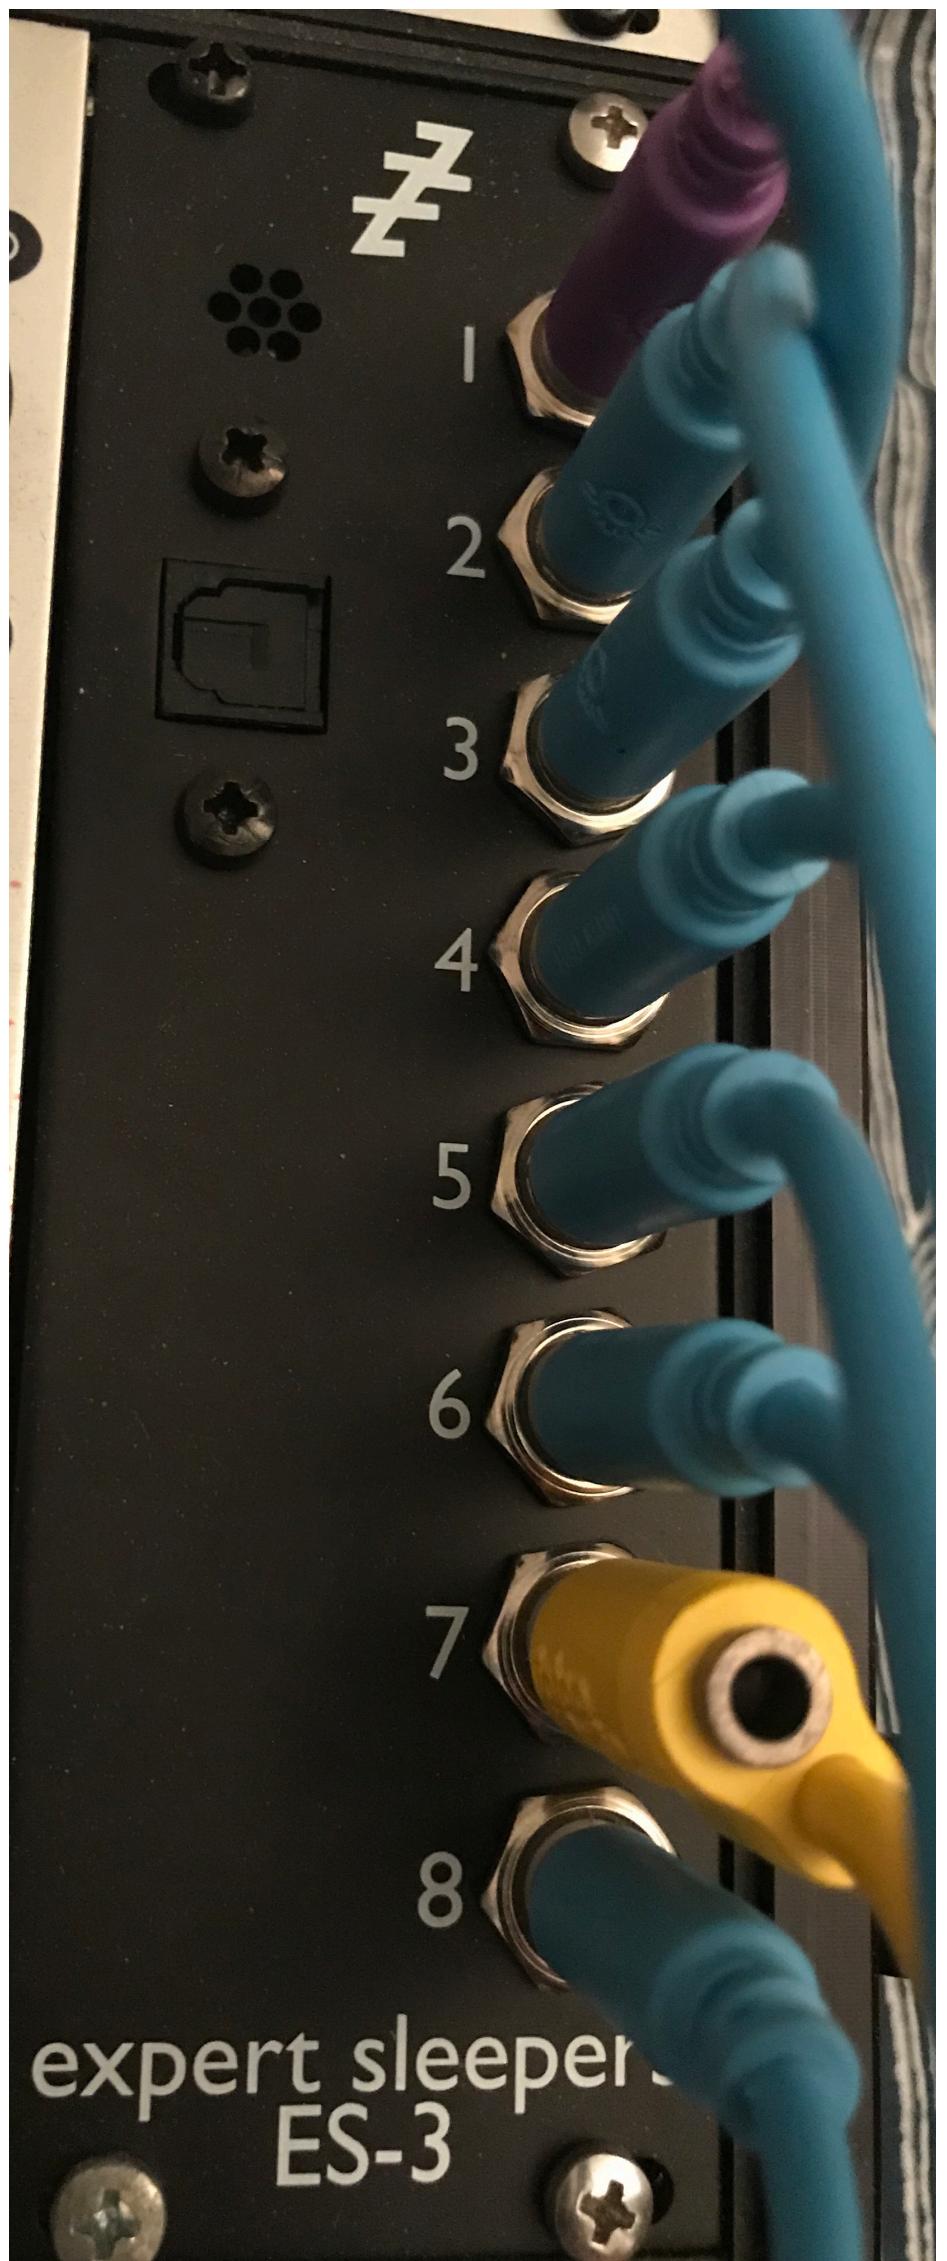
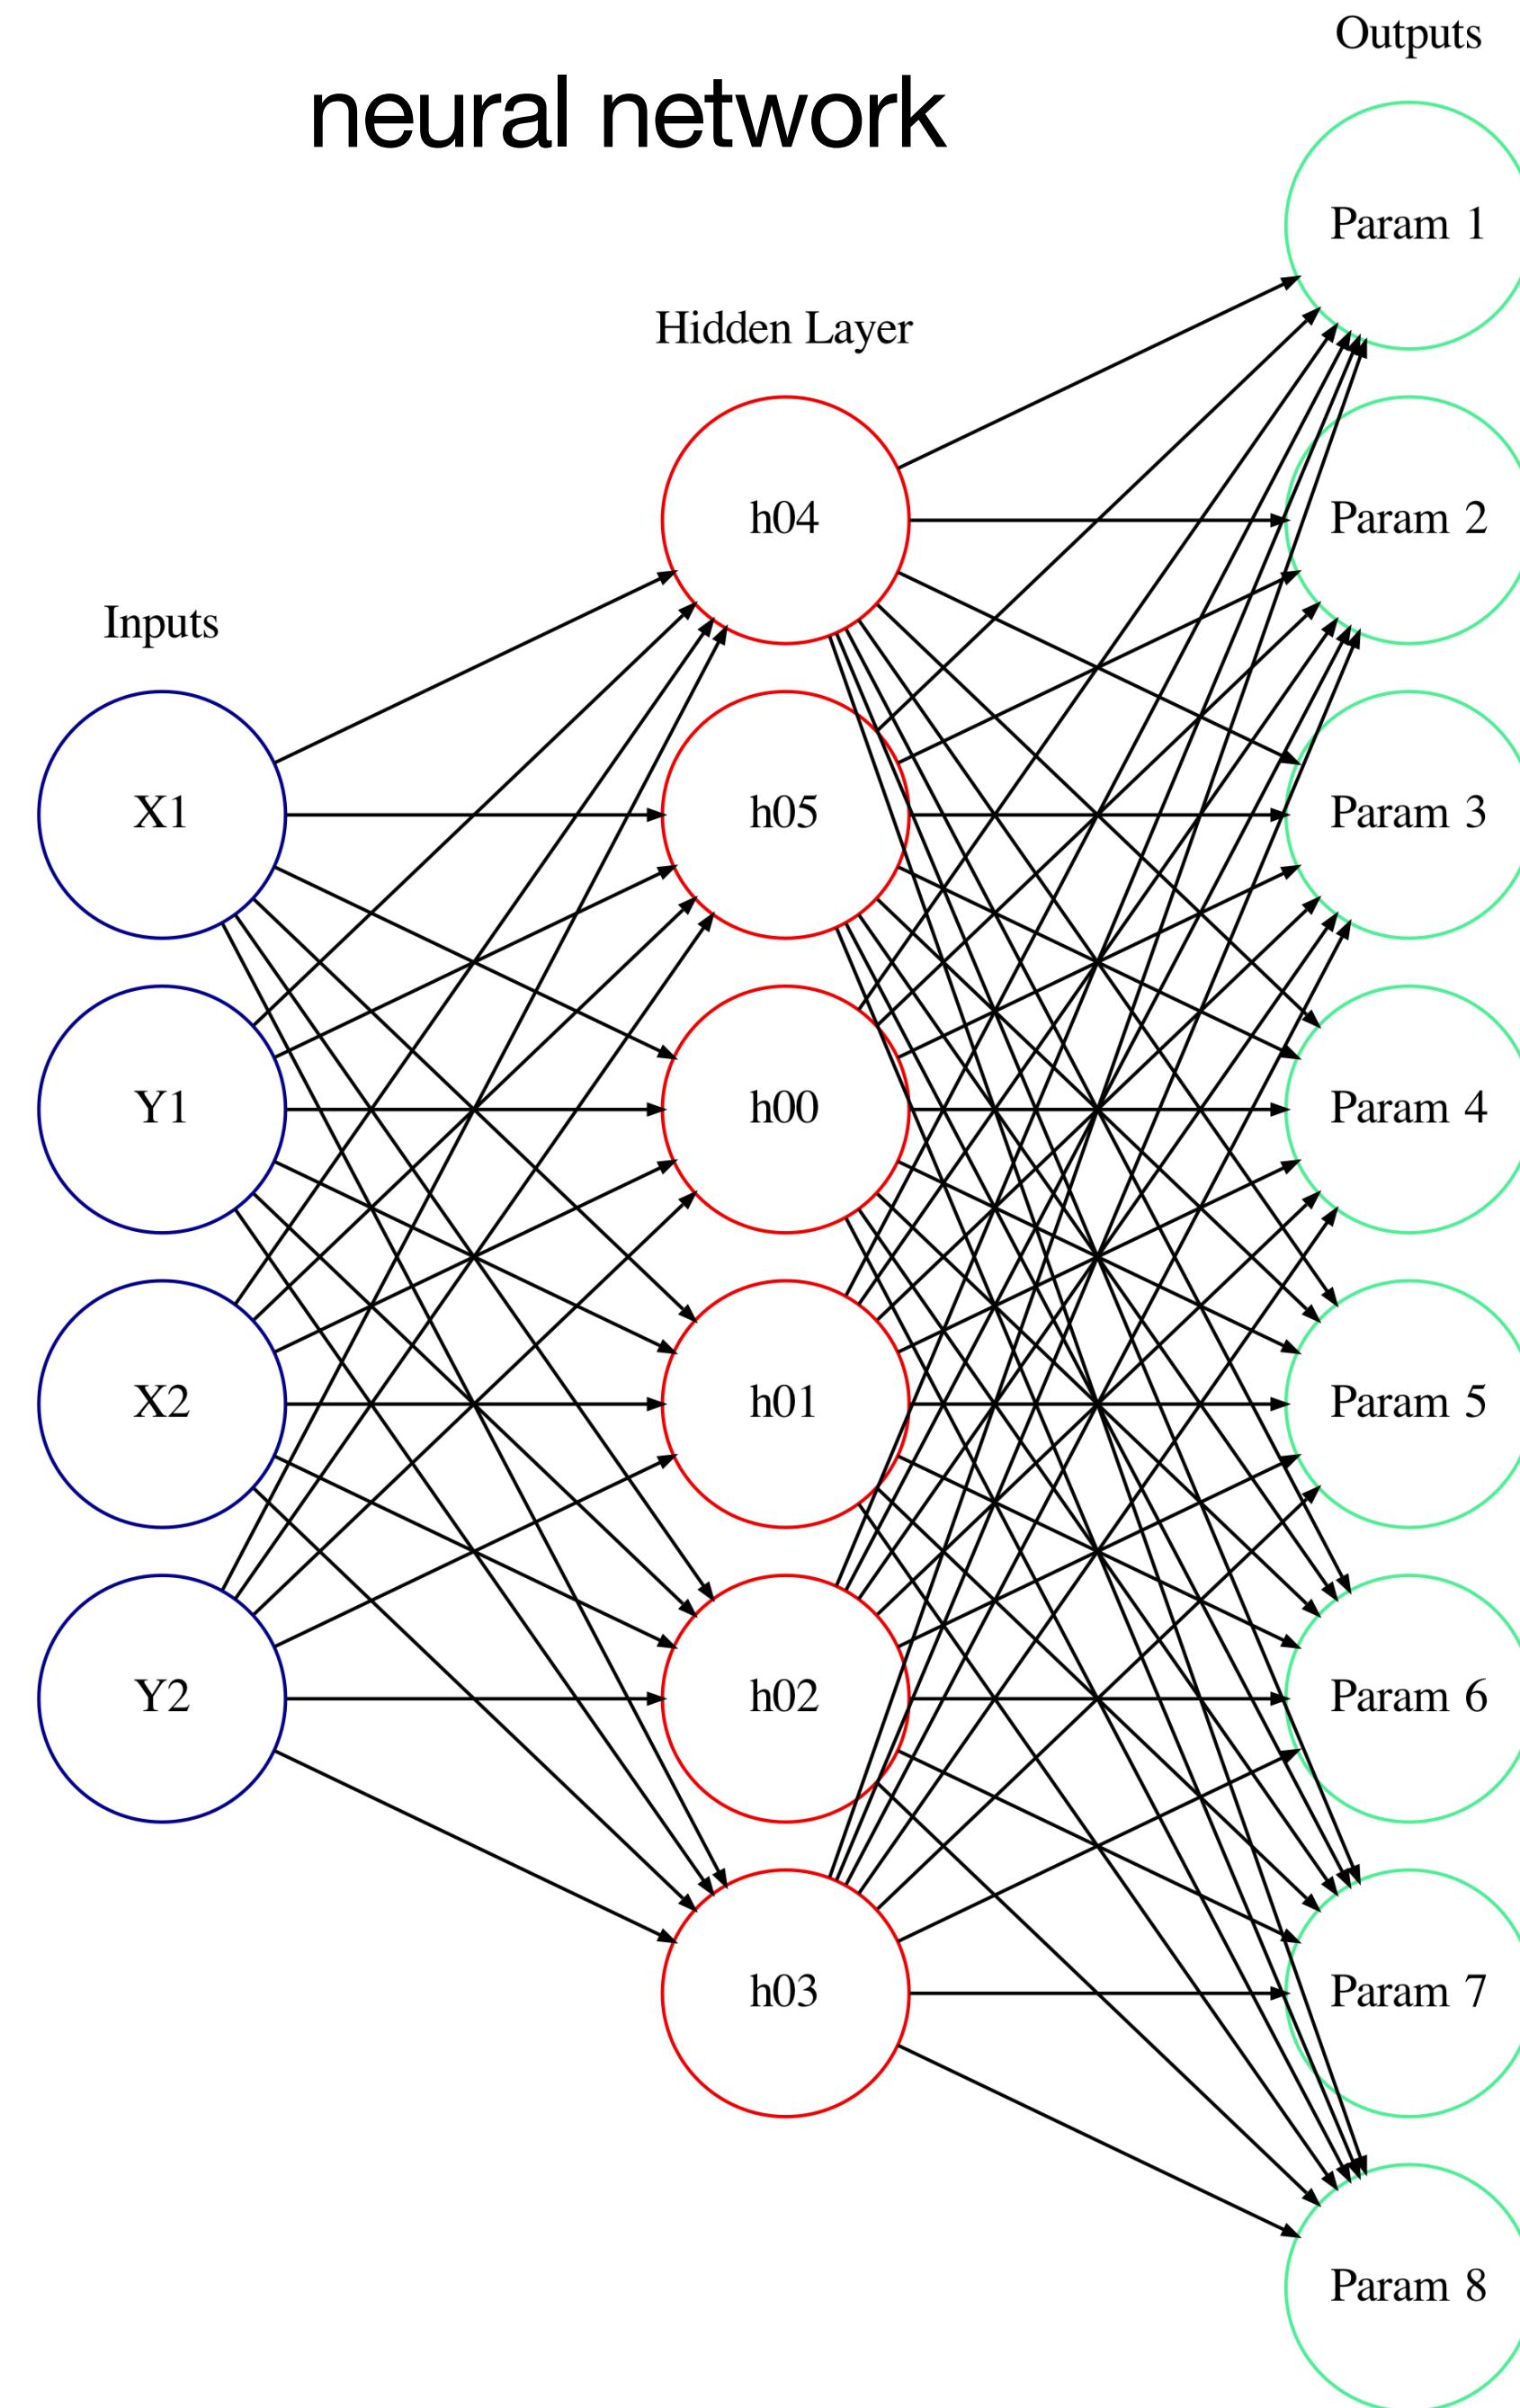
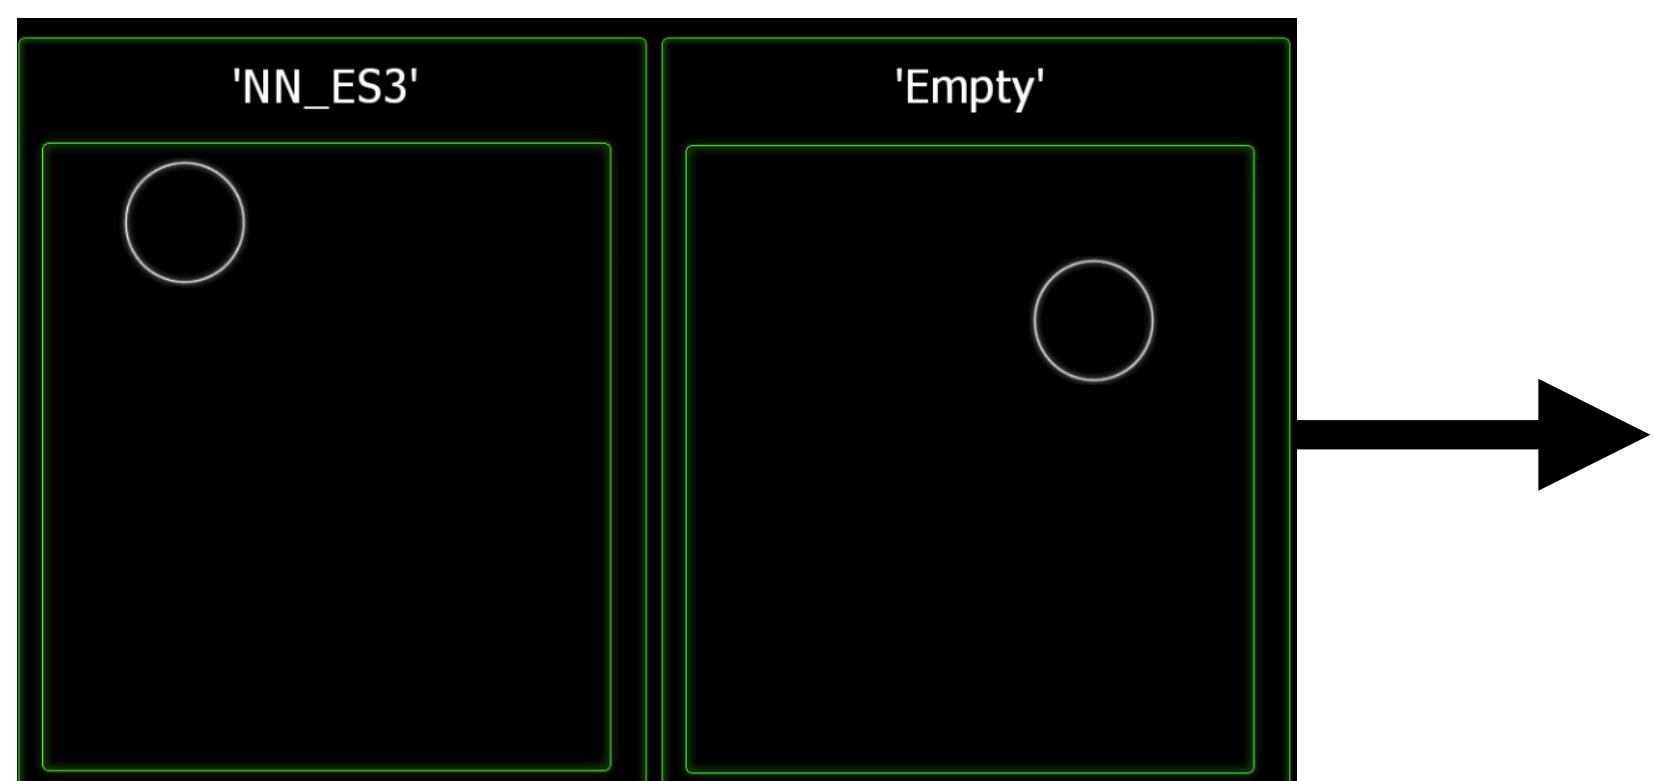
fluid.mlpregressor~







neural network



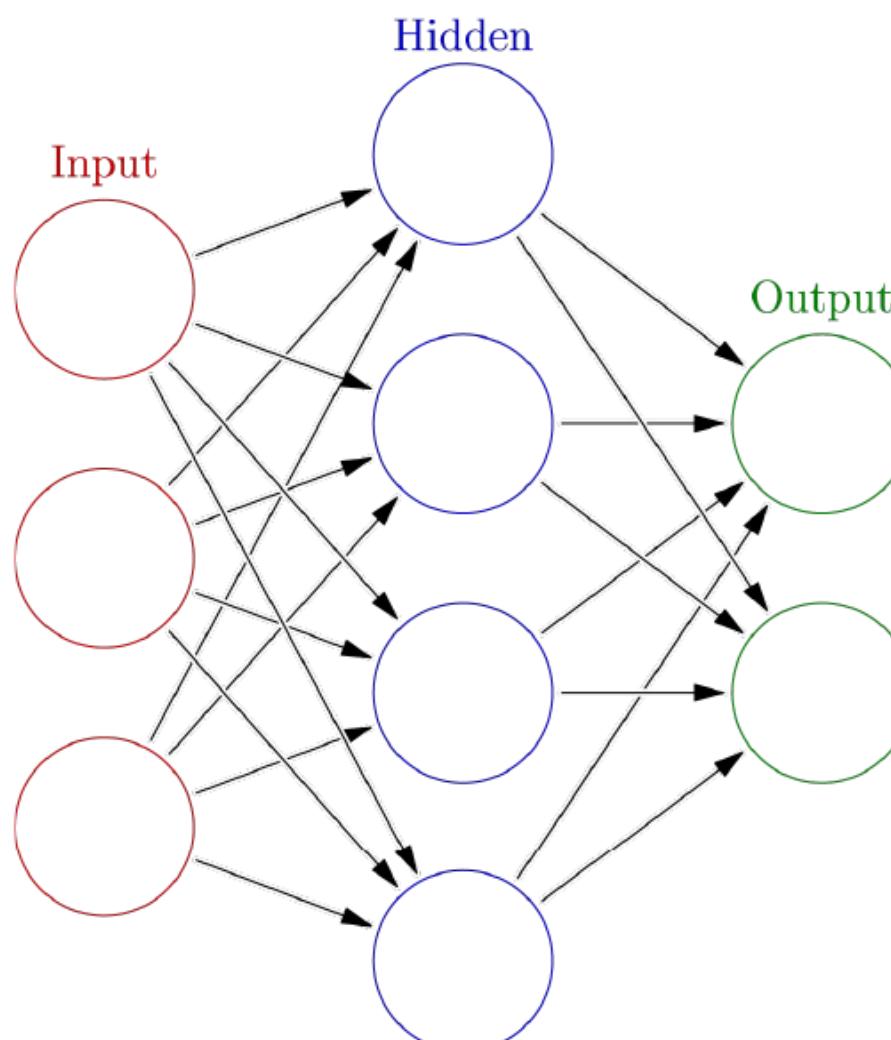
Neural Network Training a Regressor

identifier

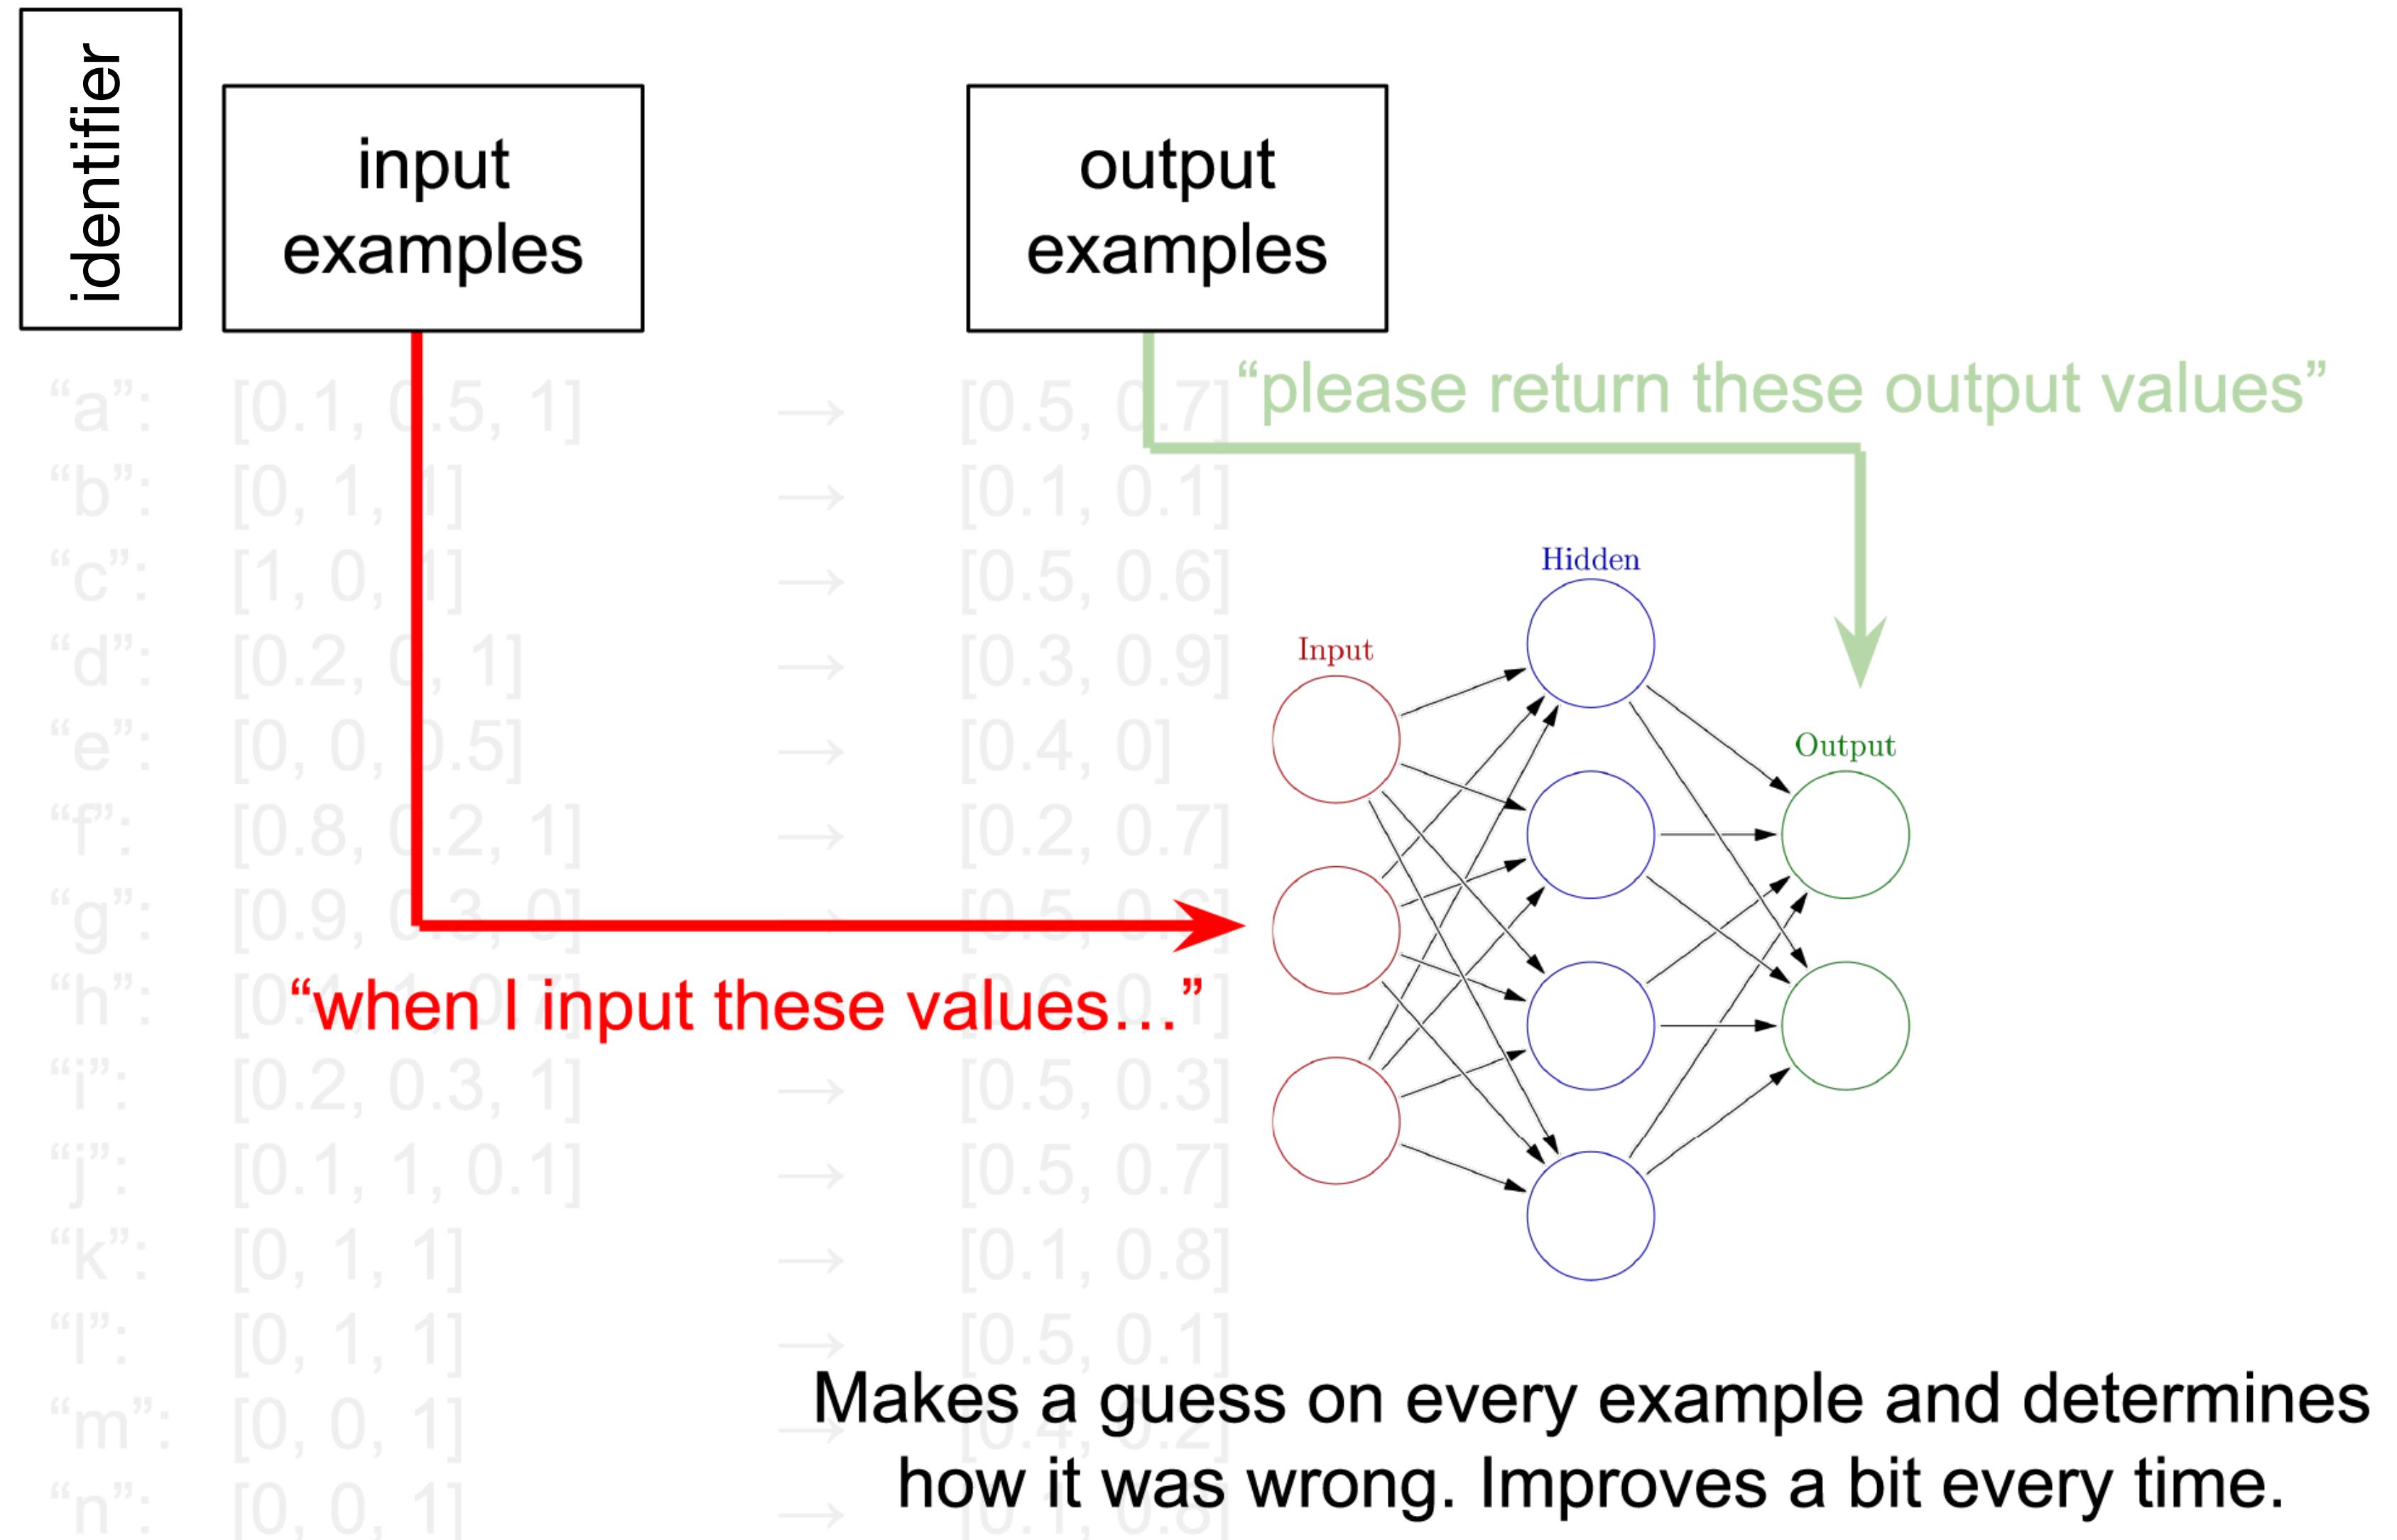
input
examples

“a”:	[0.1, 0.5, 1]	→	[0.5, 0.7]
“b”:	[0, 1, 1]	→	[0.1, 0.1]
“c”:	[1, 0, 1]	→	[0.5, 0.6]
“d”:	[0.2, 0, 1]	→	[0.3, 0.9]
“e”:	[0, 0, 0.5]	→	[0.4, 0]
“f”:	[0.8, 0.2, 1]	→	[0.2, 0.7]
“g”:	[0.9, 0.3, 0]	→	[0.5, 0.6]
“h”:	[0.4, 1, 0.7]	→	[0.6, 0.1]
“i”:	[0.2, 0.3, 1]	→	[0.5, 0.3]
“j”:	[0.1, 1, 0.1]	→	[0.5, 0.7]
“k”:	[0, 1, 1]	→	[0.1, 0.8]
“l”:	[0, 1, 1]	→	[0.5, 0.1]
“m”:	[0, 0, 1]	→	[0.4, 0.2]
“n”:	[0, 0, 1]	→	[0.1, 0.8]

output
examples



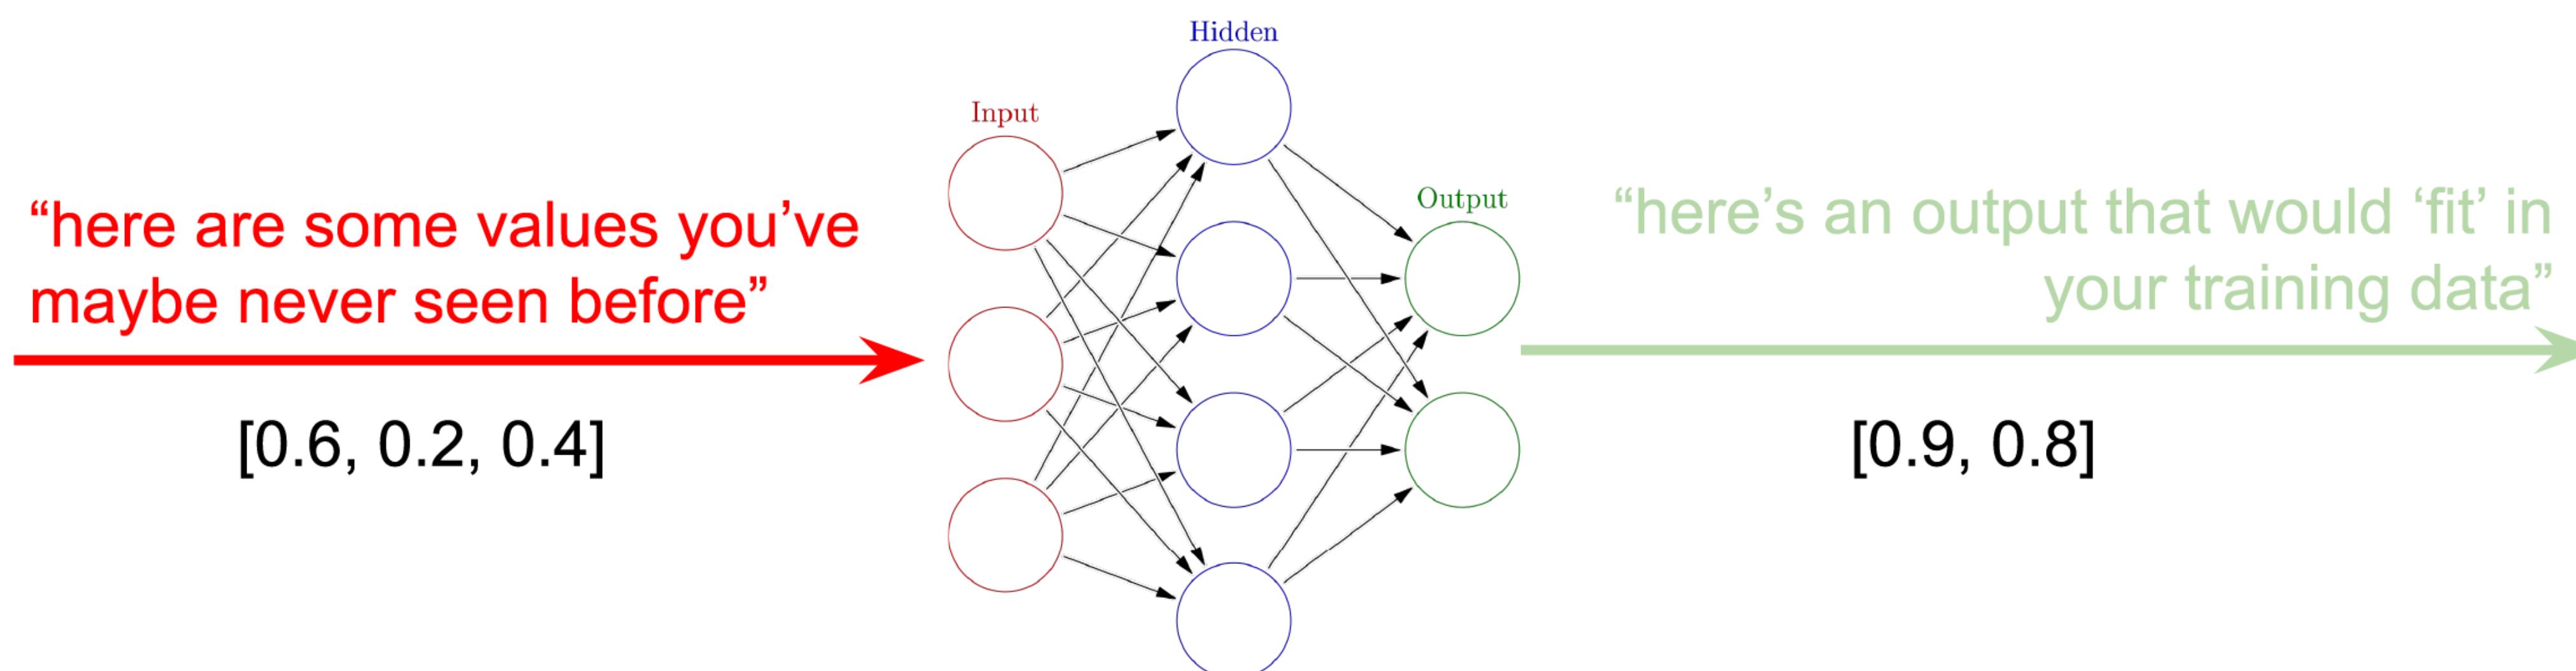
Neural Network *Training* a Regressor



Neural Network *Predicting with Regression*

input
examples

output
examples



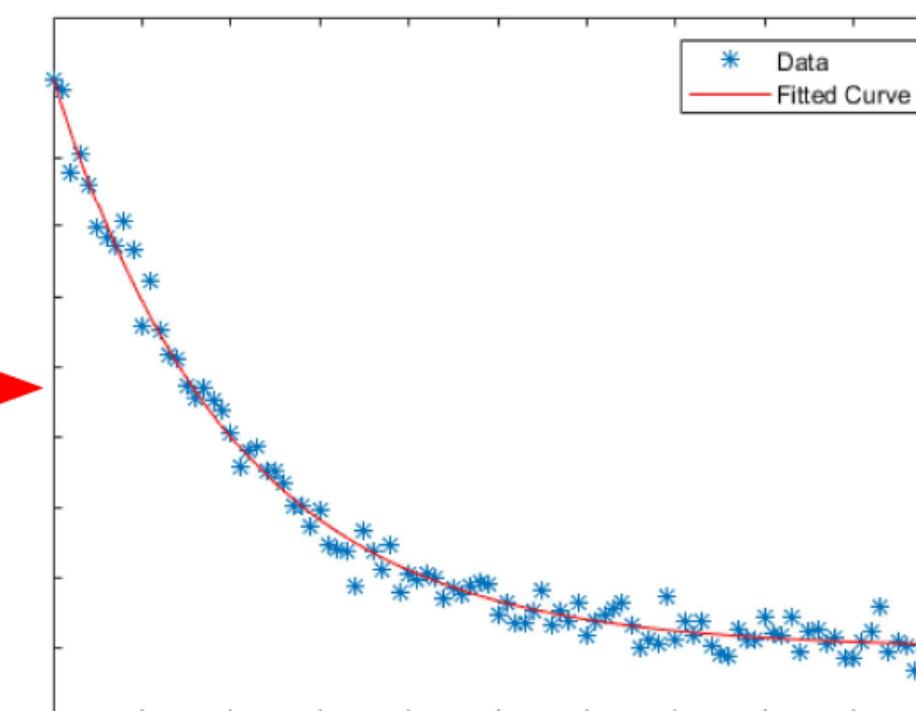
Neural Network *Predicting with Regression*

input
examples

output
examples

“here are some values you’ve
maybe never seen before”

[0.6, 0.2, 0.4]



“here’s an output that would ‘fit’ in
your training data”

[0.9, 0.8]

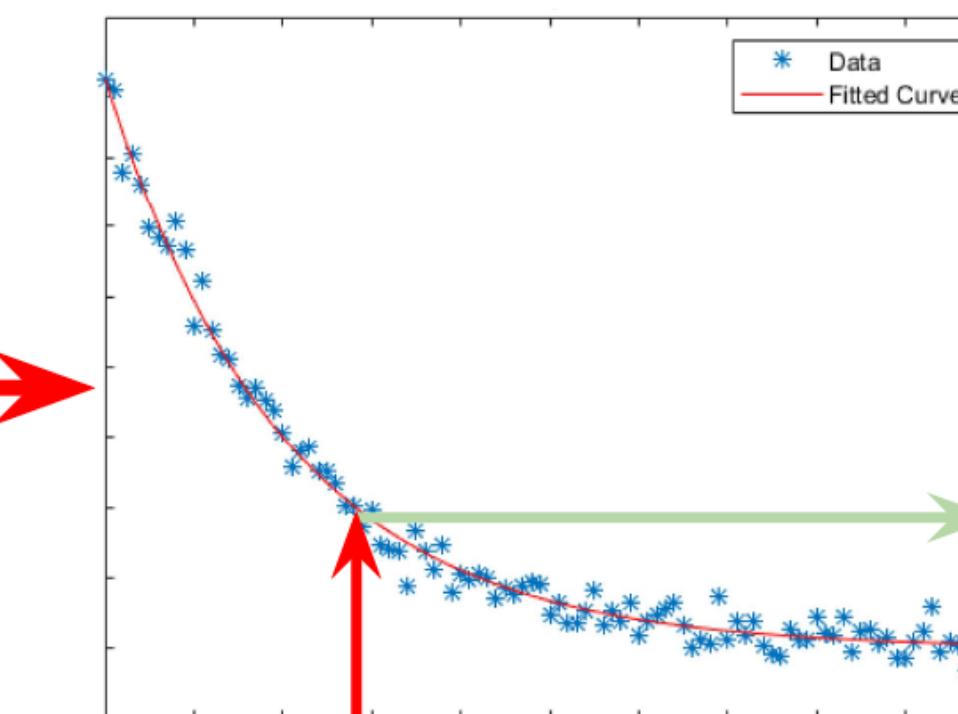
Neural Network *Predicting with Regression*

input
examples

output
examples

“here are some values you’ve
maybe never seen before”

[0.6, 0.2, 0.4]

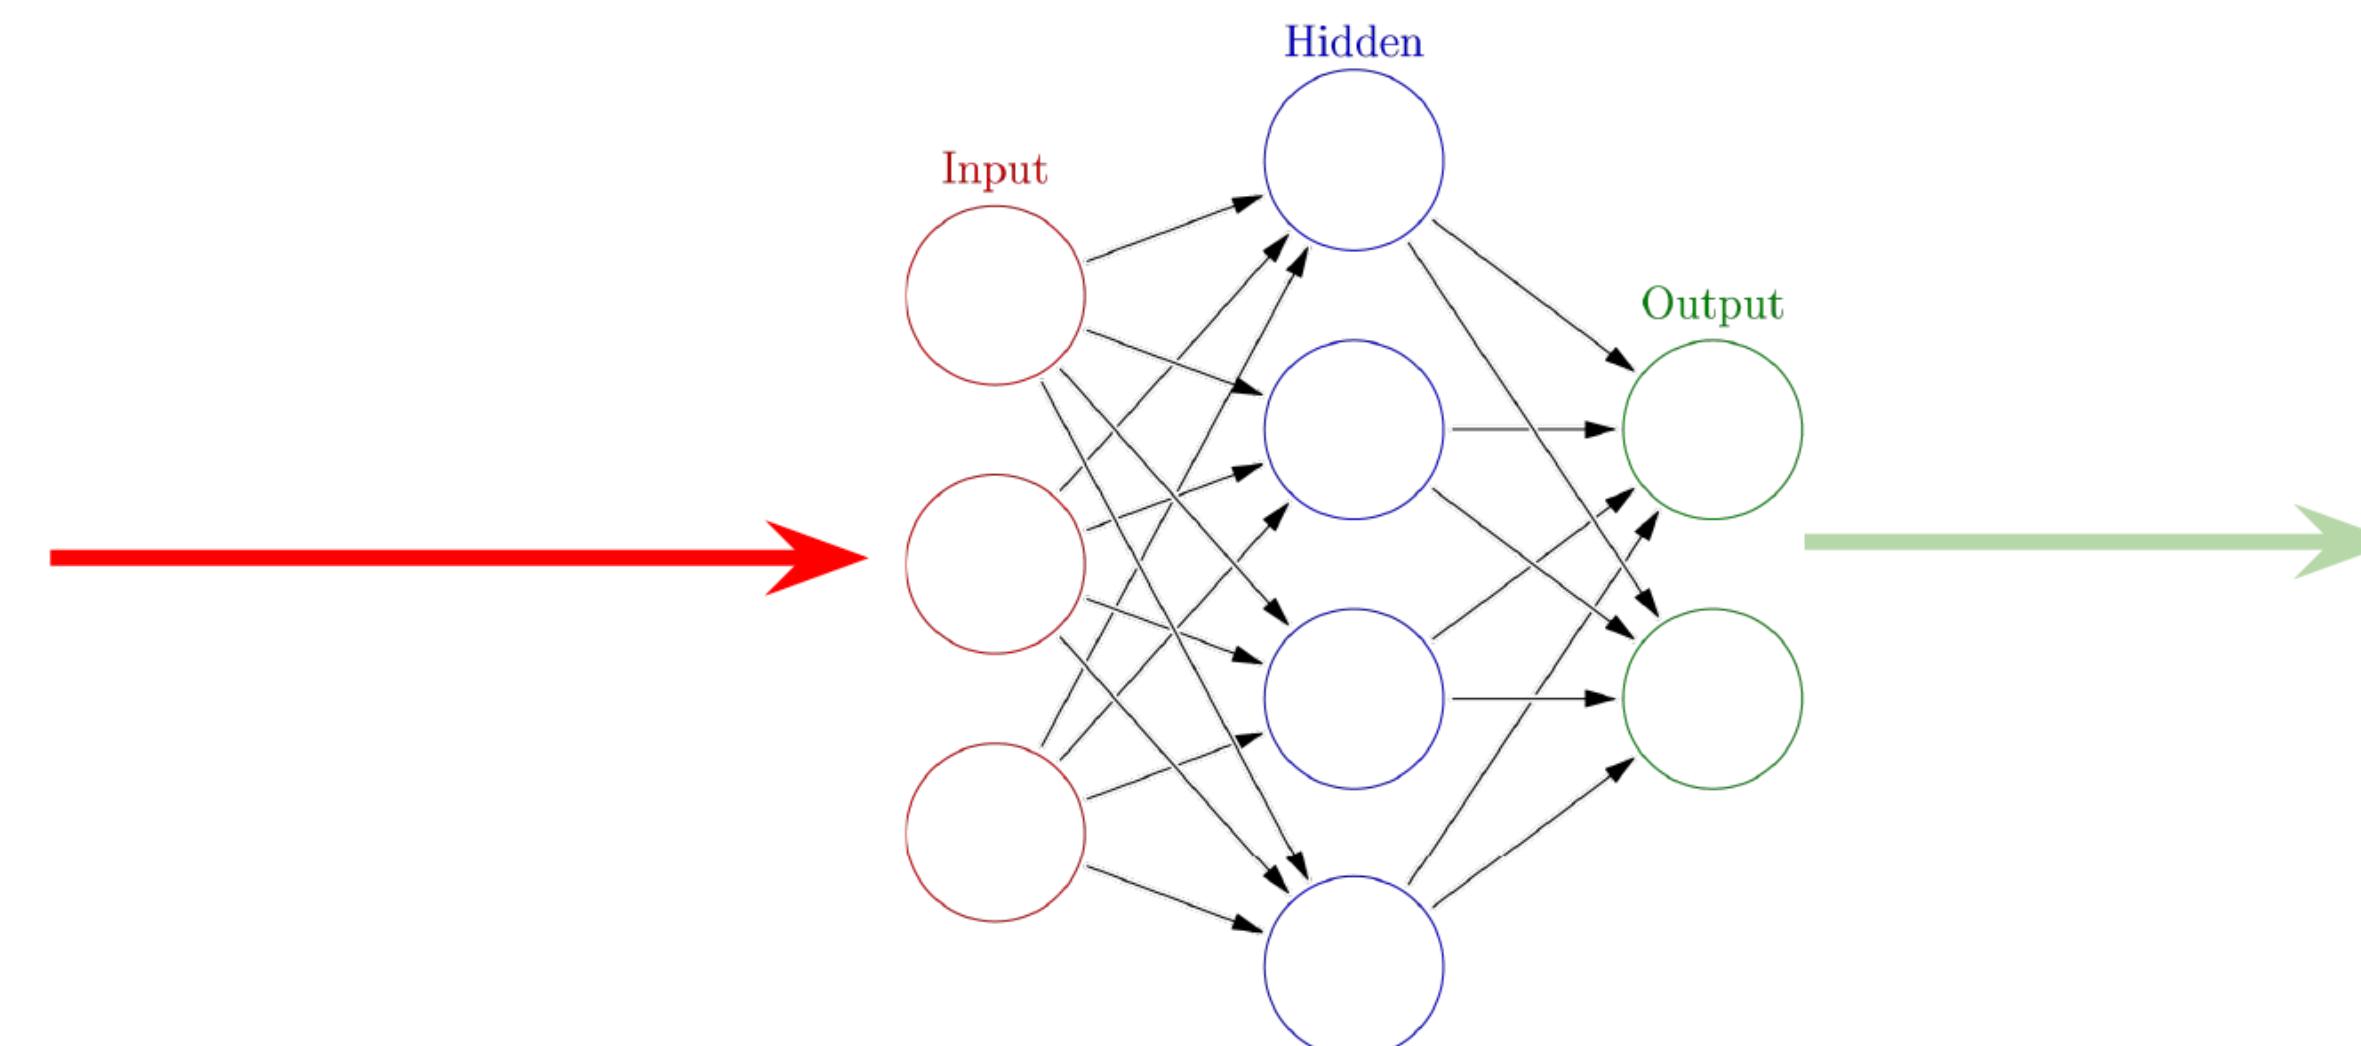


“here’s an output that would ‘fit’ in
your training data”

[0.9, 0.8]

Neural Network Predicting with Regression

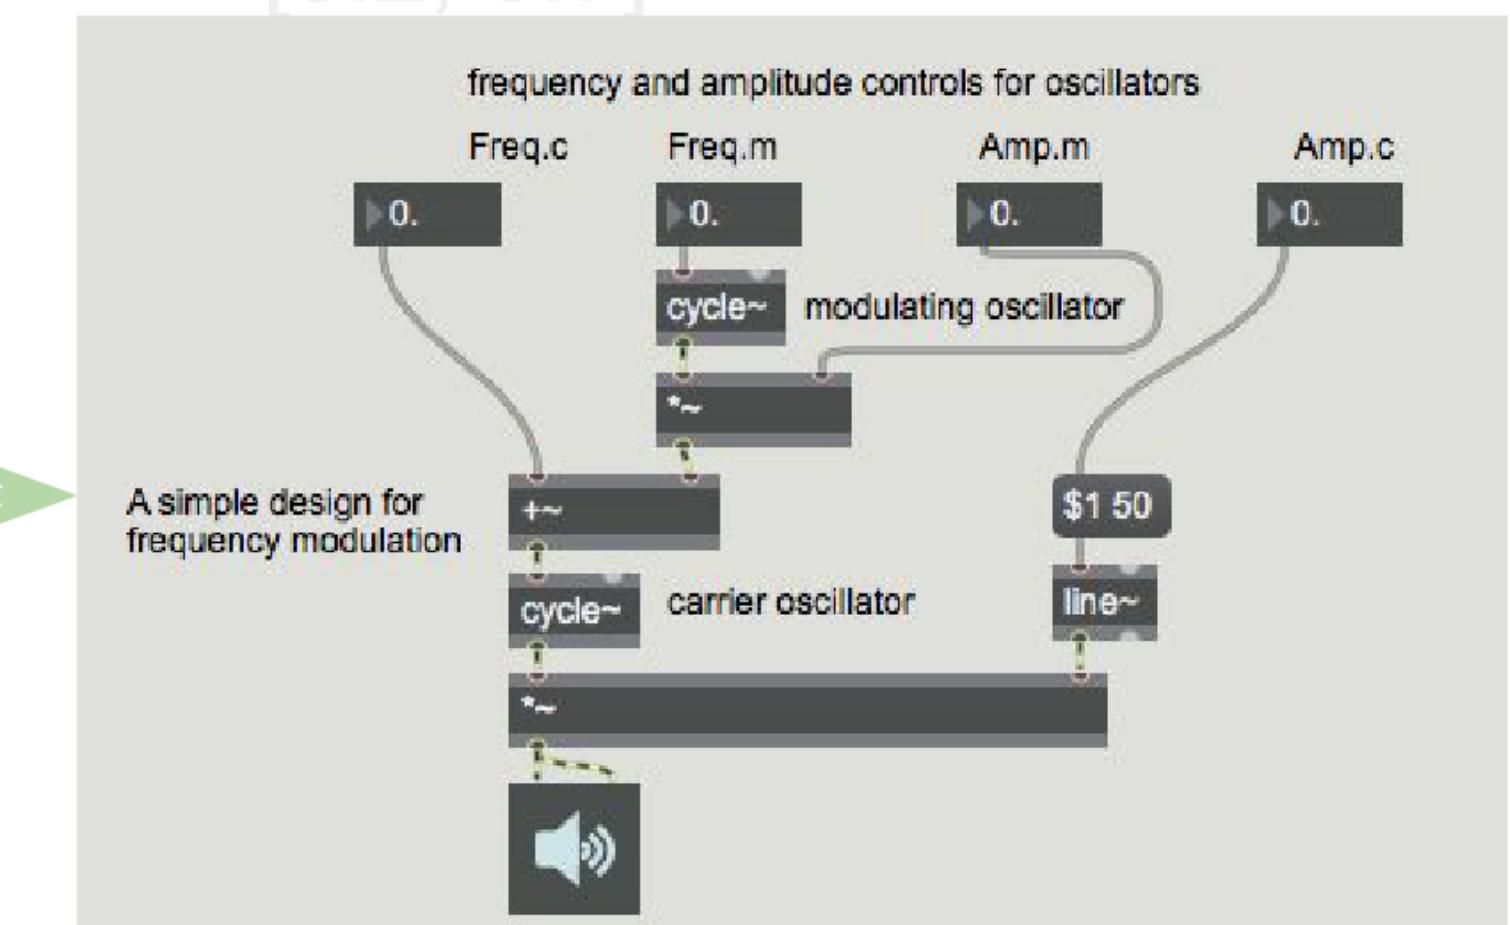
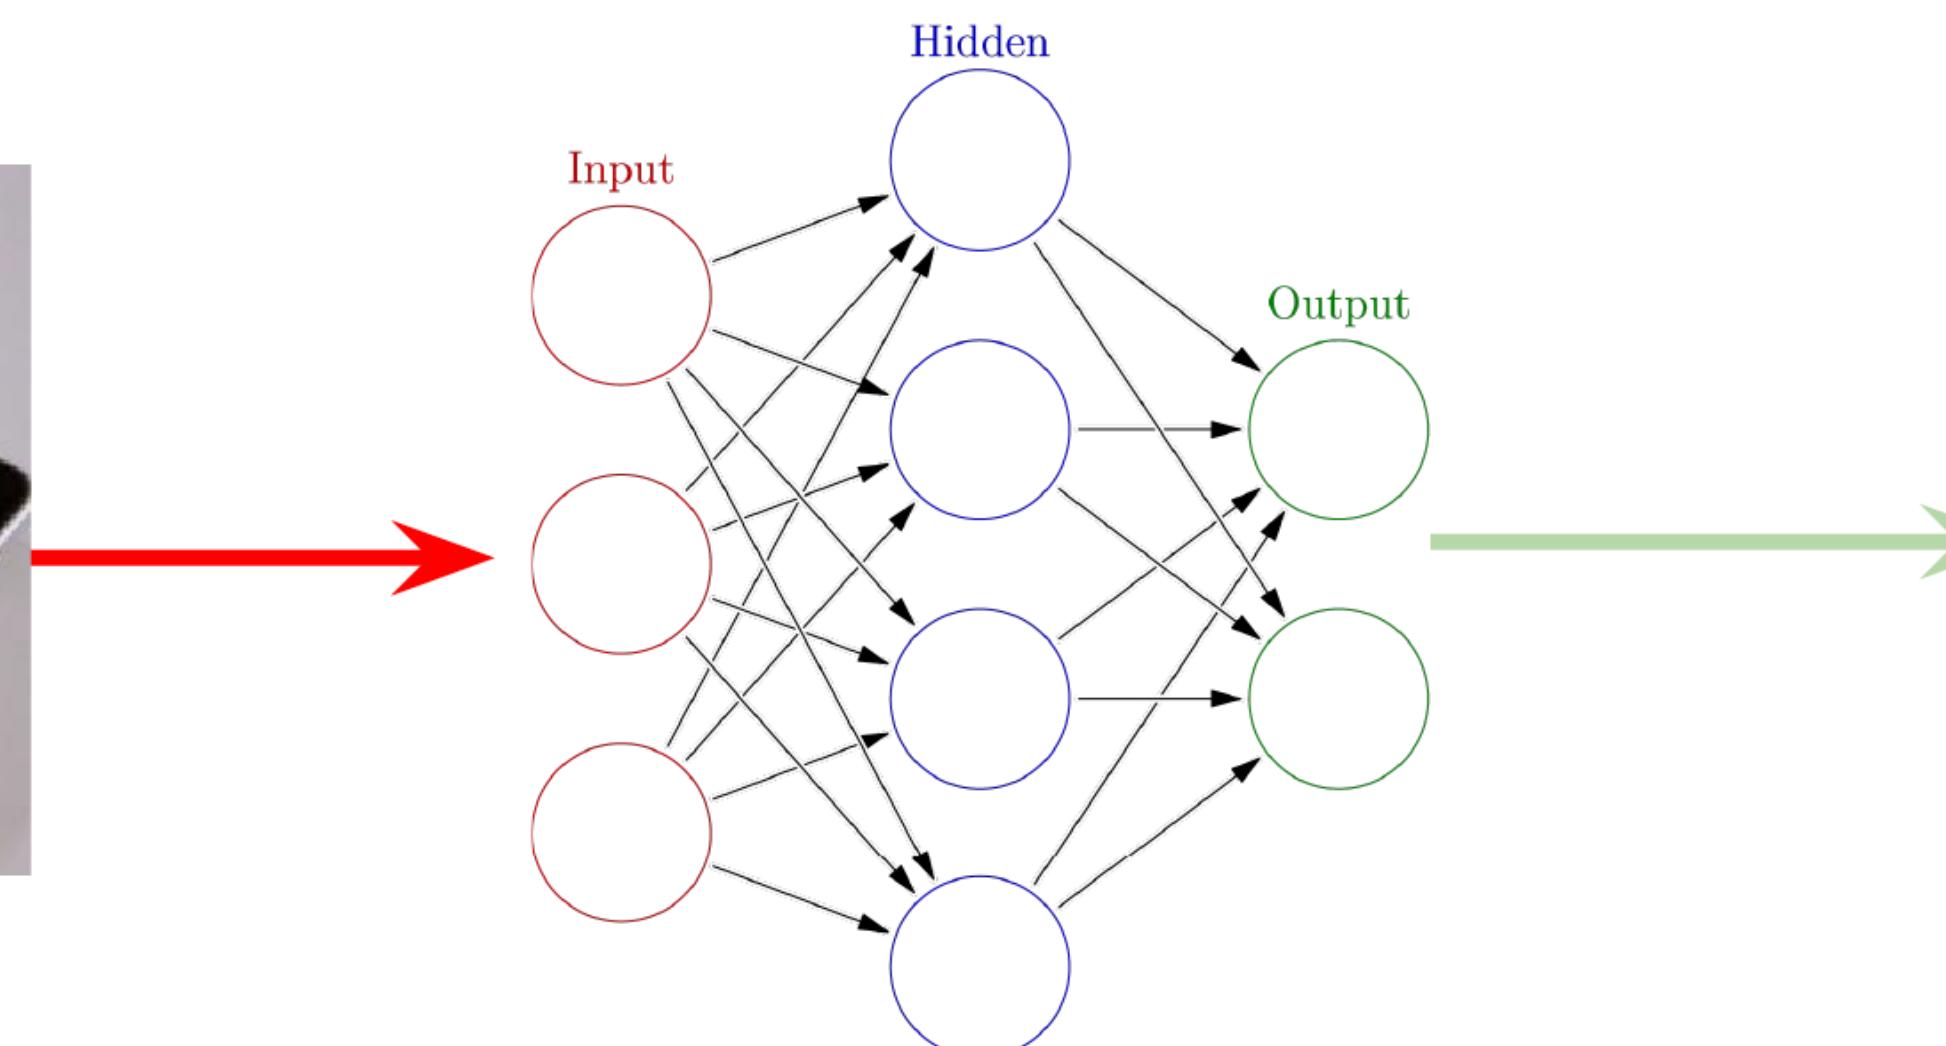
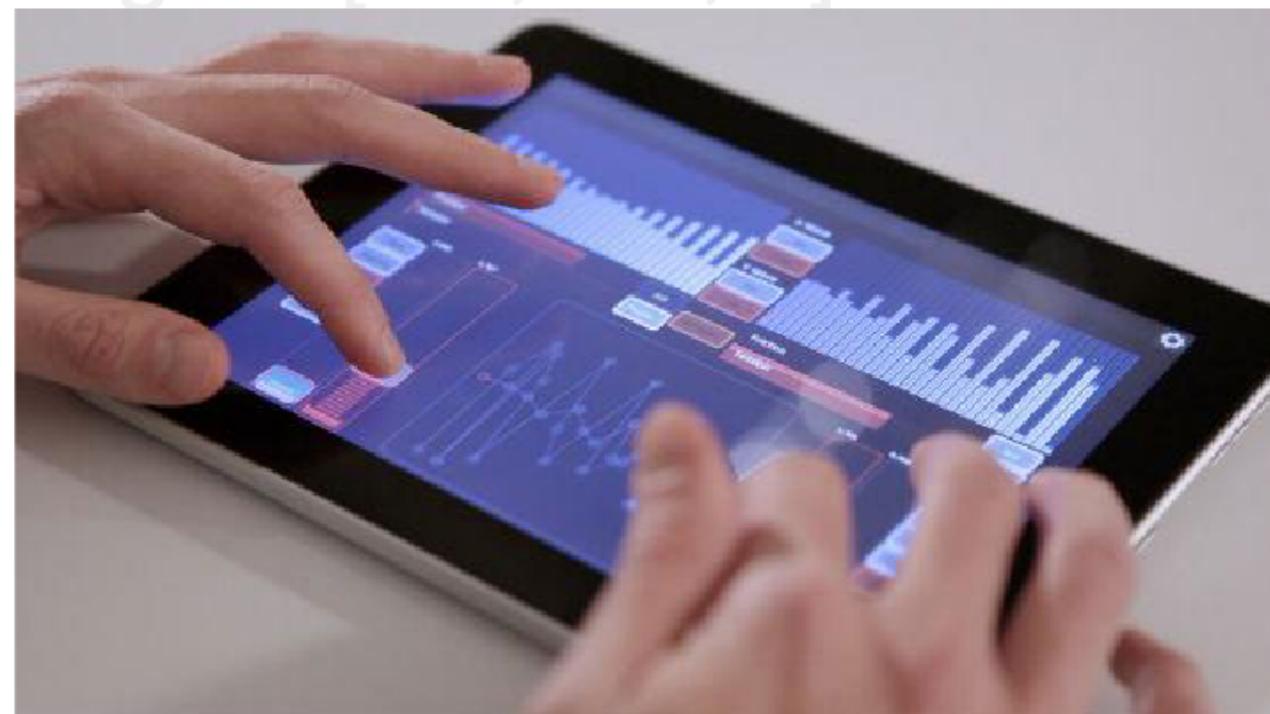
“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“g”: [0.9, 0.3, 0]
“h”: [0.4, 1, 0.7]
“i”: [0.2, 0.3, 1]
“j”: [0.1, 1, 0.1]
“k”: [0, 1, 1]
“l”: [0, 1, 1]
“m”: [0, 0, 1]
“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]



[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0]
[0.2, 0.7]
[0.5, 0.6]
[0.6, 0.1]
[0.5, 0.3]
[0.5, 0.7]
[0.1, 0.8]
[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0.3]
[0.2, 0.7]

Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“g”: [0.9, 0.3, 0]

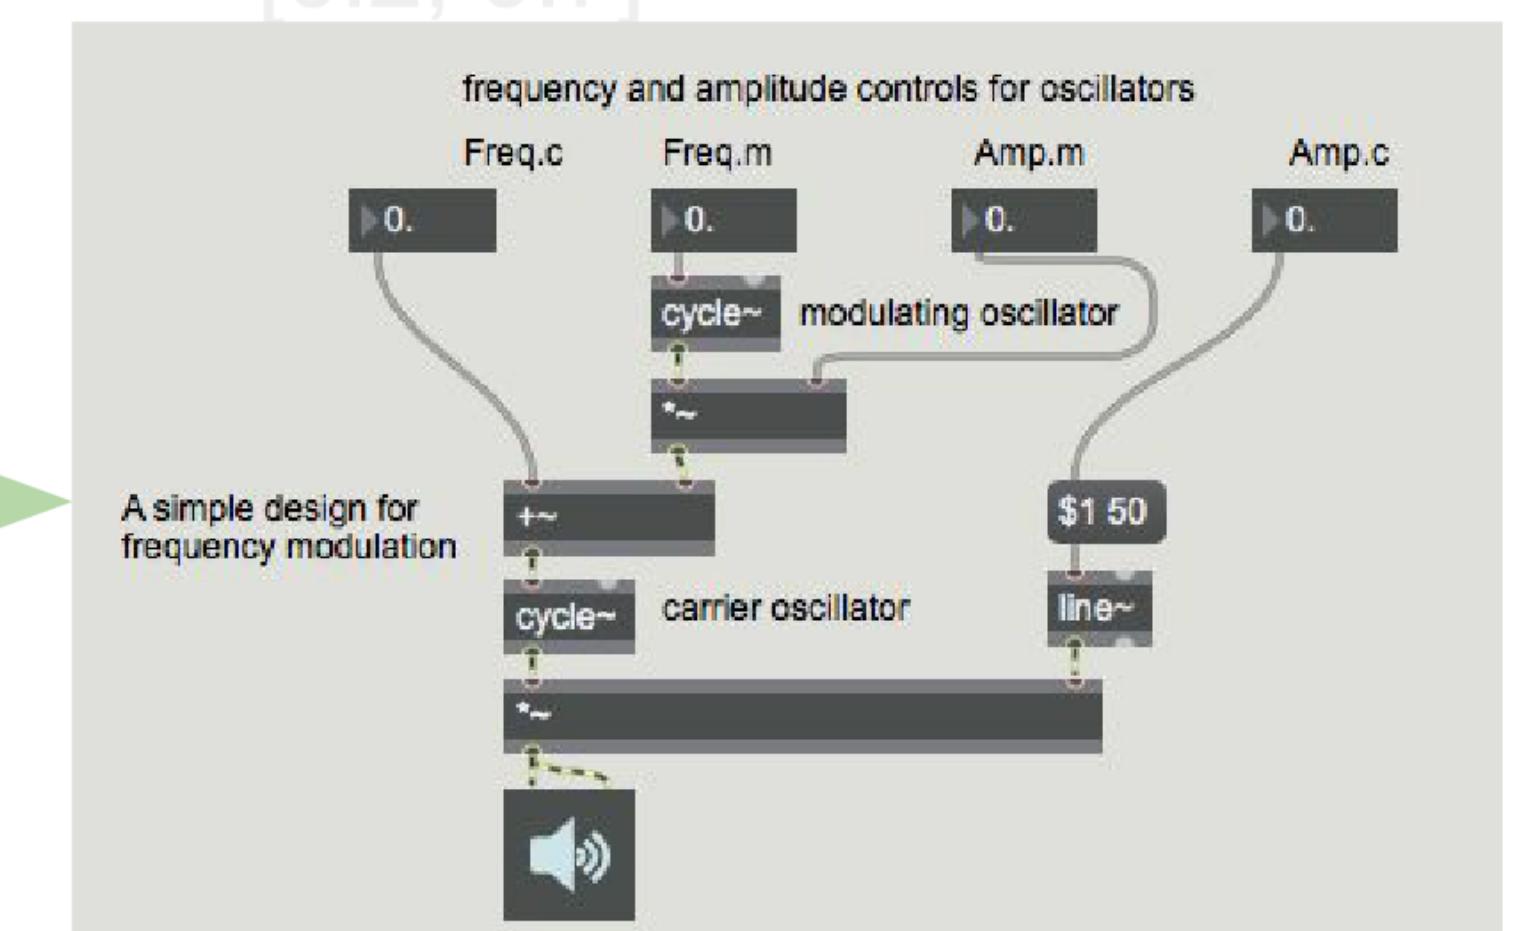
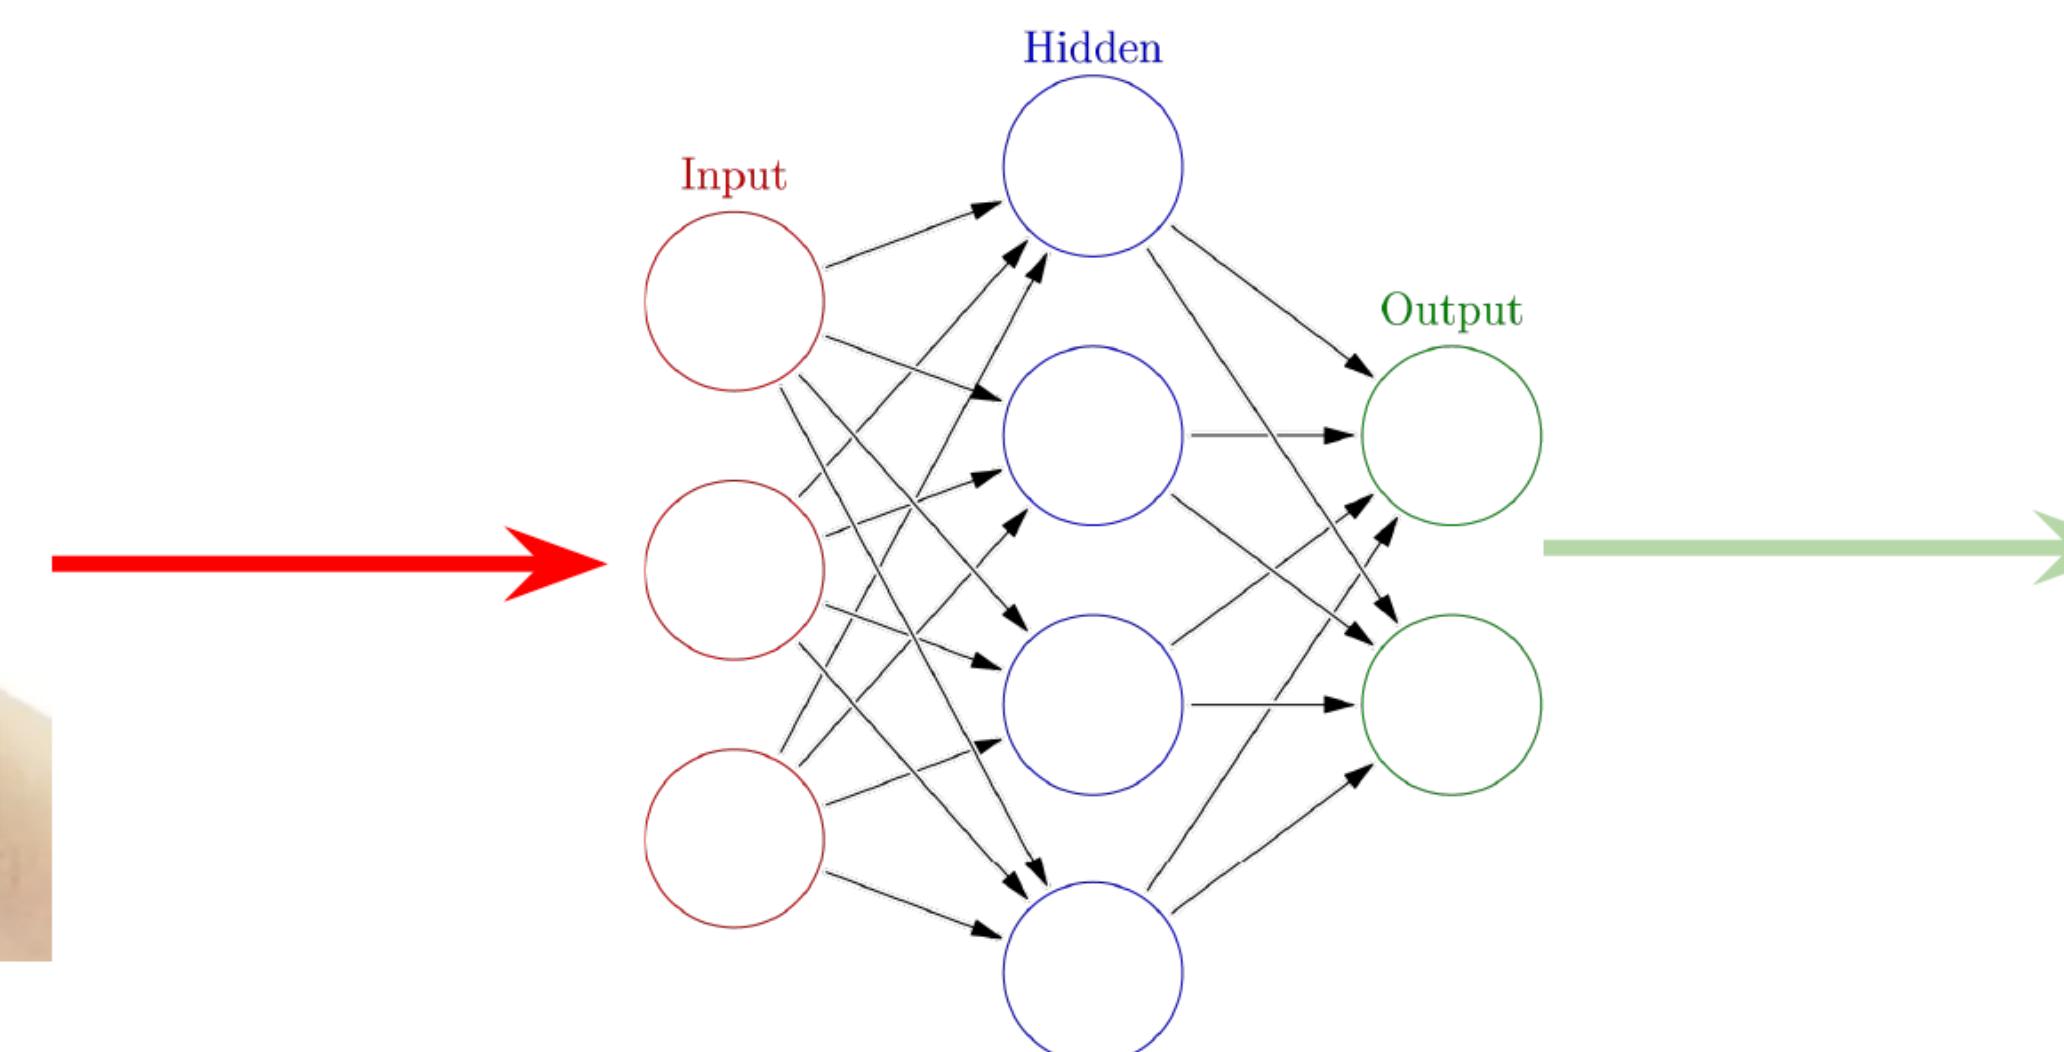
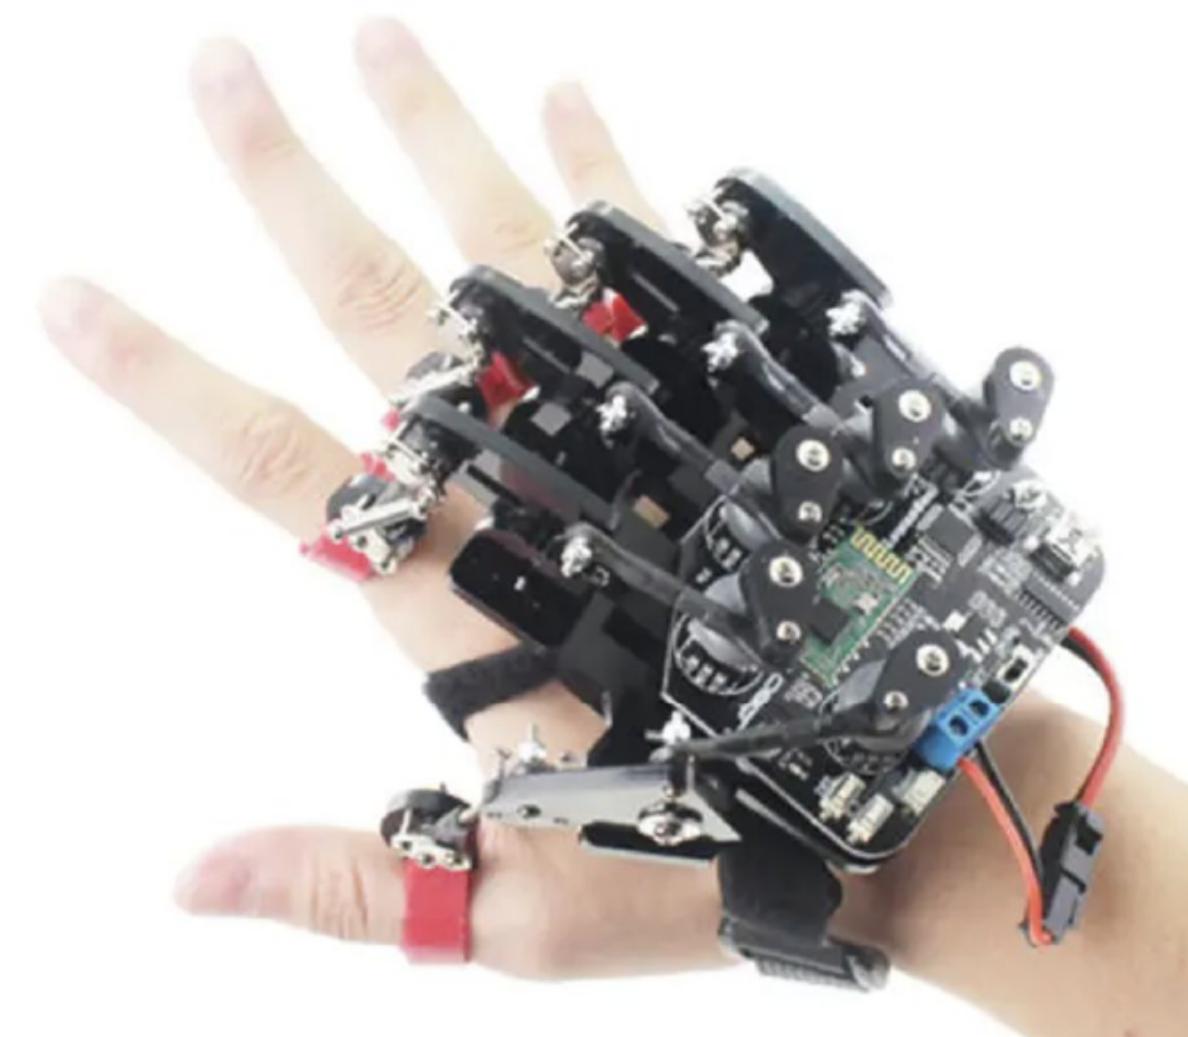


“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]

[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0]
[0.2, 0.7]

[0.5, 0.7]
[0.1, 0.8]
[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]

Neural Network Predicting with Regression



Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]

“b”: [0, 1, 1]

“c”: [1, 0, 1]

“d”: [0.2, 0, 1]

“e”: [0, 0, 0.5]

“f”: [0.8, 0.2, 1]



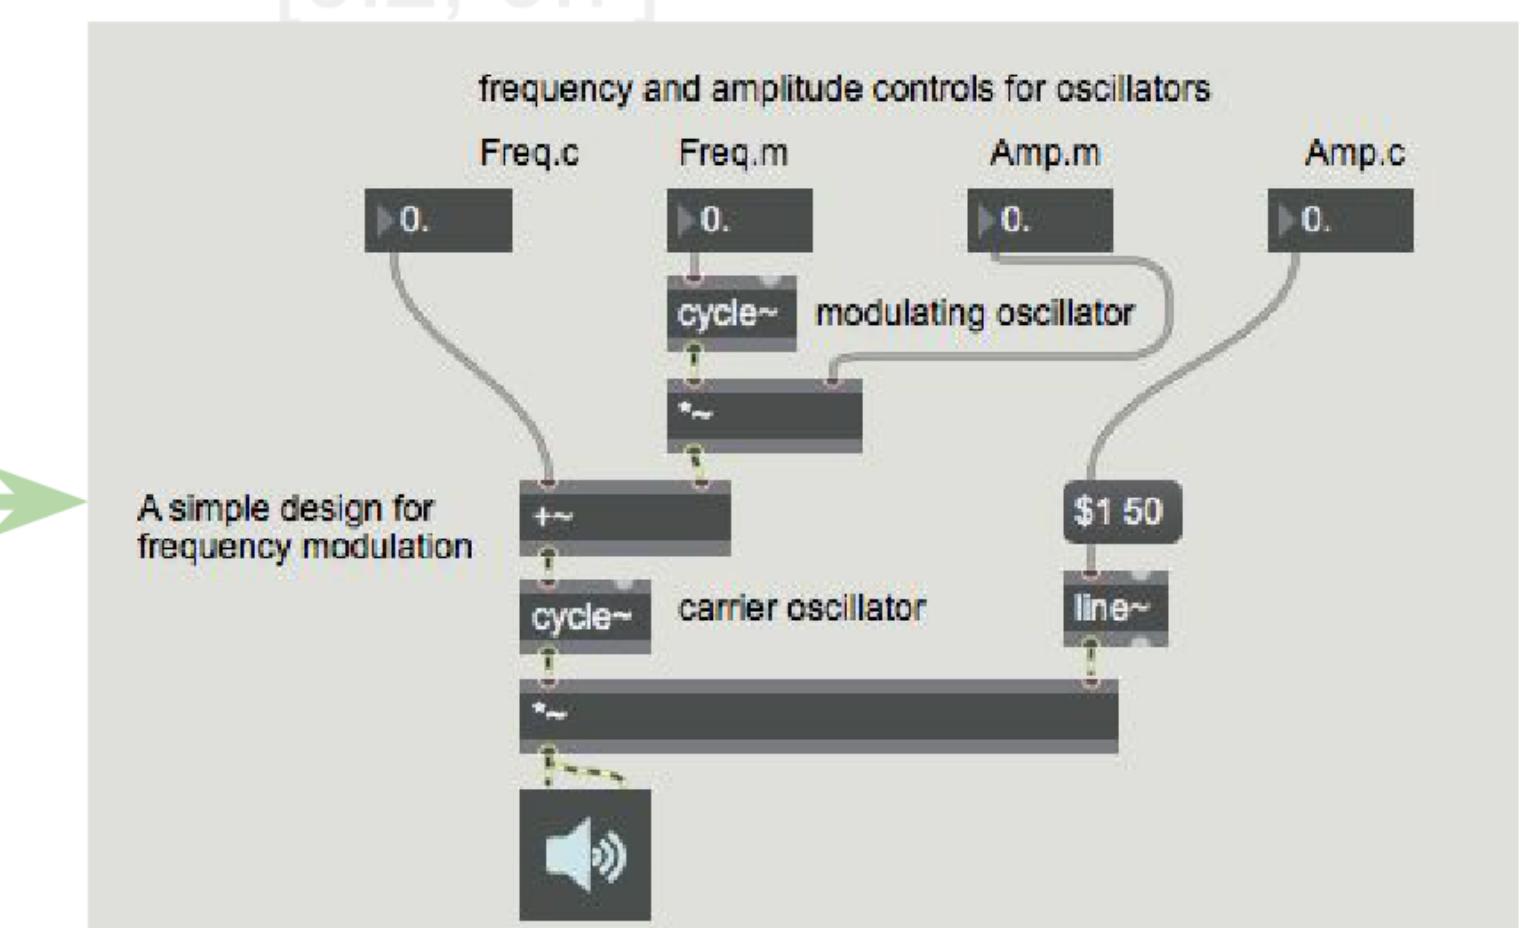
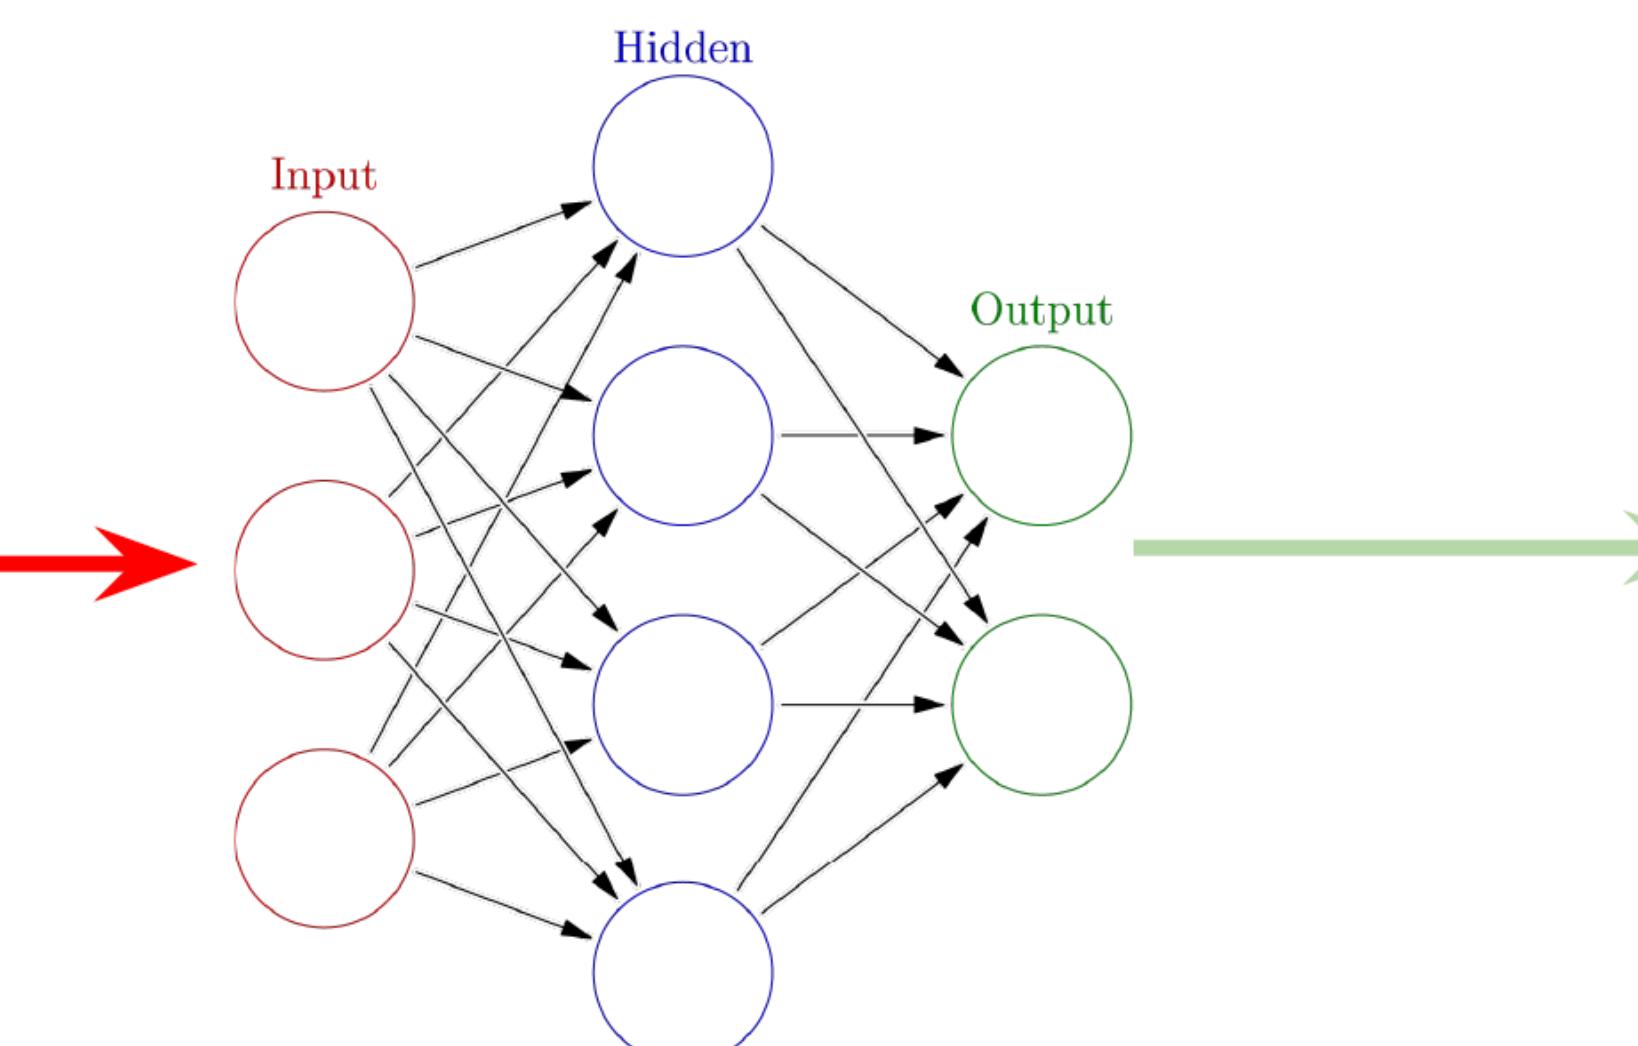
e.g.:

pitch

loudness

spectral centroid

spectral flatness



“g”: [0, 1, 1]

“h”: [0, 0, 1]

“i”: [0.4, 1, 0.7]

“j”: [0.1, 0.3, 0.1]

“k”: [0, 1, 1]

“l”: [0, 1, 1]

[0.5, 0.7]

[0.1, 0.1]

[0.5, 0.6]

[0.3, 0.9]

[0.4, 0]

[0.2, 0.7]

[0.5, 0.7]

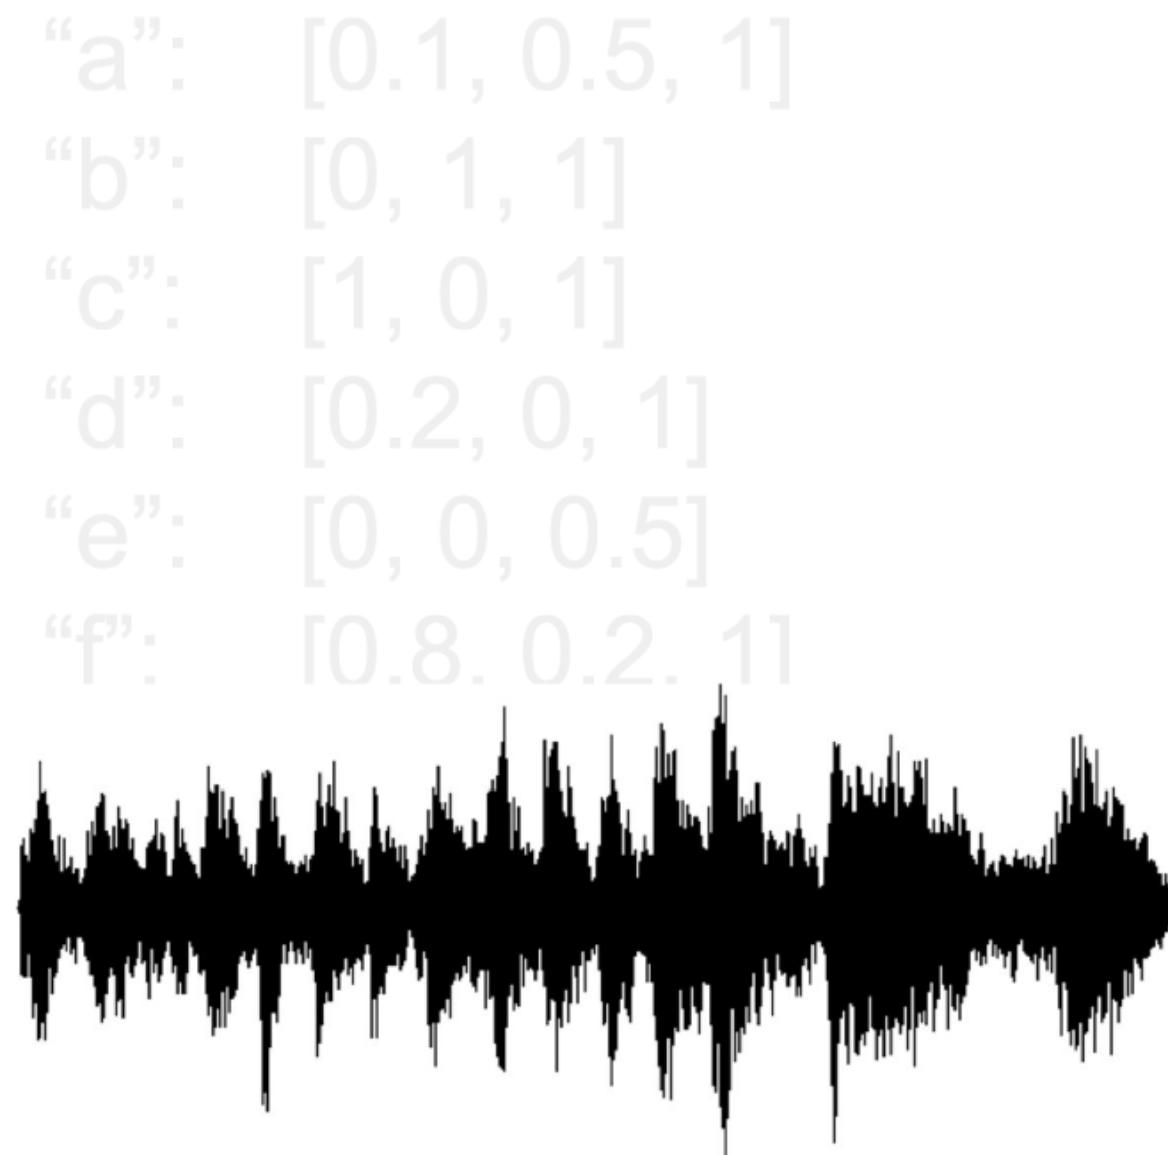
[0.1, 0.8]

[0.5, 0.1]

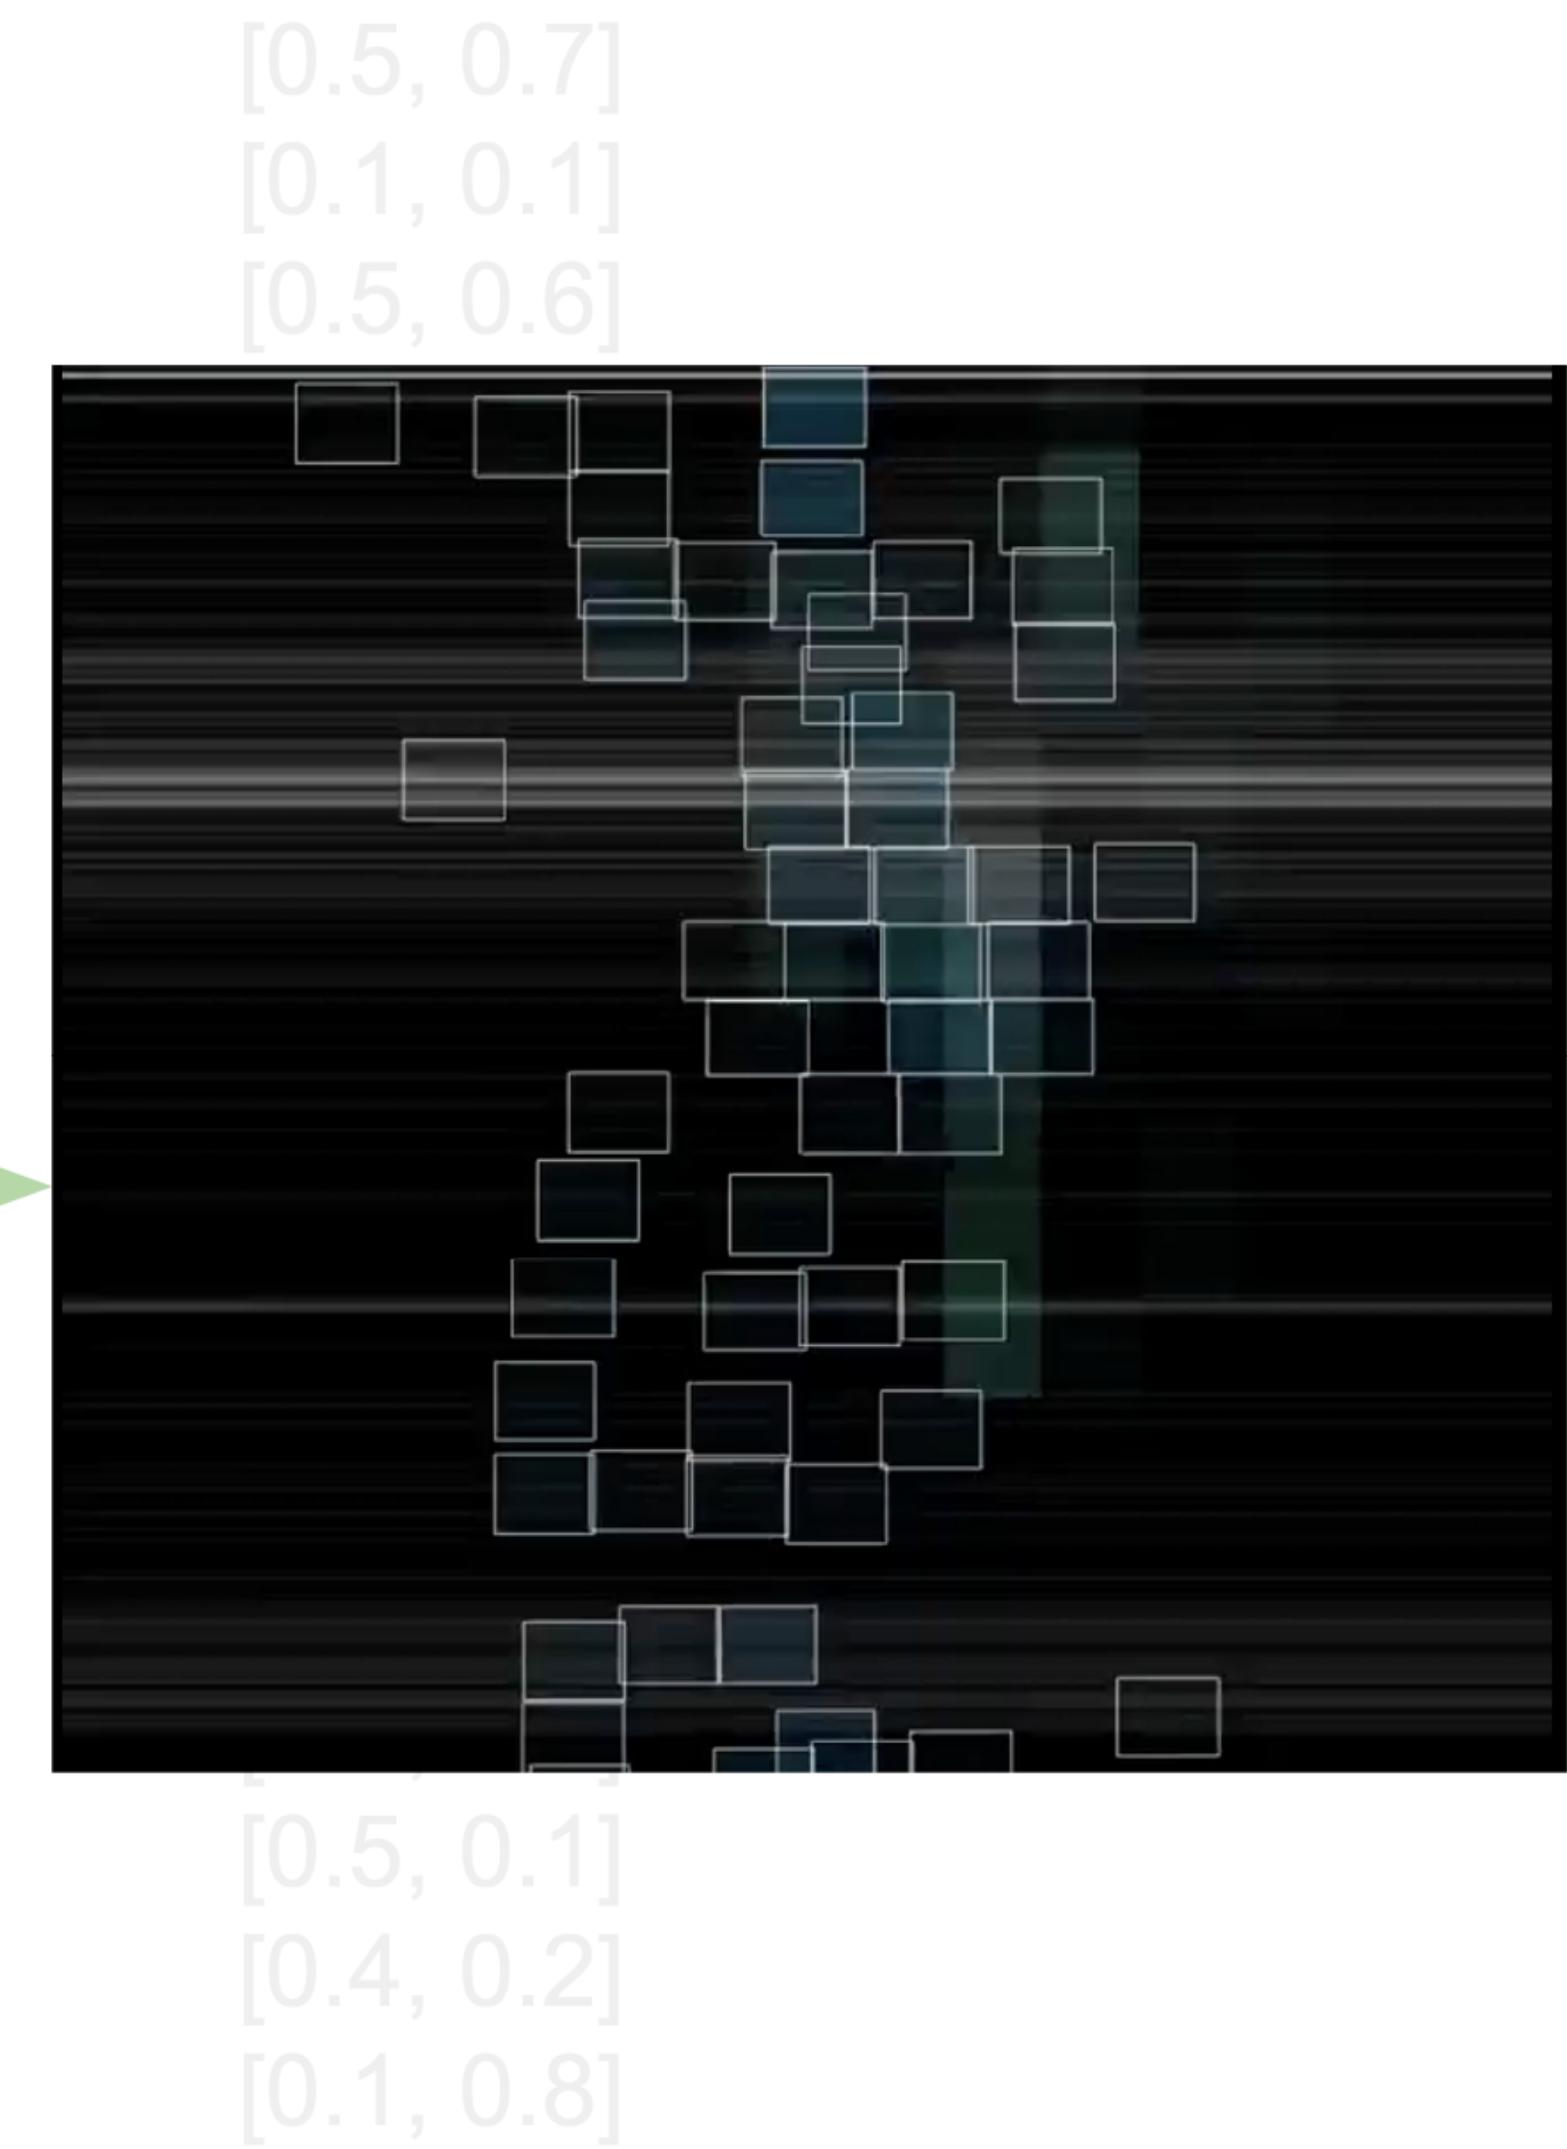
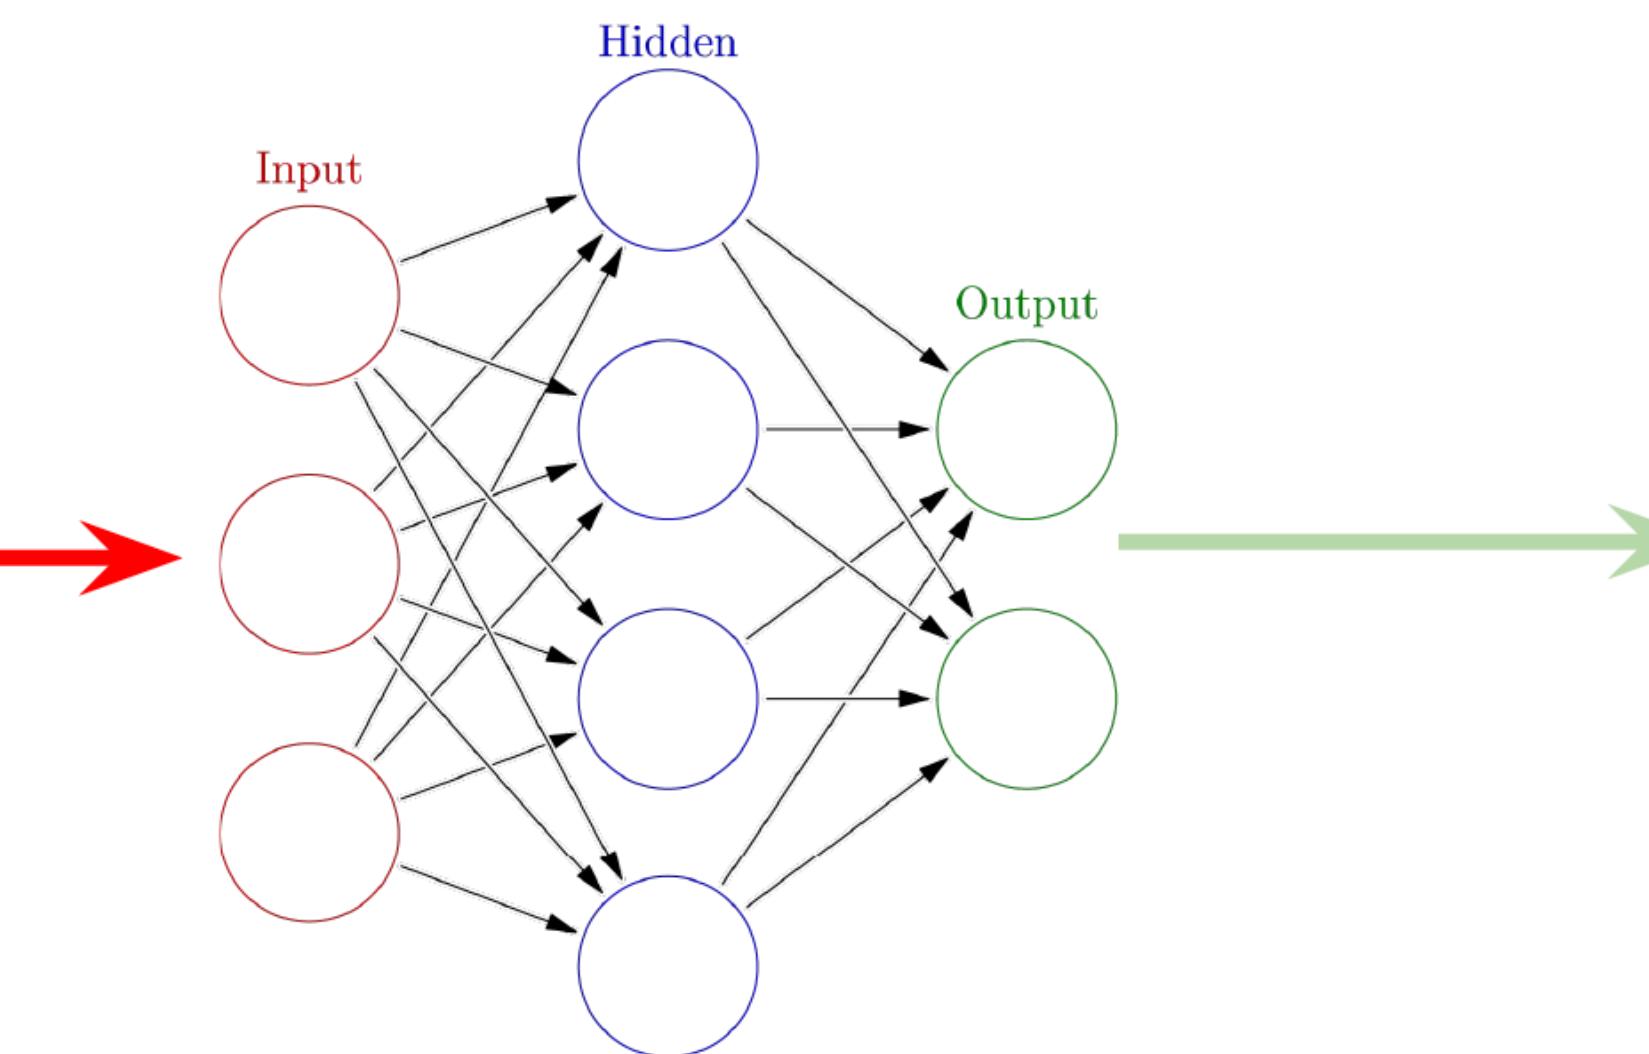
[0.4, 0.2]

[0.1, 0.8]

Neural Network Predicting with Regression

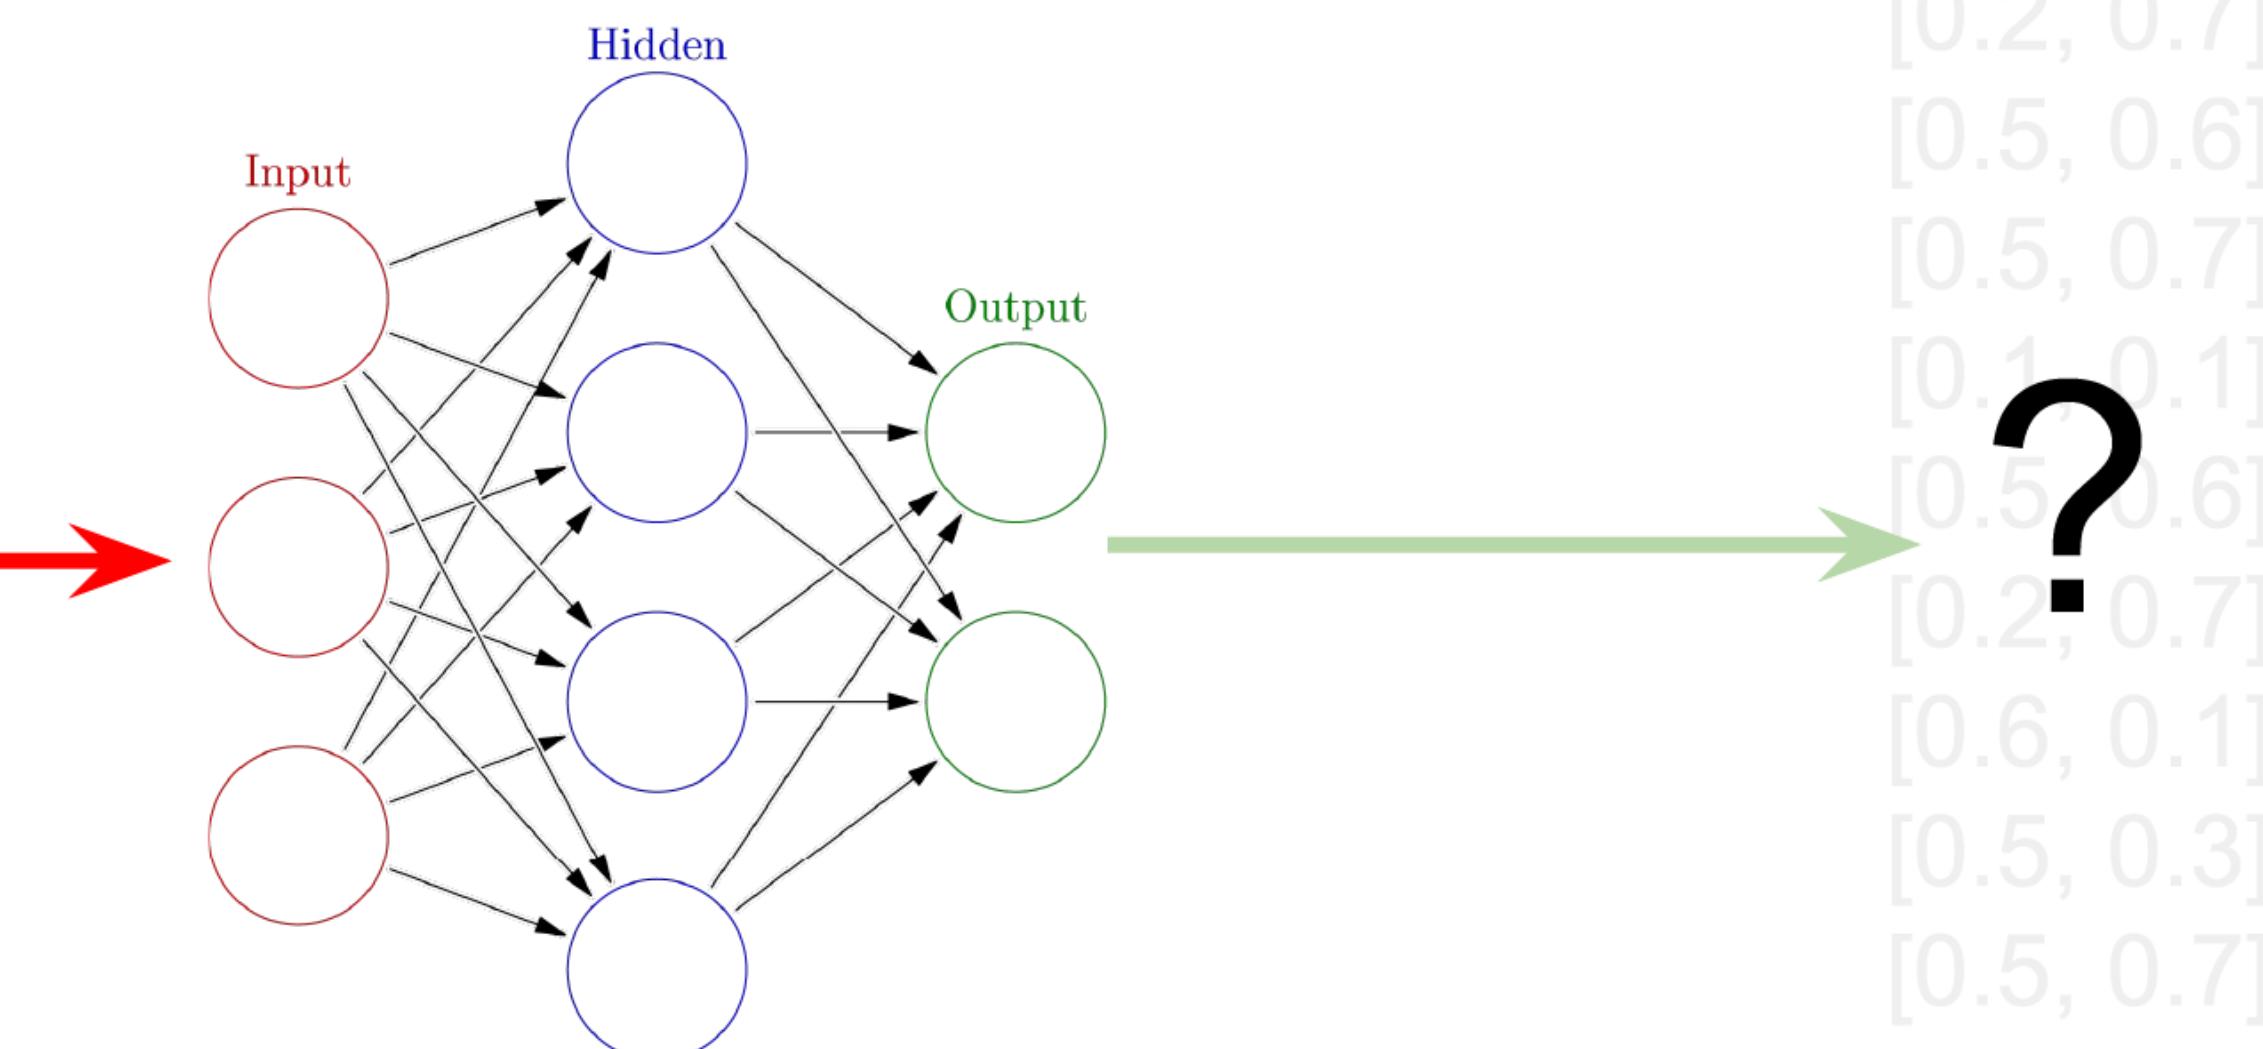


“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“j”: [0, 0, 0.5]
“k”: [0, 1, 1]
“l”: [0, 1, 1]
“m”: [0, 0, 1]
“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]



Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“g”: [0.9, 0.3, 0]
“h”: [0.4, 1, 0.7]
“i”: [0.2, 0.3, 1]
“j”: [0.1, 0.1, 0.1]
“k”: [0, 1, 1]
“l”: [0, 1, 1]
“m”: [0, 0, 1]
“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]



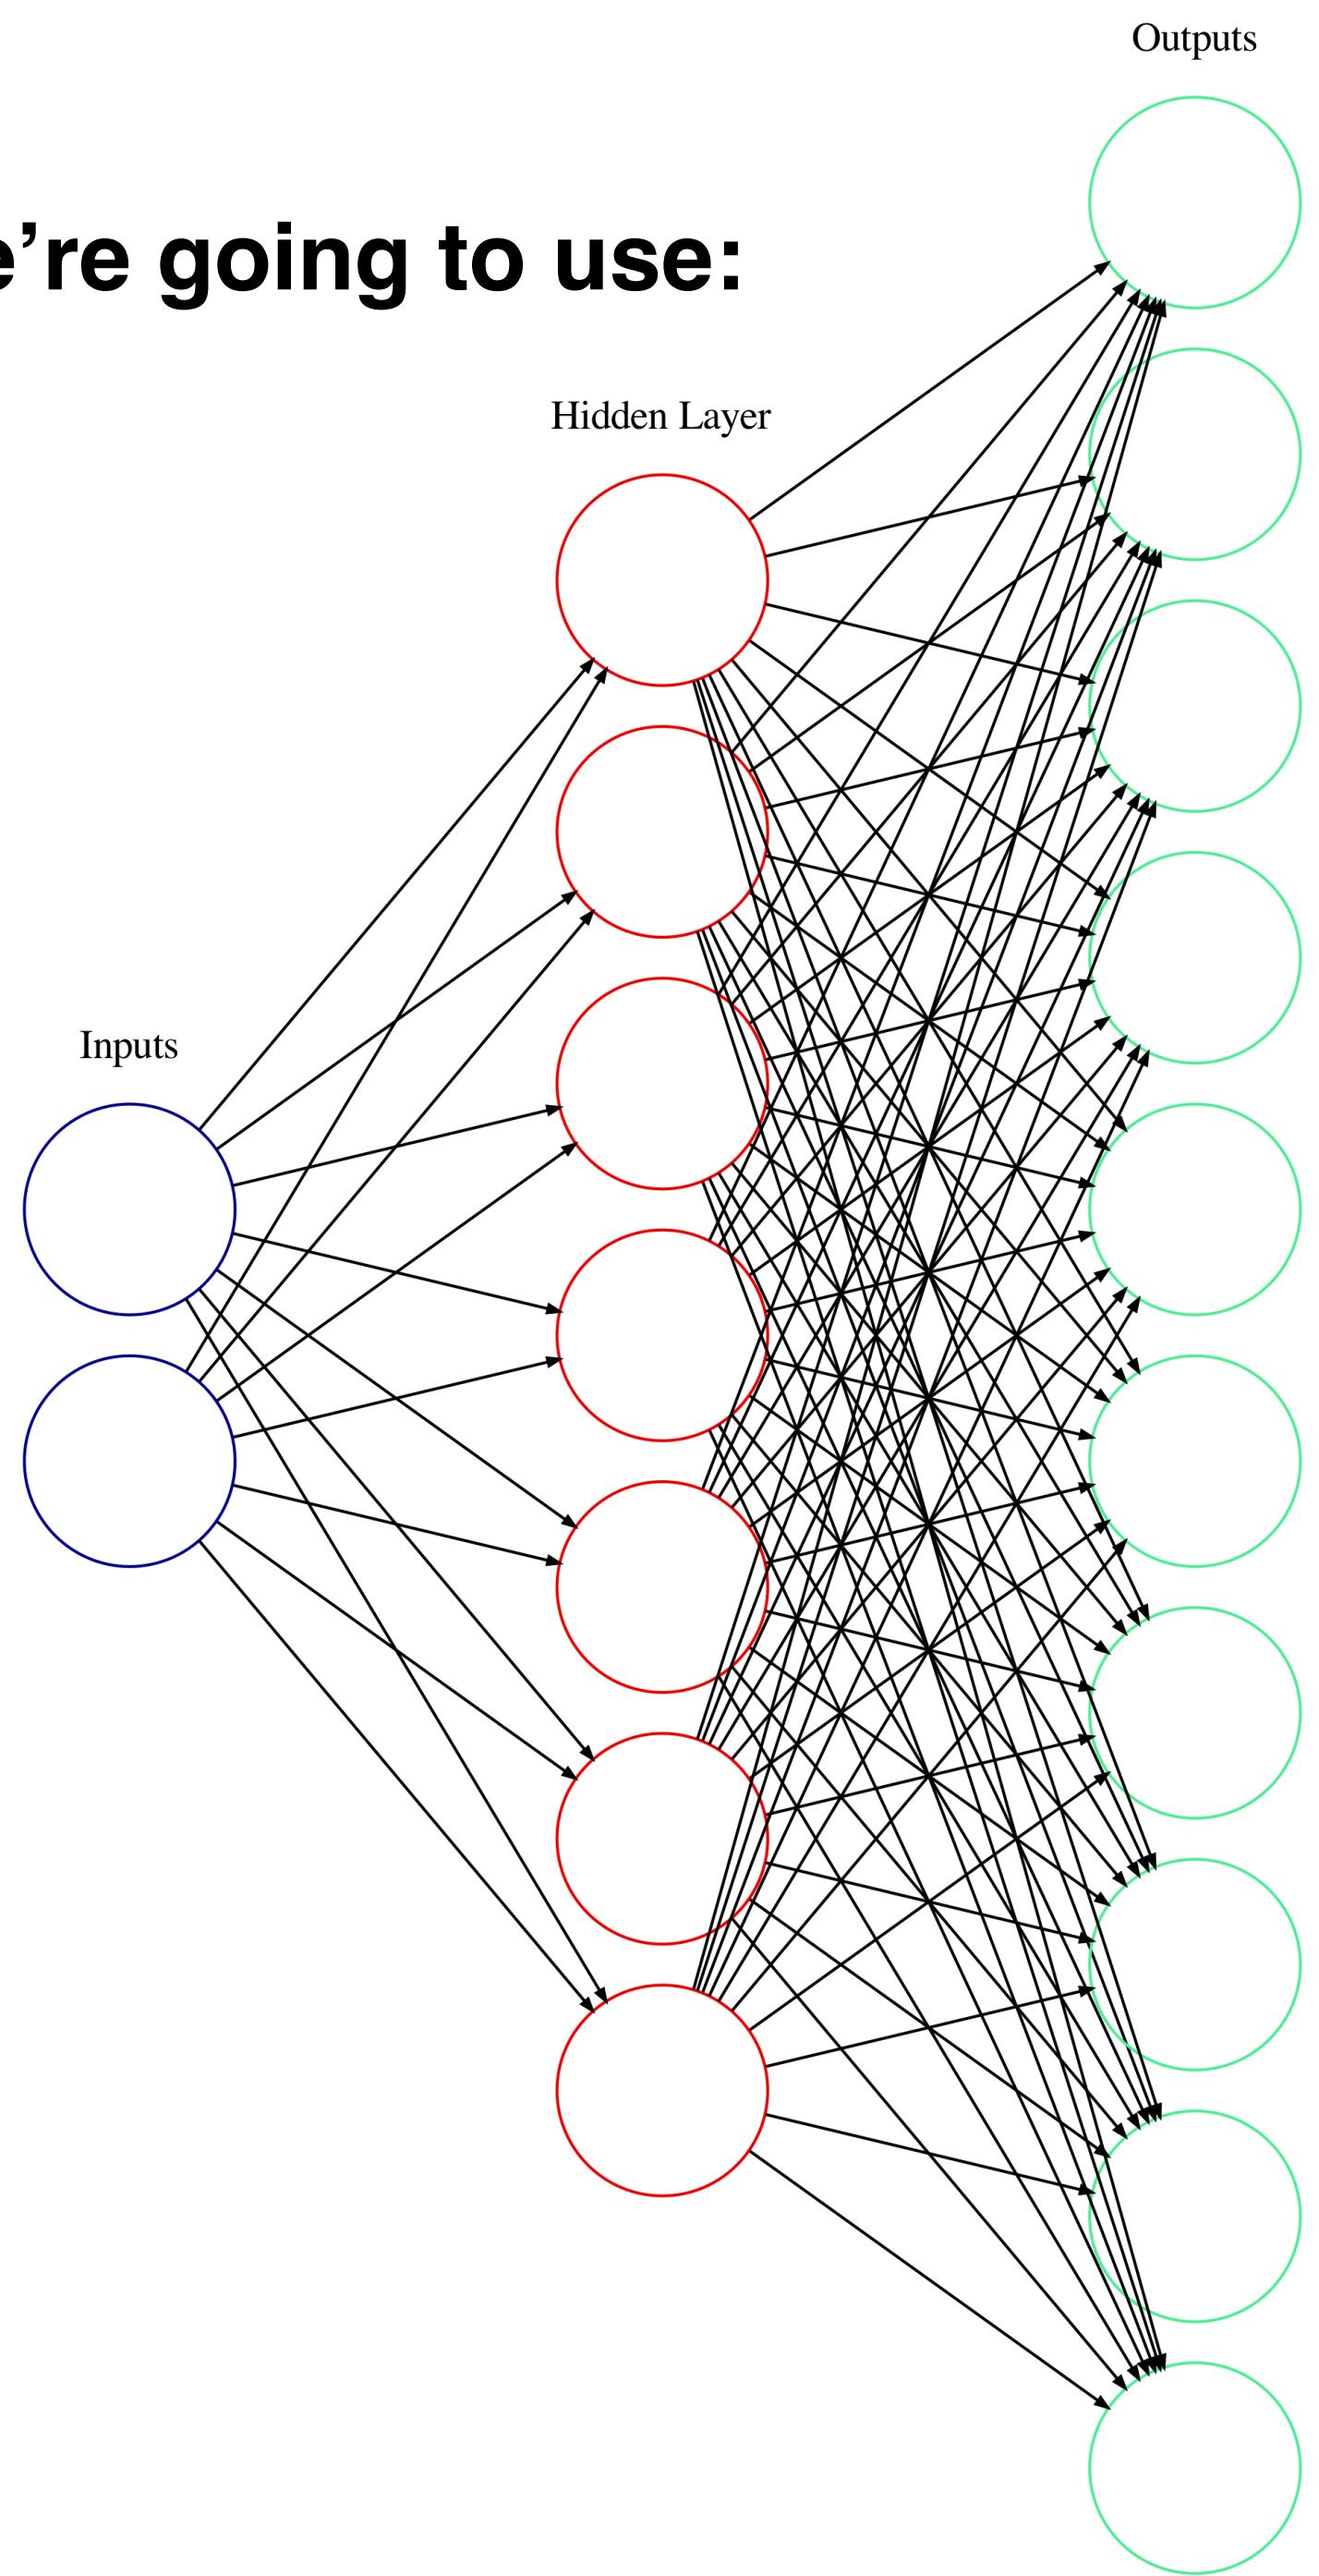
[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0]
[0.2, 0.7]
[0.5, 0.6]
[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.2, 0.7]
[0.6, 0.1]
[0.5, 0.3]
[0.5, 0.7]
[0.1, 0.8]
[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]

Structure of the neural network we're going to use:

2 inputs

1 hidden layer of 7 nodes

10 outputs

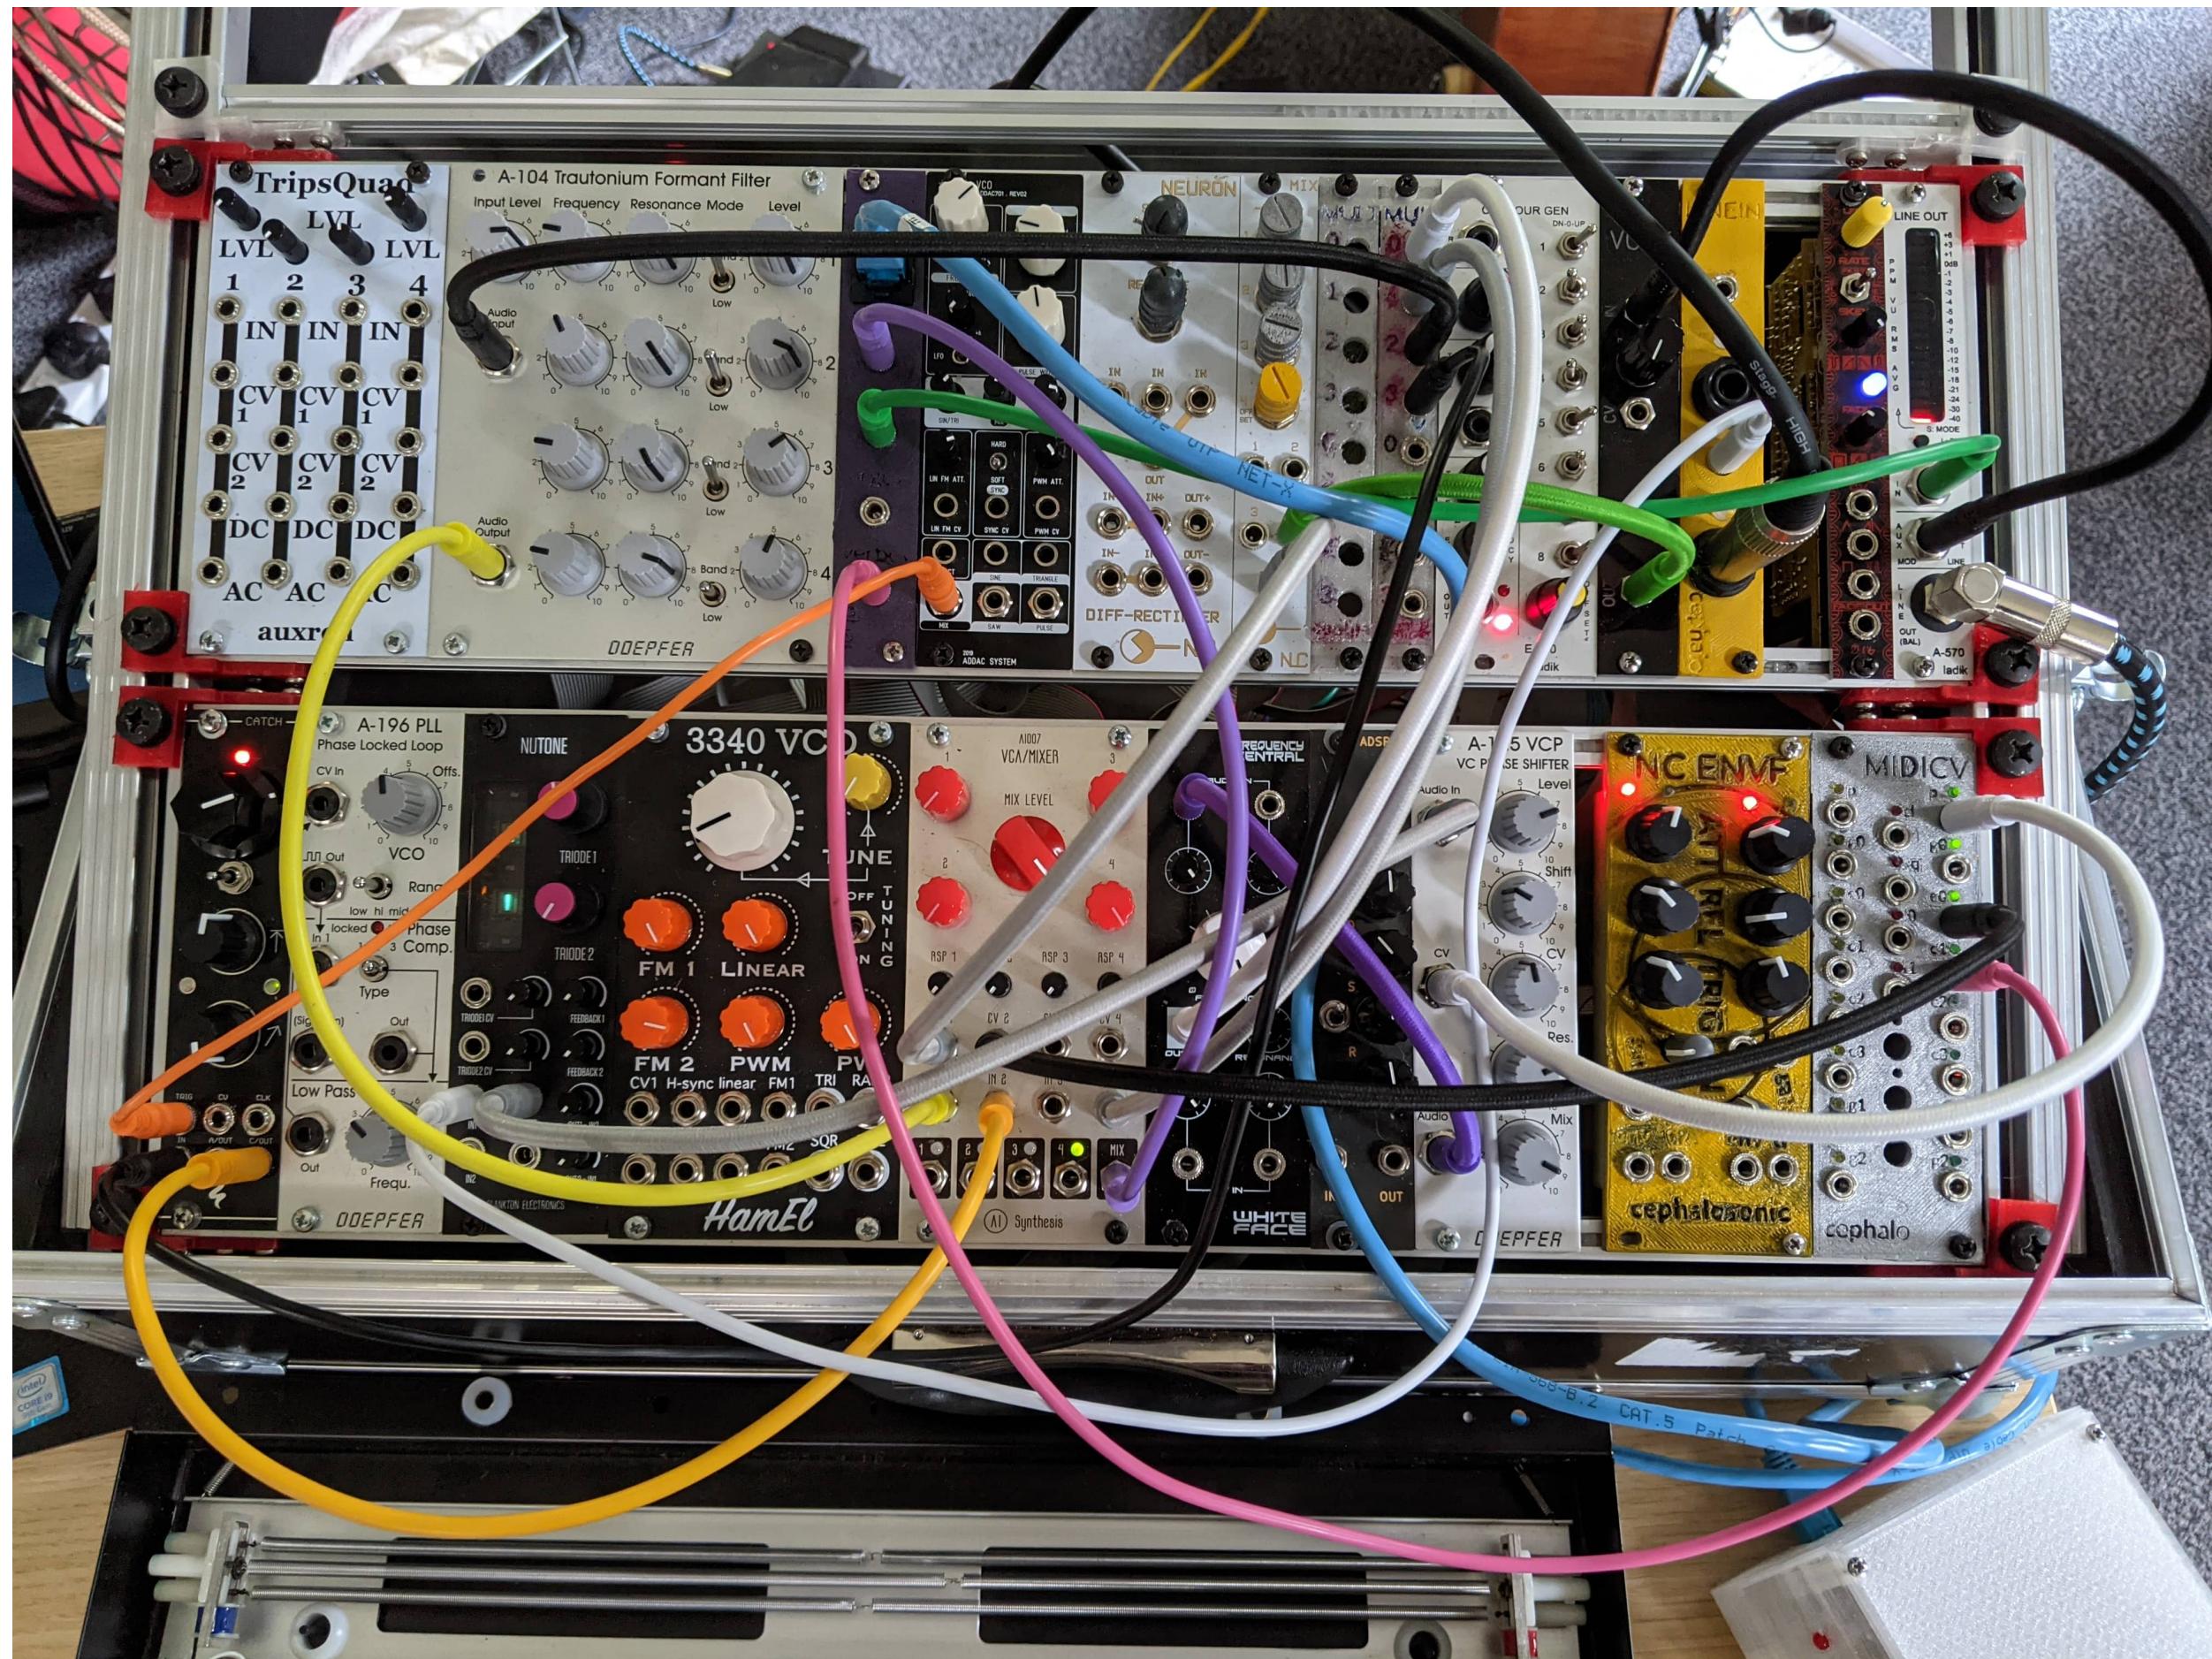


fluid.mlpregressor~



There are certain things that we care about, as musicians for example, that are really hard to articulate in code. It's hard for me to talk about what kind of quality of sound I want and then translate that into a set of filter coefficients. It's hard for me to talk about how I want a performer to move on stage and then translate that into some sort of mathematical equation for their trajectory. But it's a lot easier for me to either find examples of sounds that have a particular quality or to give examples of movements or if I'm using other types of modalities, often curating or creating examples are just way easier for us as people. And this relates to the types of tacit knowledge and embodied knowledge we bring to creative practices.

-Rebecca Fiebrink



FeedbackFeedforward

by Alice Eldridge and
Chris Kiefer

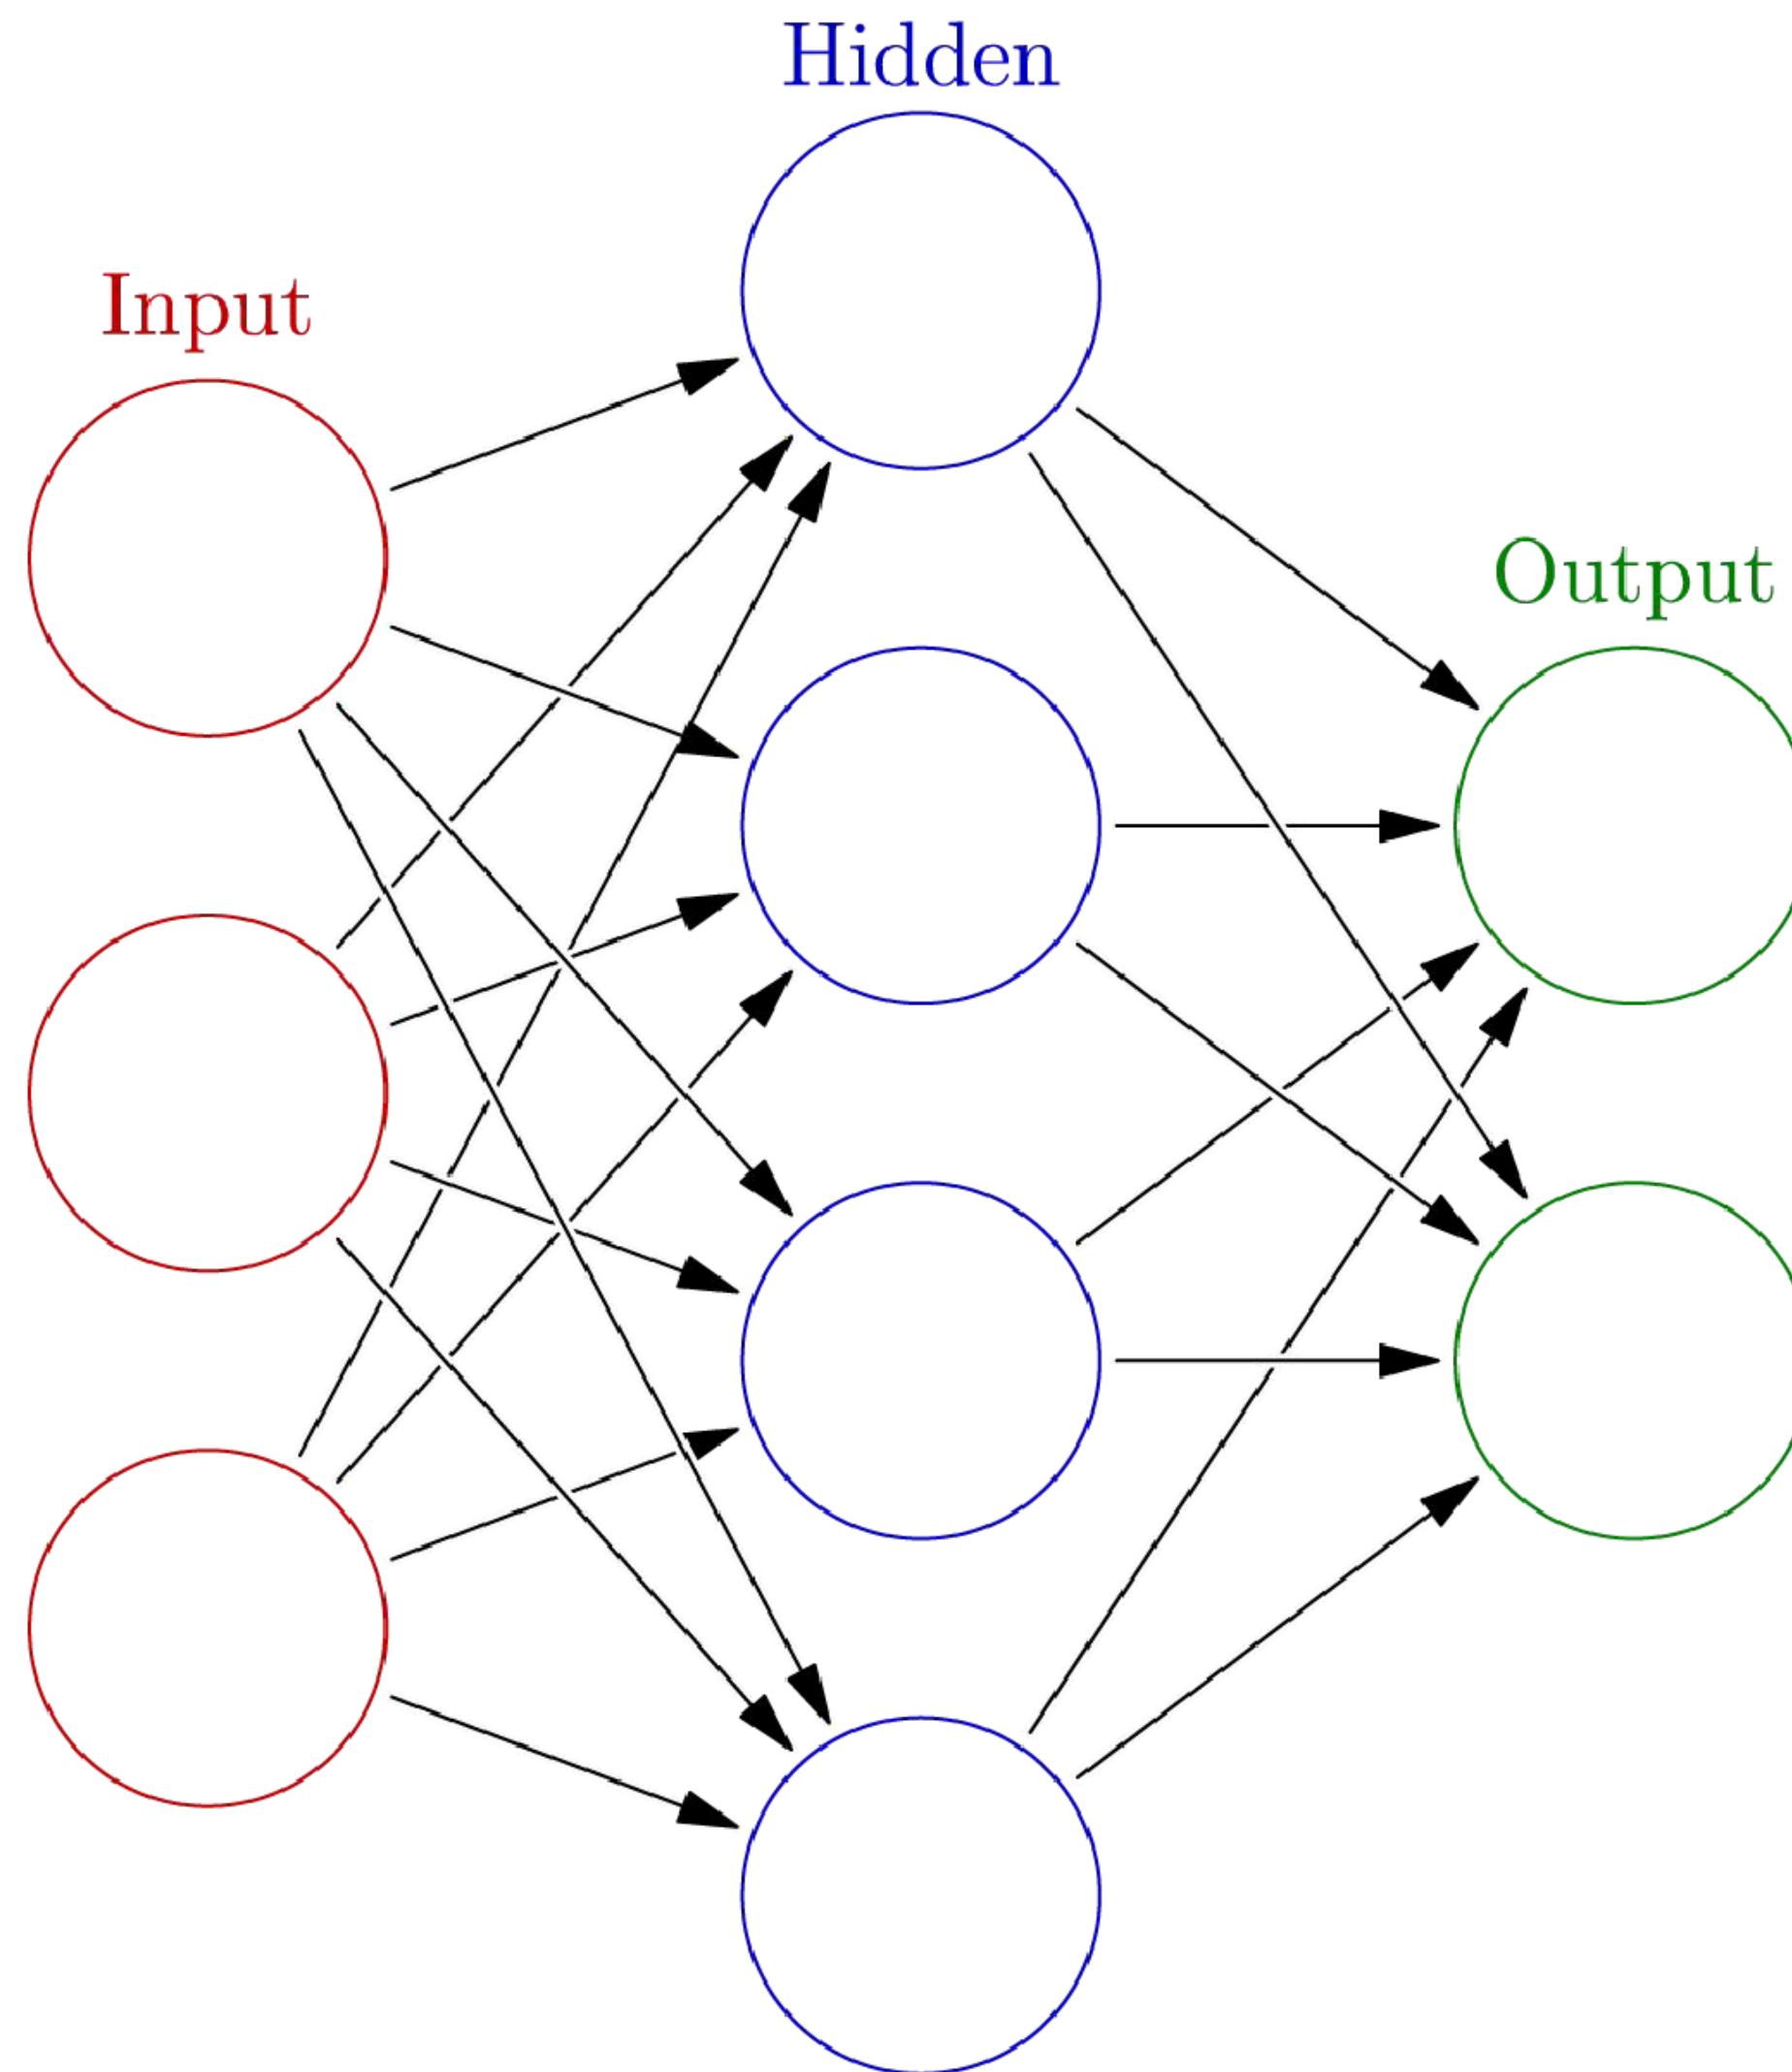
learn.flucoma.org/explore/eldridge-kiefer

FluidMLPClassifier

classify sounds by
timbre in real-time



Neural Network (Multi-Layer Perceptron)

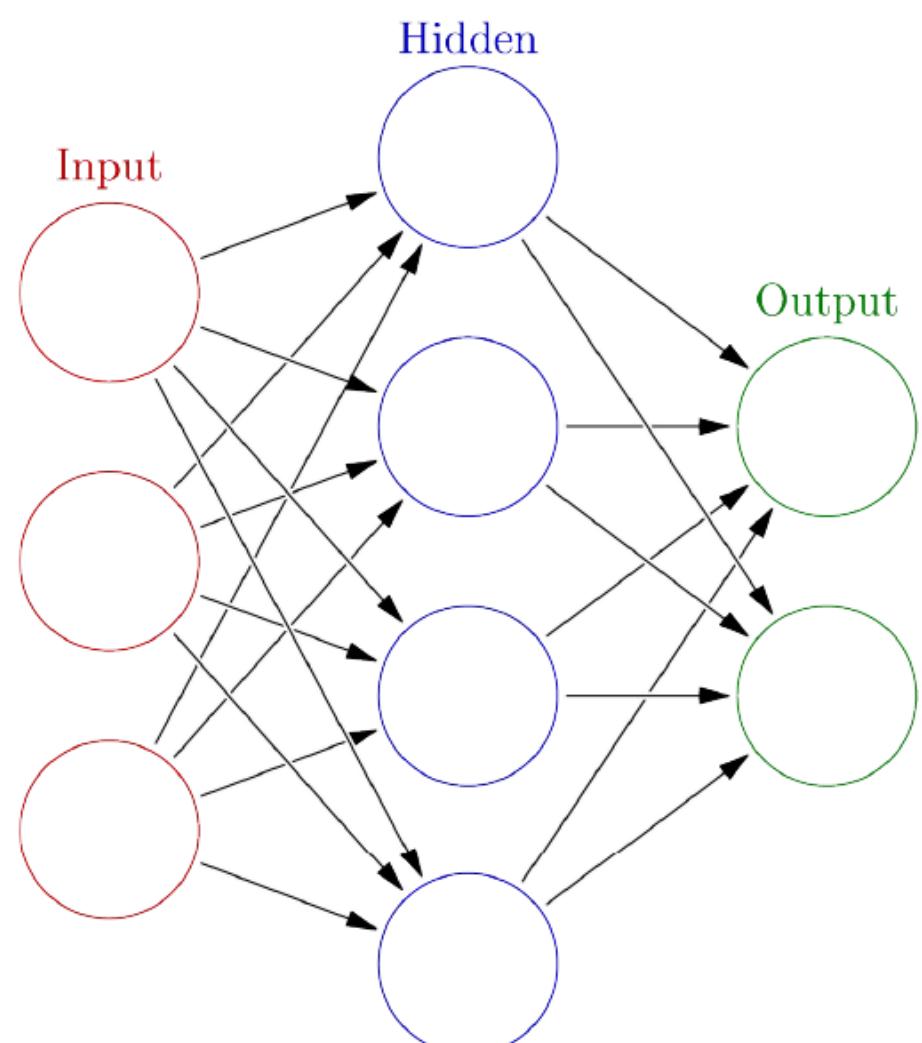


Classification:
a neural network predicts
**which category (or “class”)
an input belongs to**

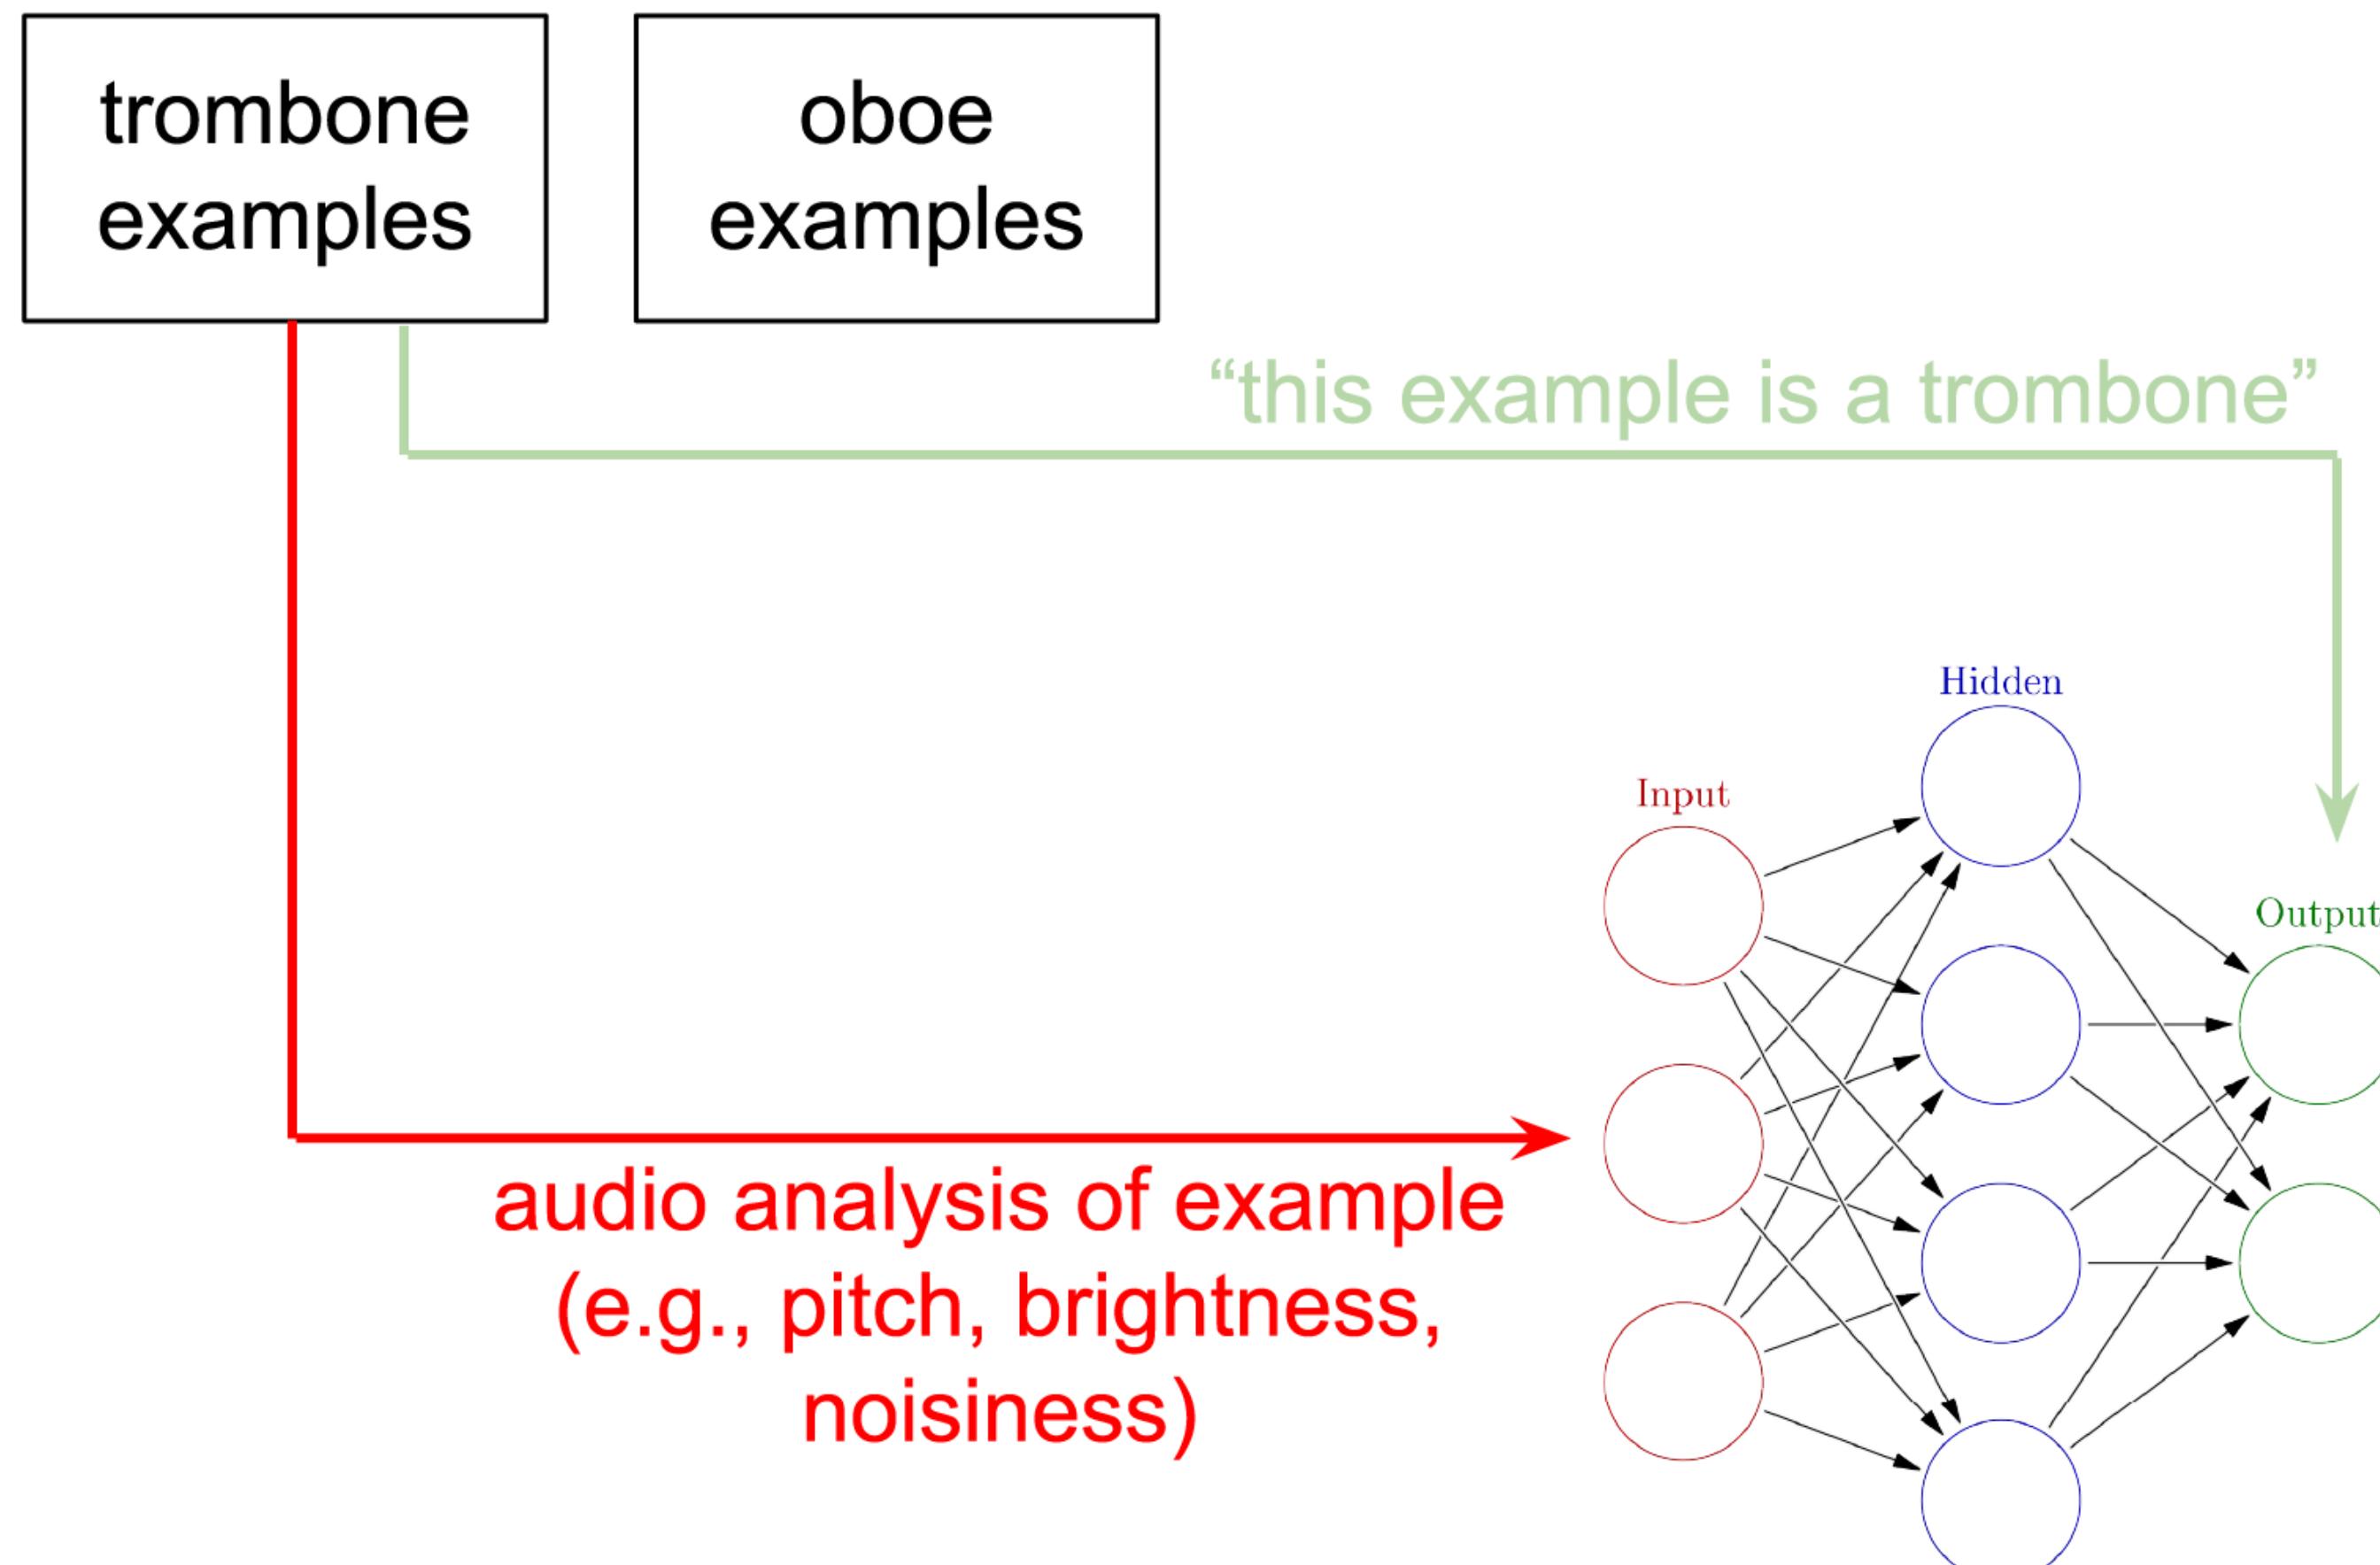
Neural Network ***Training a Classifier***

trombone
examples

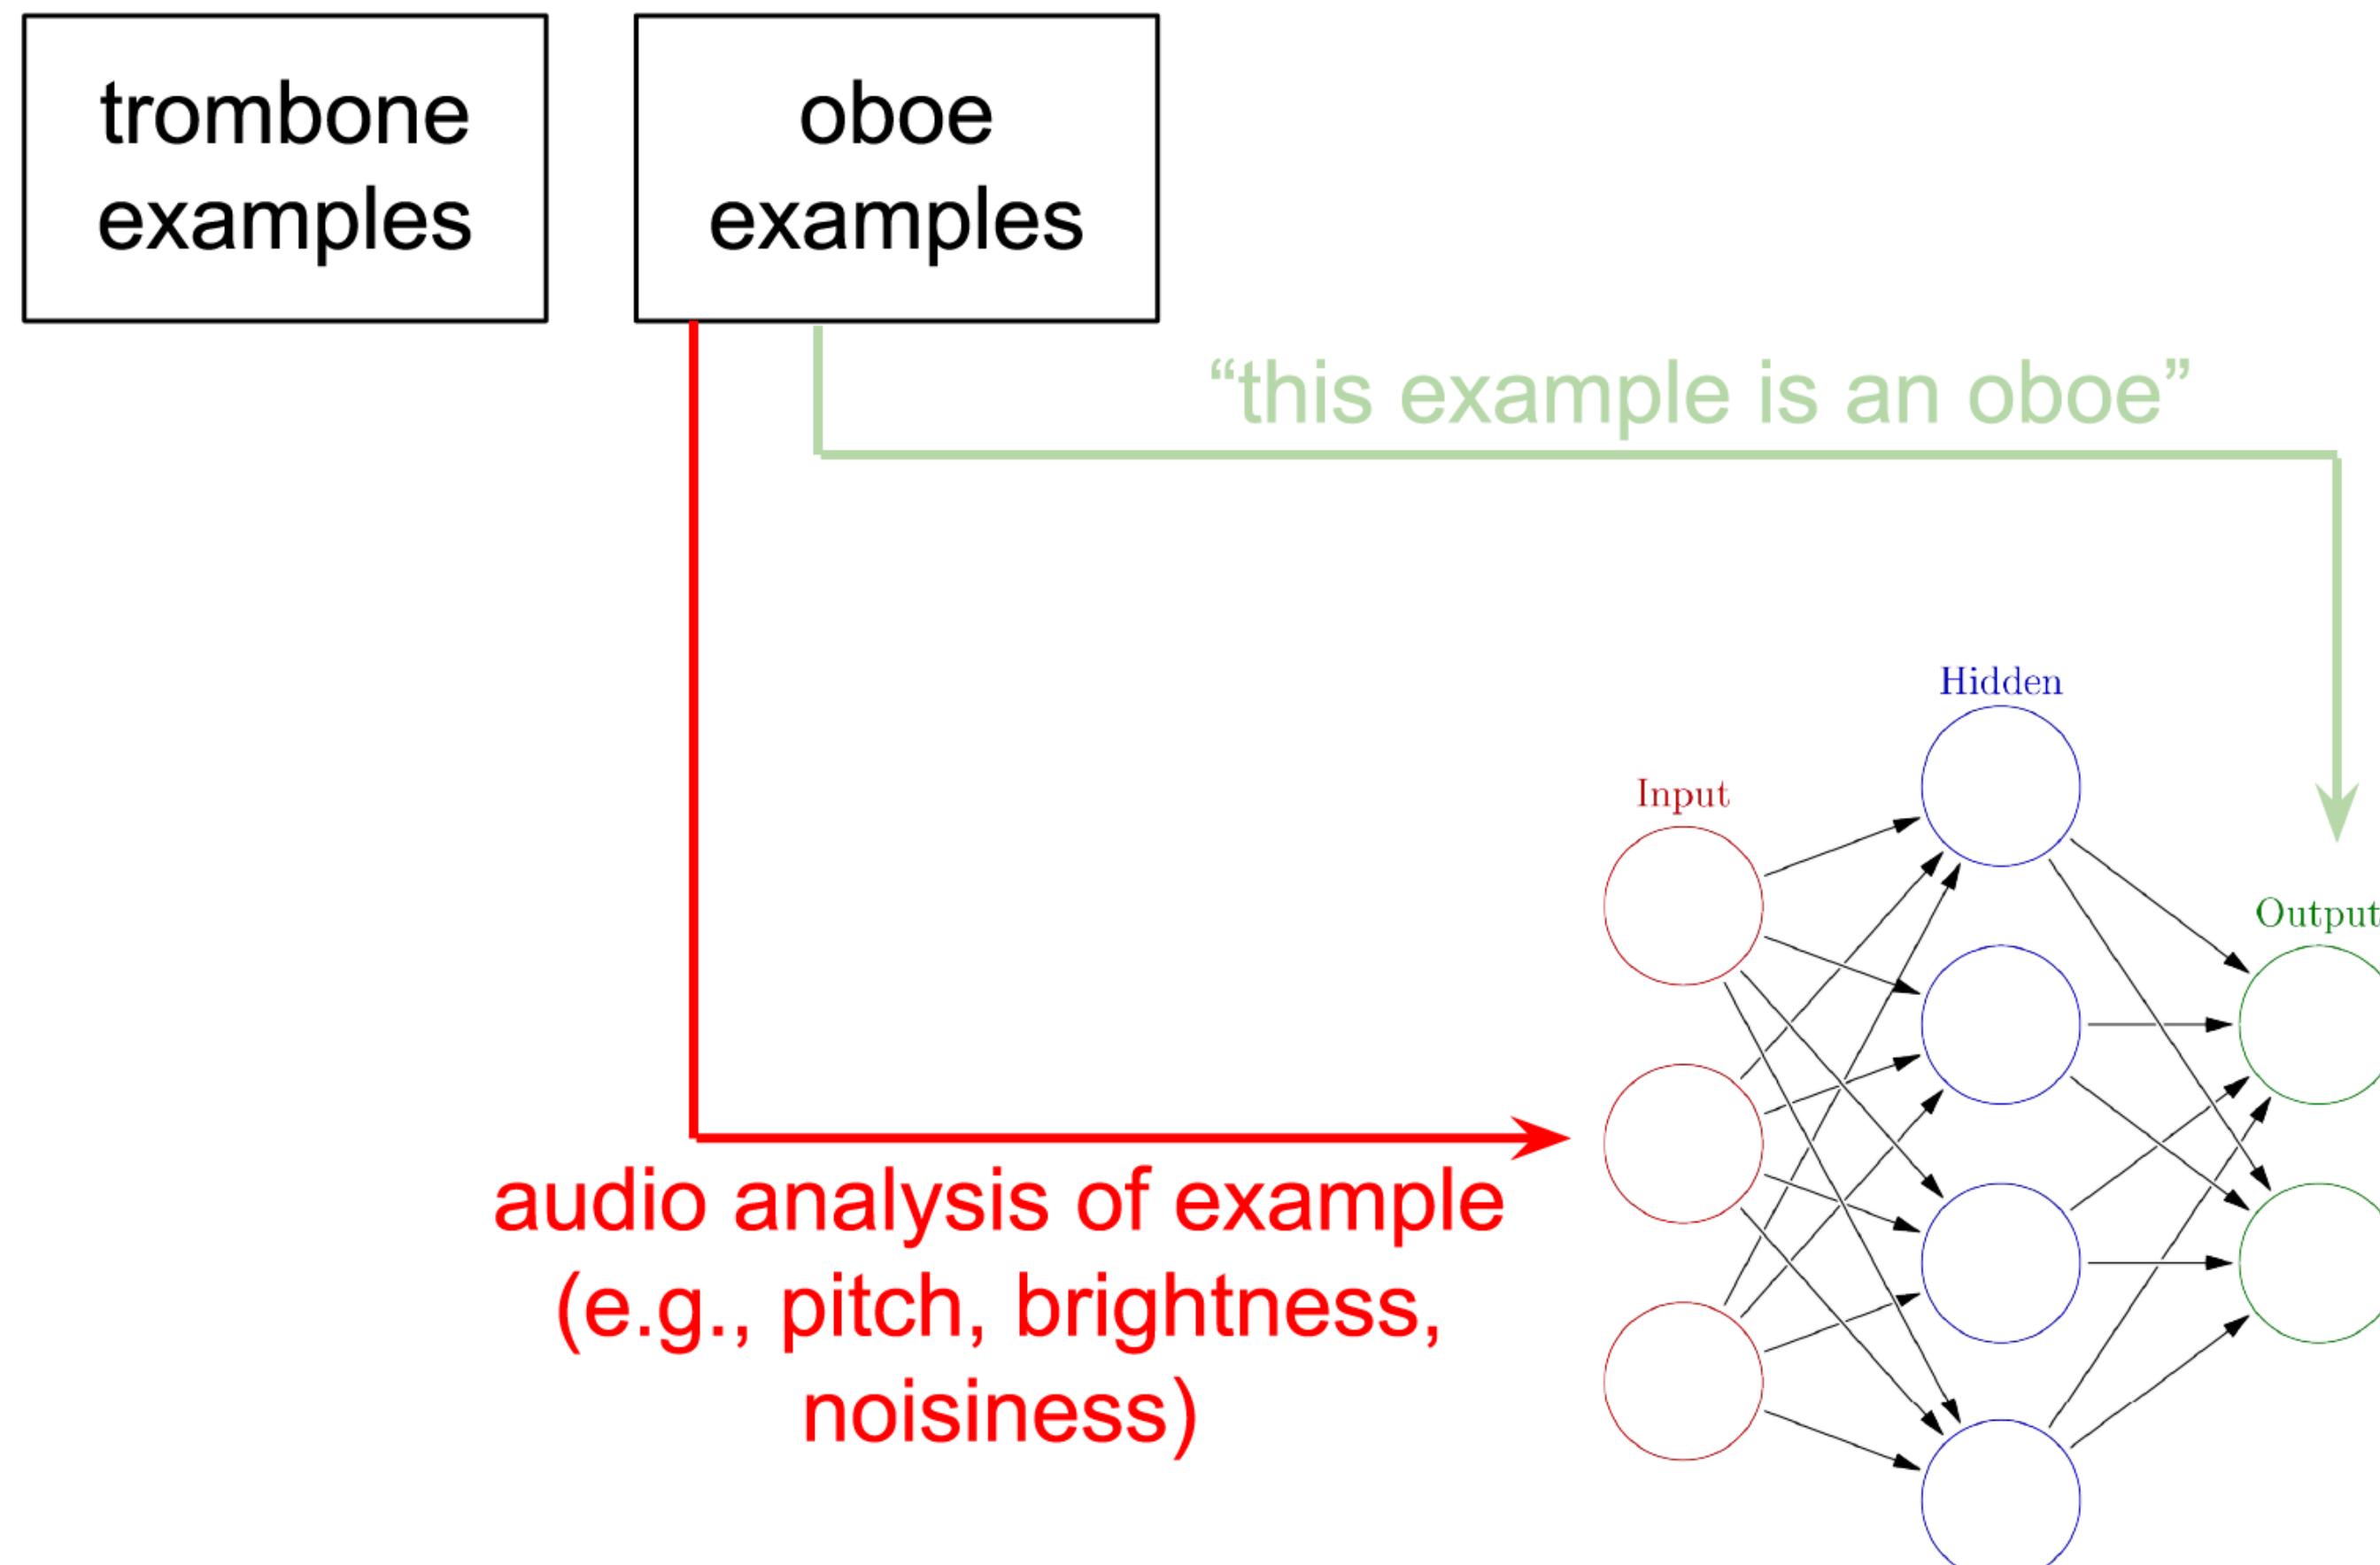
oboe
examples



Neural Network **Training a Classifier**



Neural Network **Training a Classifier**



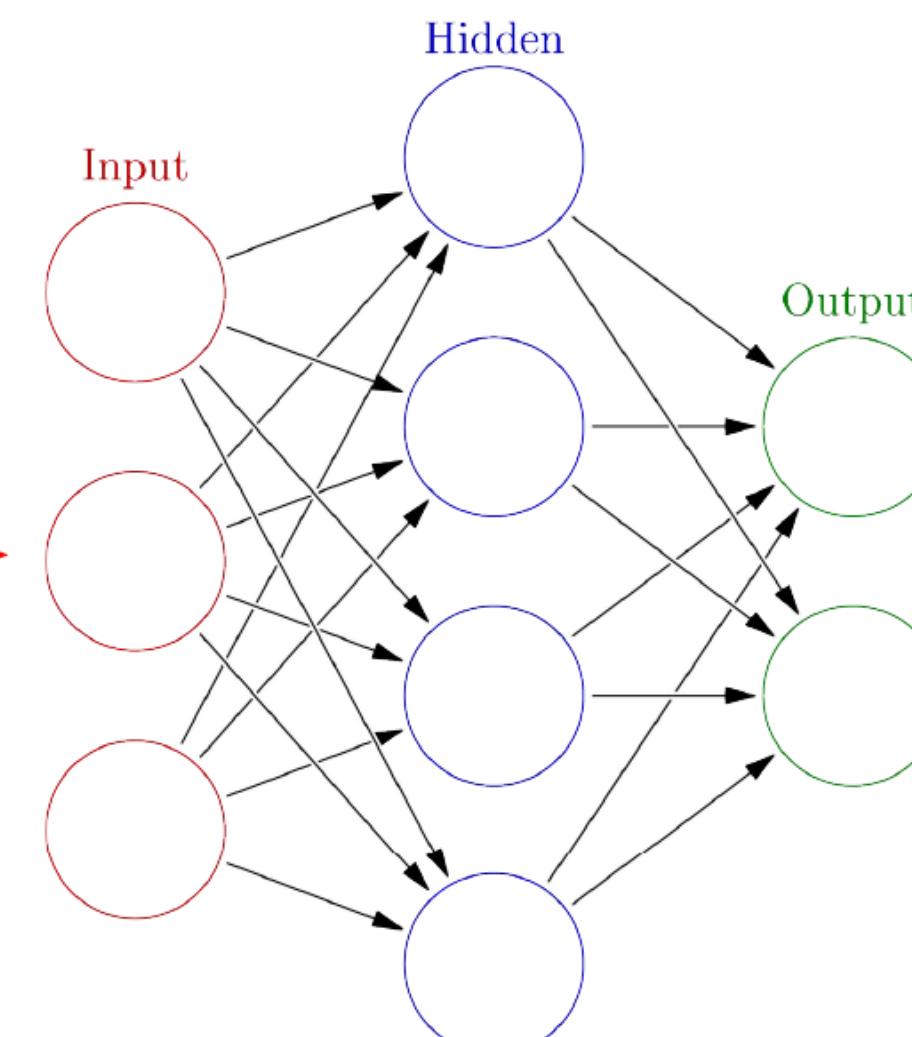
Neural Network Predicting a Classification

trombone
examples

oboe
examples

new example it has
never seen before

audio analysis of example
(e.g., pitch, brightness,
noisiness)



“this new example is most
like the oboe examples you
showed me before”
(or trombone...)

fluid.mlpclassifier

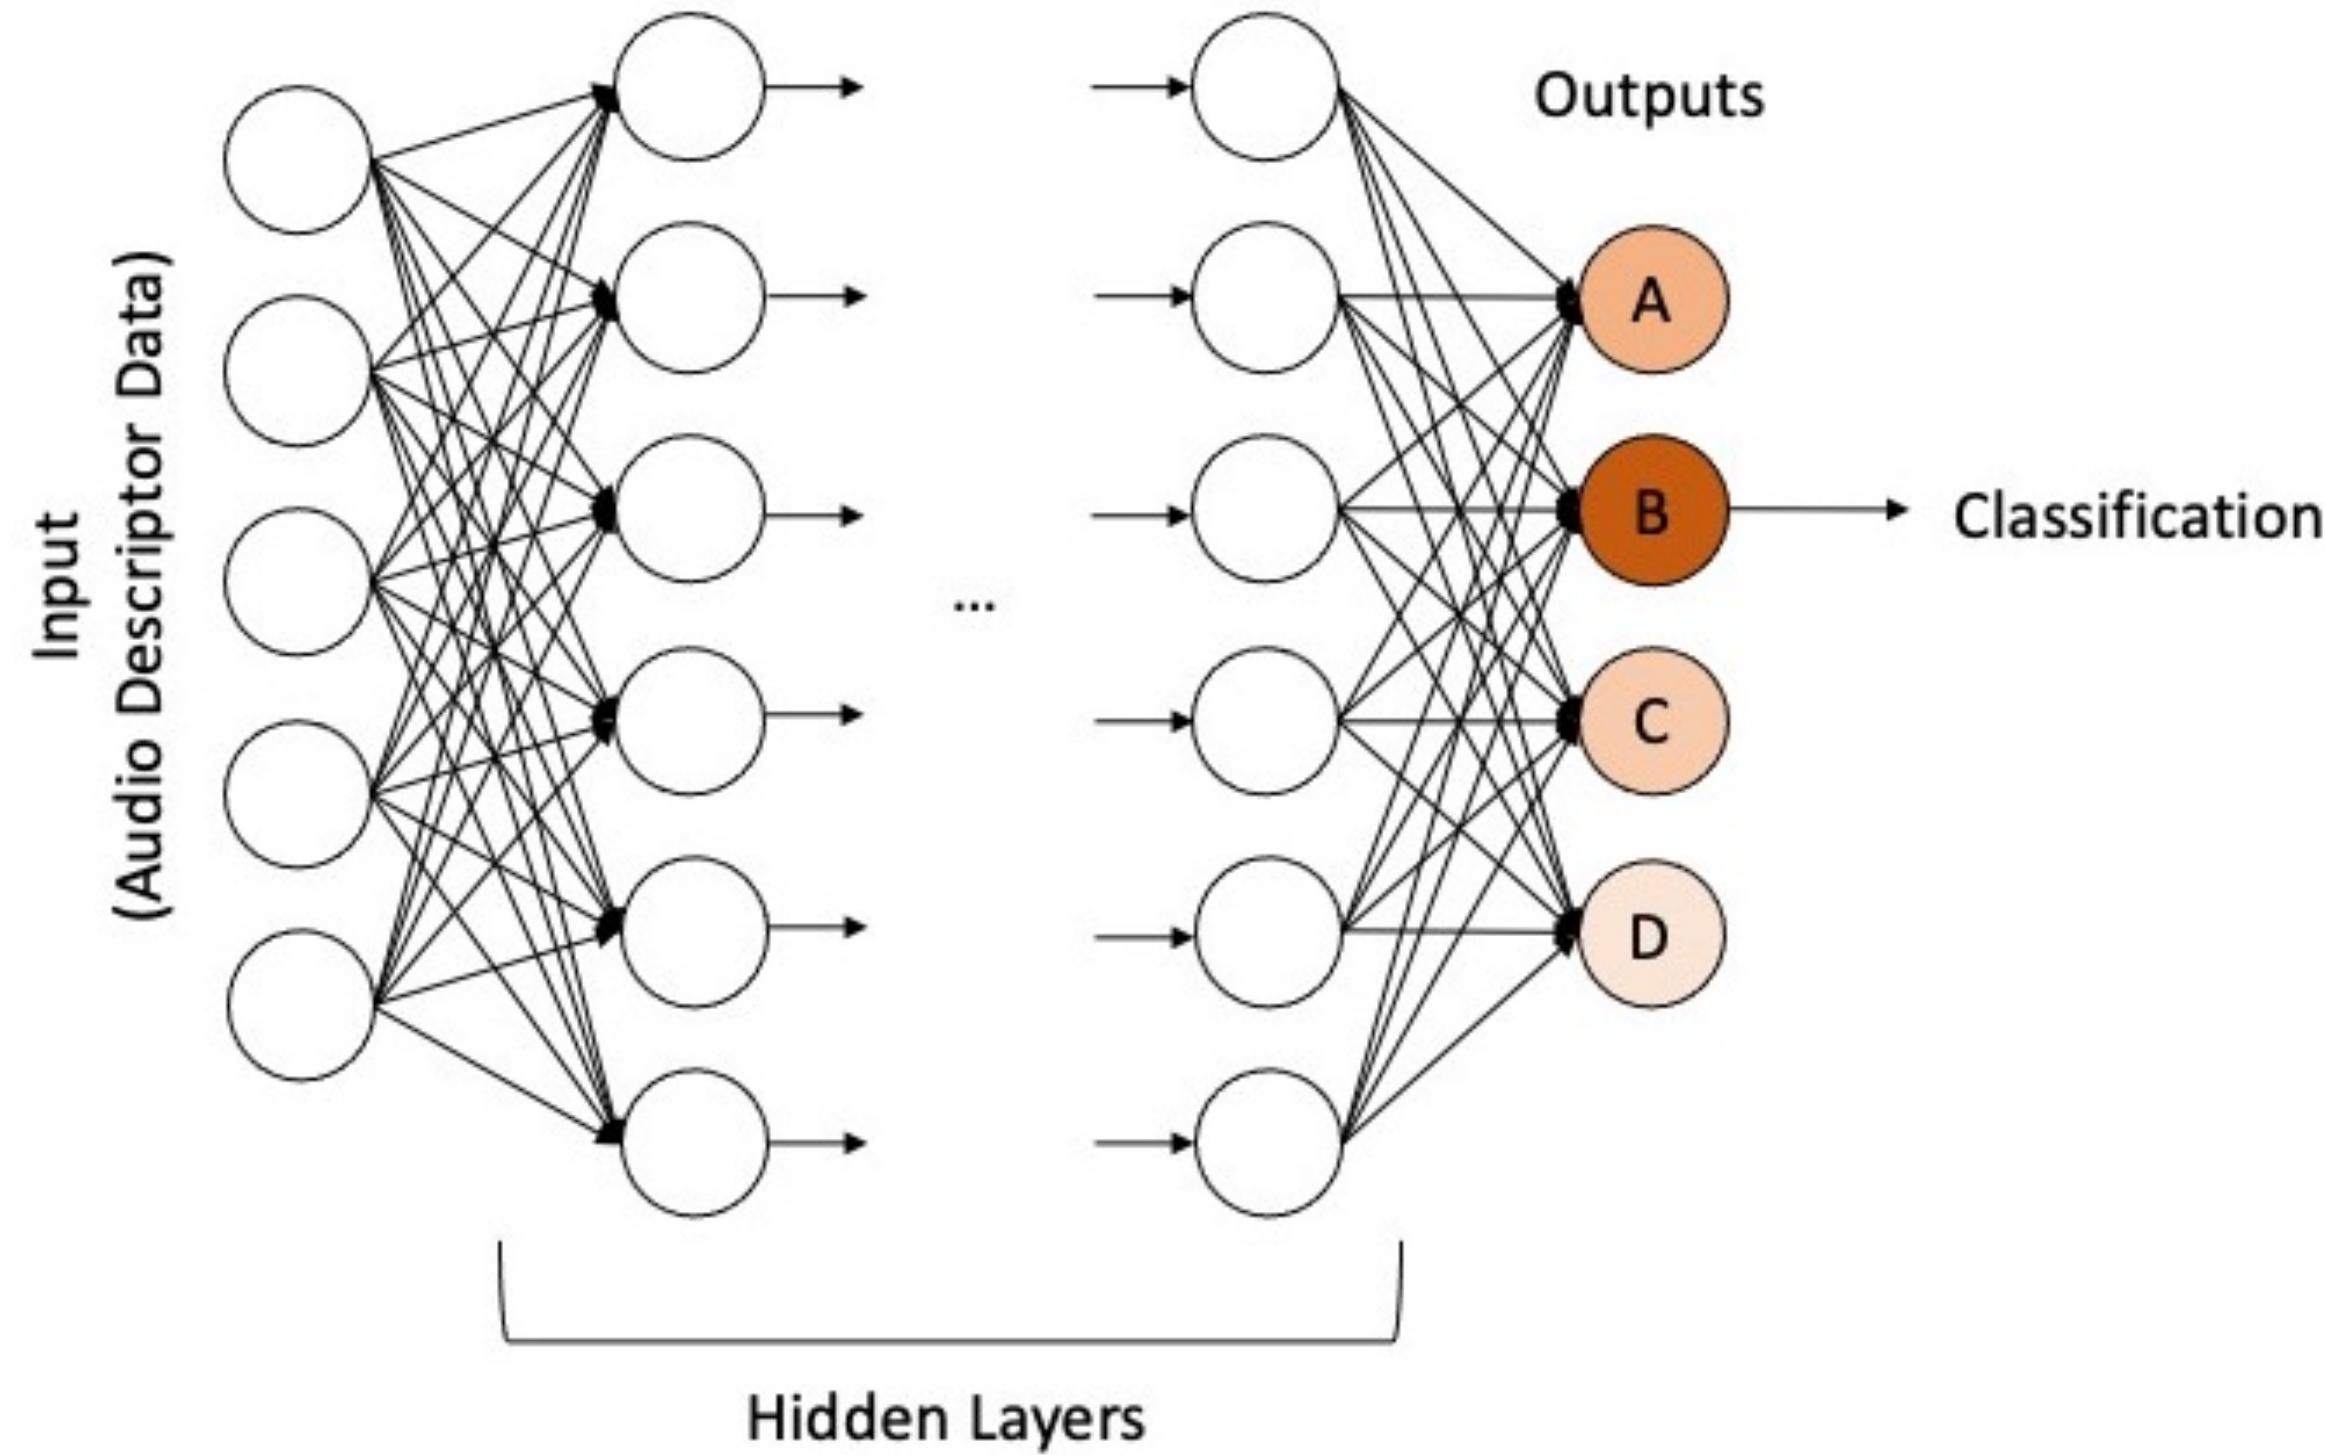


fluid.mlpclassifier

validation



- **overfitting**
 - excellent on training data, poor on anything else
- **underfitting**
 - poor on everything...keep training!



Drift Shadow by Alex Harker

learn.flucoma.org/explore/harker