

<https://bit.ly/flucoma-upenn>



# FluCoMa



**Ted Moore**

[tedmooremusic.com](http://tedmooremusic.com)

[ted@tedmooremusic.com](mailto:ted@tedmooremusic.com)

# FluCoMa: Fluid Corpus Manipulation



- “enable techno-fluent musicians to use machine listening and machine learning in their creative practices”
- Integrating Machine Listening and Machine Learning in...
- Max, SuperCollider & Pure Data
- Learning Resources ([learn.flucoma.org](https://learn.flucoma.org))
- Discourse Community ([discourse.flucoma.org](https://discourse.flucoma.org))



# FluCoMa: Fluid Corpus Manipulation



## Slice Audio

- onset slice
- transient slice
- novelty slice
- amplitude slice
- amplitude gate

## Decompose Audio

- extract transients
- harmonic/percussive separation
- model as sine waves
- non-negative matrix factorisation

## Analyse Audio

- pitch
- loudness
- mel-bands
- mel-frequency cepstral coefficients
- spectral centroid
- spectral flatness
- chromagram

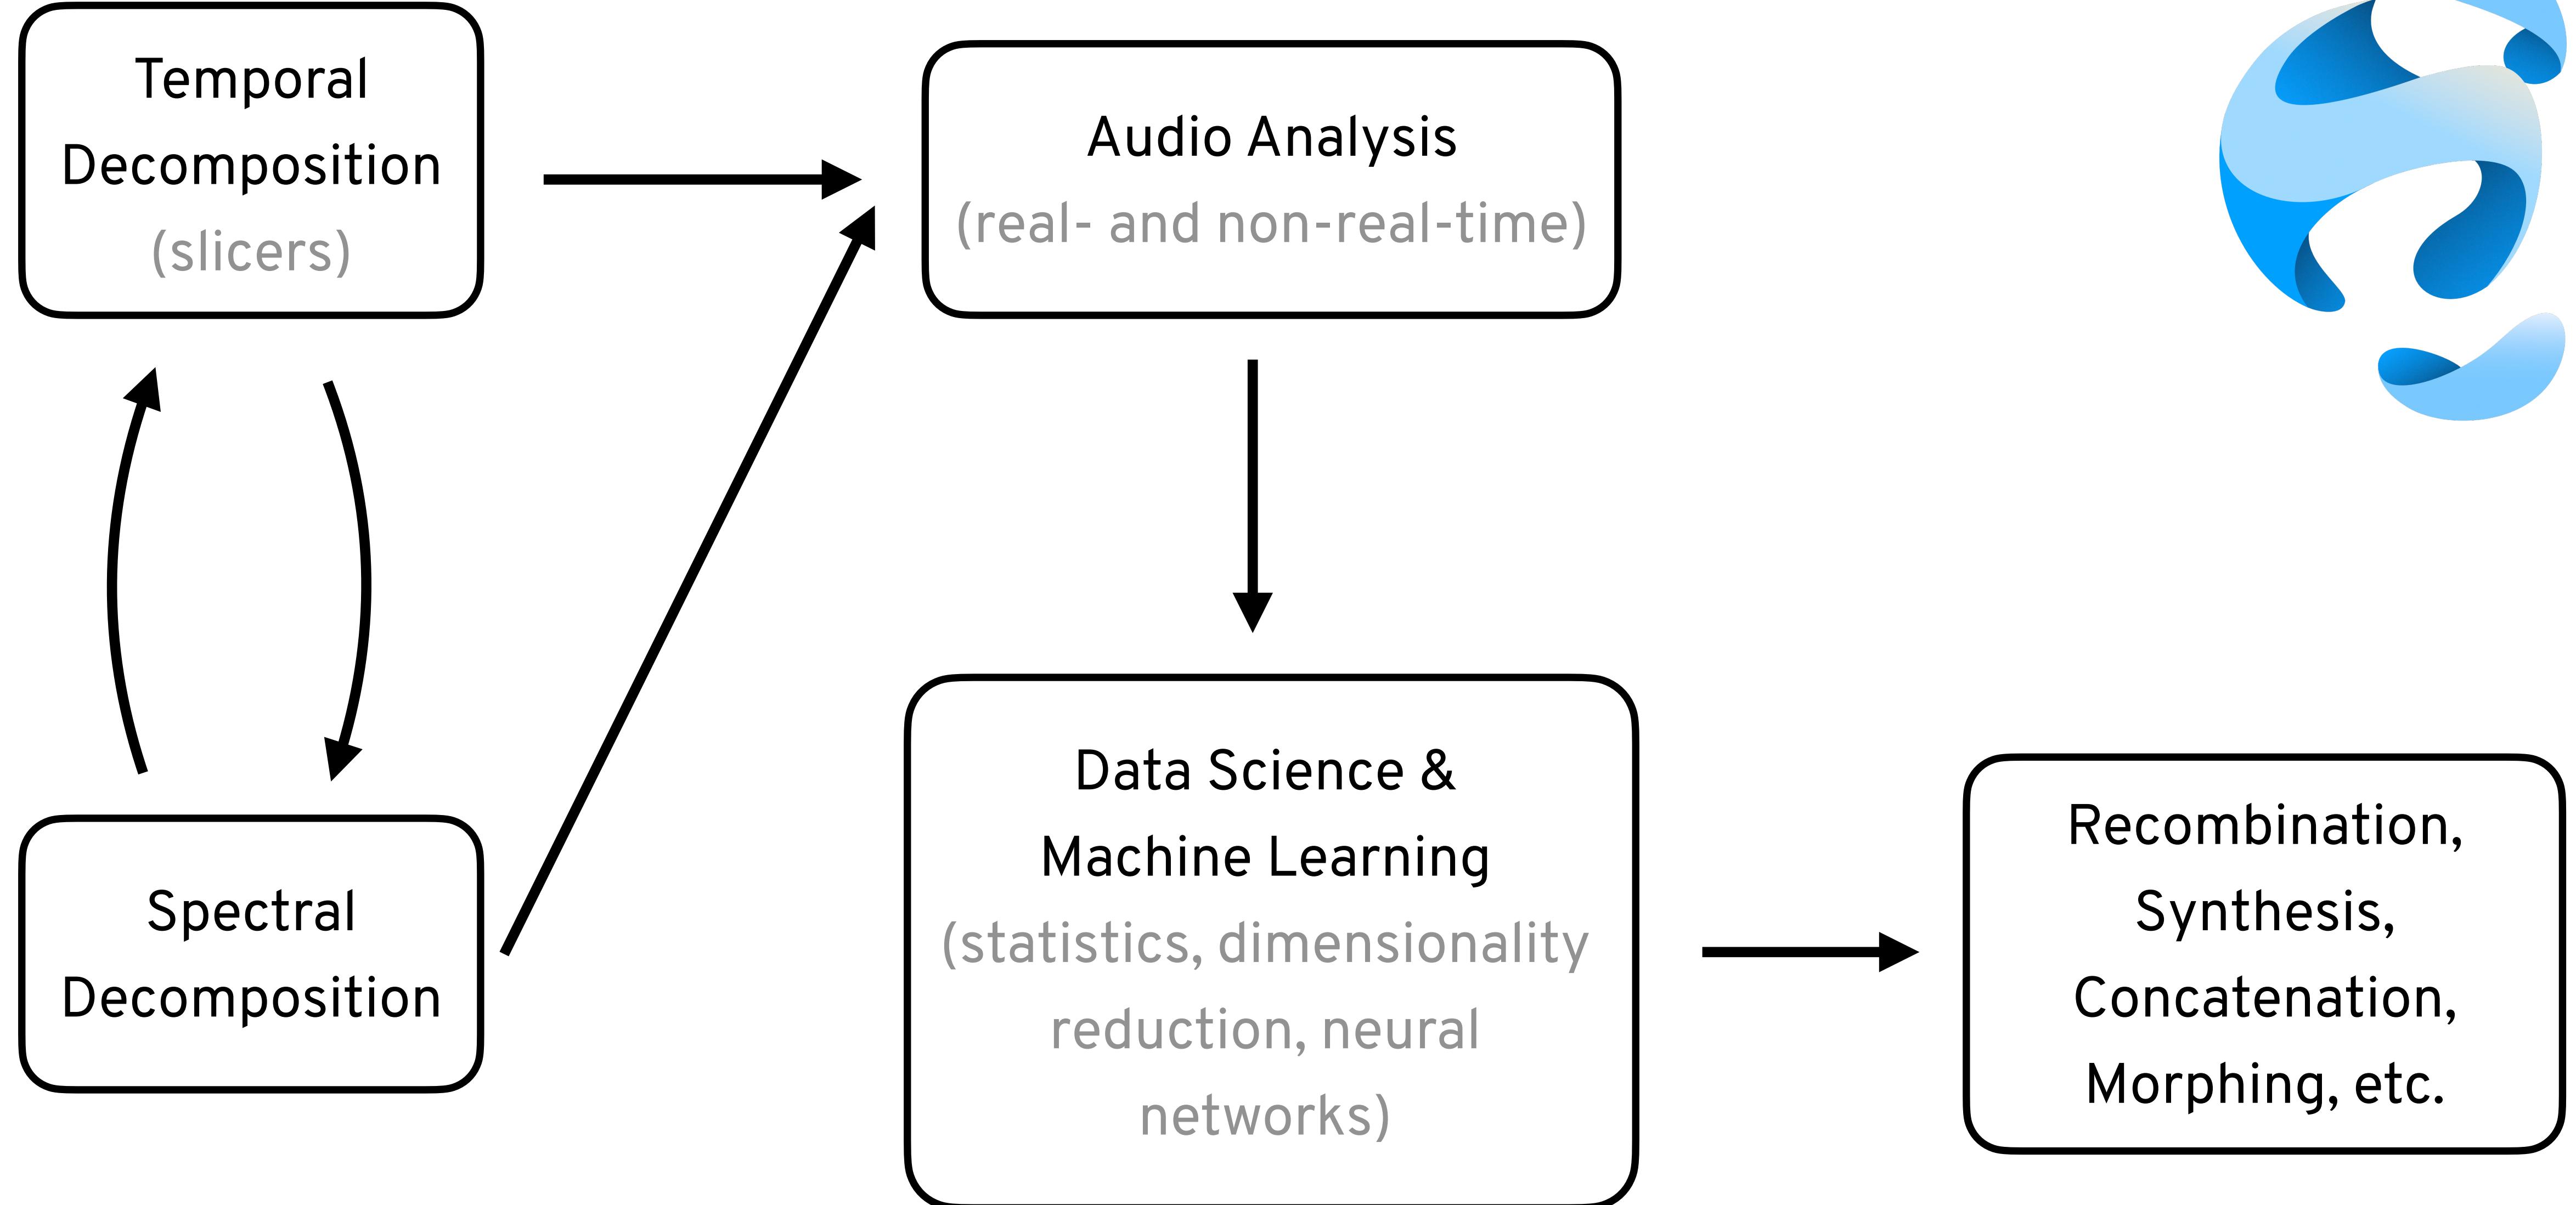
## Transform Audio

- audio transport
- non-negative matrix factorisation filters & morphing

## Analyse Data

- datasets
- labelsets
- statistical analyses
- normalization
- standardization
- robust scaler
- principal component analysis
- MDS
- KDTree
- K Nearest Neighbours
- neural networks
- SQL-type query
- KMeans
- UMAP
- grid

... and more



# Plan for the Week...



# FluidMLPClassifier

classify sounds by  
timbre in real-time



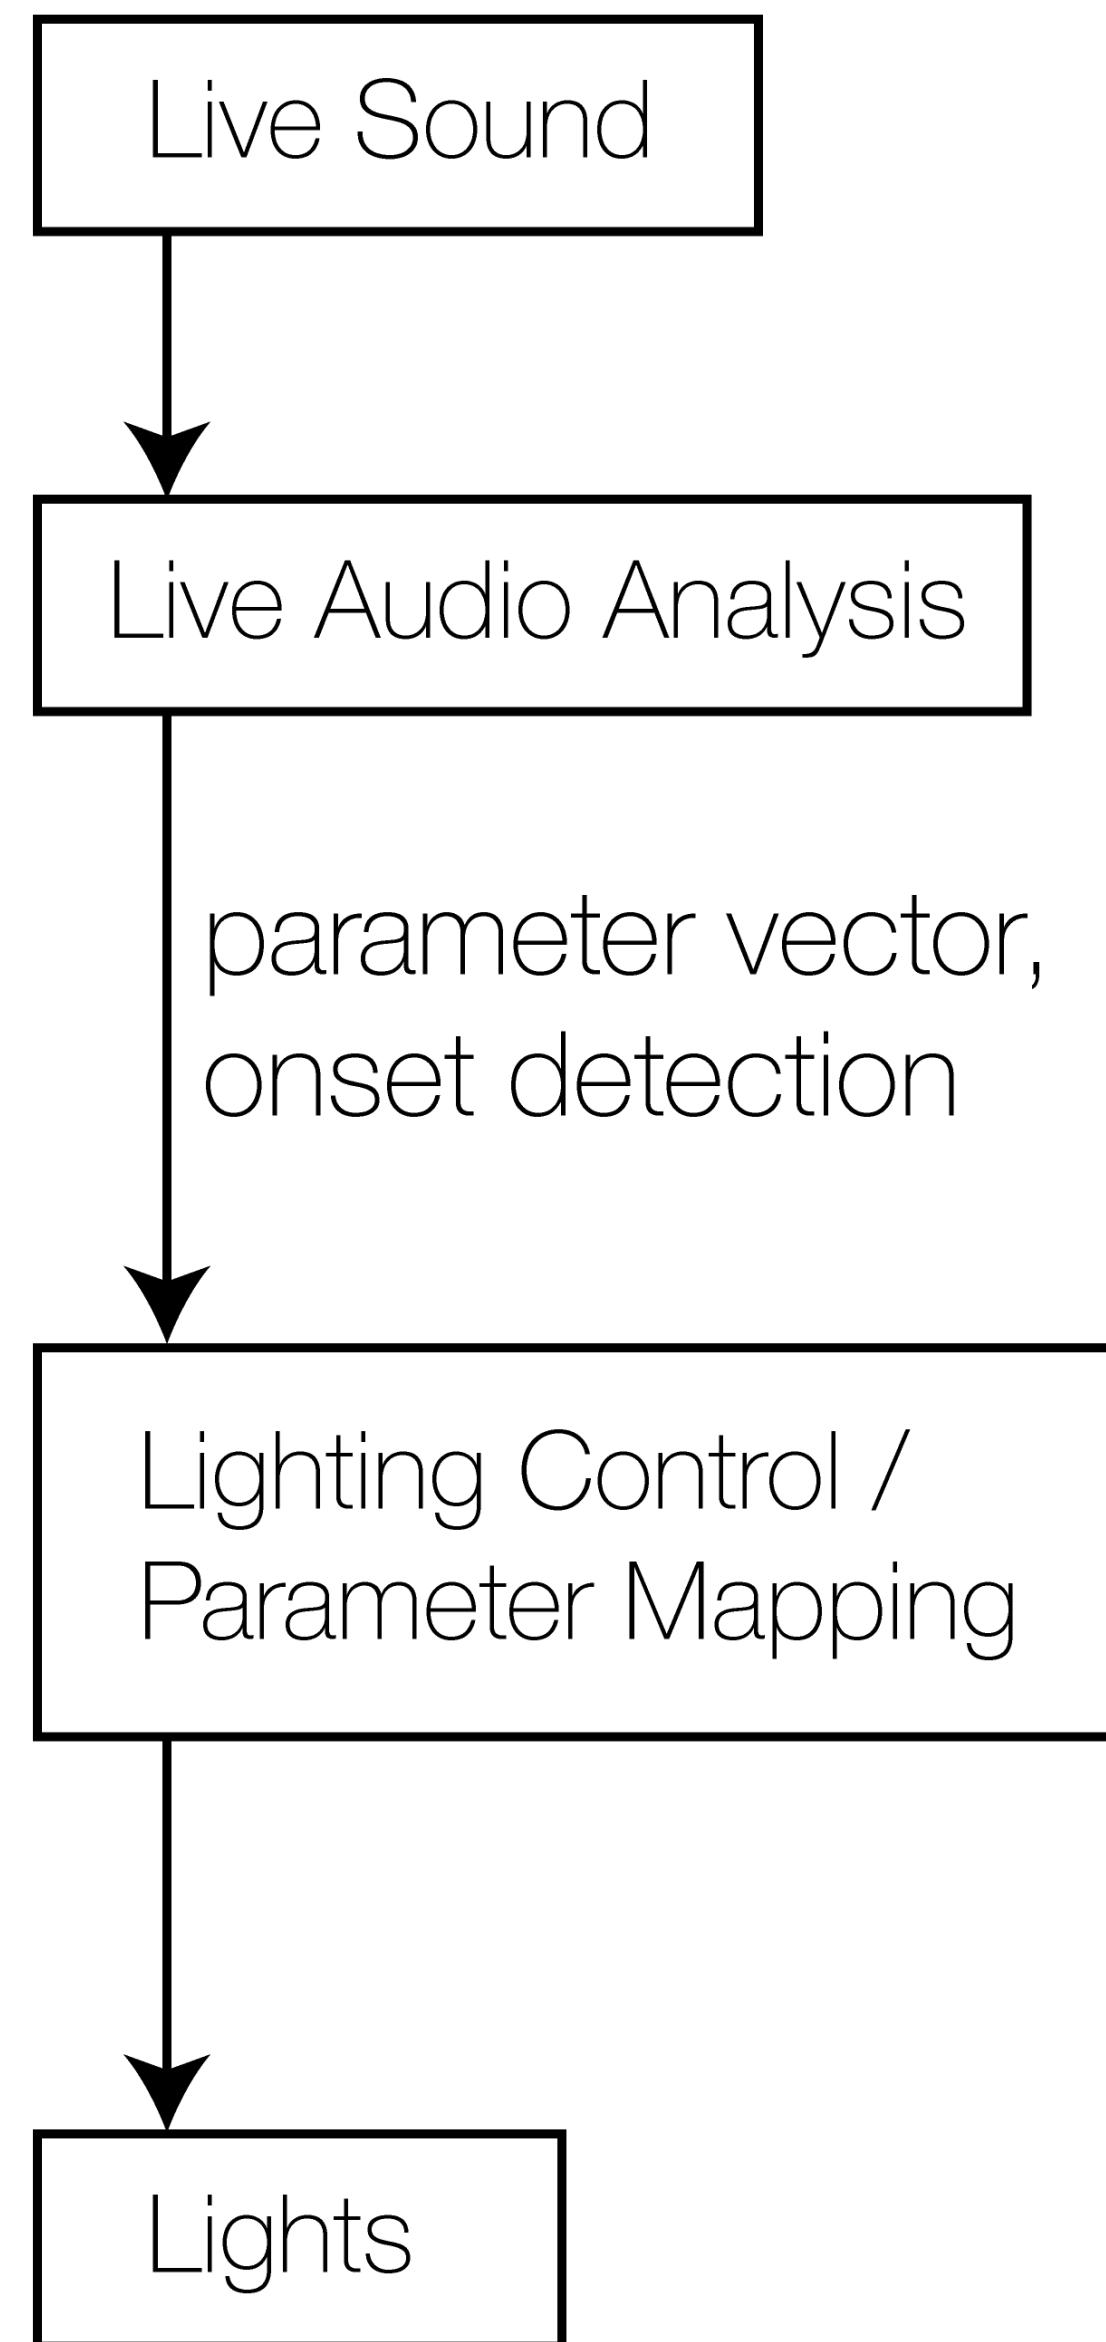
# **feed**

using a neural network for real-time  
audio classification

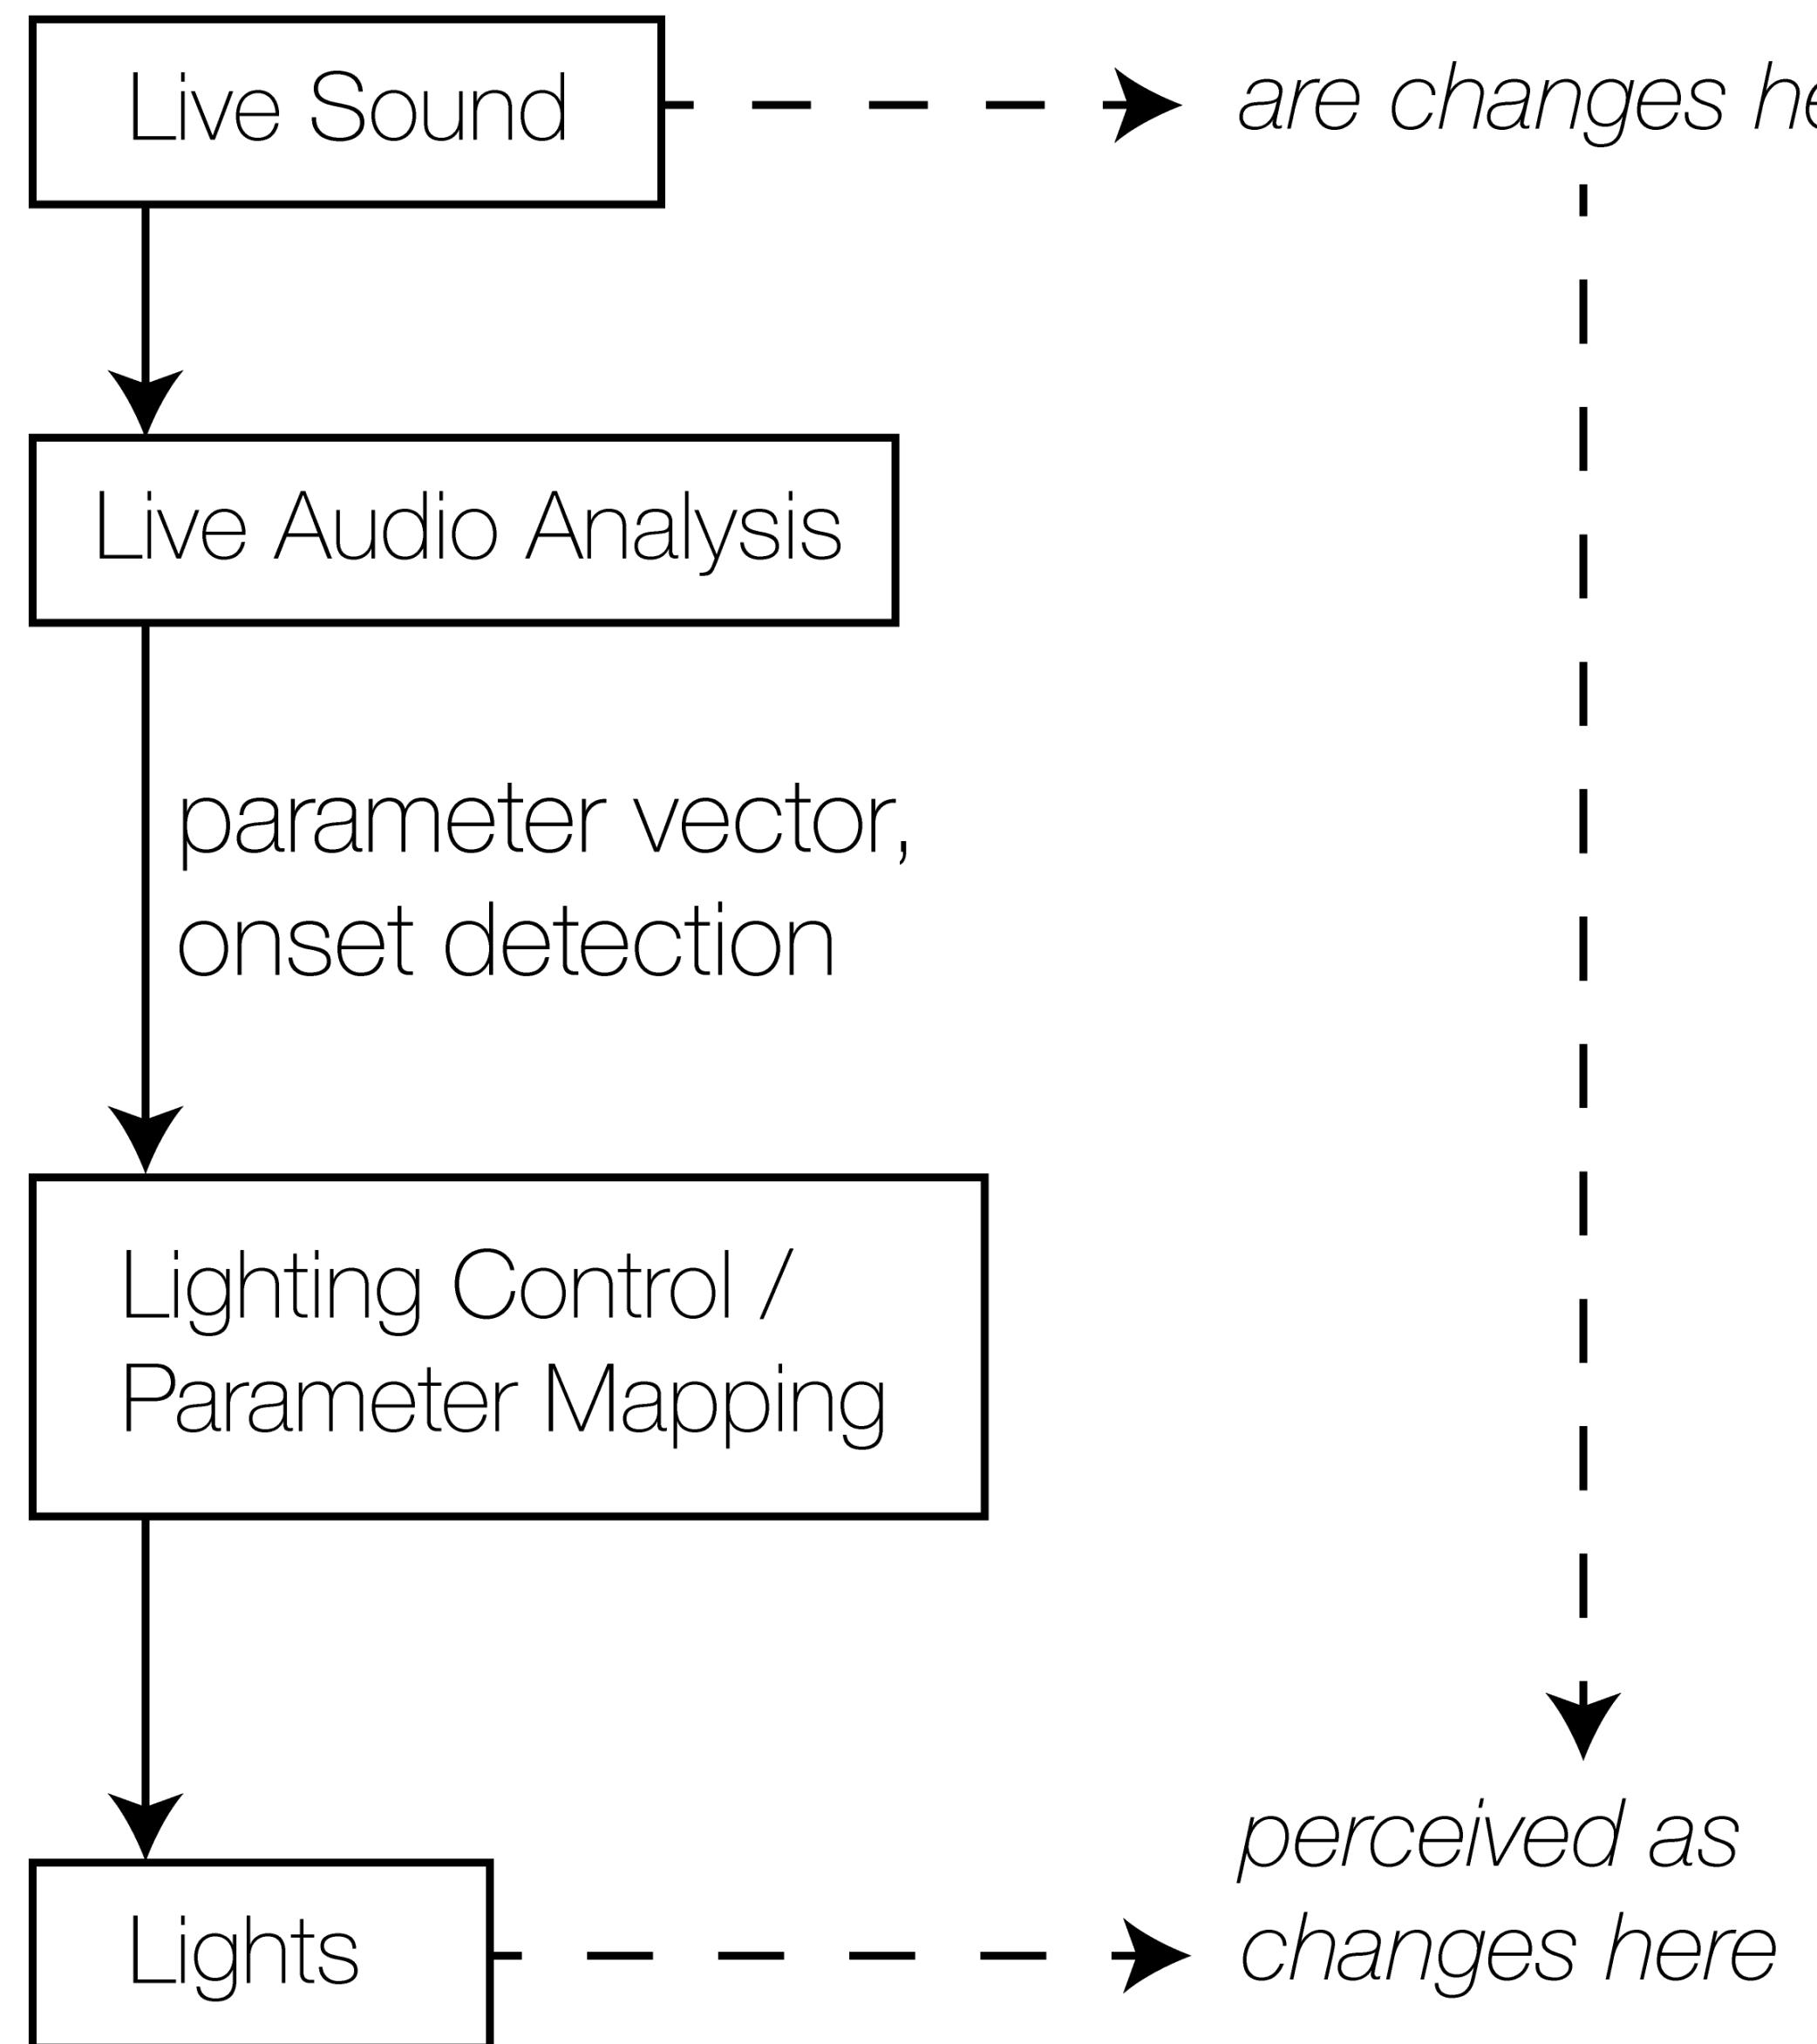
training a machine to hear the way I hear



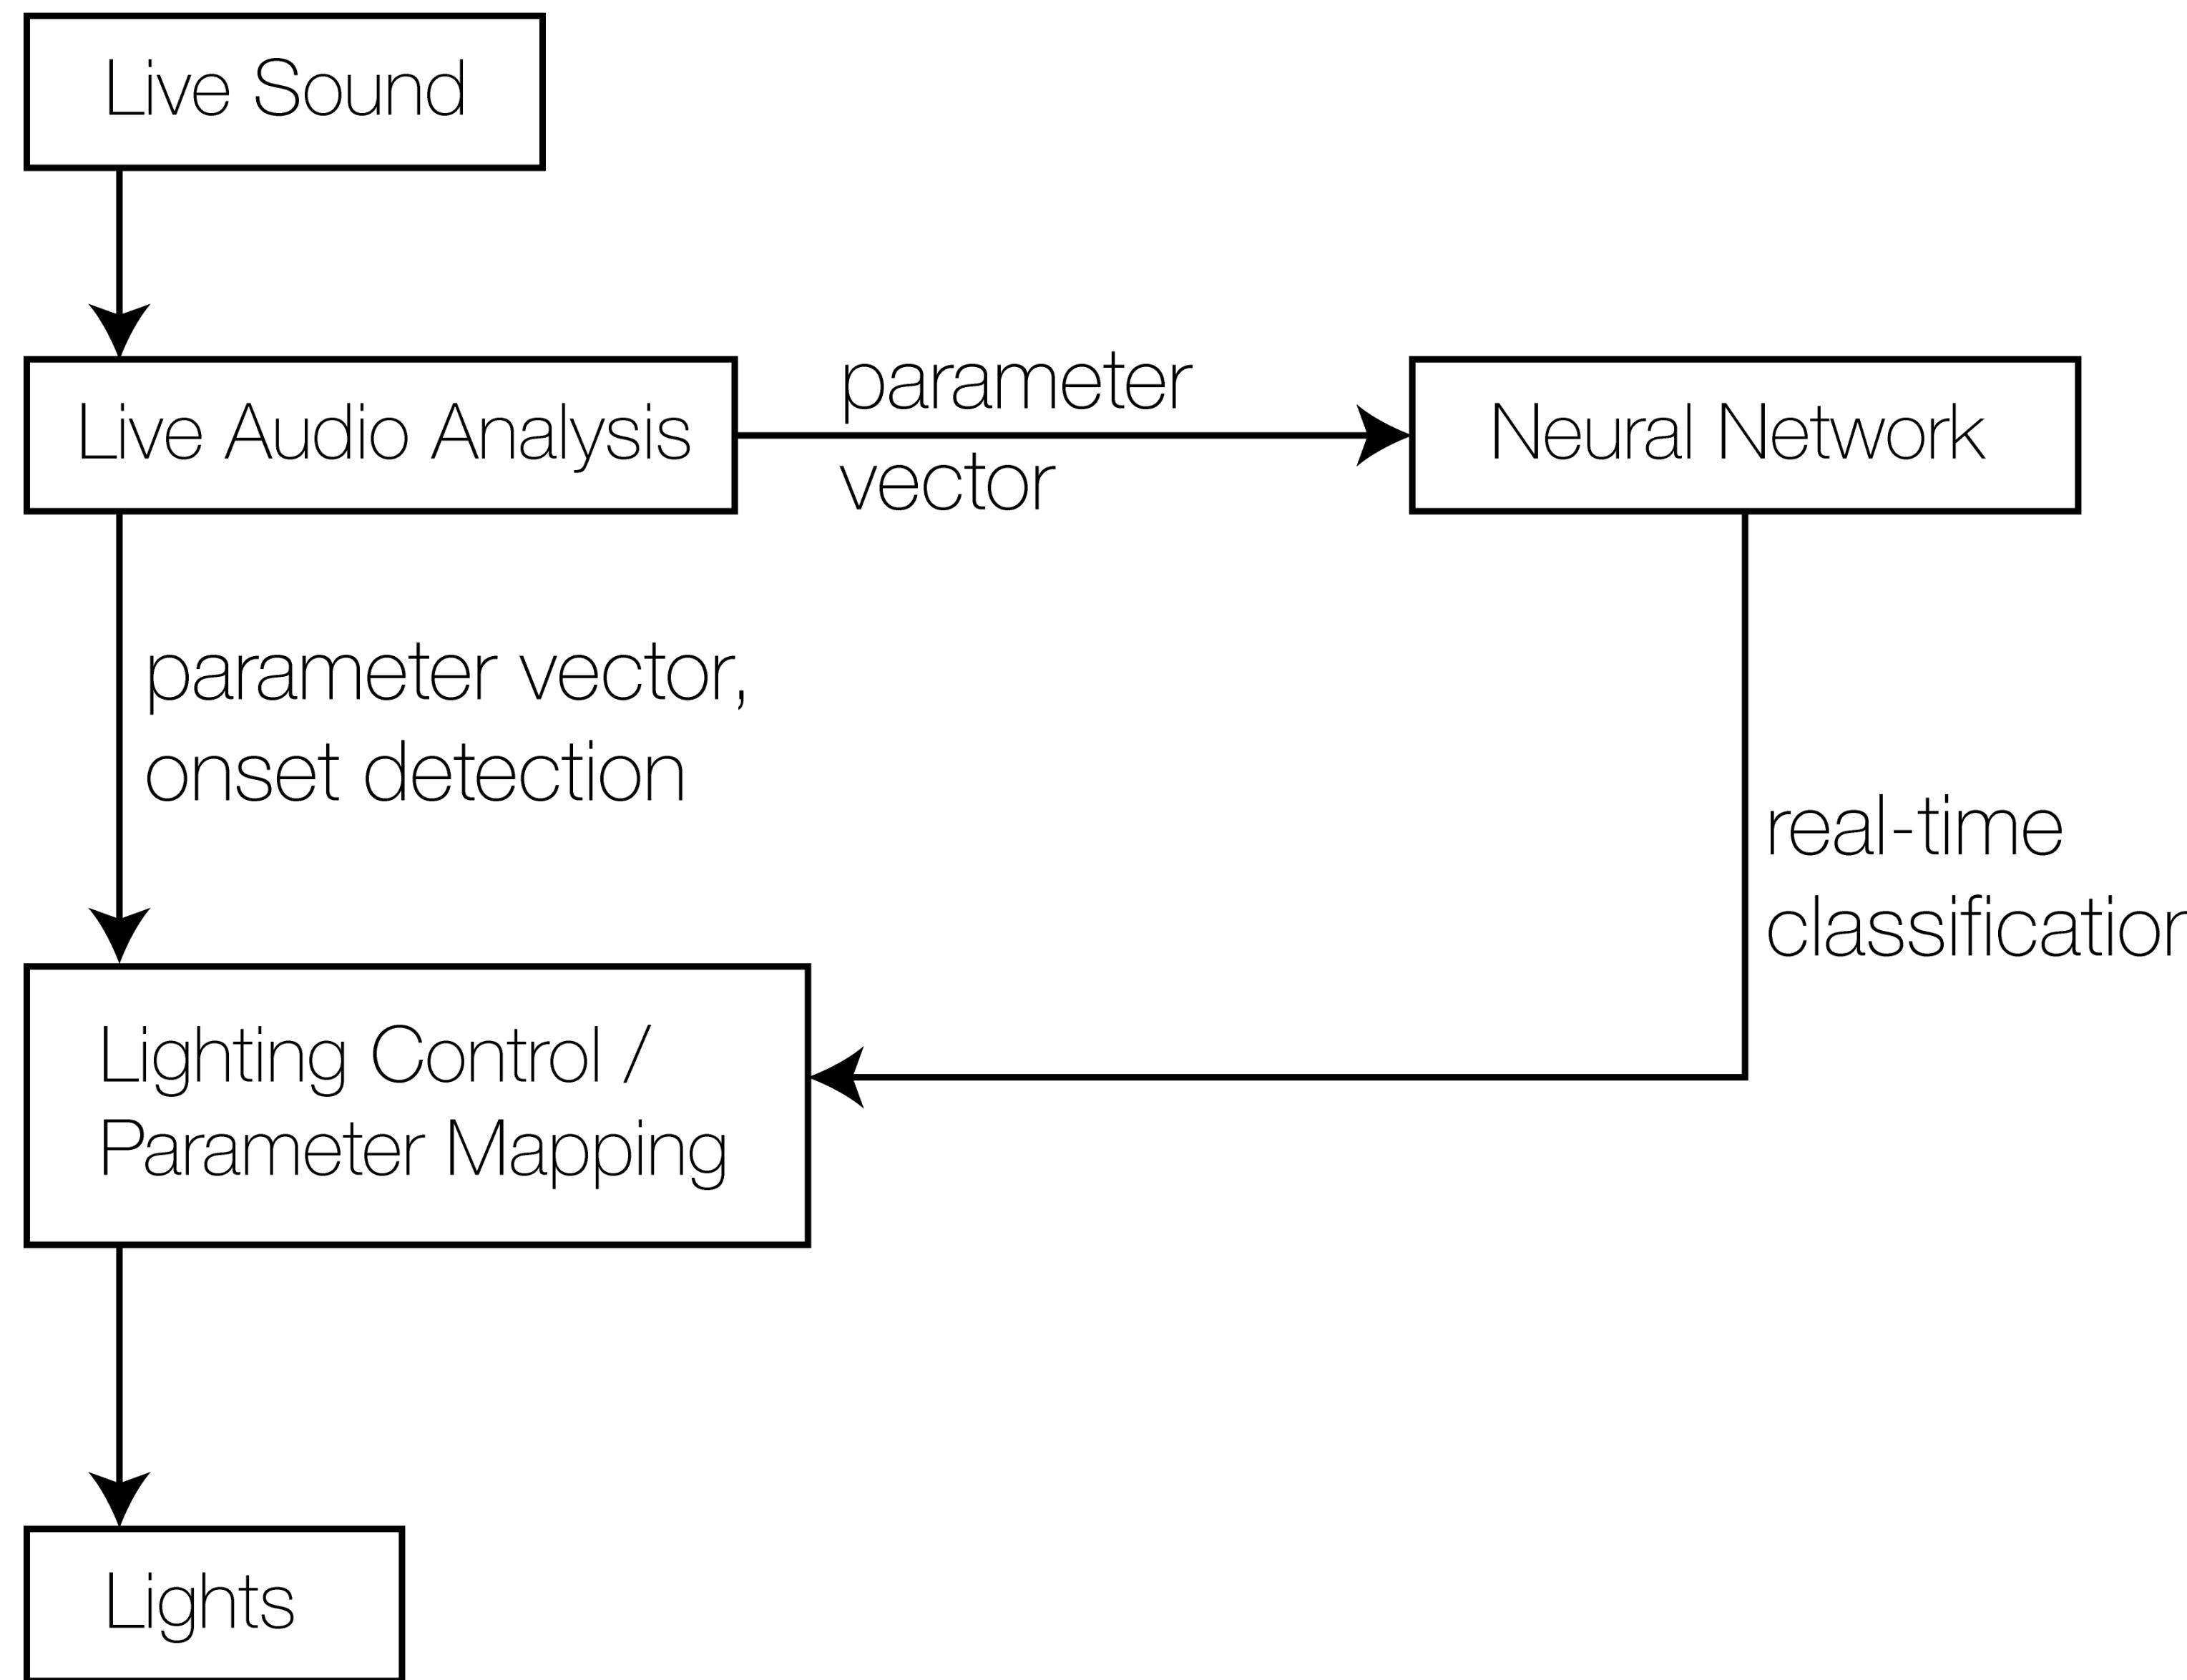
# Machine Listening System



# Machine Listening System



# Machine Learning System





distorted\_noise



high\_squeal



low\_impulses



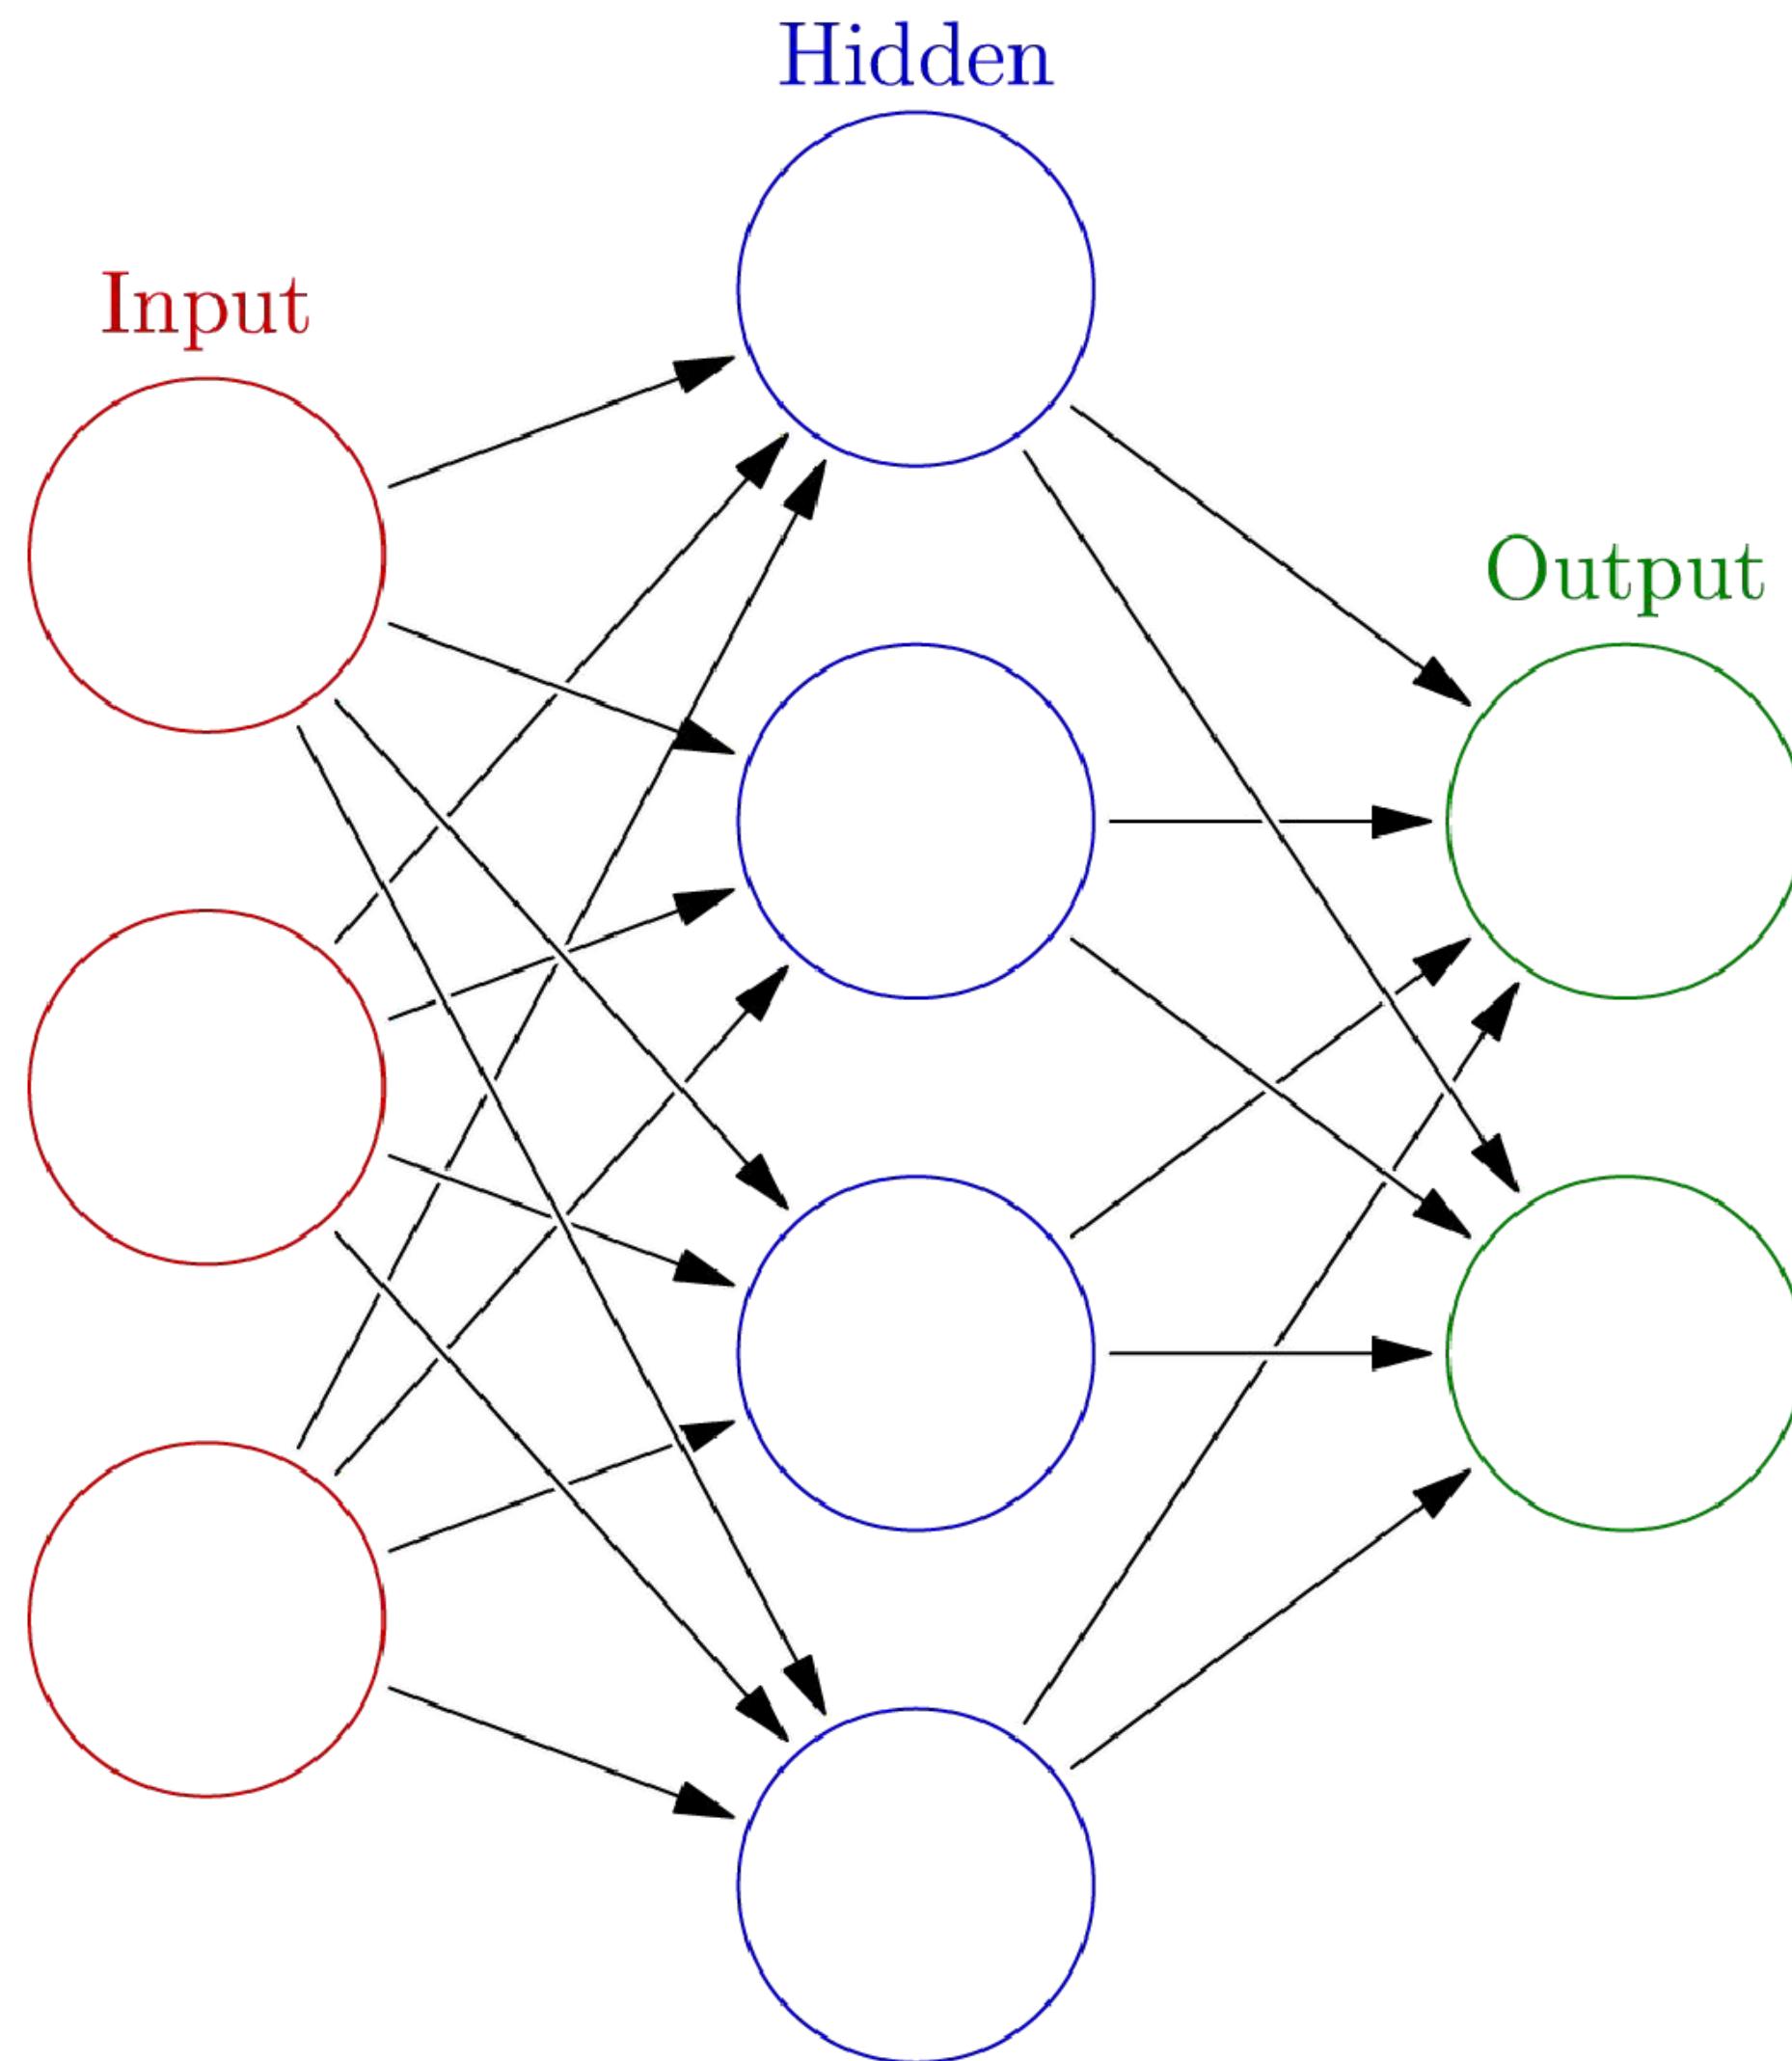
sus\_noise\_quiet

**flashing lights warning**





# *Neural Network (Multi-Layer Perceptron)*

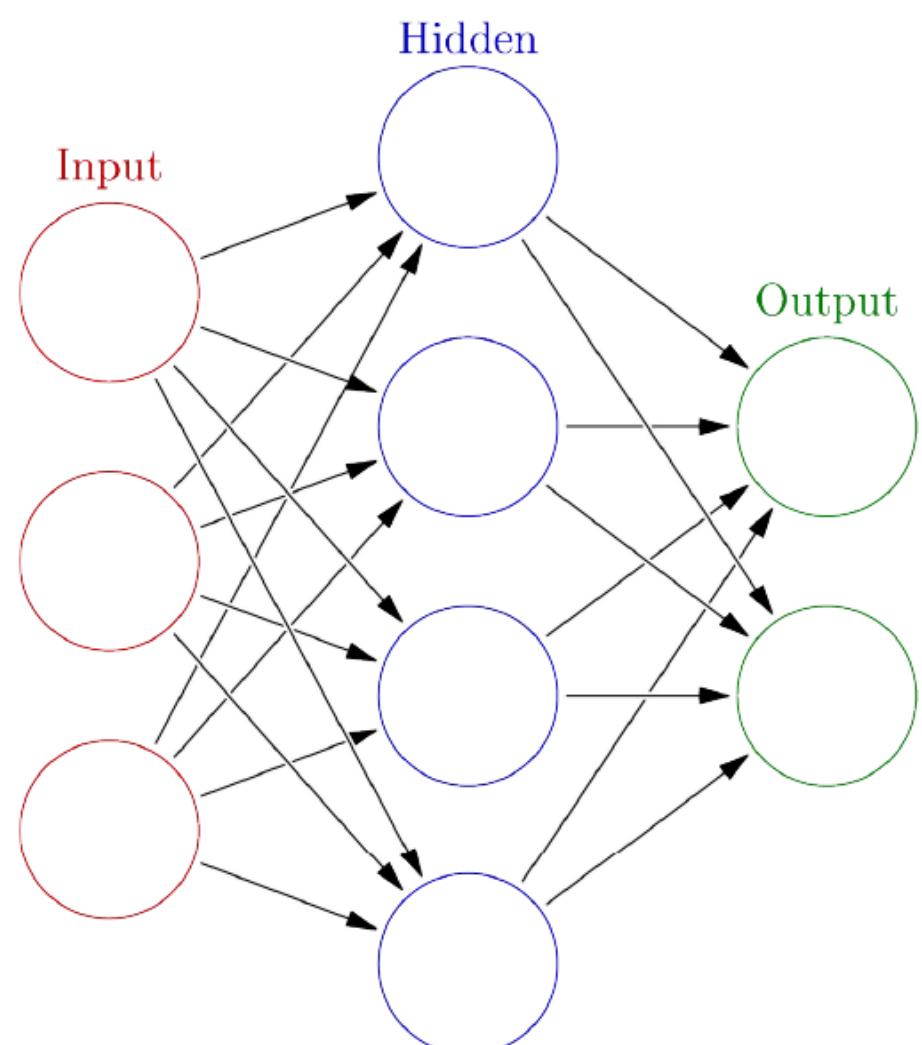


**Classification:**  
a neural network predicts  
which category (or “class”)  
an input belongs to

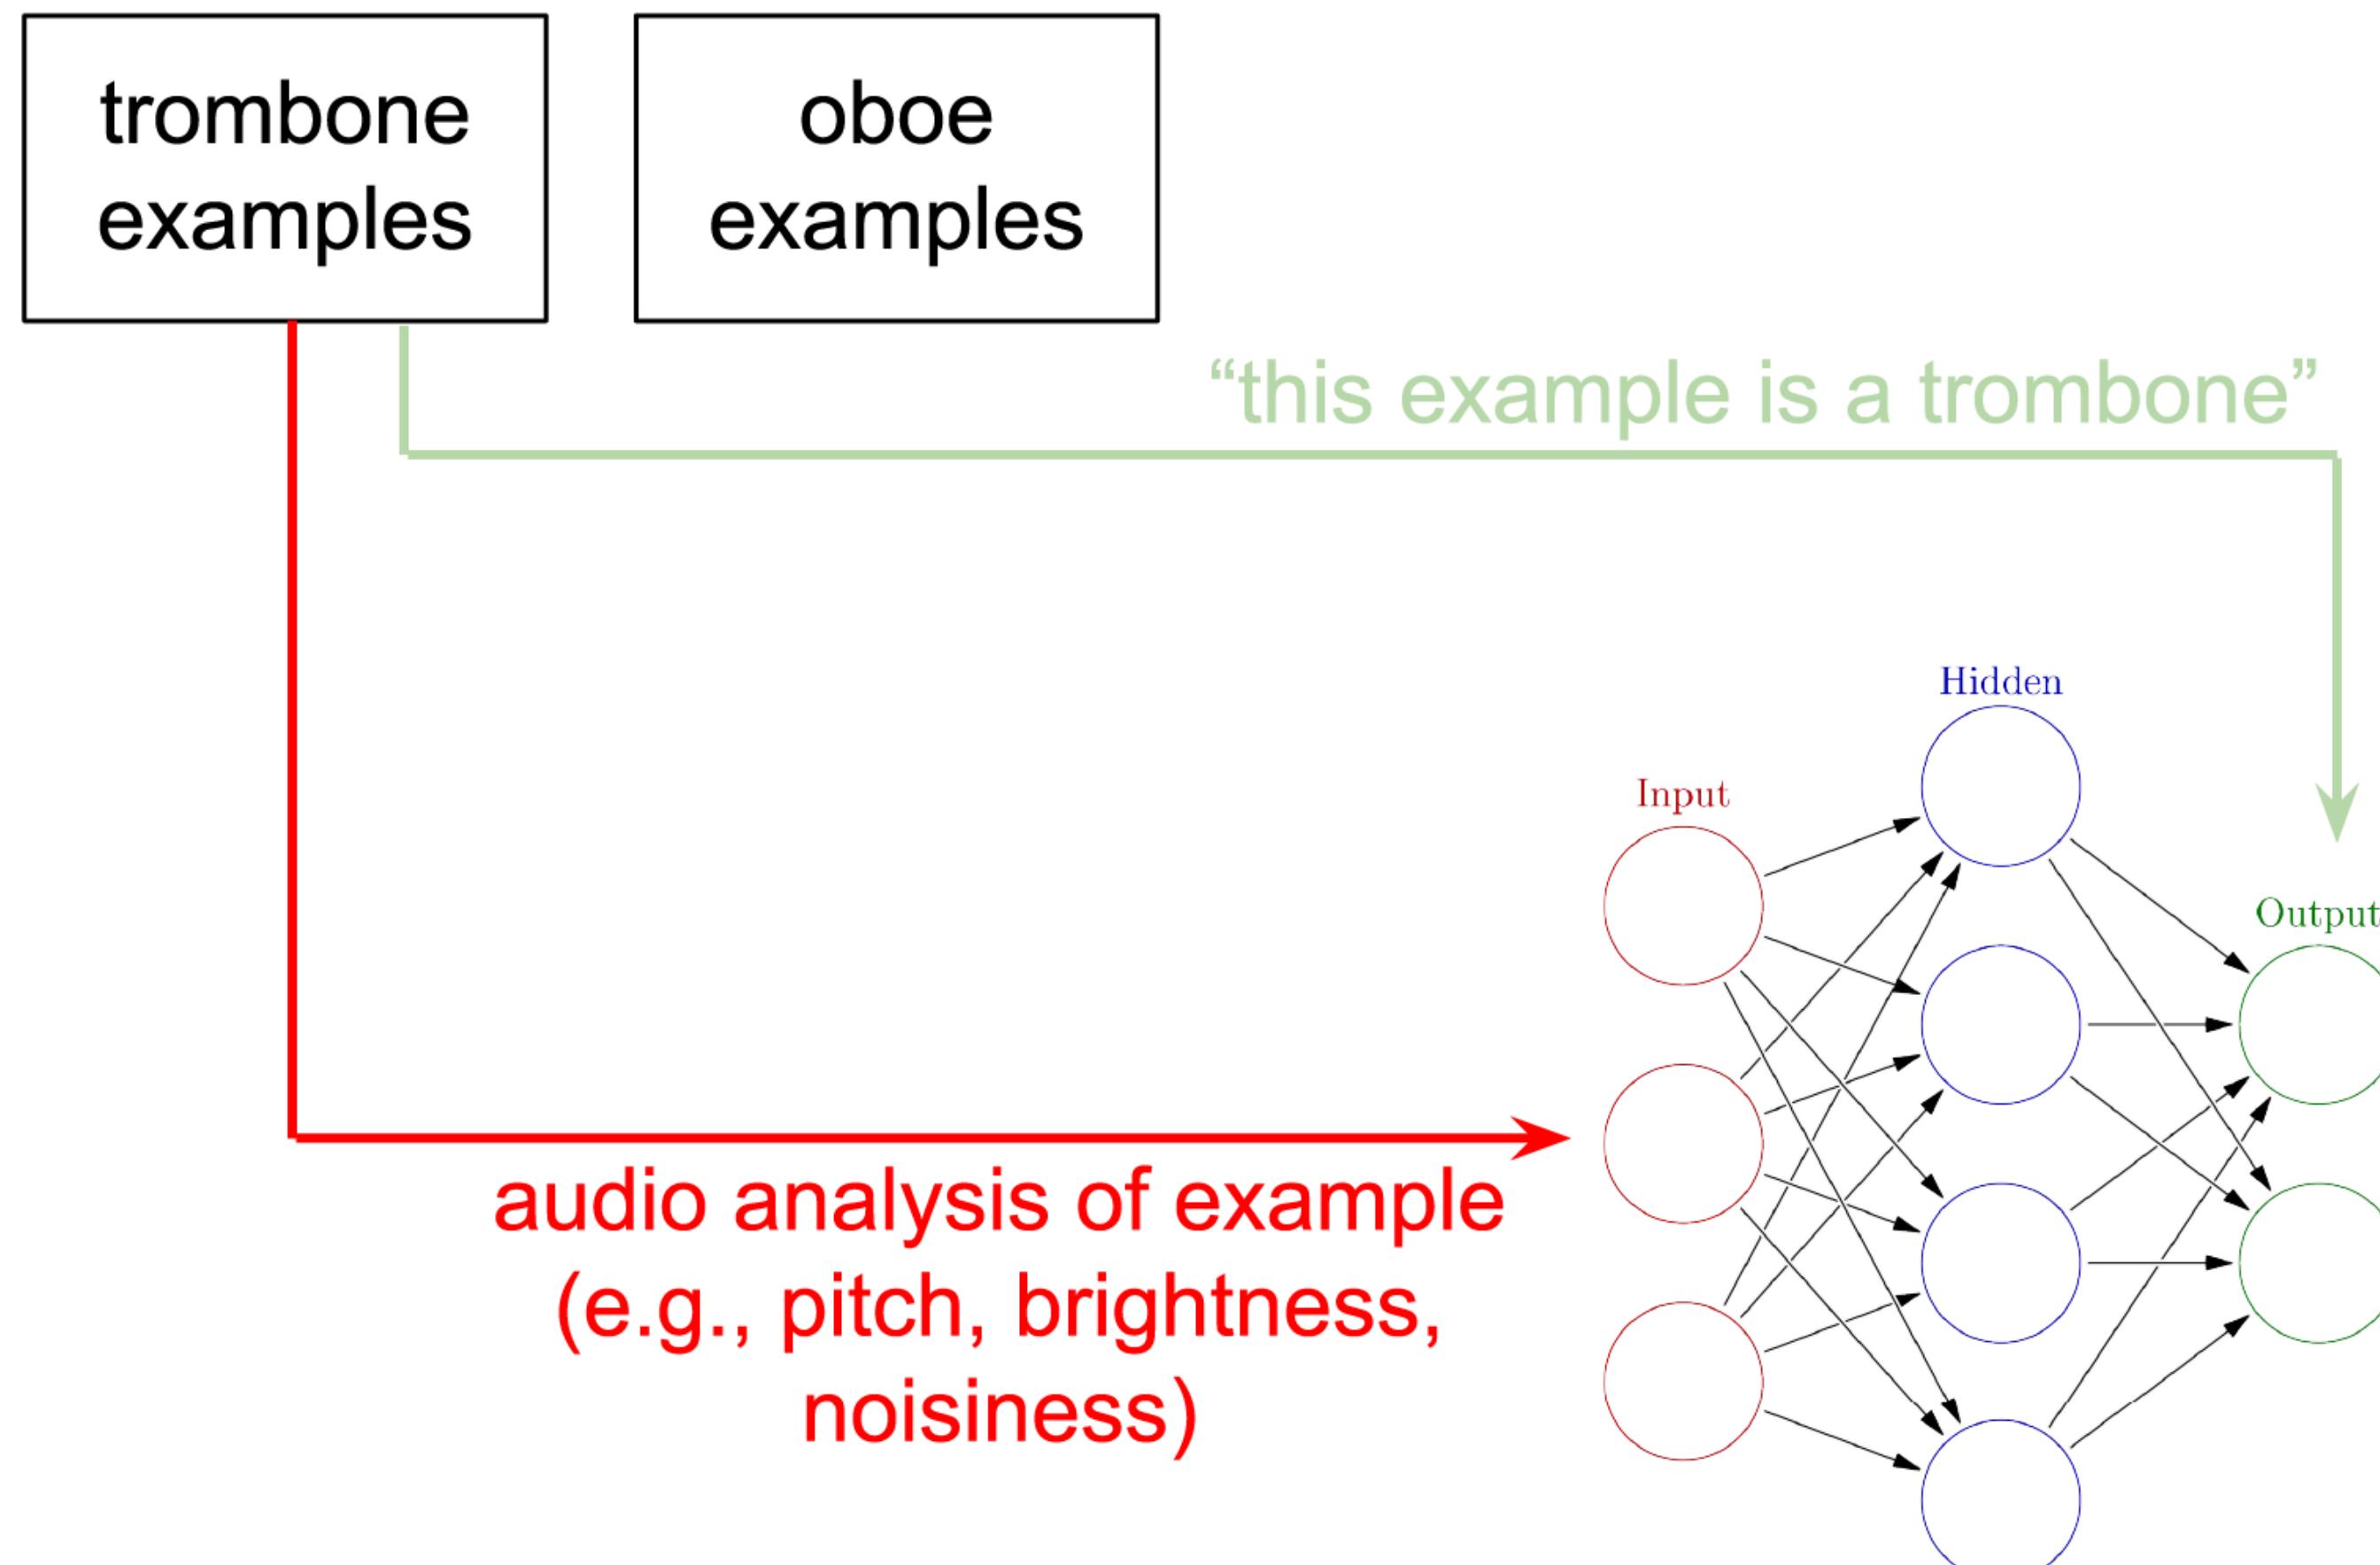
# *Neural Network* ***Training a Classifier***

trombone  
examples

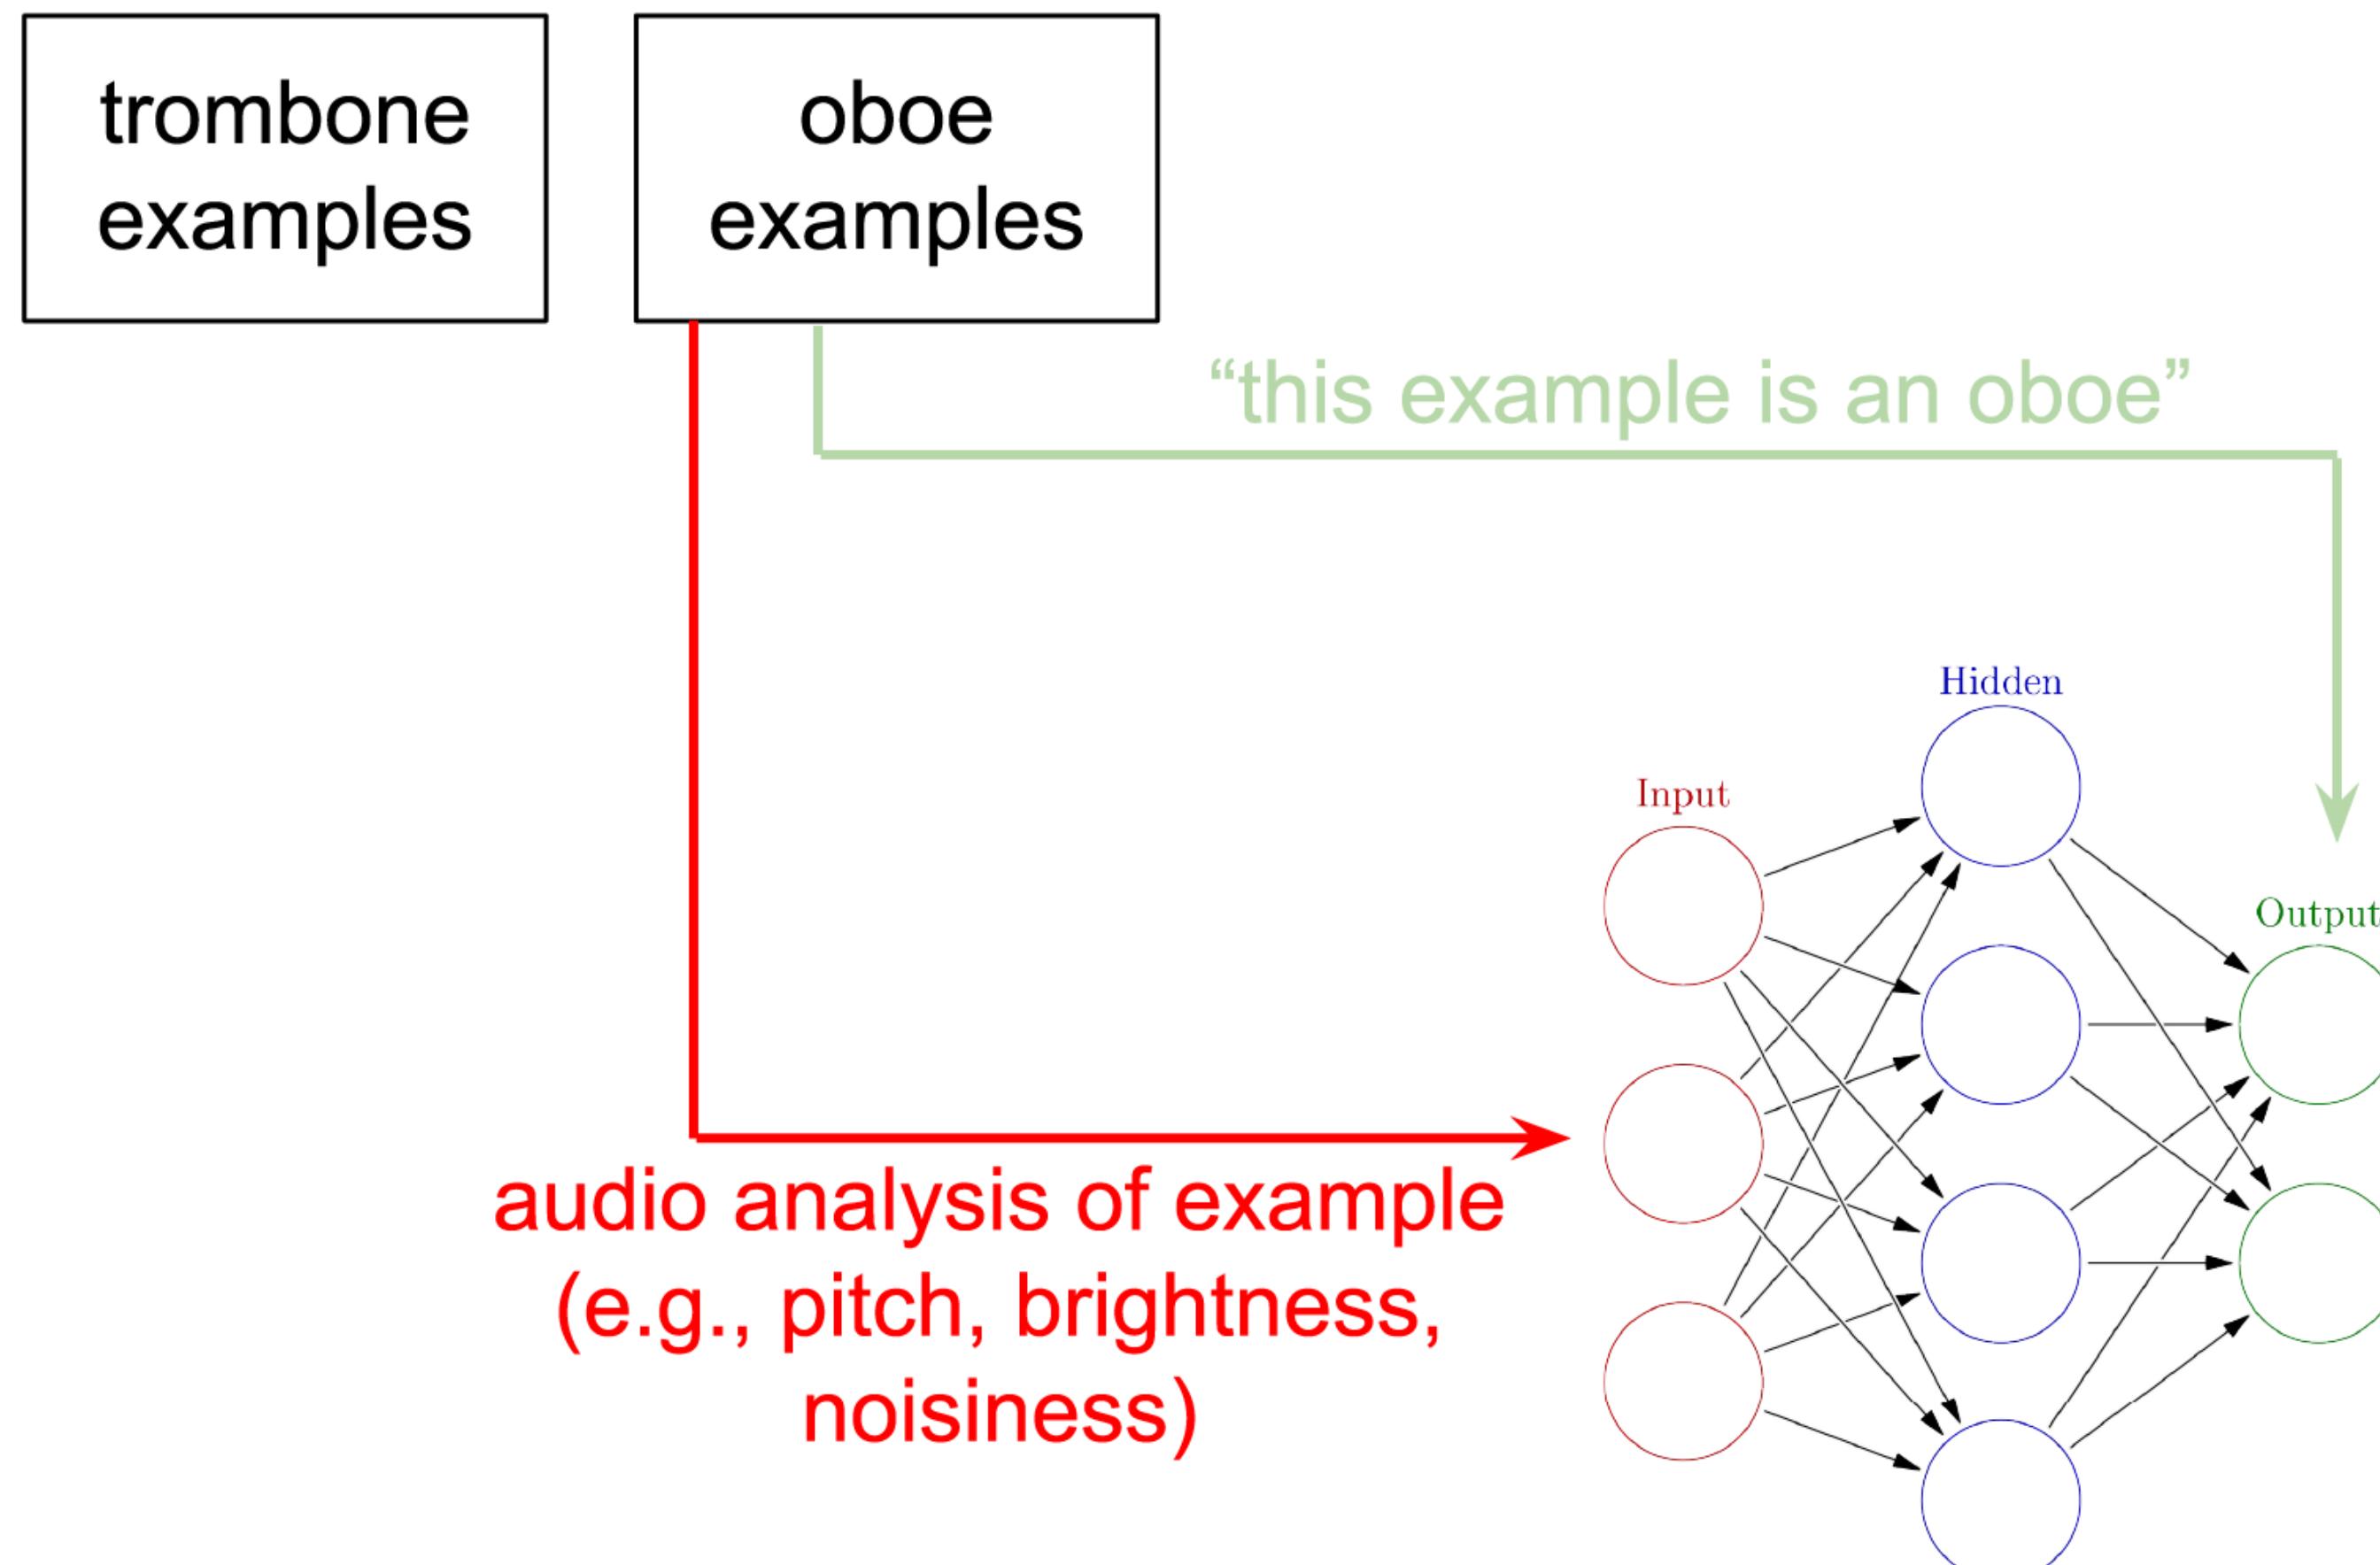
oboe  
examples



# *Neural Network* **Training a Classifier**



# *Neural Network* **Training a Classifier**



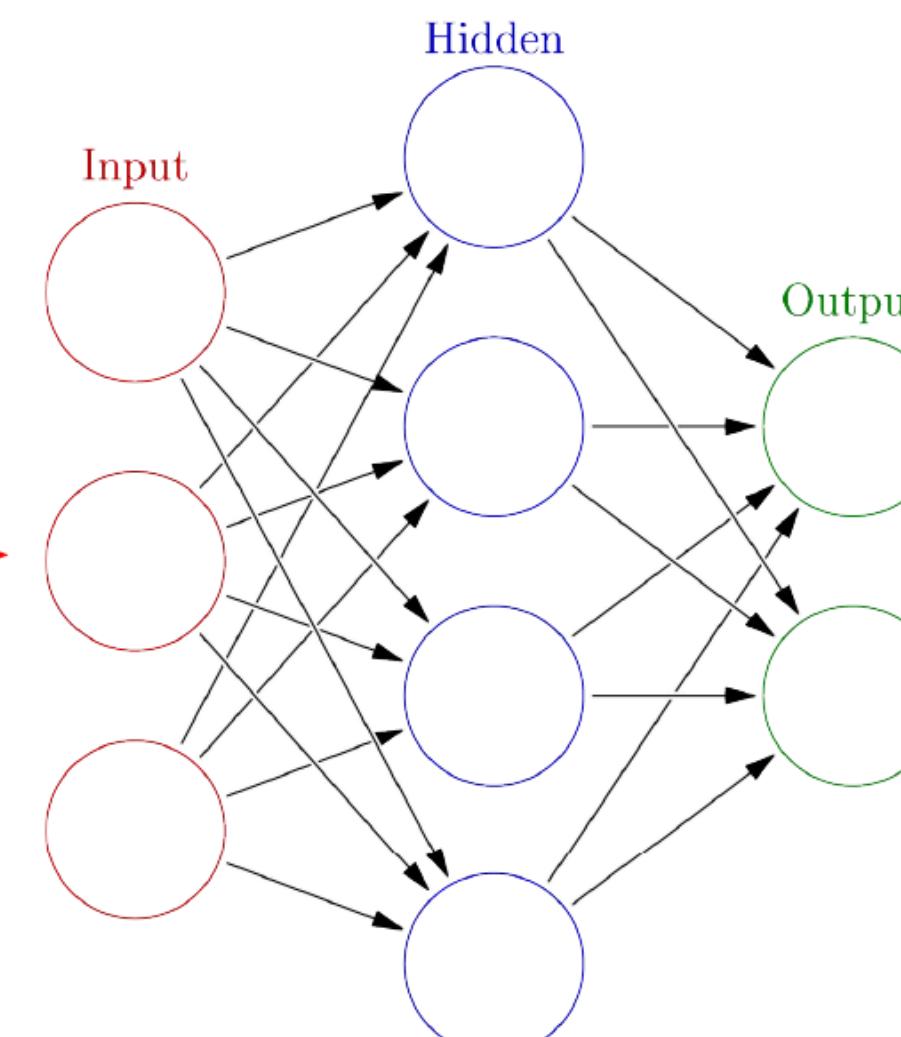
# Neural Network Predicting a Classification

trombone  
examples

oboe  
examples

new example it has  
never seen before

audio analysis of example  
(e.g., pitch, brightness,  
noisiness)



“this new example is most  
like the oboe examples you  
showed me before”  
(or trombone...)

*fluid.mlpclassifier~*

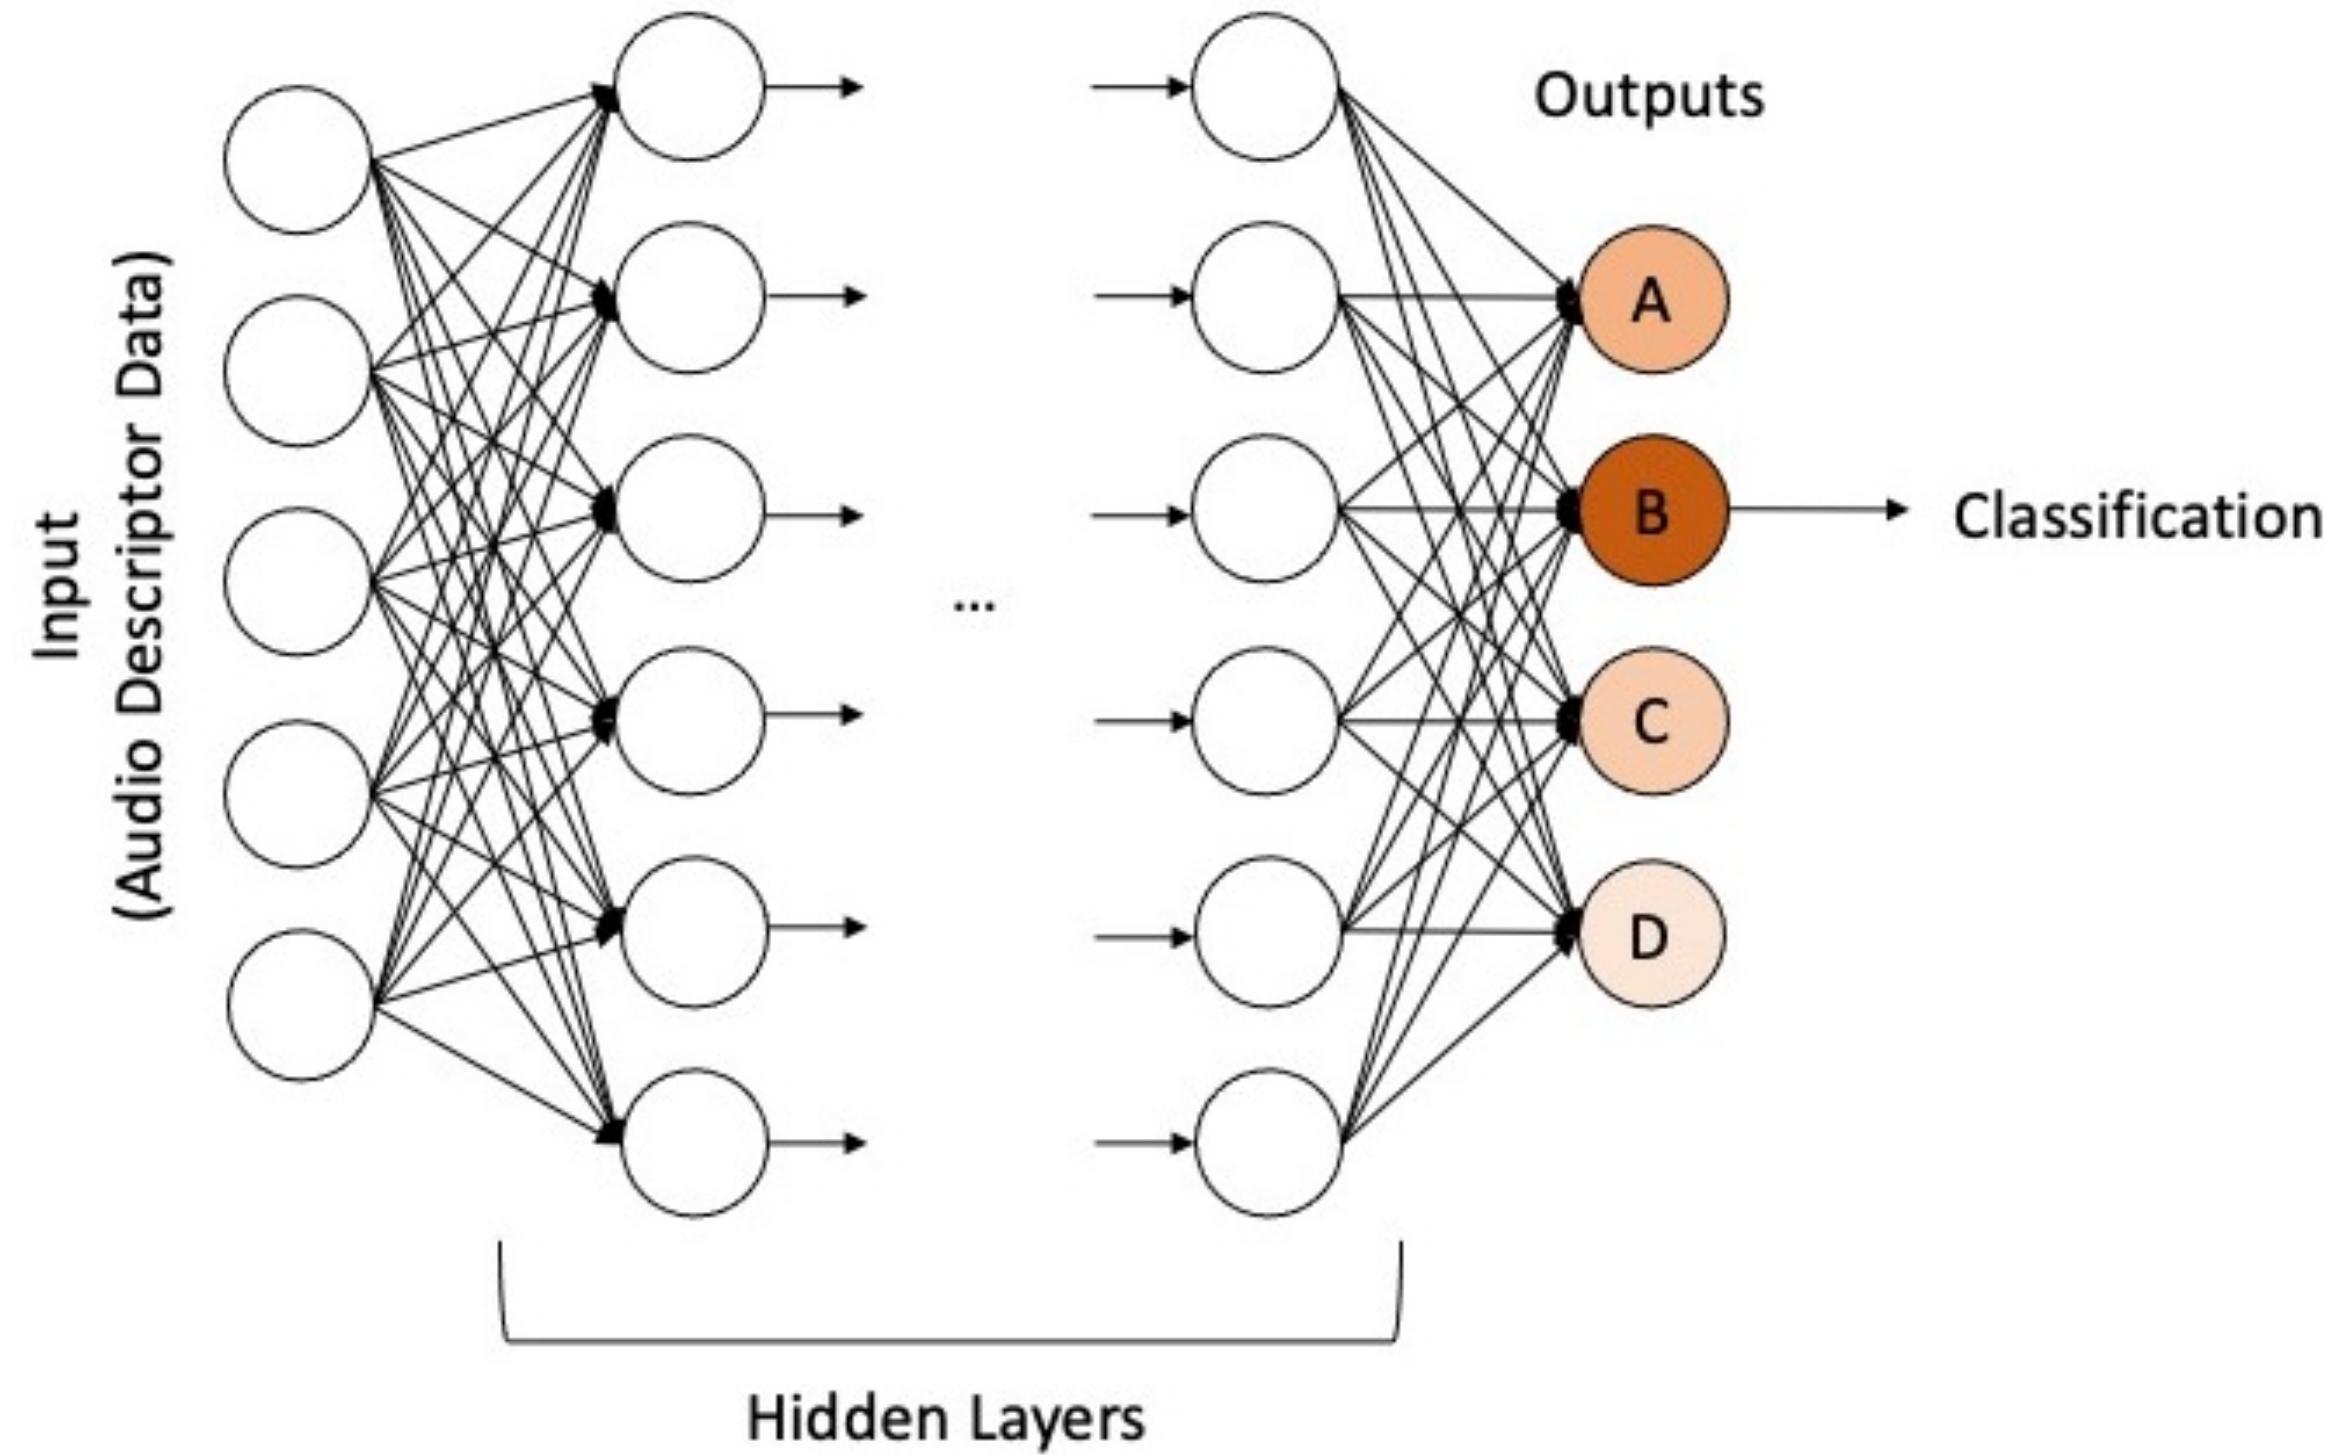


***fluid.mlpclassifier~***

validation



- **overfitting**
  - excellent on training data, poor on anything else
- **underfitting**
  - poor on everything...keep training!



# *Drift Shadow* by Alex Harker

[learn.flucoma.org/explore/harker](http://learn.flucoma.org/explore/harker)

# MLPRegressor

control *many* synthesizer  
parameters from a smaller  
control space







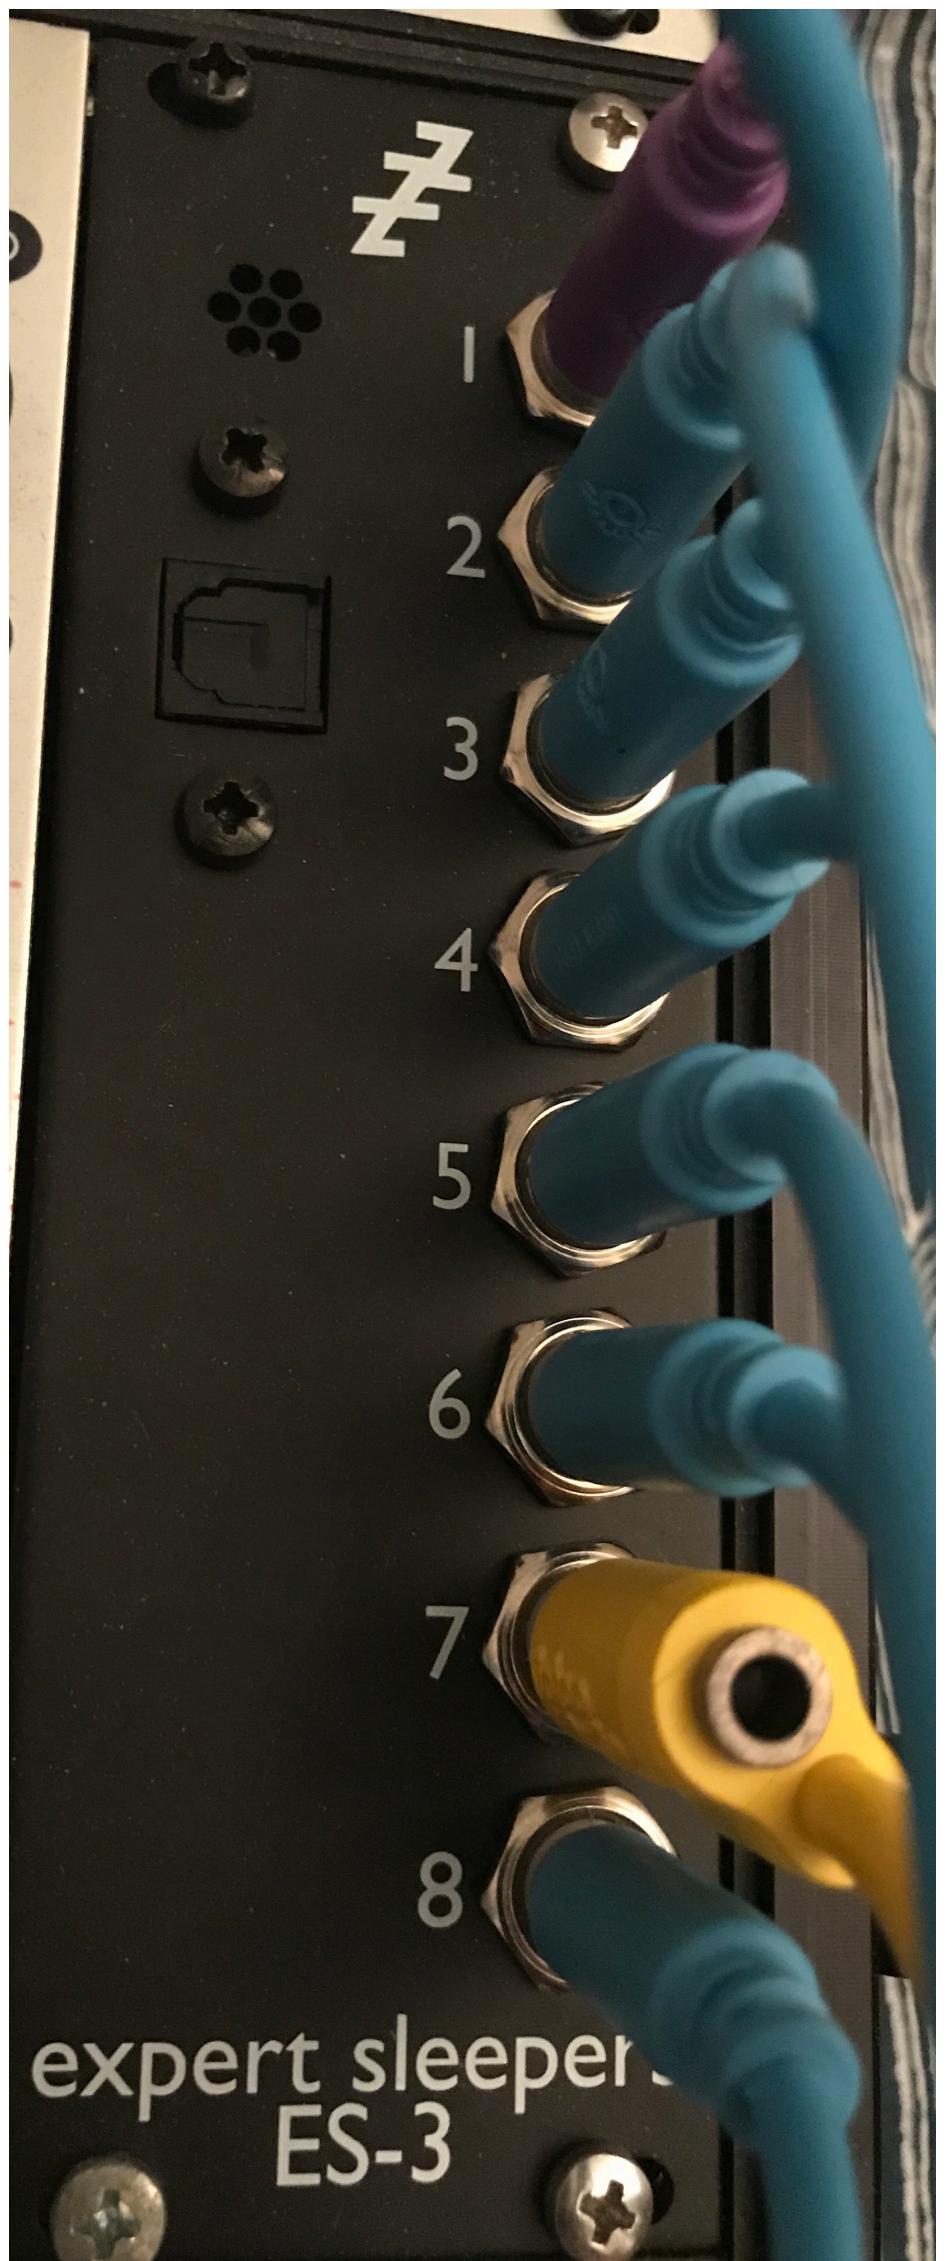
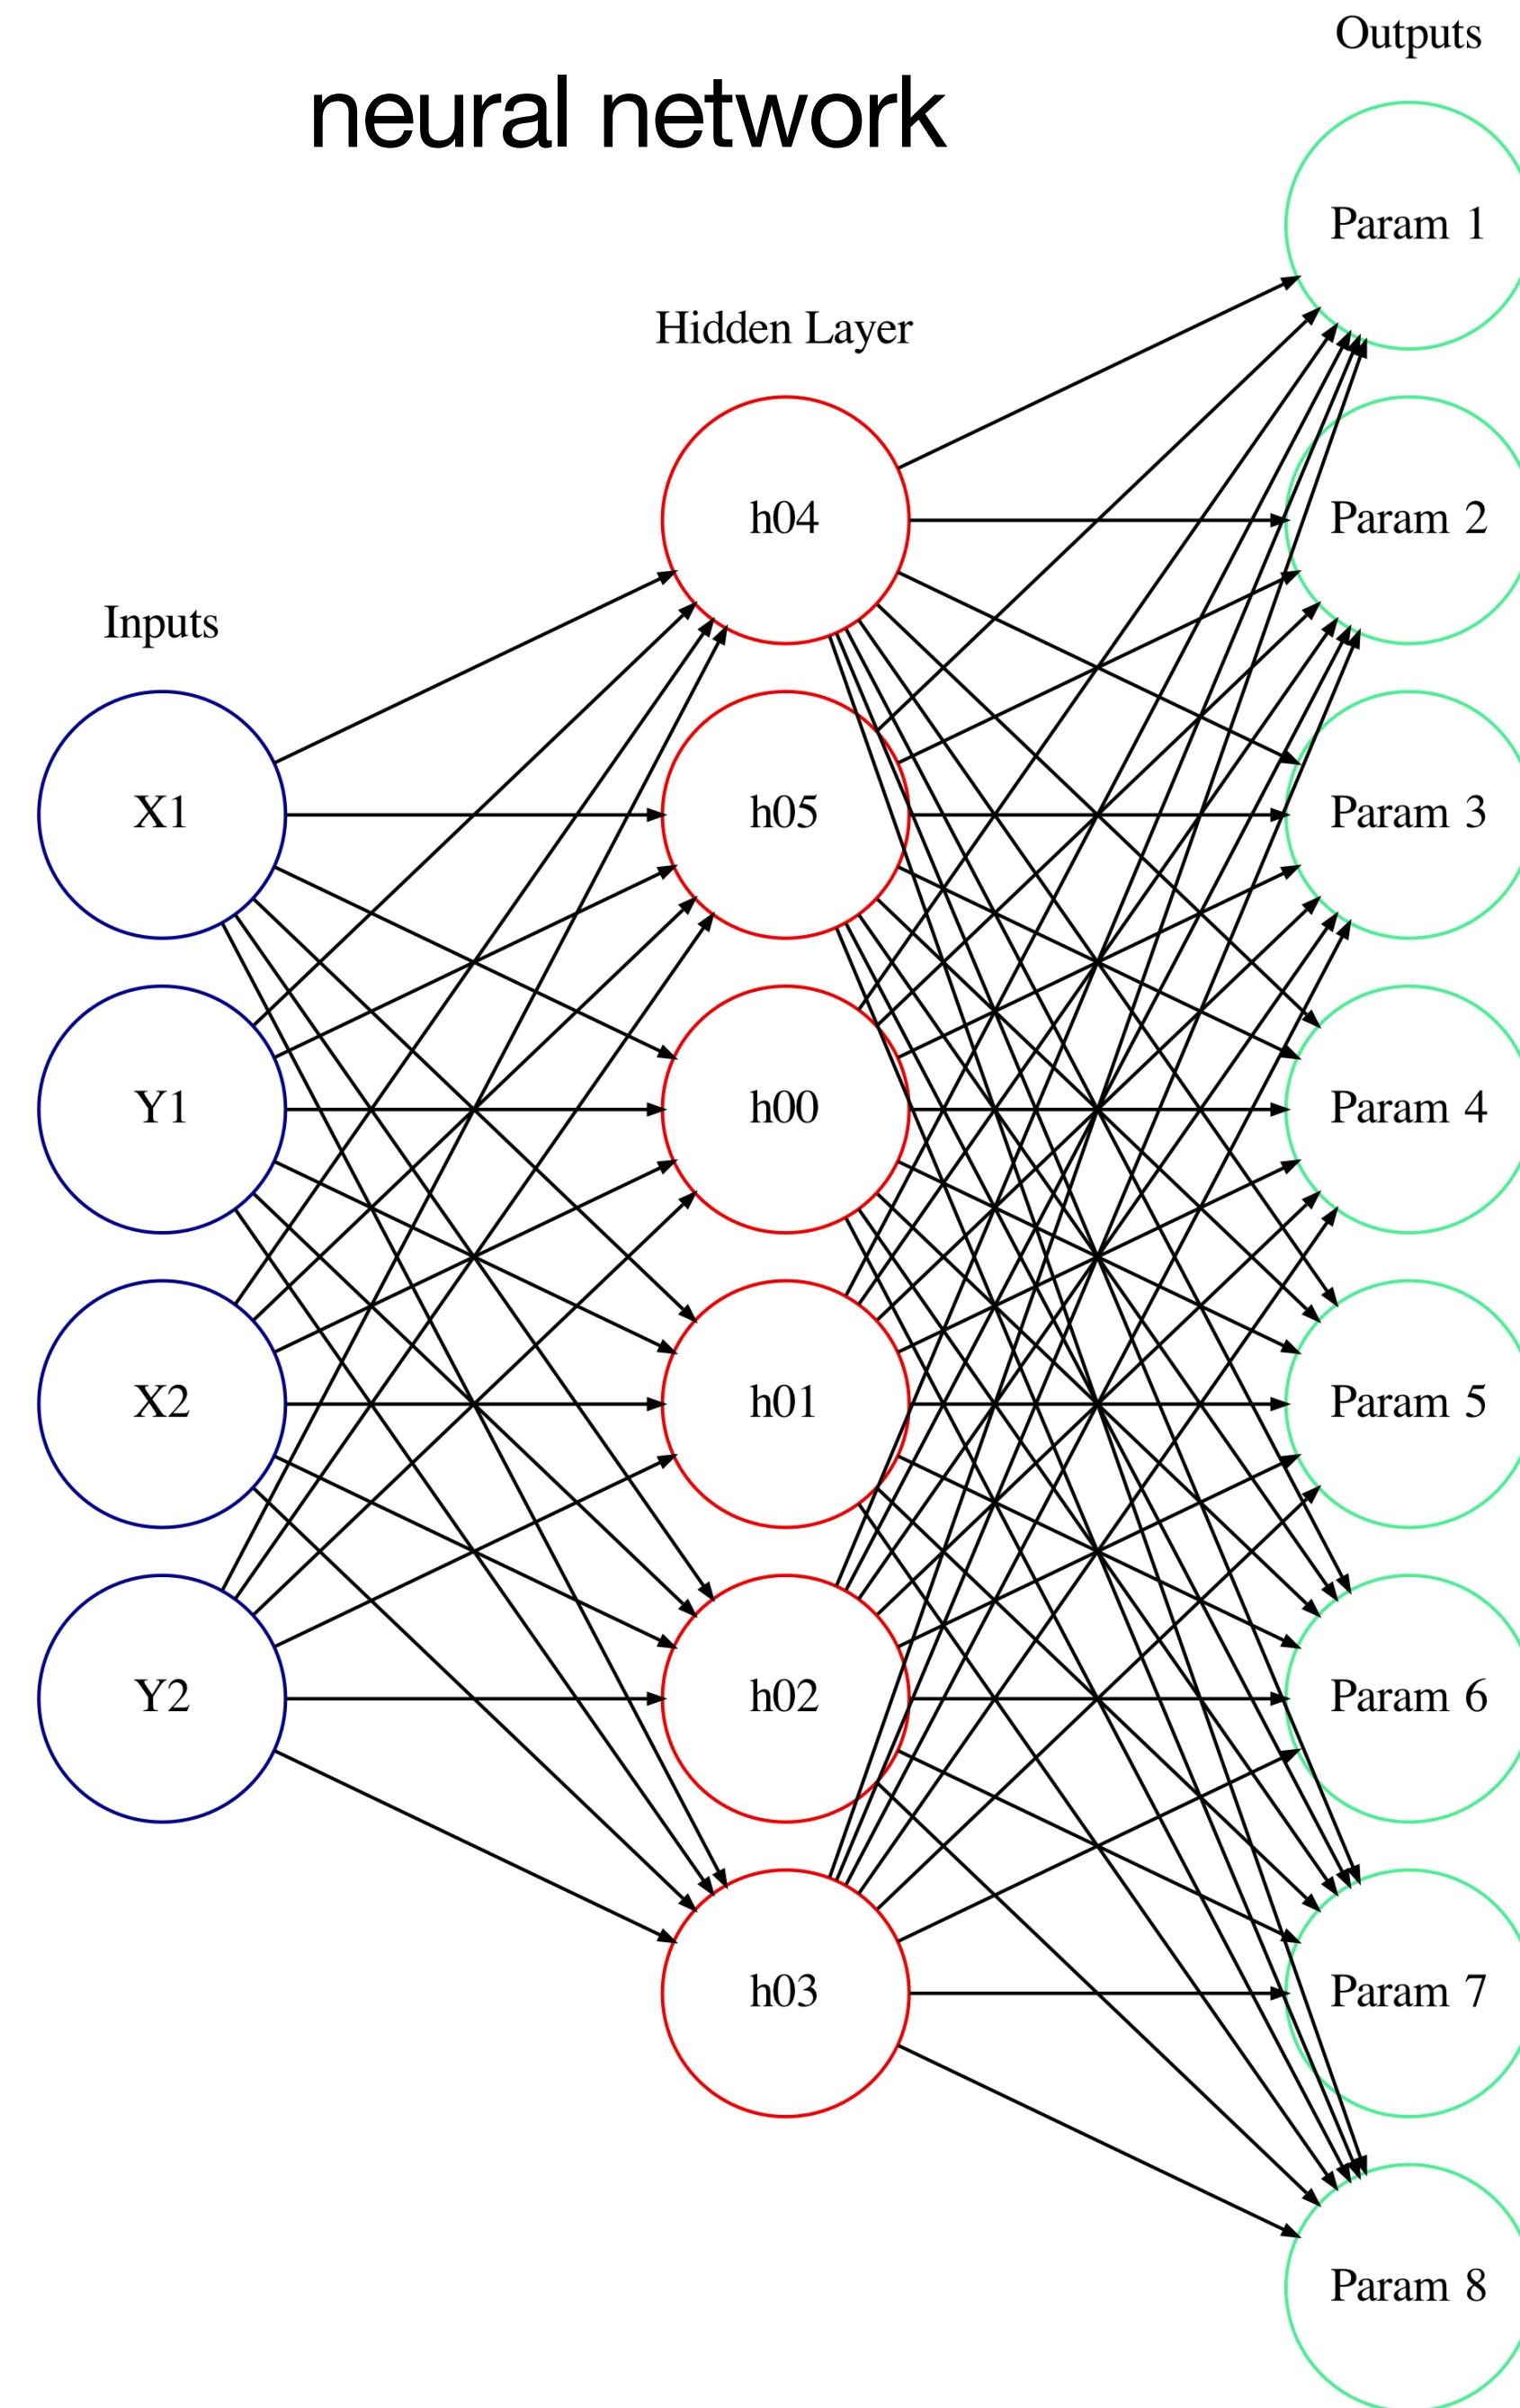
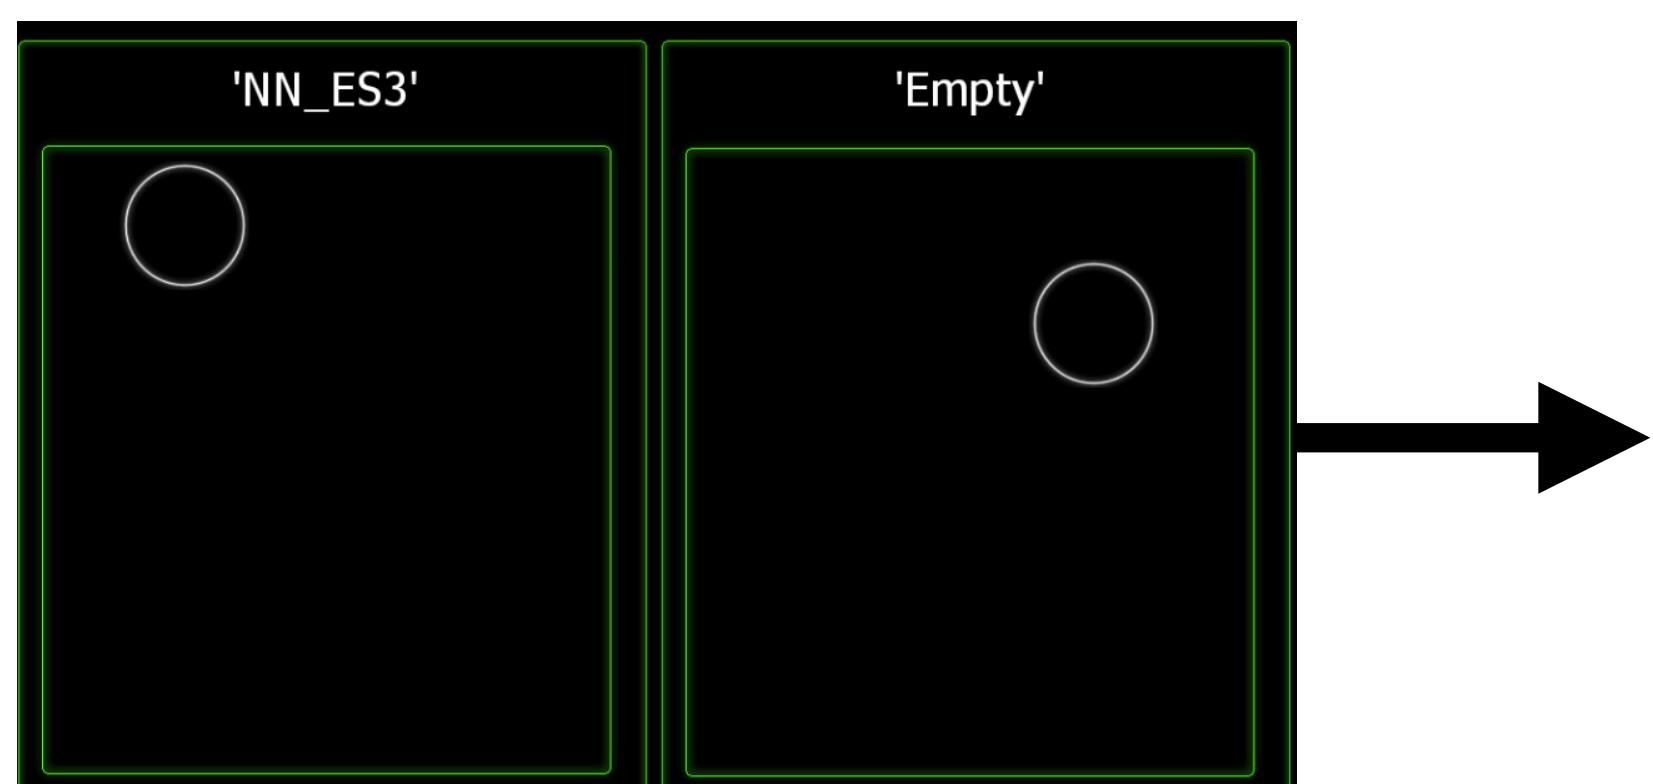
The QuNeo interface displays a signal flow graph with four tracks (0, 1, 2, 3). The graph consists of various processing blocks connected by lines:

- Track 0:** SPEARPlayer → PShiftDel → Benjolin.
- Track 1:** Random Delay → SpecSmear.
- Track 2:** AmpMod → PitchShift.
- Track 3:** (empty)

Nodes in the graph include:  
- SPEARPlayer (blue)  
- PShiftDel (blue)  
- Benjolin (green)  
- Random Delay (cyan)  
- SpecSmear (cyan)  
- AmpMod (green)  
- PitchShift (cyan)



# neural network





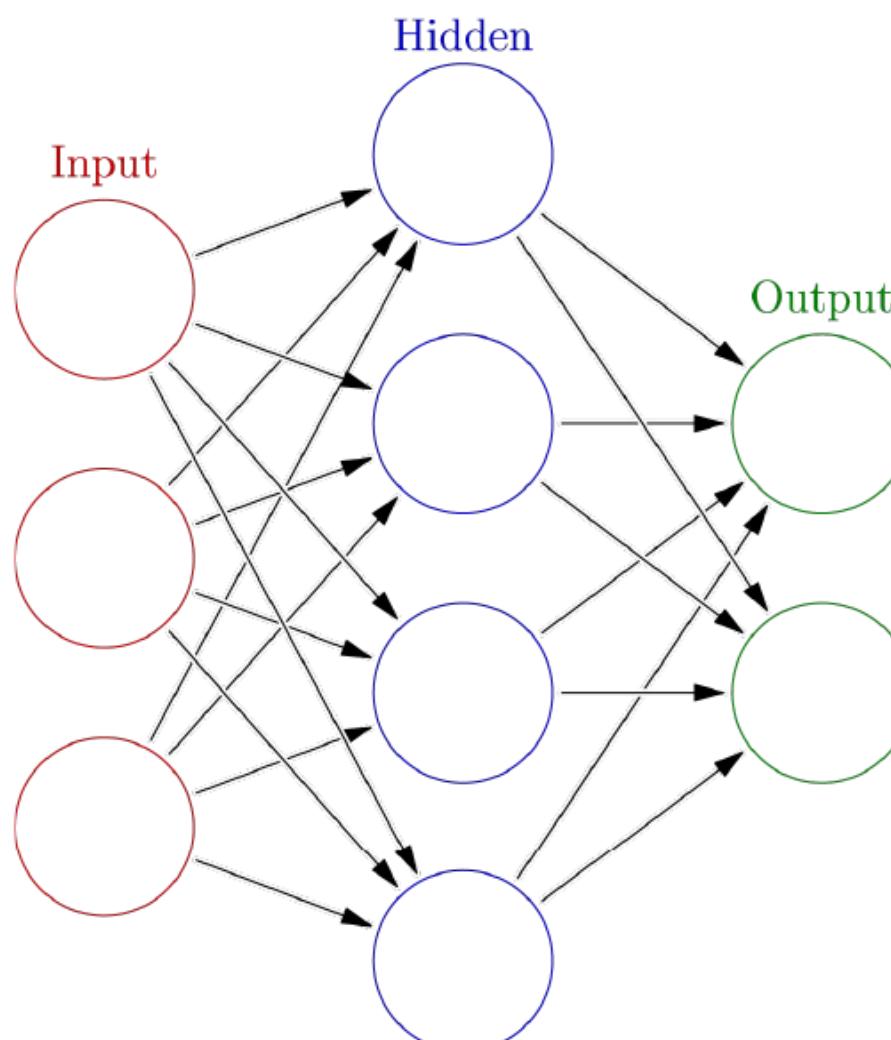
# *Neural Network Training a Regressor*

identifier

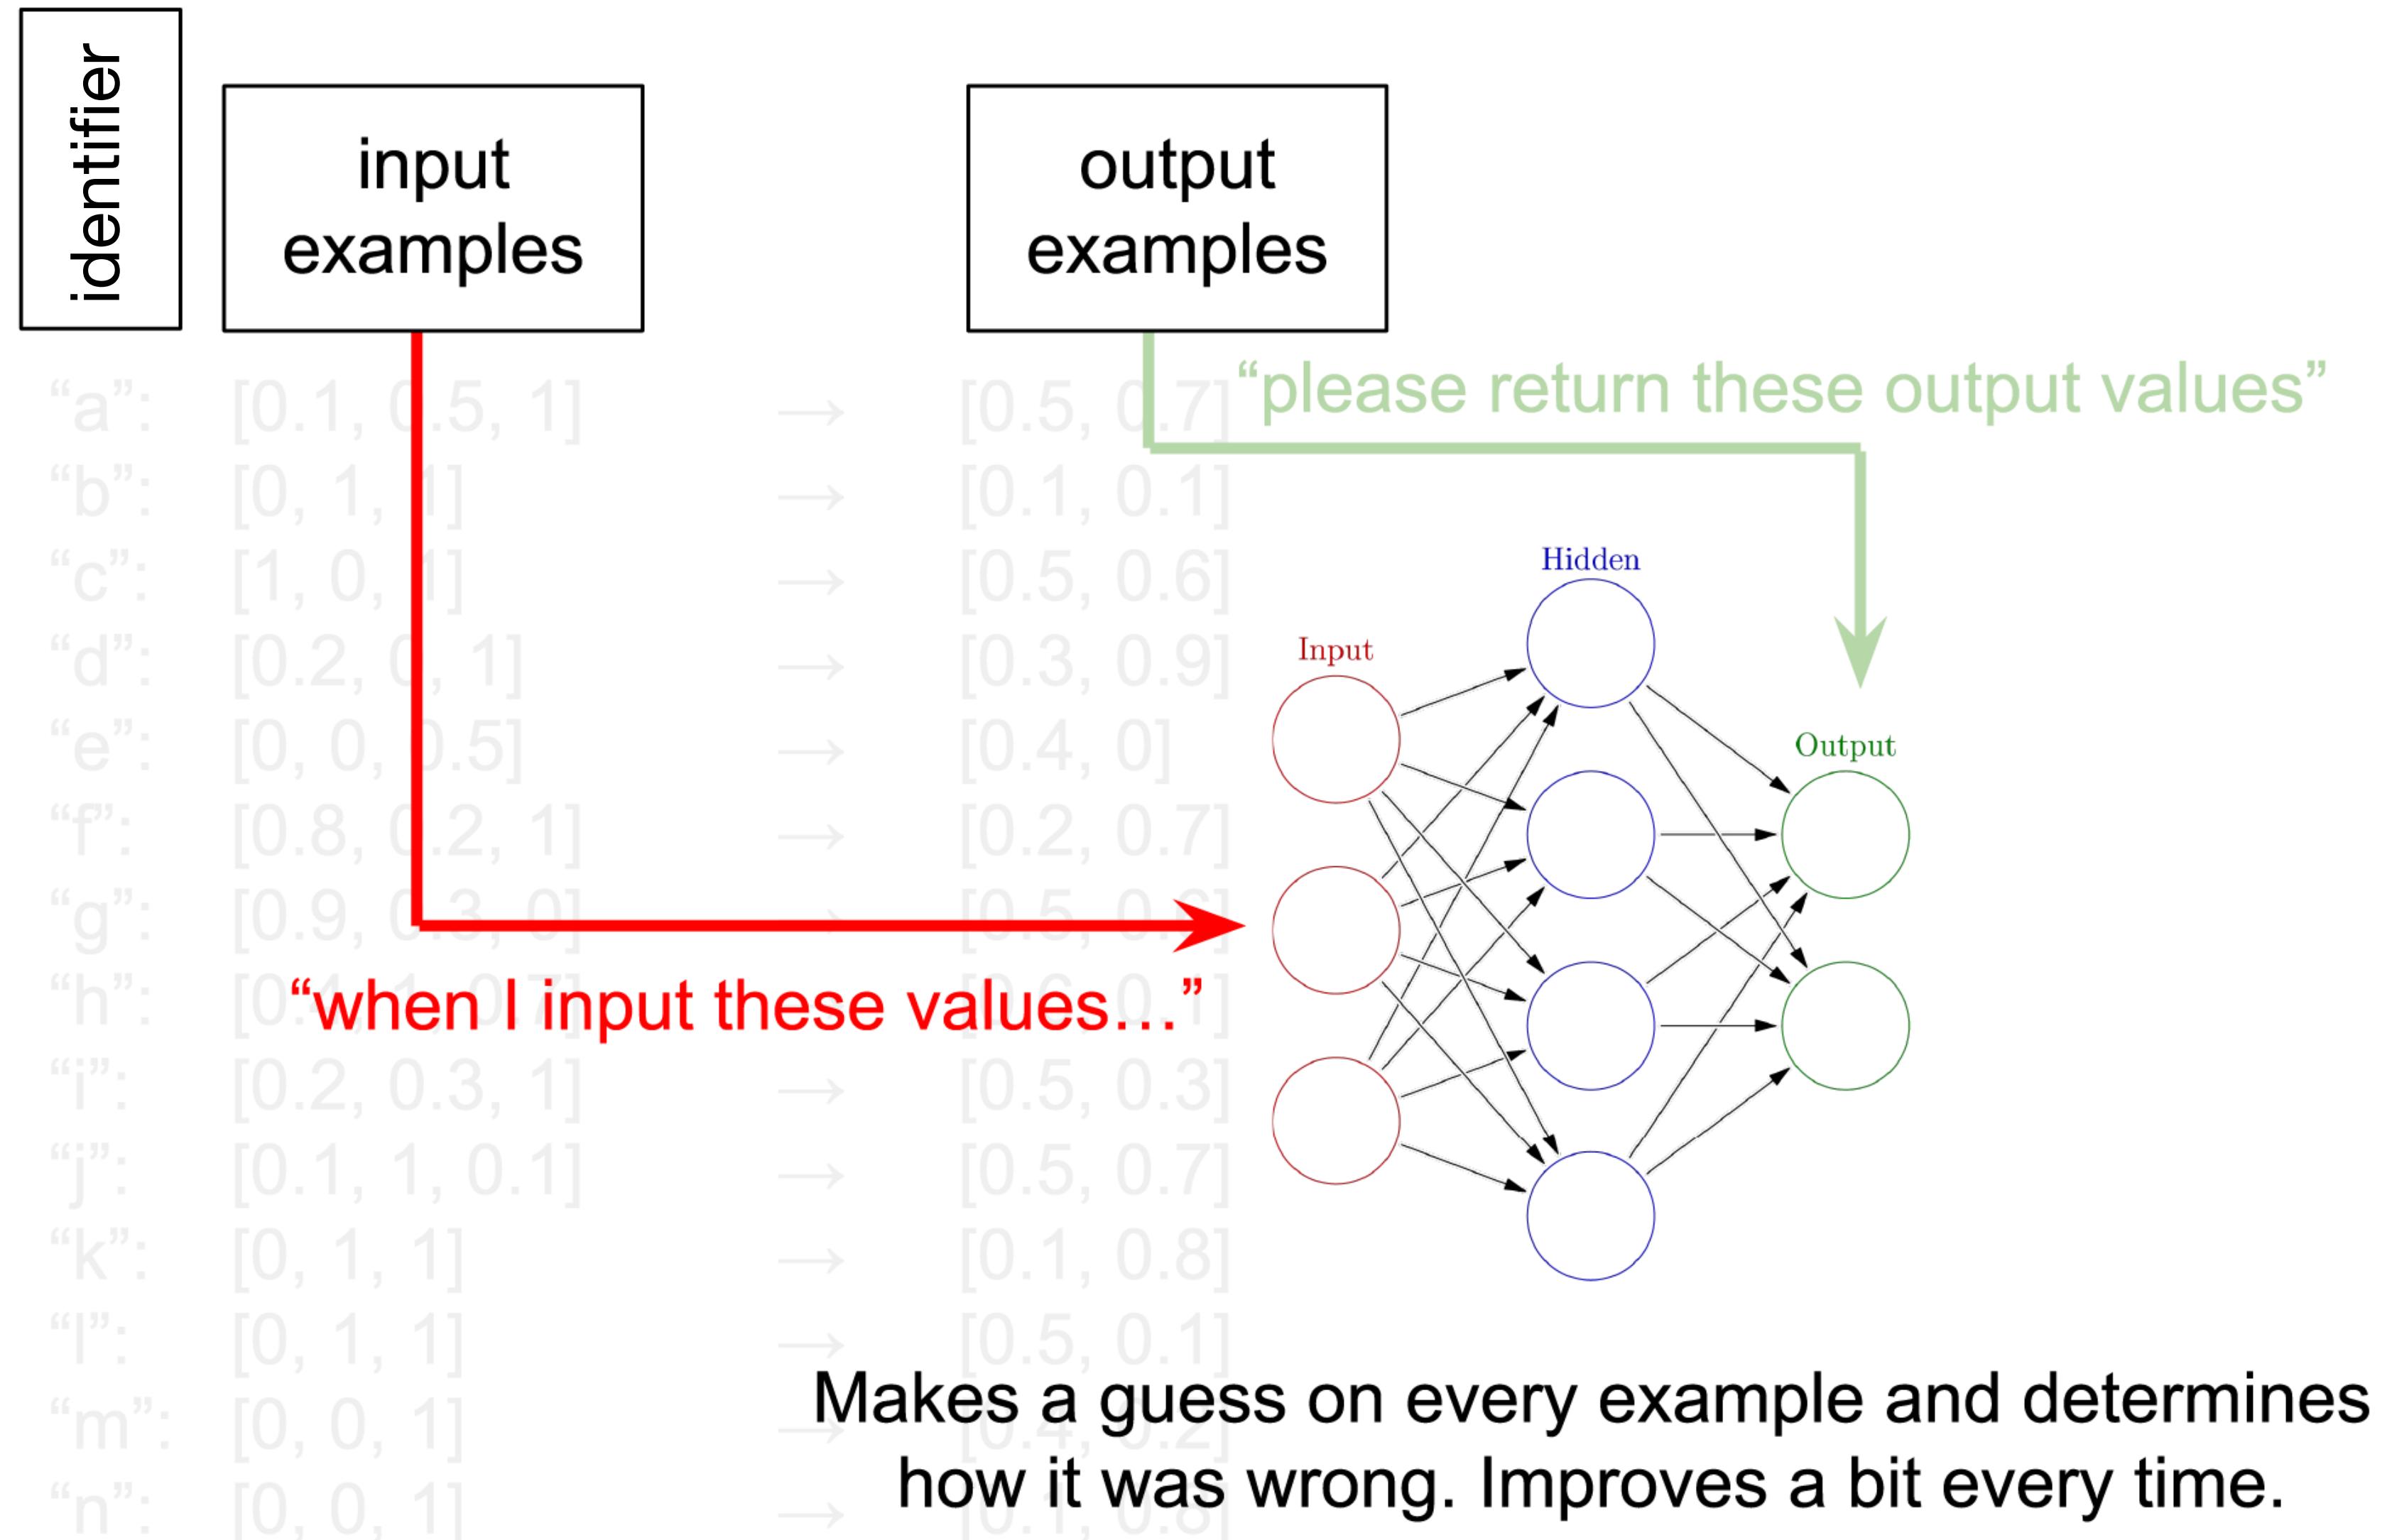
input  
examples

“a”:	[0.1, 0.5, 1]	→	[0.5, 0.7]
“b”:	[0, 1, 1]	→	[0.1, 0.1]
“c”:	[1, 0, 1]	→	[0.5, 0.6]
“d”:	[0.2, 0, 1]	→	[0.3, 0.9]
“e”:	[0, 0, 0.5]	→	[0.4, 0]
“f”:	[0.8, 0.2, 1]	→	[0.2, 0.7]
“g”:	[0.9, 0.3, 0]	→	[0.5, 0.6]
“h”:	[0.4, 1, 0.7]	→	[0.6, 0.1]
“i”:	[0.2, 0.3, 1]	→	[0.5, 0.3]
“j”:	[0.1, 1, 0.1]	→	[0.5, 0.7]
“k”:	[0, 1, 1]	→	[0.1, 0.8]
“l”:	[0, 1, 1]	→	[0.5, 0.1]
“m”:	[0, 0, 1]	→	[0.4, 0.2]
“n”:	[0, 0, 1]	→	[0.1, 0.8]

output  
examples



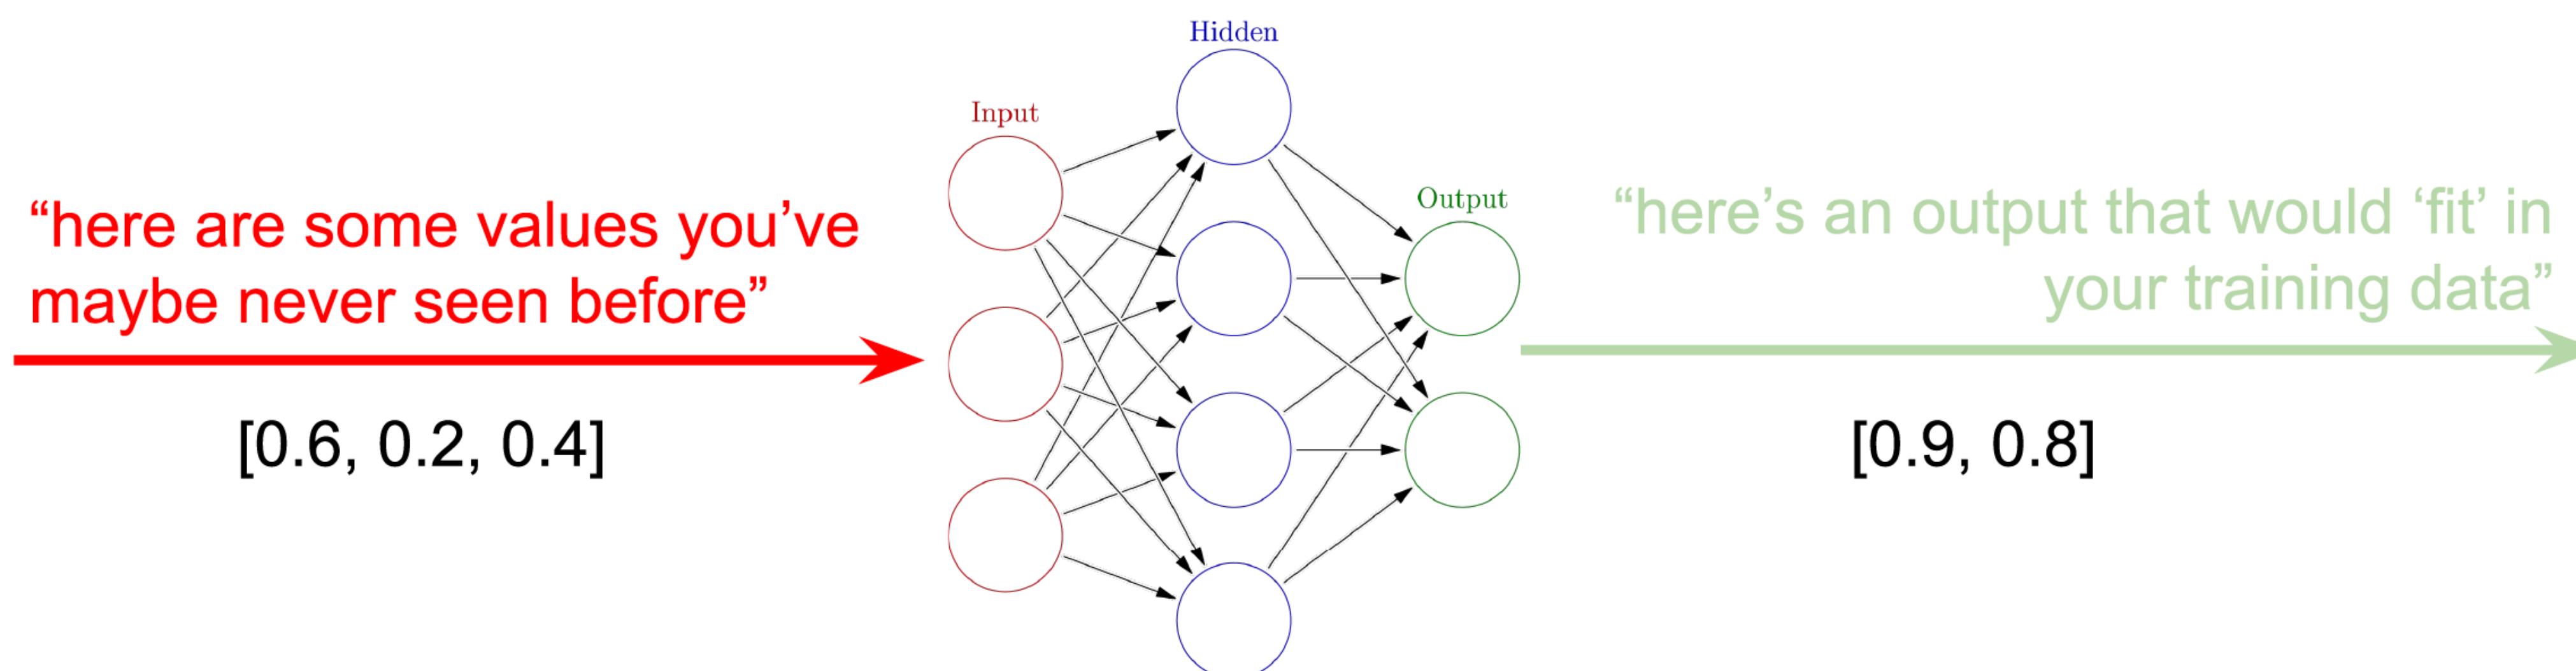
# Neural Network *Training* a Regressor



# Neural Network *Predicting with Regression*

input  
examples

output  
examples



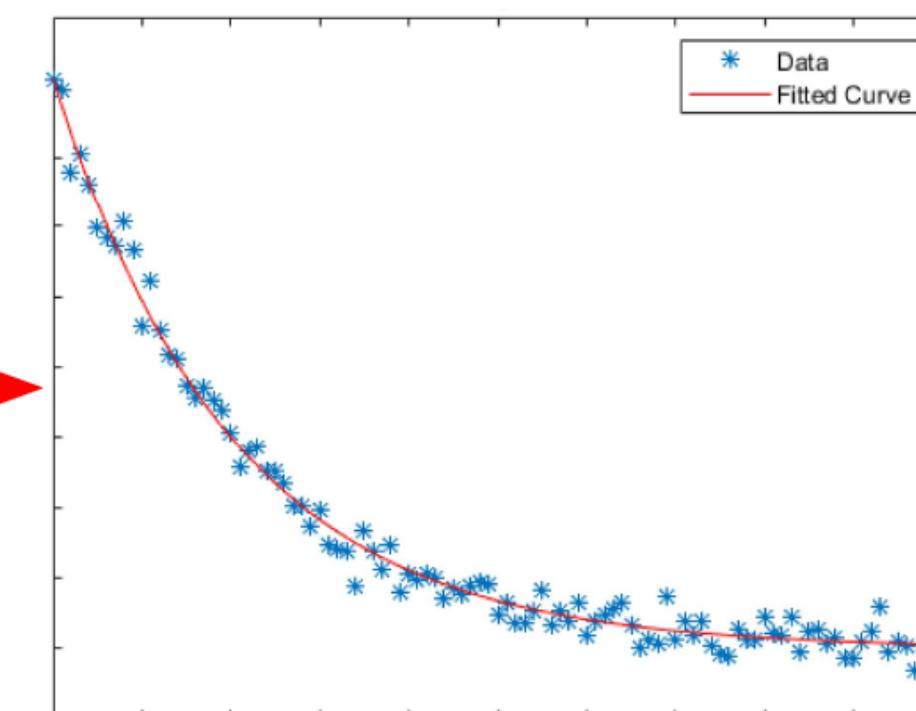
# Neural Network *Predicting with Regression*

input  
examples

output  
examples

“here are some values you’ve  
maybe never seen before”

[0.6, 0.2, 0.4]



“here’s an output that would ‘fit’ in  
your training data”

[0.9, 0.8]

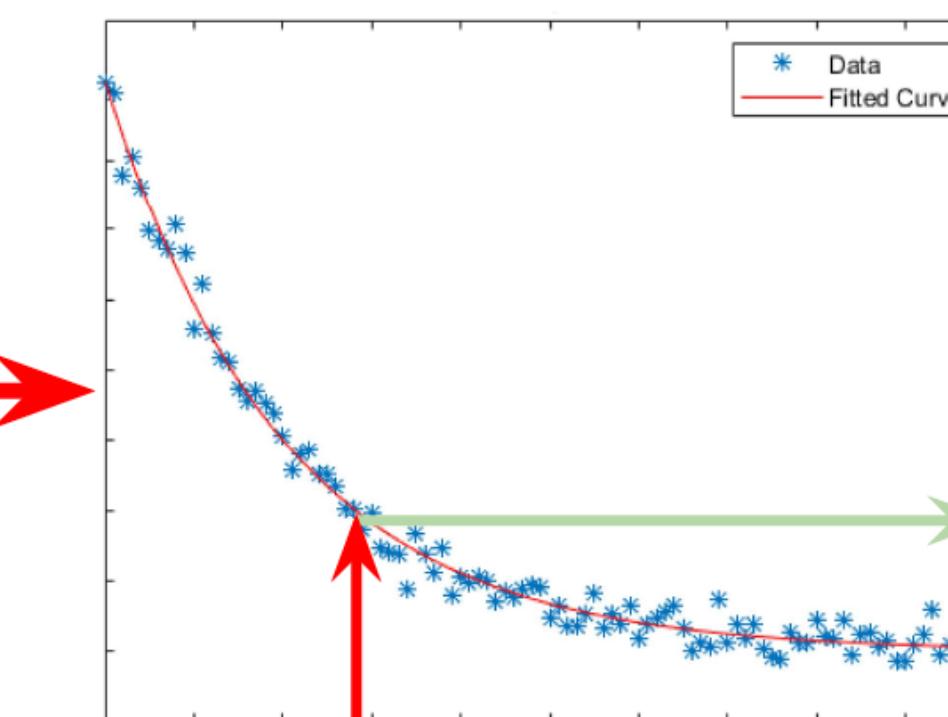
# Neural Network *Predicting with Regression*

input  
examples

output  
examples

“here are some values you’ve  
maybe never seen before”

[0.6, 0.2, 0.4]

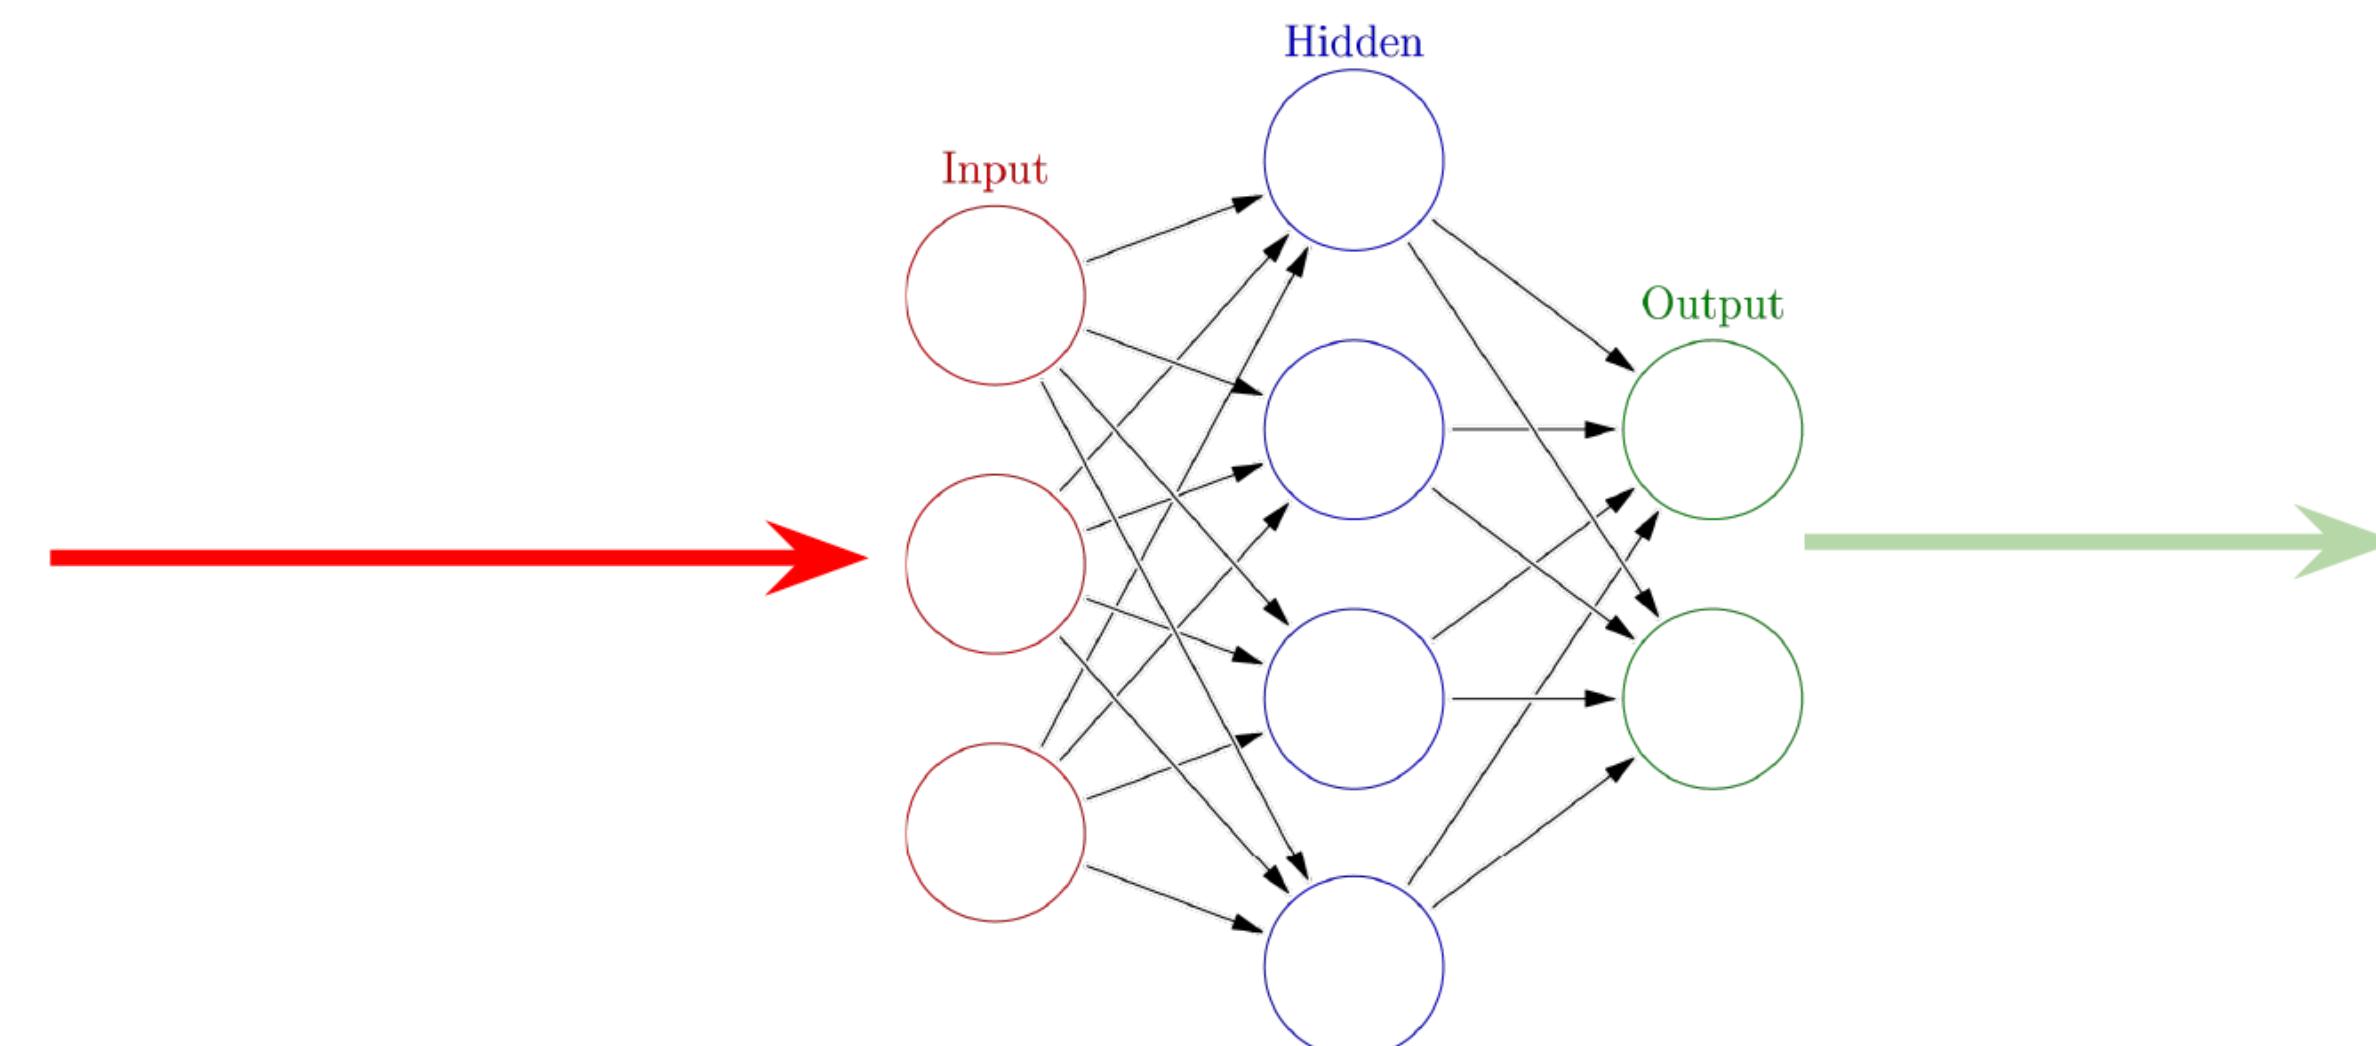


“here’s an output that would ‘fit’ in  
your training data”

[0.9, 0.8]

# Neural Network Predicting with Regression

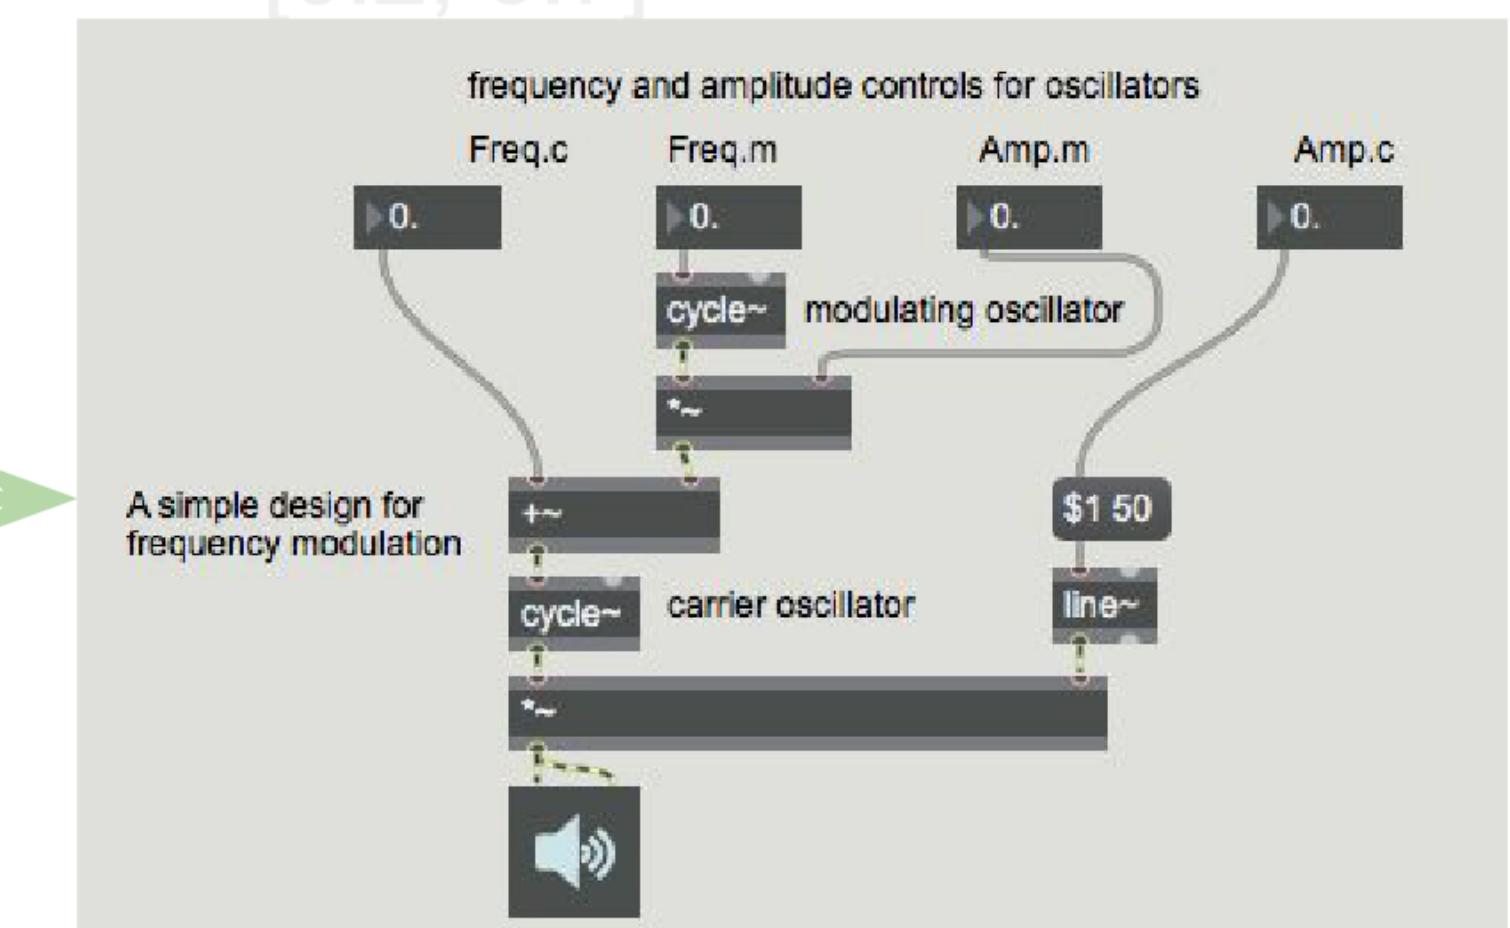
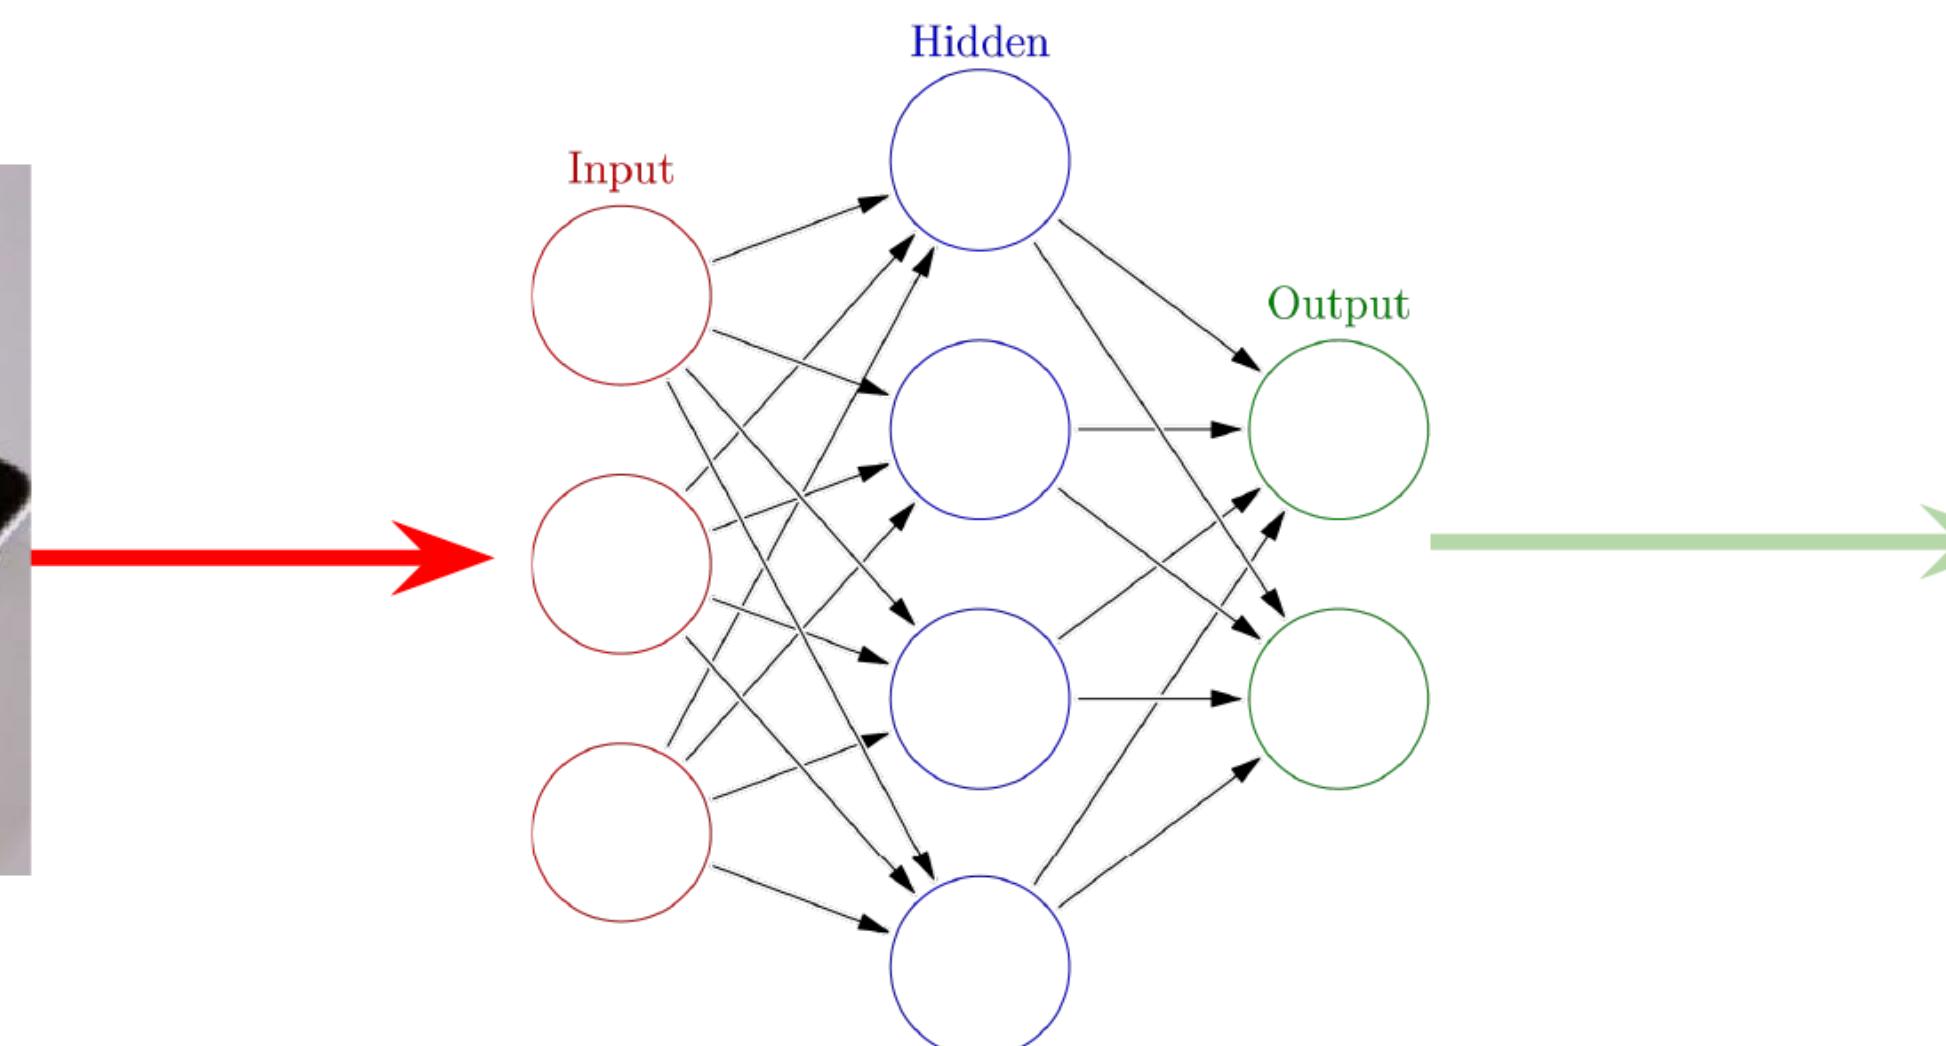
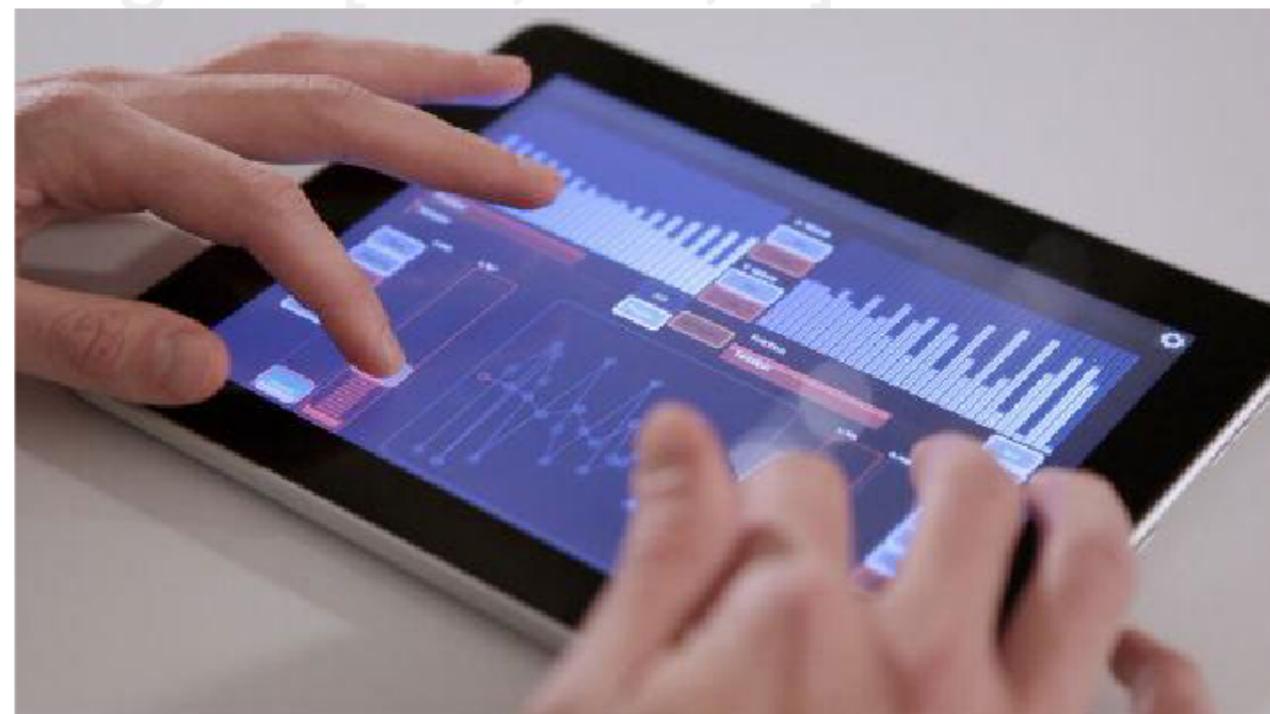
“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“g”: [0.9, 0.3, 0]  
“h”: [0.4, 1, 0.7]  
“i”: [0.2, 0.3, 1]  
“j”: [0.1, 1, 0.1]  
“k”: [0, 1, 1]  
“l”: [0, 1, 1]  
“m”: [0, 0, 1]  
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]



[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]  
[0.5, 0.6]  
[0.6, 0.1]  
[0.5, 0.3]  
[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0.3]  
[0.2, 0.7]

# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“g”: [0.9, 0.3, 0]



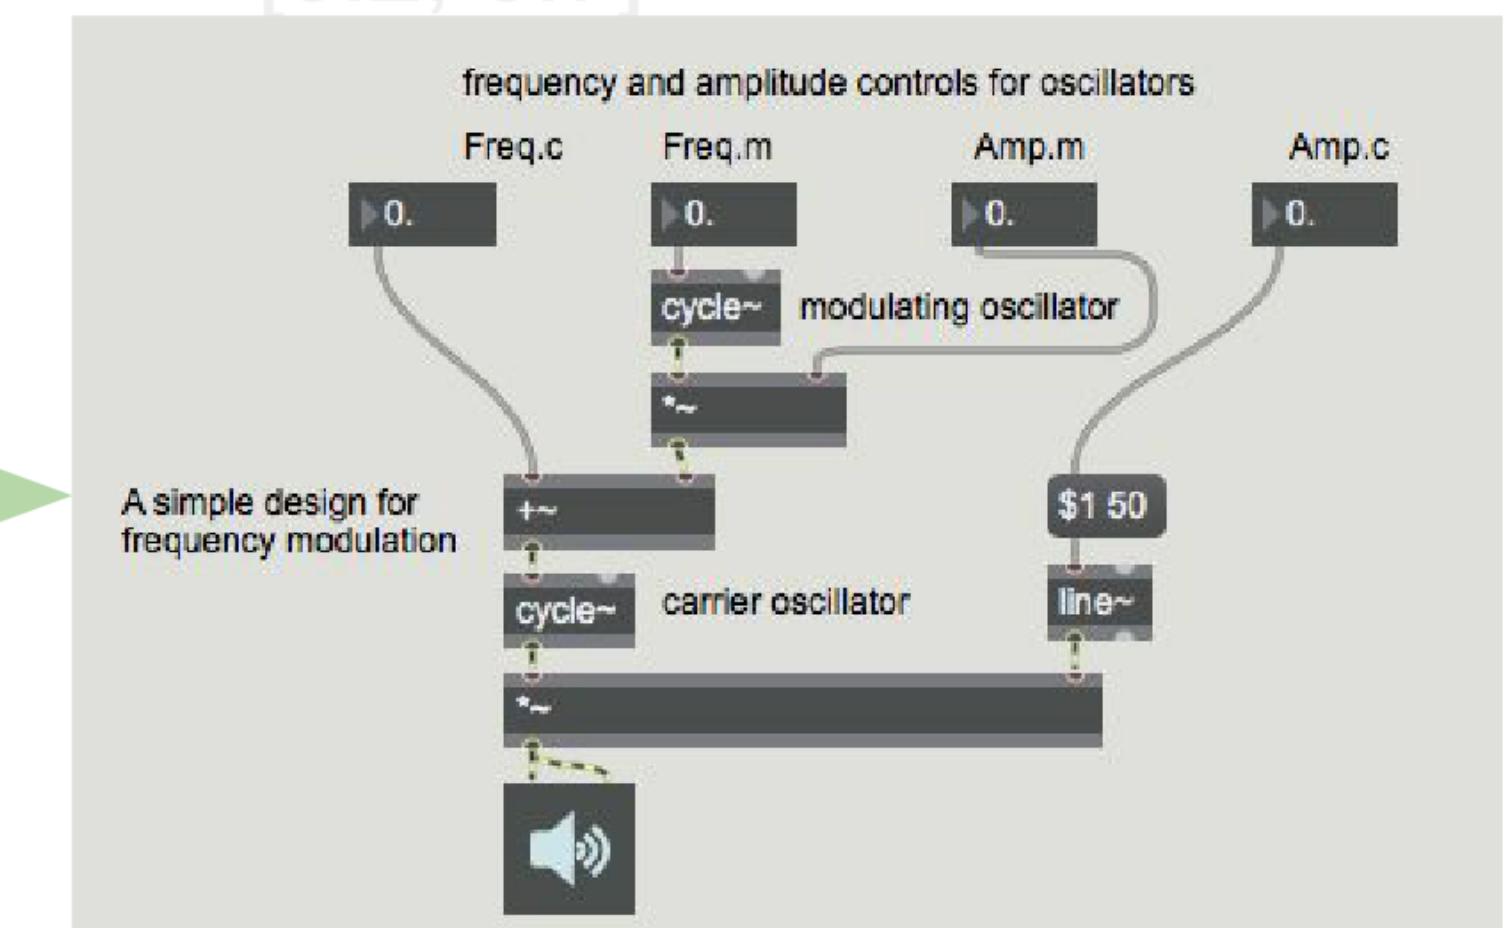
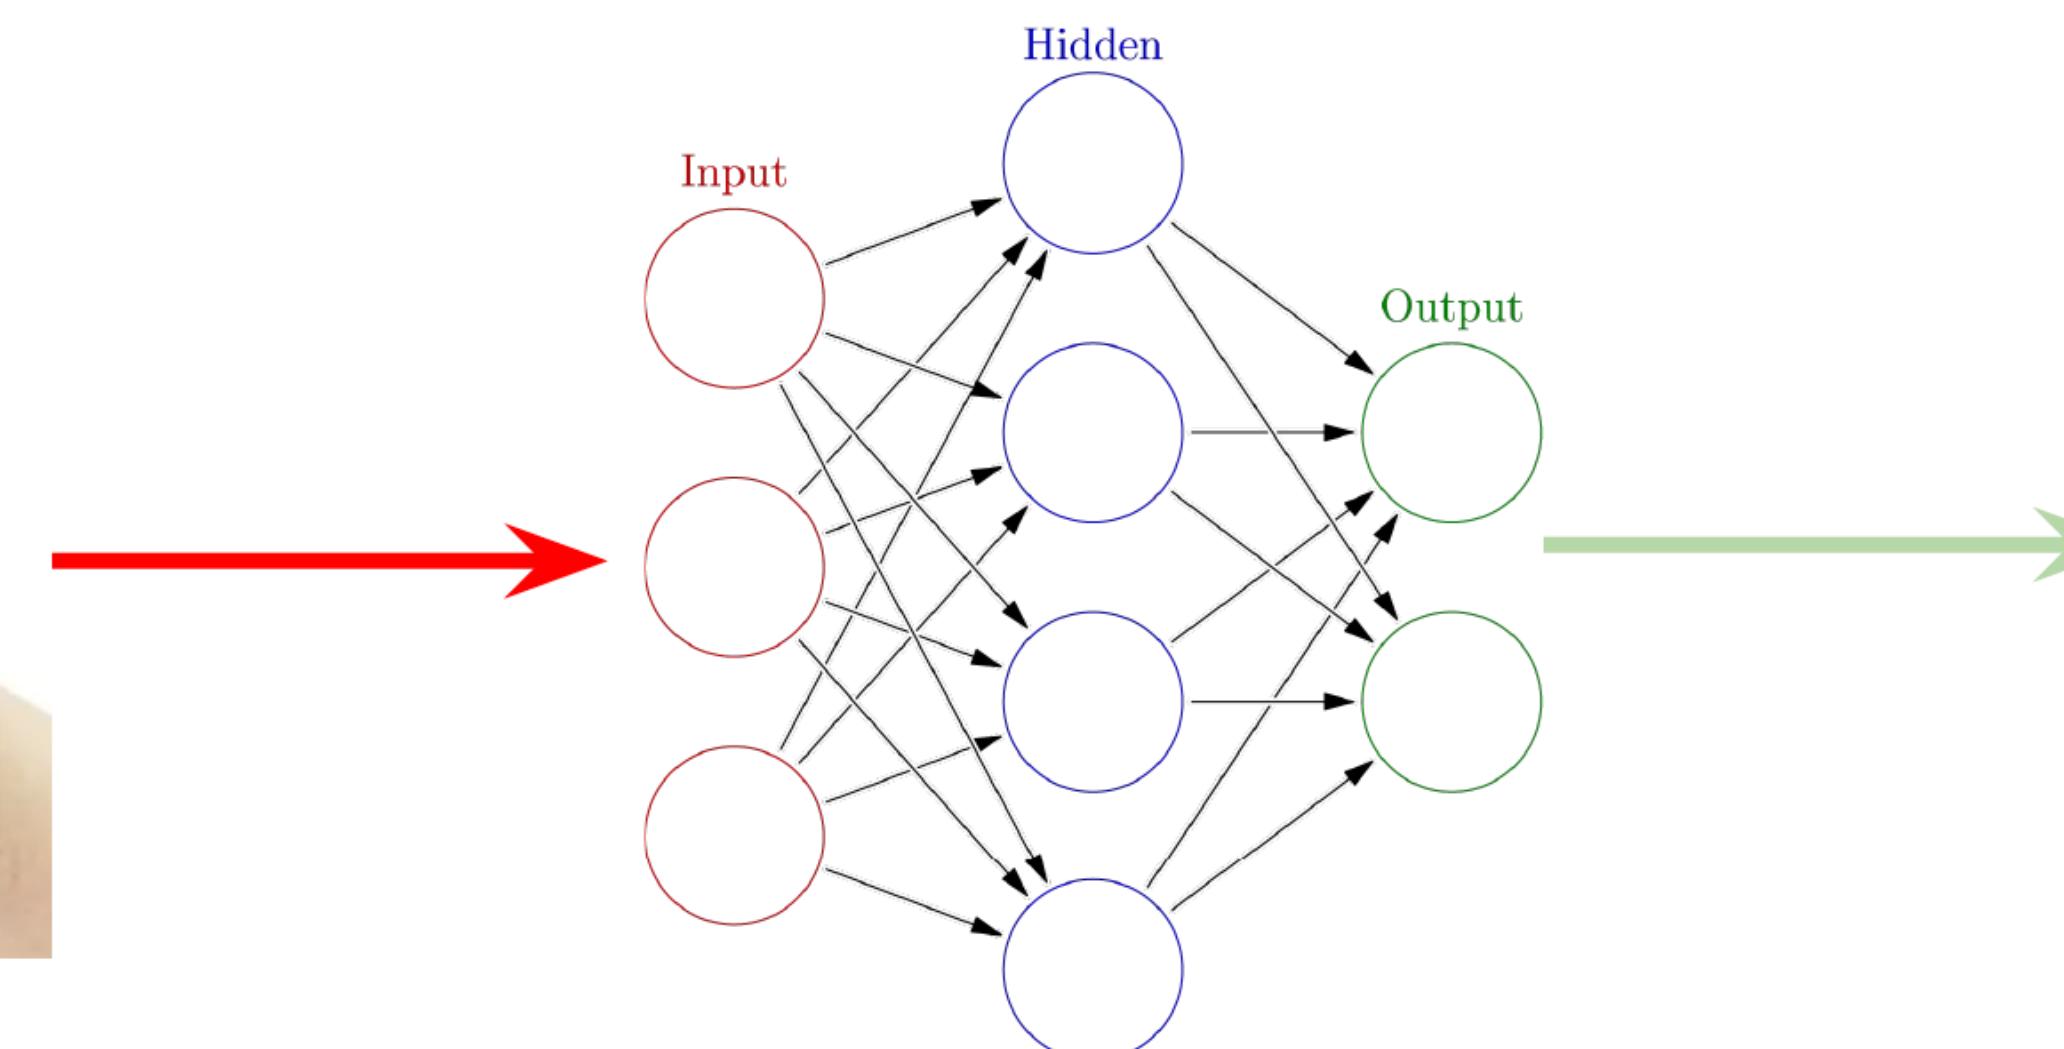
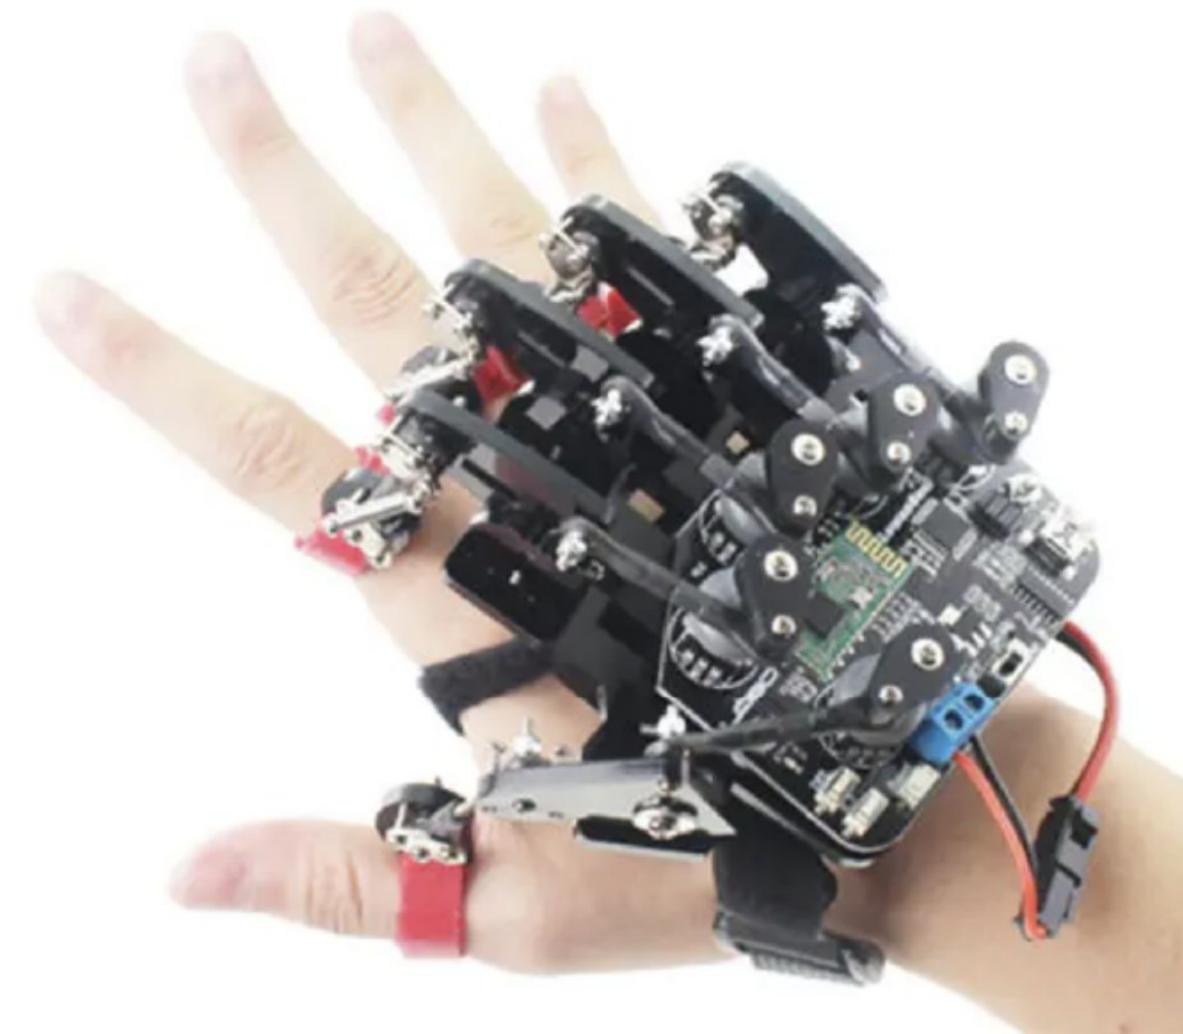
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]

[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]

[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]

# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]



“f”: [0.1, 1, 0.1]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]

[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]

[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]

# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]

“b”: [0, 1, 1]

“c”: [1, 0, 1]

“d”: [0.2, 0, 1]

“e”: [0, 0, 0.5]

“f”: [0.8, 0.2, 1]



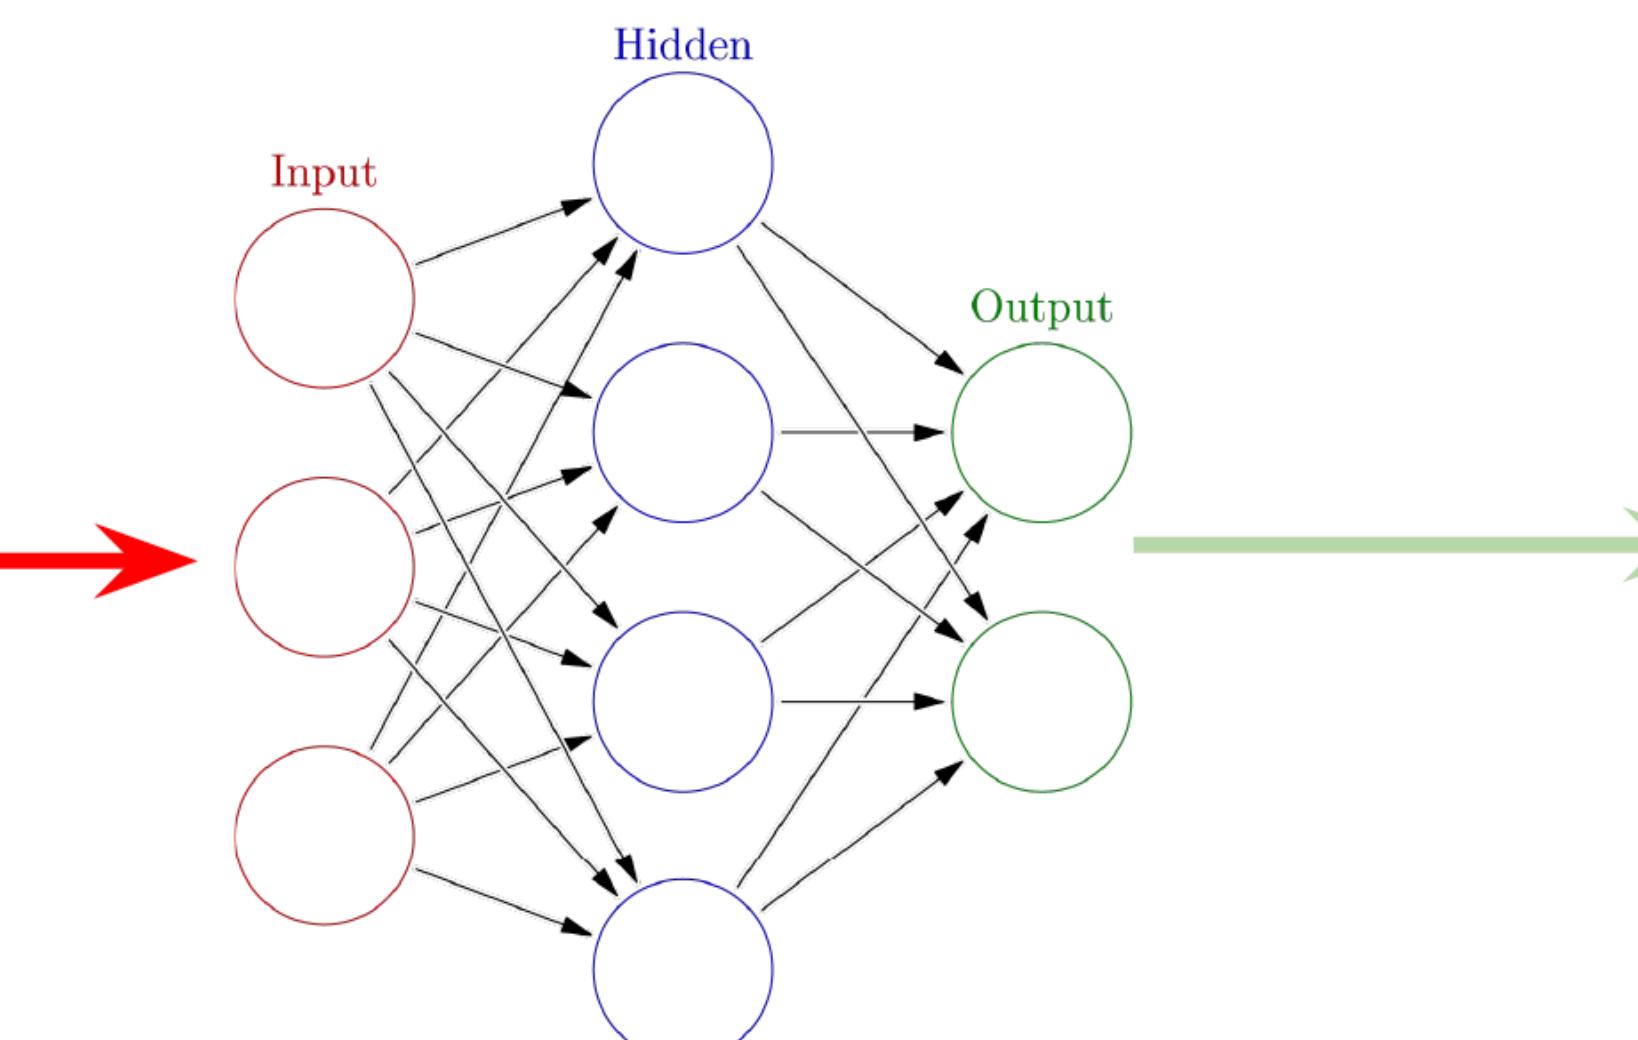
e.g.:

**pitch**

**loudness**

**spectral centroid**

**spectral flatness**



[0.5, 0.7]

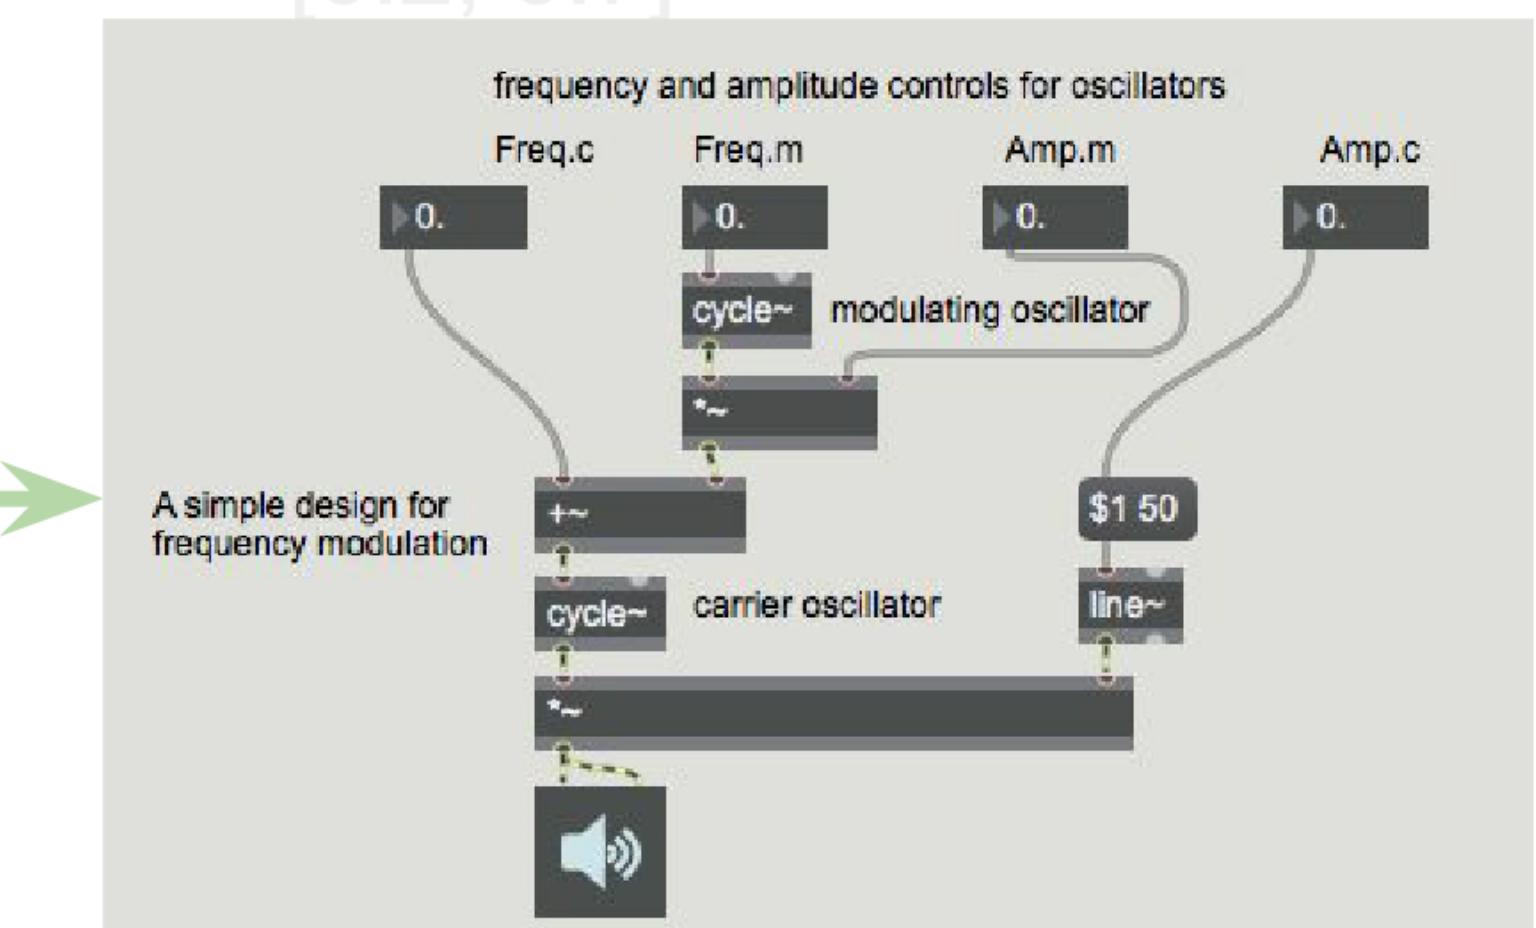
[0.1, 0.1]

[0.5, 0.6]

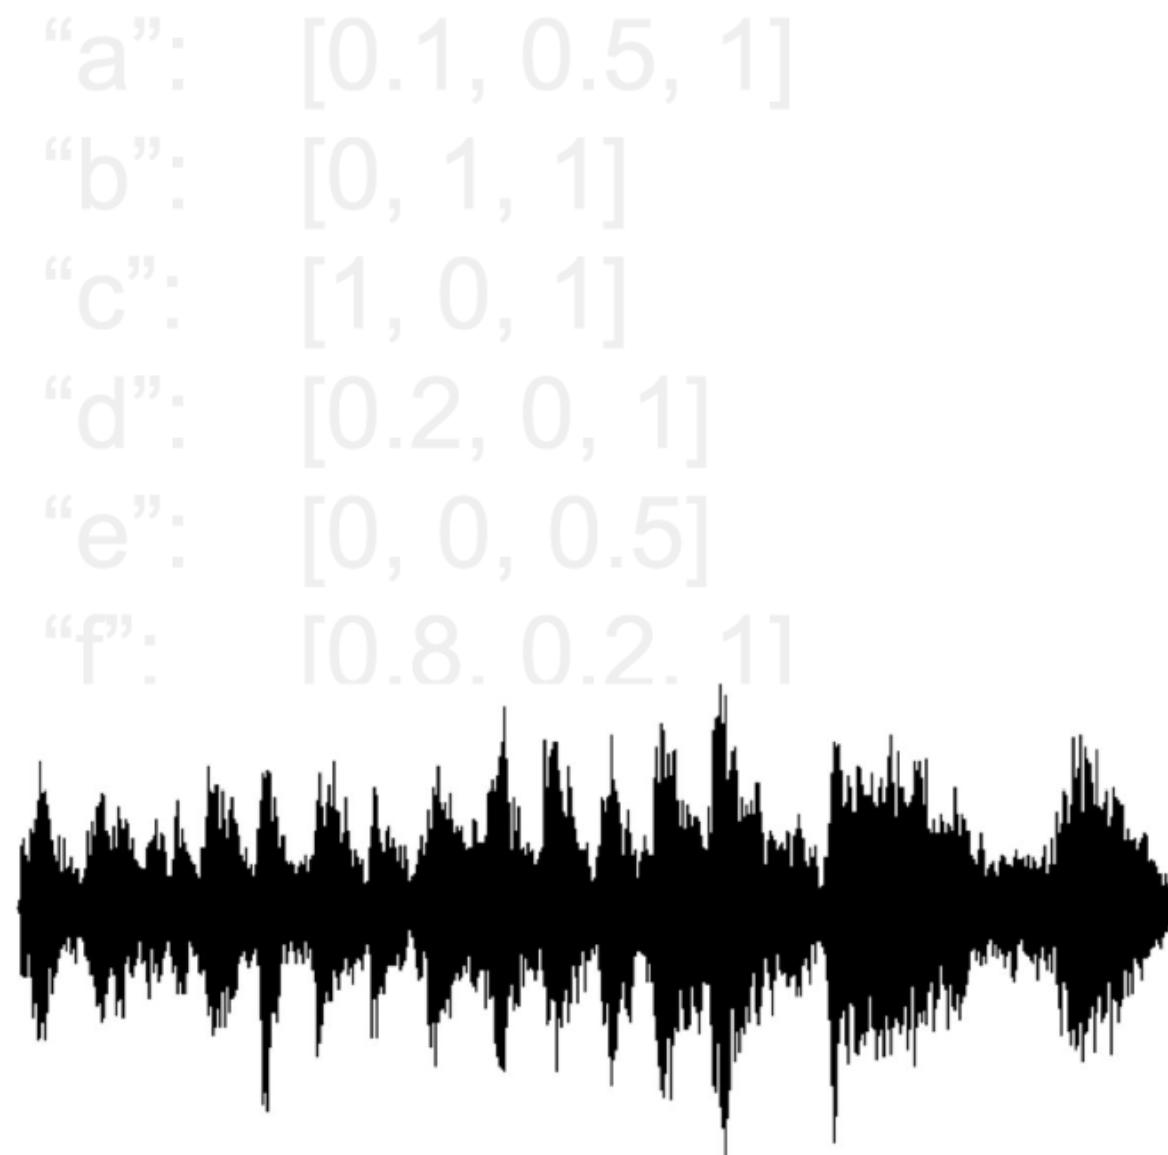
[0.3, 0.9]

[0.4, 0]

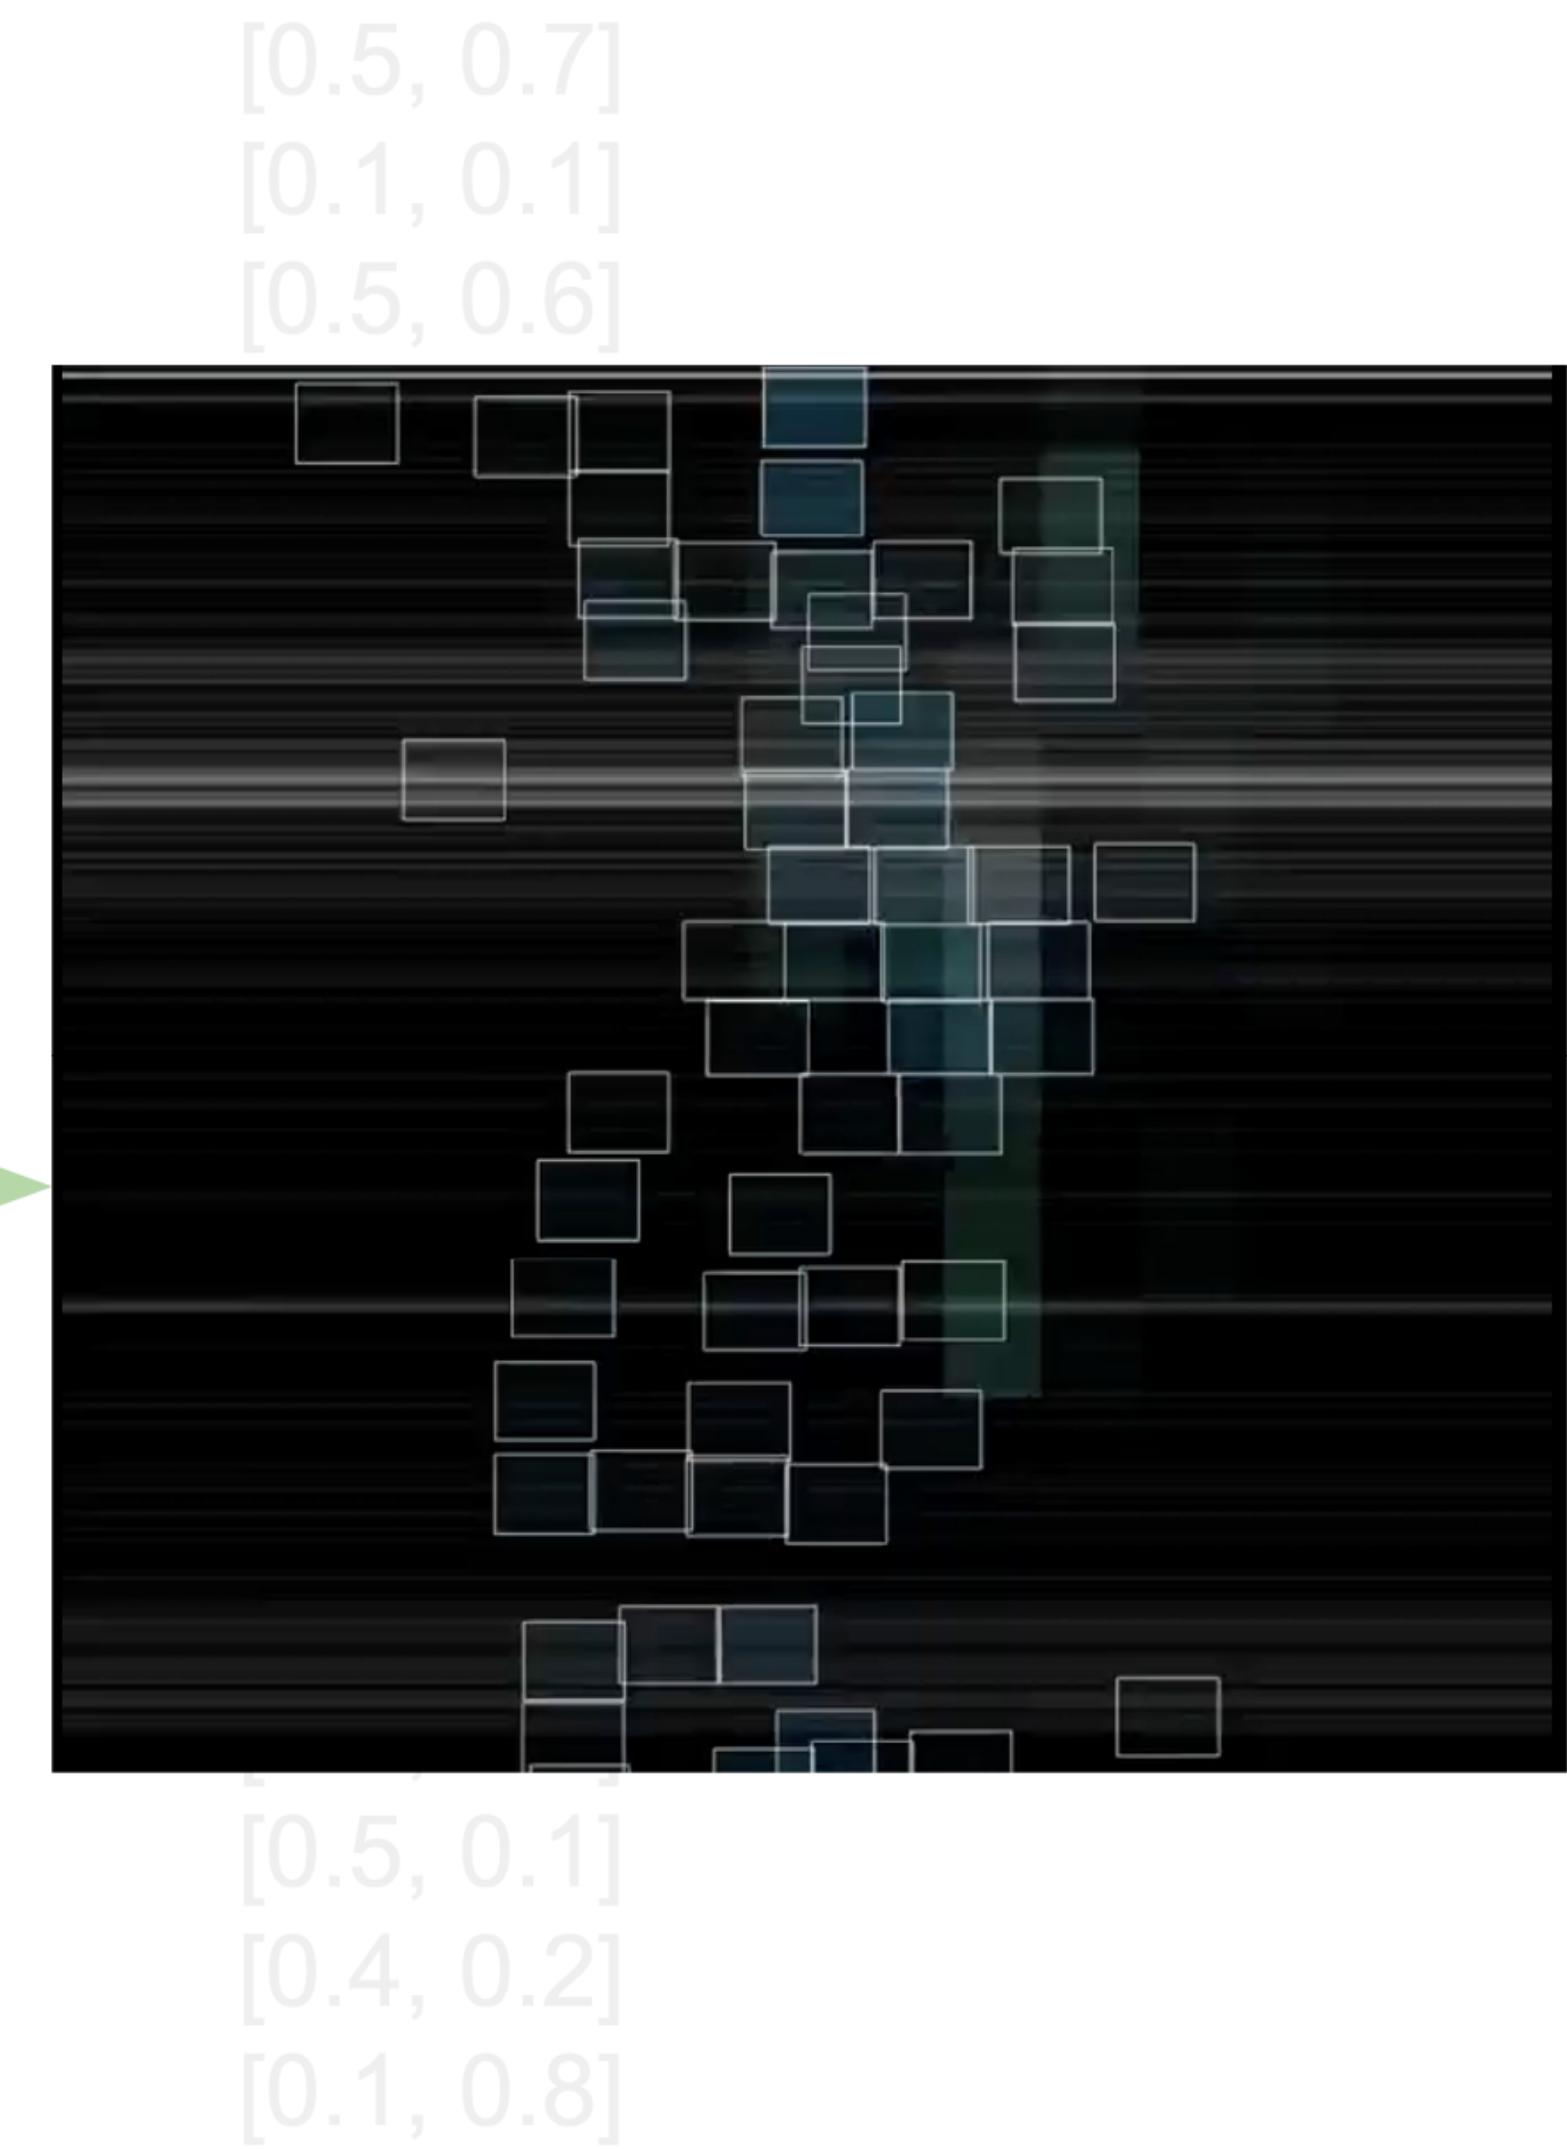
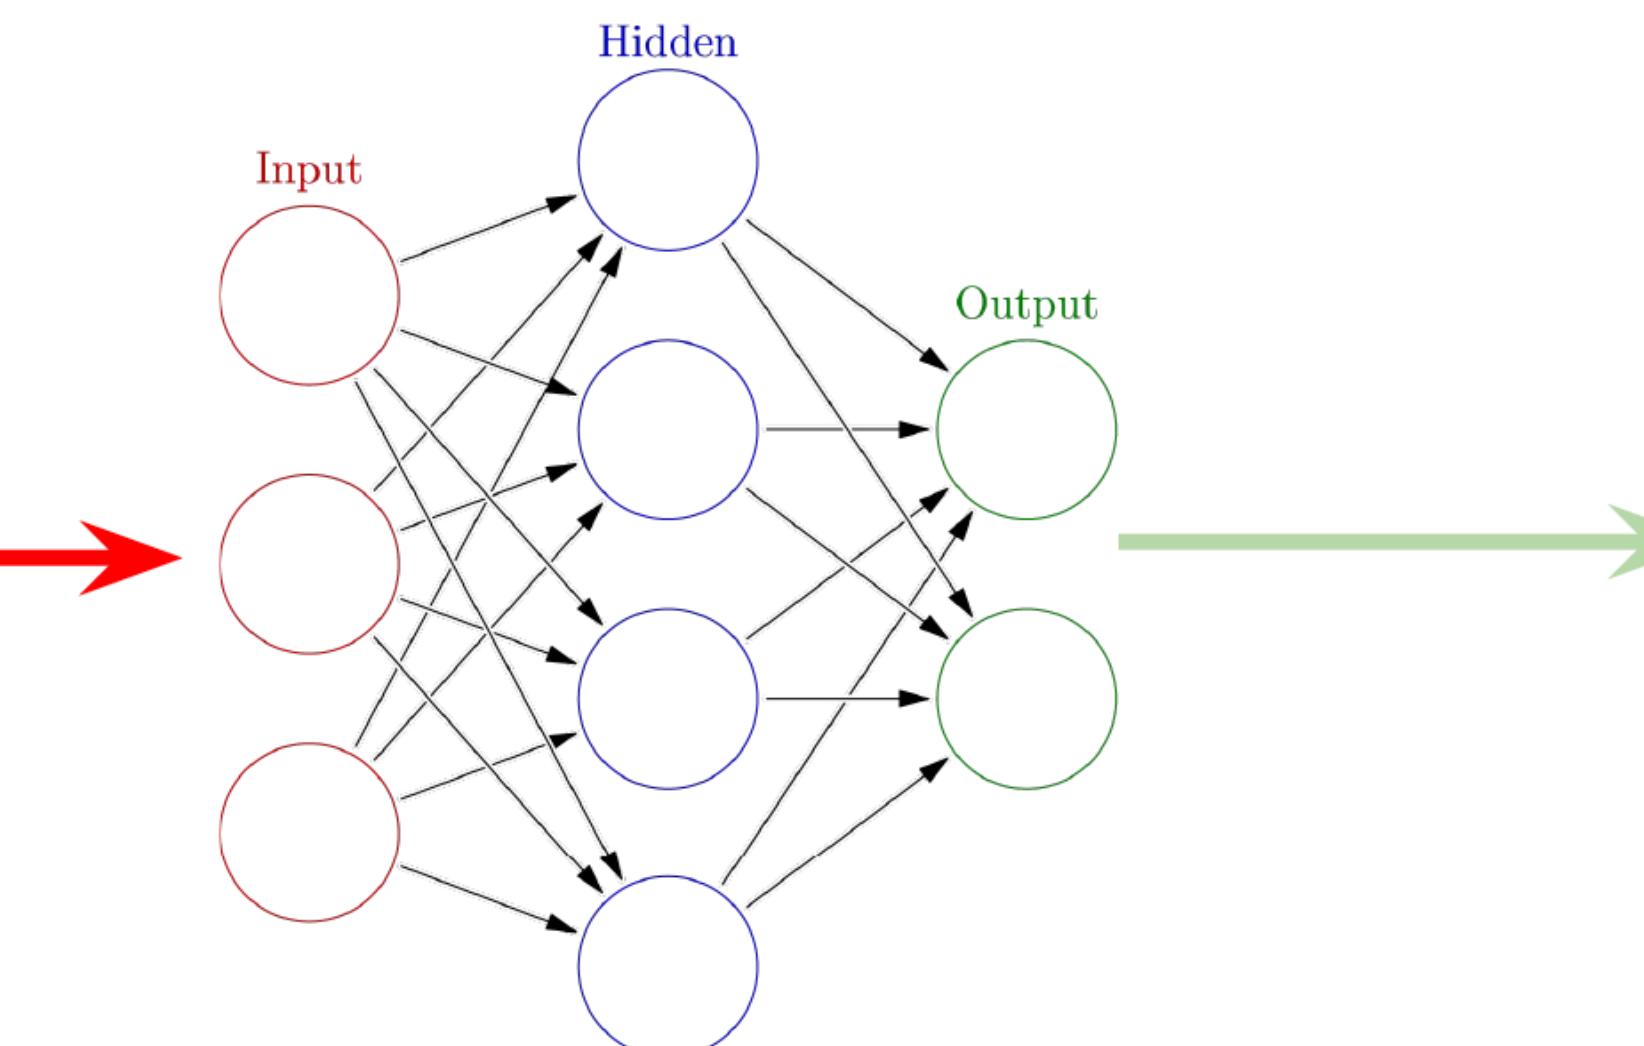
[0.2, 0.7]



# Neural Network Predicting with Regression

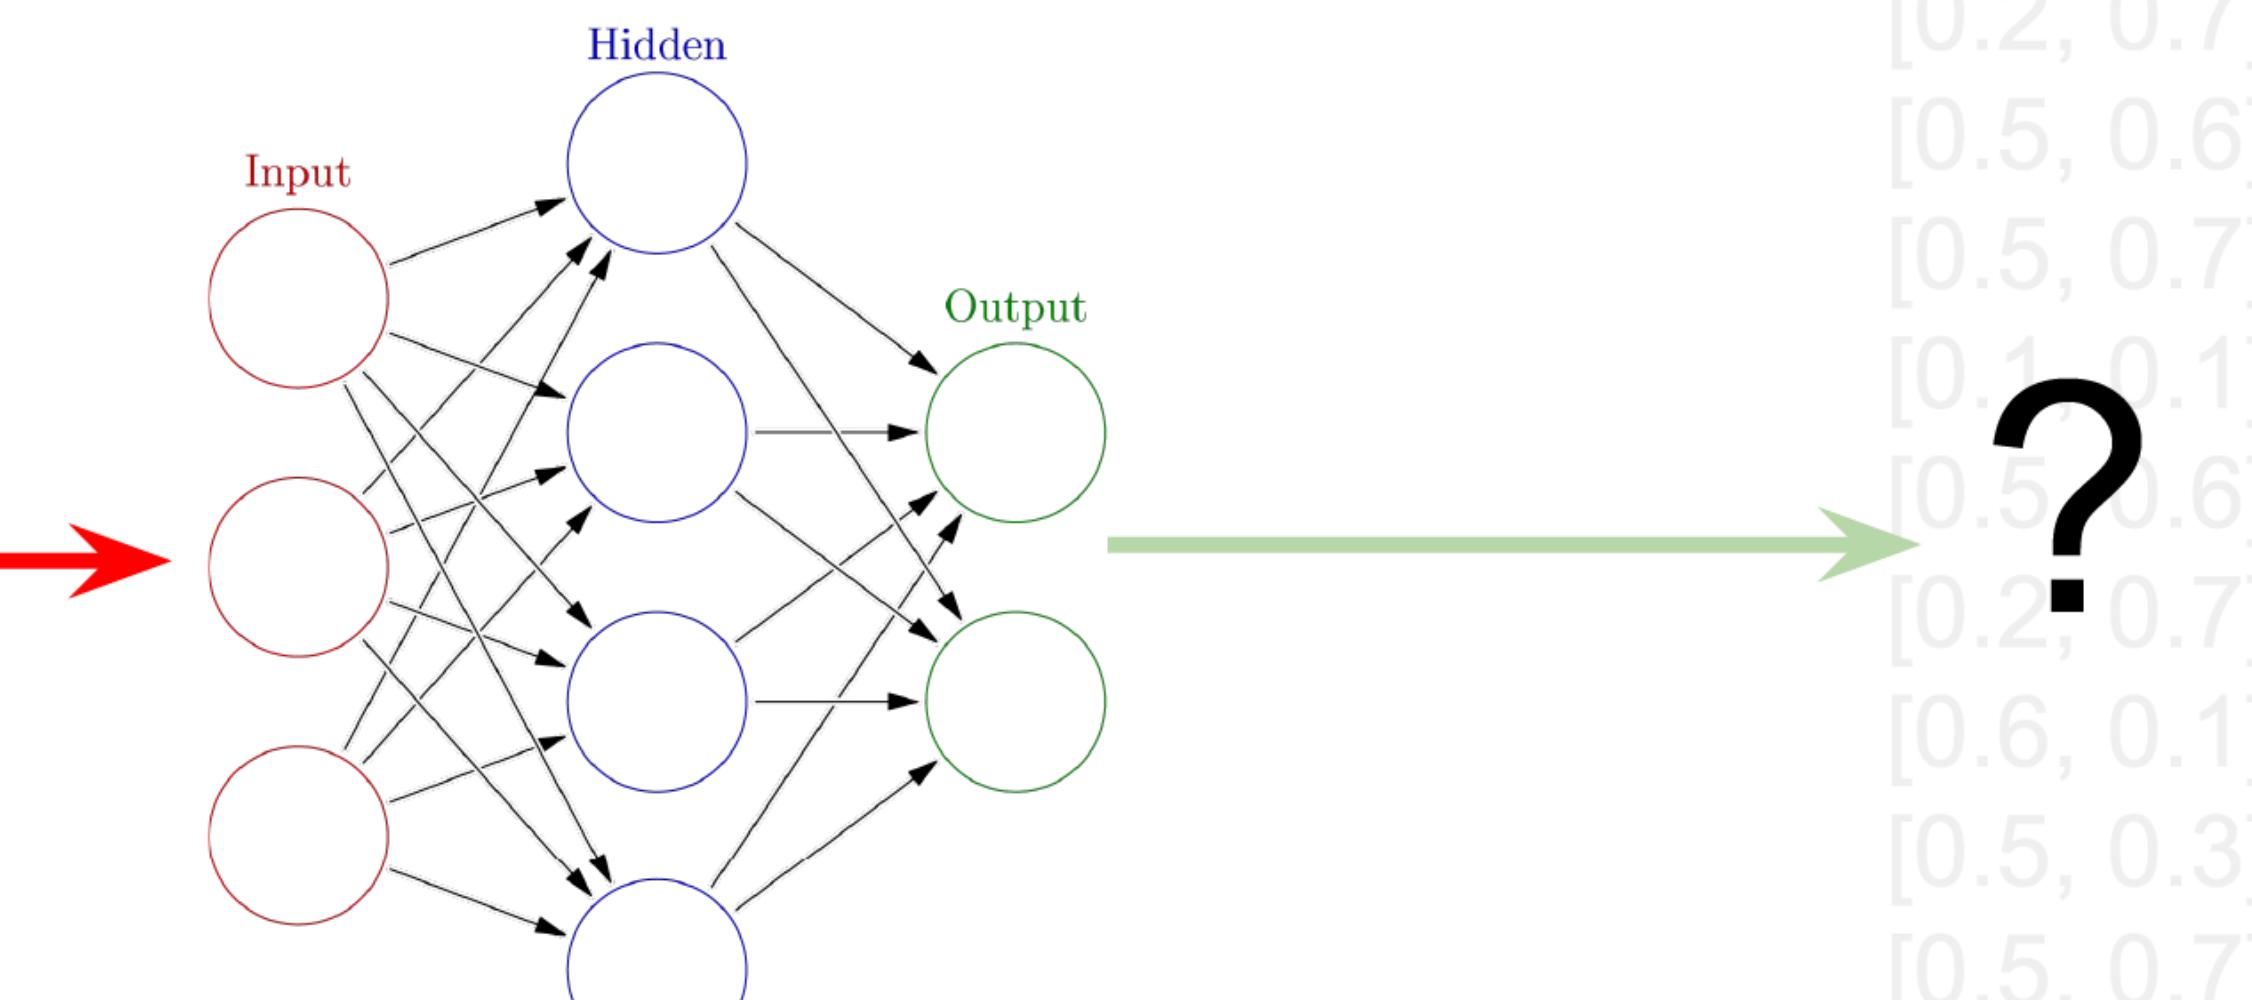


“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“j”: [0, 0, 0.5]  
“k”: [0, 1, 1]  
“l”: [0, 1, 1]  
“m”: [0, 0, 1]  
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]



# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“g”: [0.9, 0.3, 0]  
“h”: [0.4, 1, 0.7]  
“i”: [0.2, 0.3, 1]  
“j”: [0.1, 0.1, 0.1]  
“k”: [0, 1, 1]  
“l”: [0, 1, 1]  
“m”: [0, 0, 1]  
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]



[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]  
[0.5, 0.6]  
[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.2, 0.7]  
[0.6, 0.1]  
[0.5, 0.3]  
[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]

*There are certain things that we care about, as musicians for example, that are really hard to articulate in code. It's hard for me to talk about what kind of quality of sound I want and then translate that into a set of filter coefficients. It's hard for me to talk about how I want a performer to move on stage and then translate that into some sort of mathematical equation for their trajectory. But it's a lot easier for me to either find examples of sounds that have a particular quality or to give examples of movements or if I'm using other types of modalities, often curating or creating examples are just way easier for us as people. And this relates to the types of tacit knowledge and embodied knowledge we bring to creative practices.*

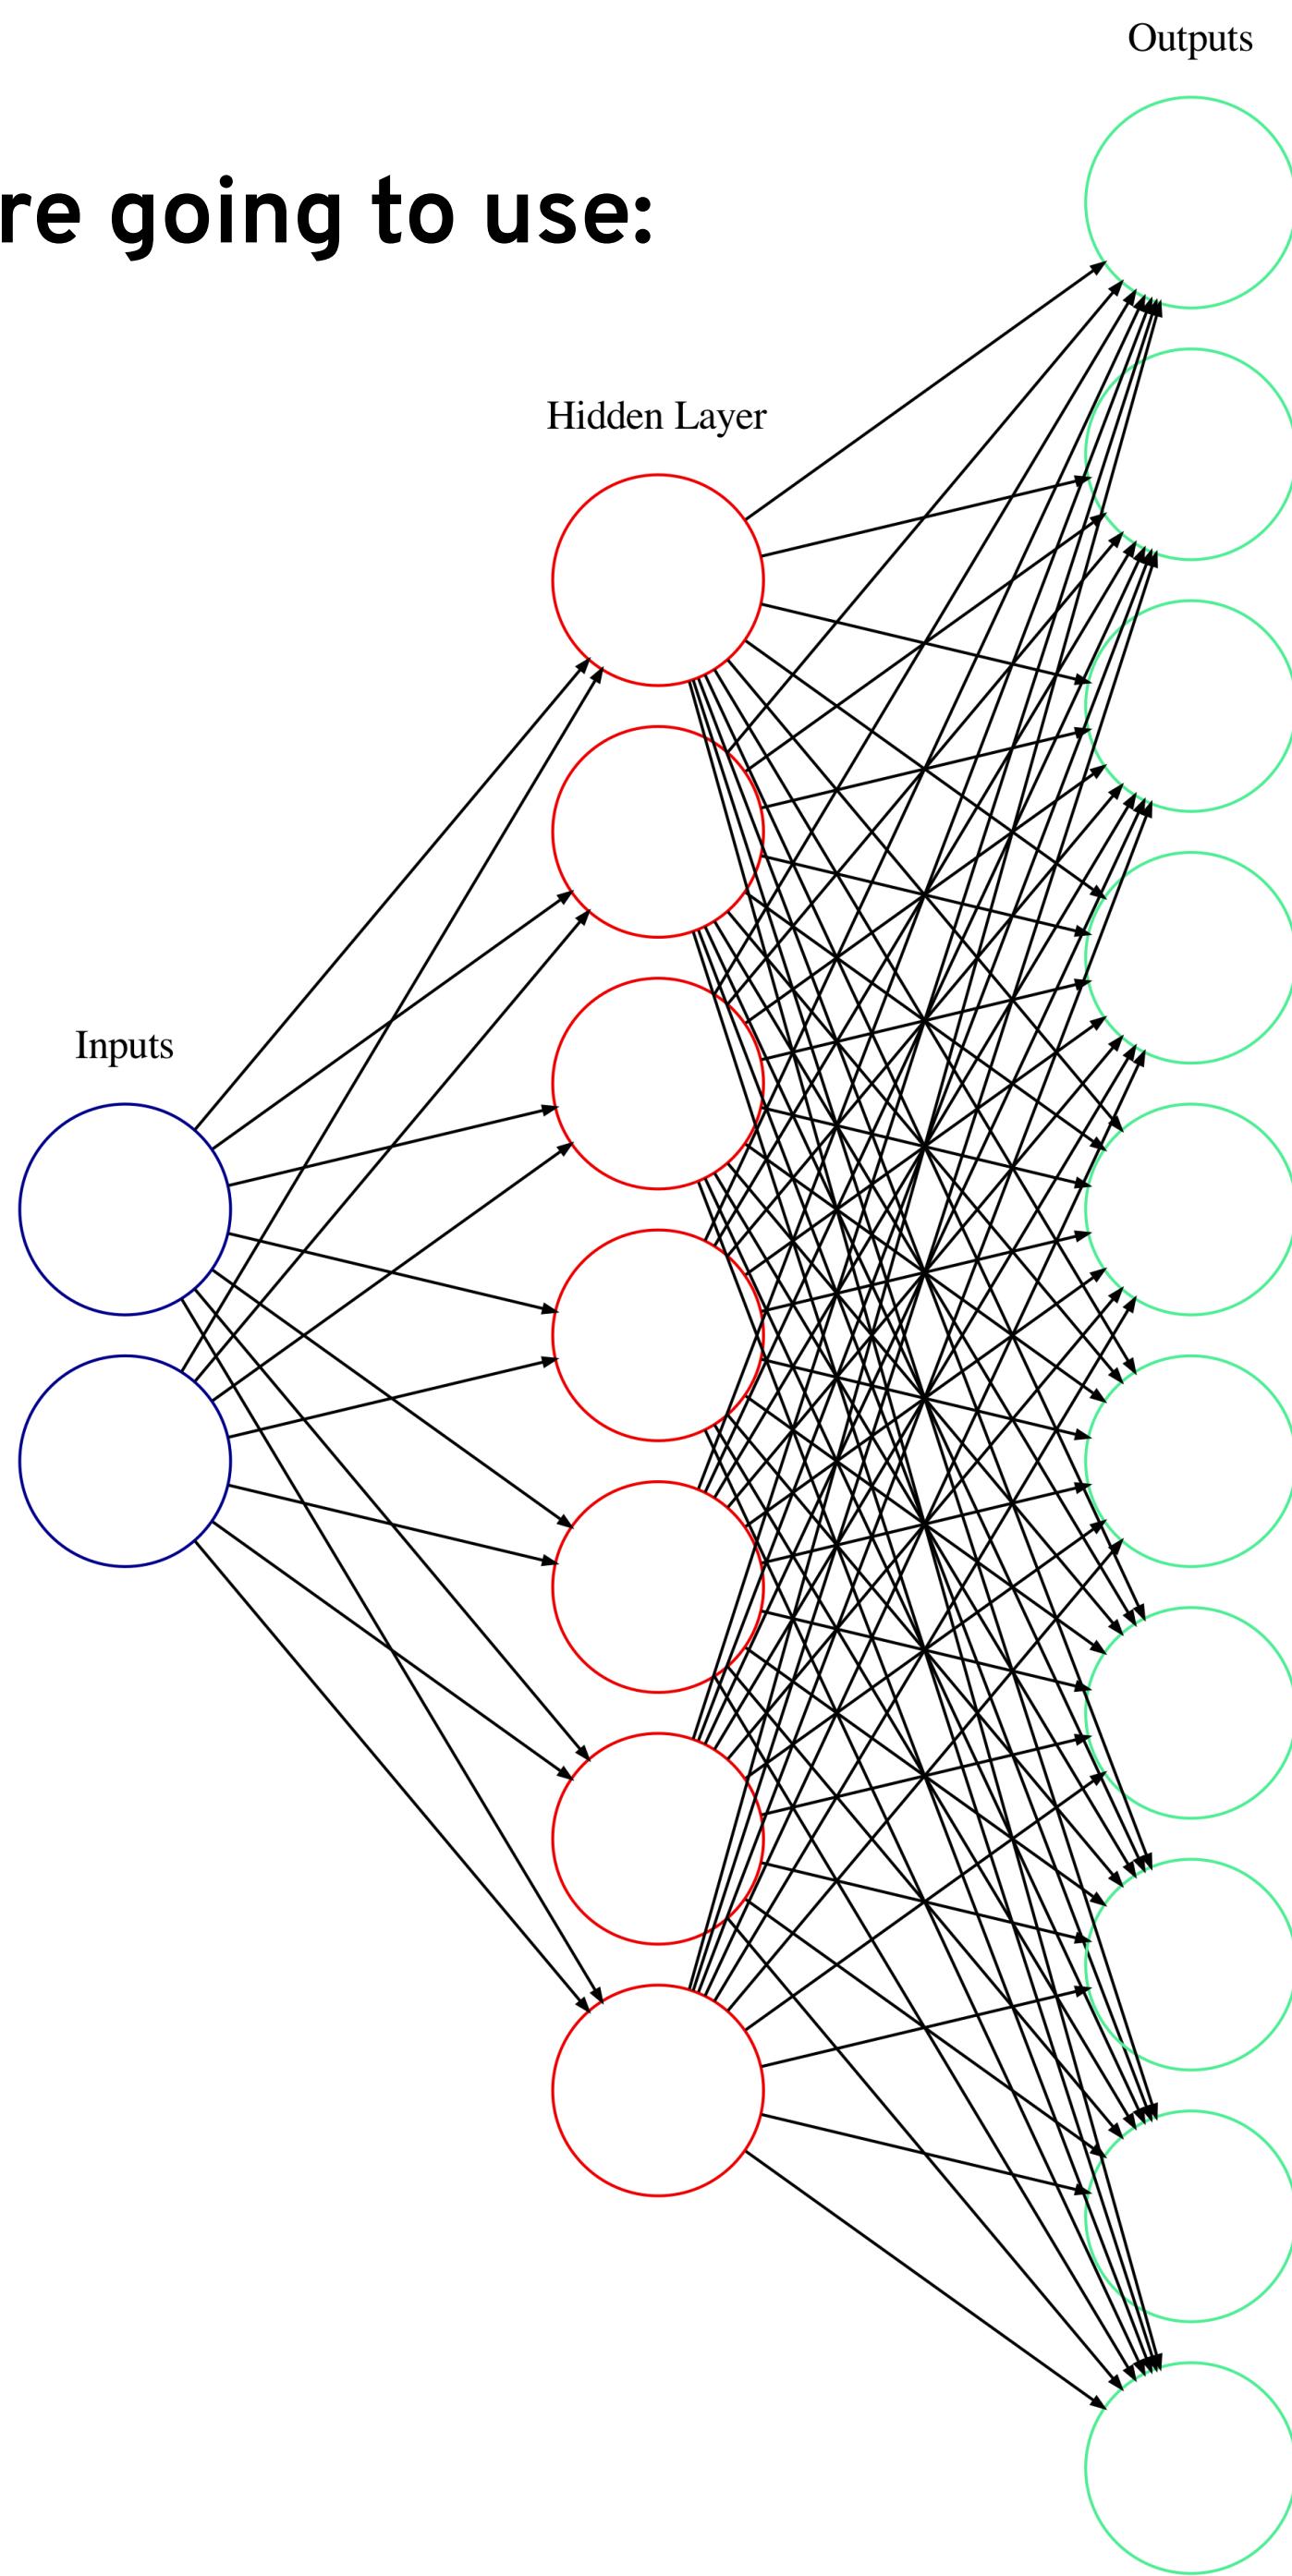
-Rebecca Fiebrink

## Structure of the neural network we're going to use:

2 inputs

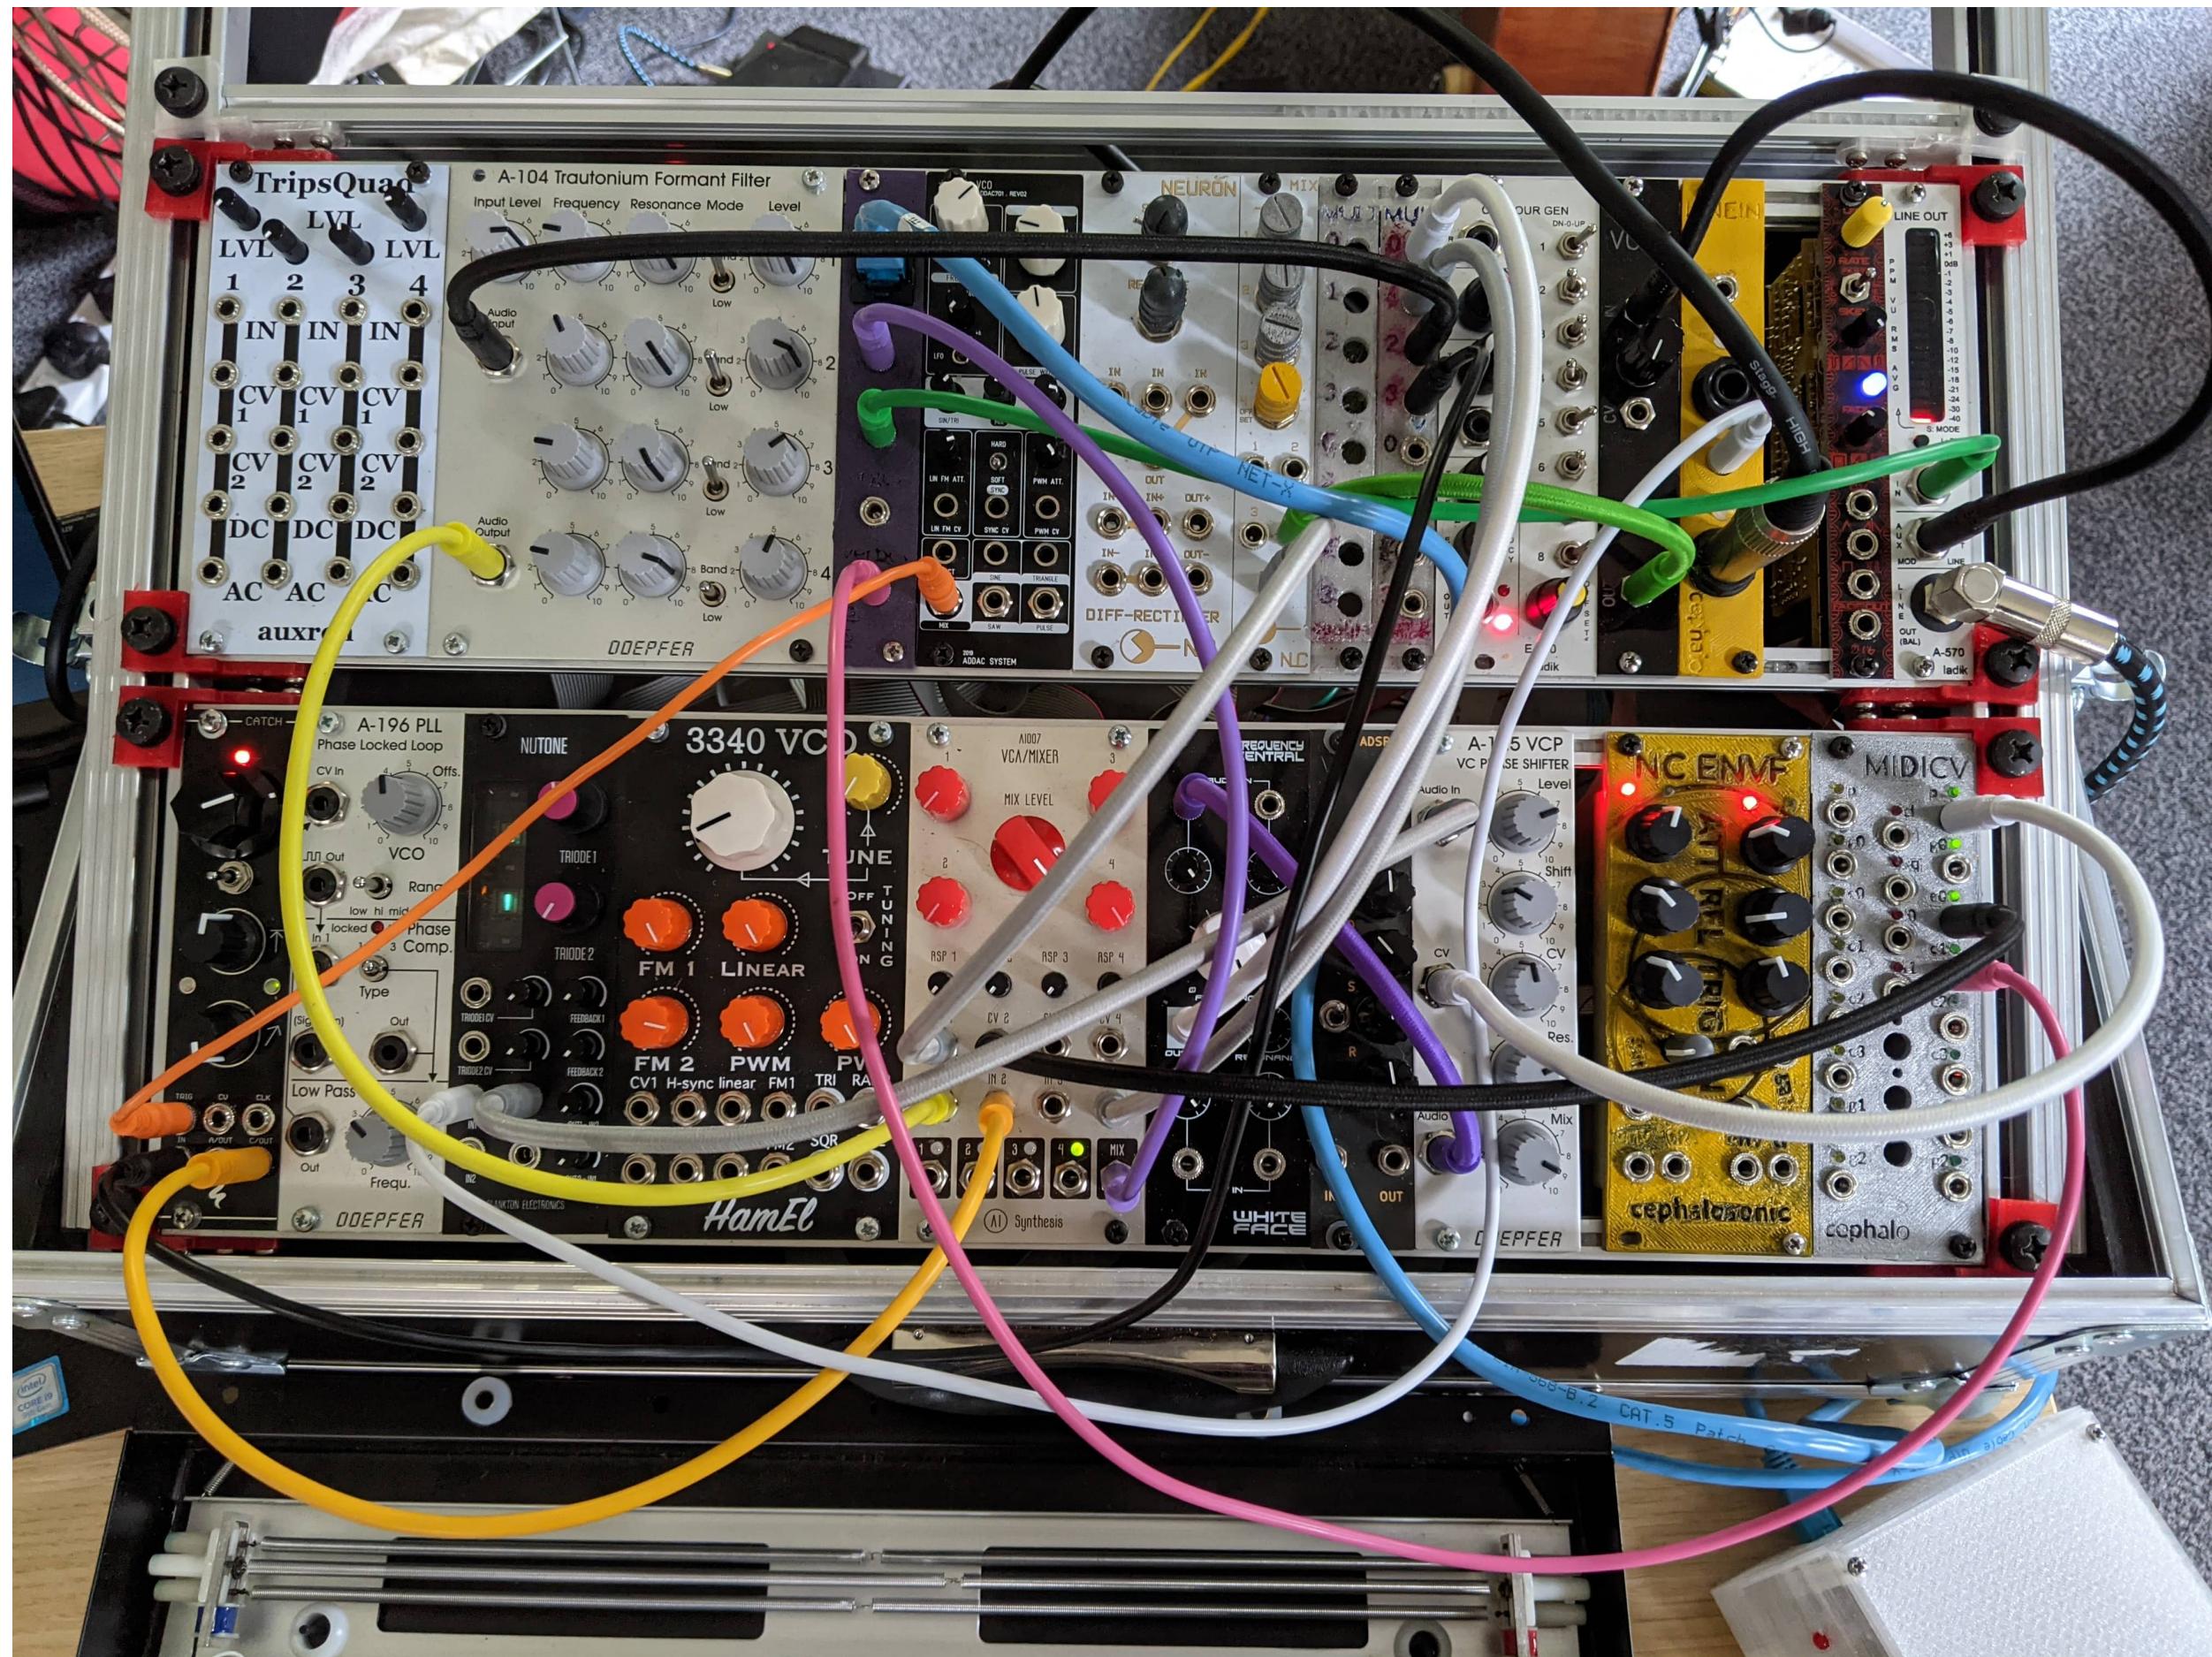
1 hidden layer of 7 nodes

10 outputs



*fluid.ml/pregressor~*





# FeedbackFeedforward

by Alice Eldridge and  
Chris Kiefer

[learn.flucoma.org/explore/eldridge-kiefer](https://learn.flucoma.org/explore/eldridge-kiefer)