

FluCoMa



Ted Moore

tedmooremusic.com

ted@tedmooremusic.com

FluCoMa: Fluid Corpus Manipulation



- “enable techno-fluent musicians to use machine listening and machine learning in their creative practices”
- Integrating Machine Listening and Machine Learning in...
- Max, SuperCollider & Pure Data
- Learning Resources (learn.flucoma.org)
- Discourse Community (discourse.flucoma.org)



FluCoMa: Fluid Corpus Manipulation



Slice Audio

- onset slice
- transient slice
- novelty slice
- amplitude slice
- amplitude gate

Decompose Audio

- extract transients
- harmonic/percussive separation
- model as sine waves
- non-negative matrix factorisation

Analyse Audio

- pitch
- loudness
- mel-bands

- mel-frequency cepstral coefficients
- spectral centroid
- spectral flatness
- chromagram

Transform Audio

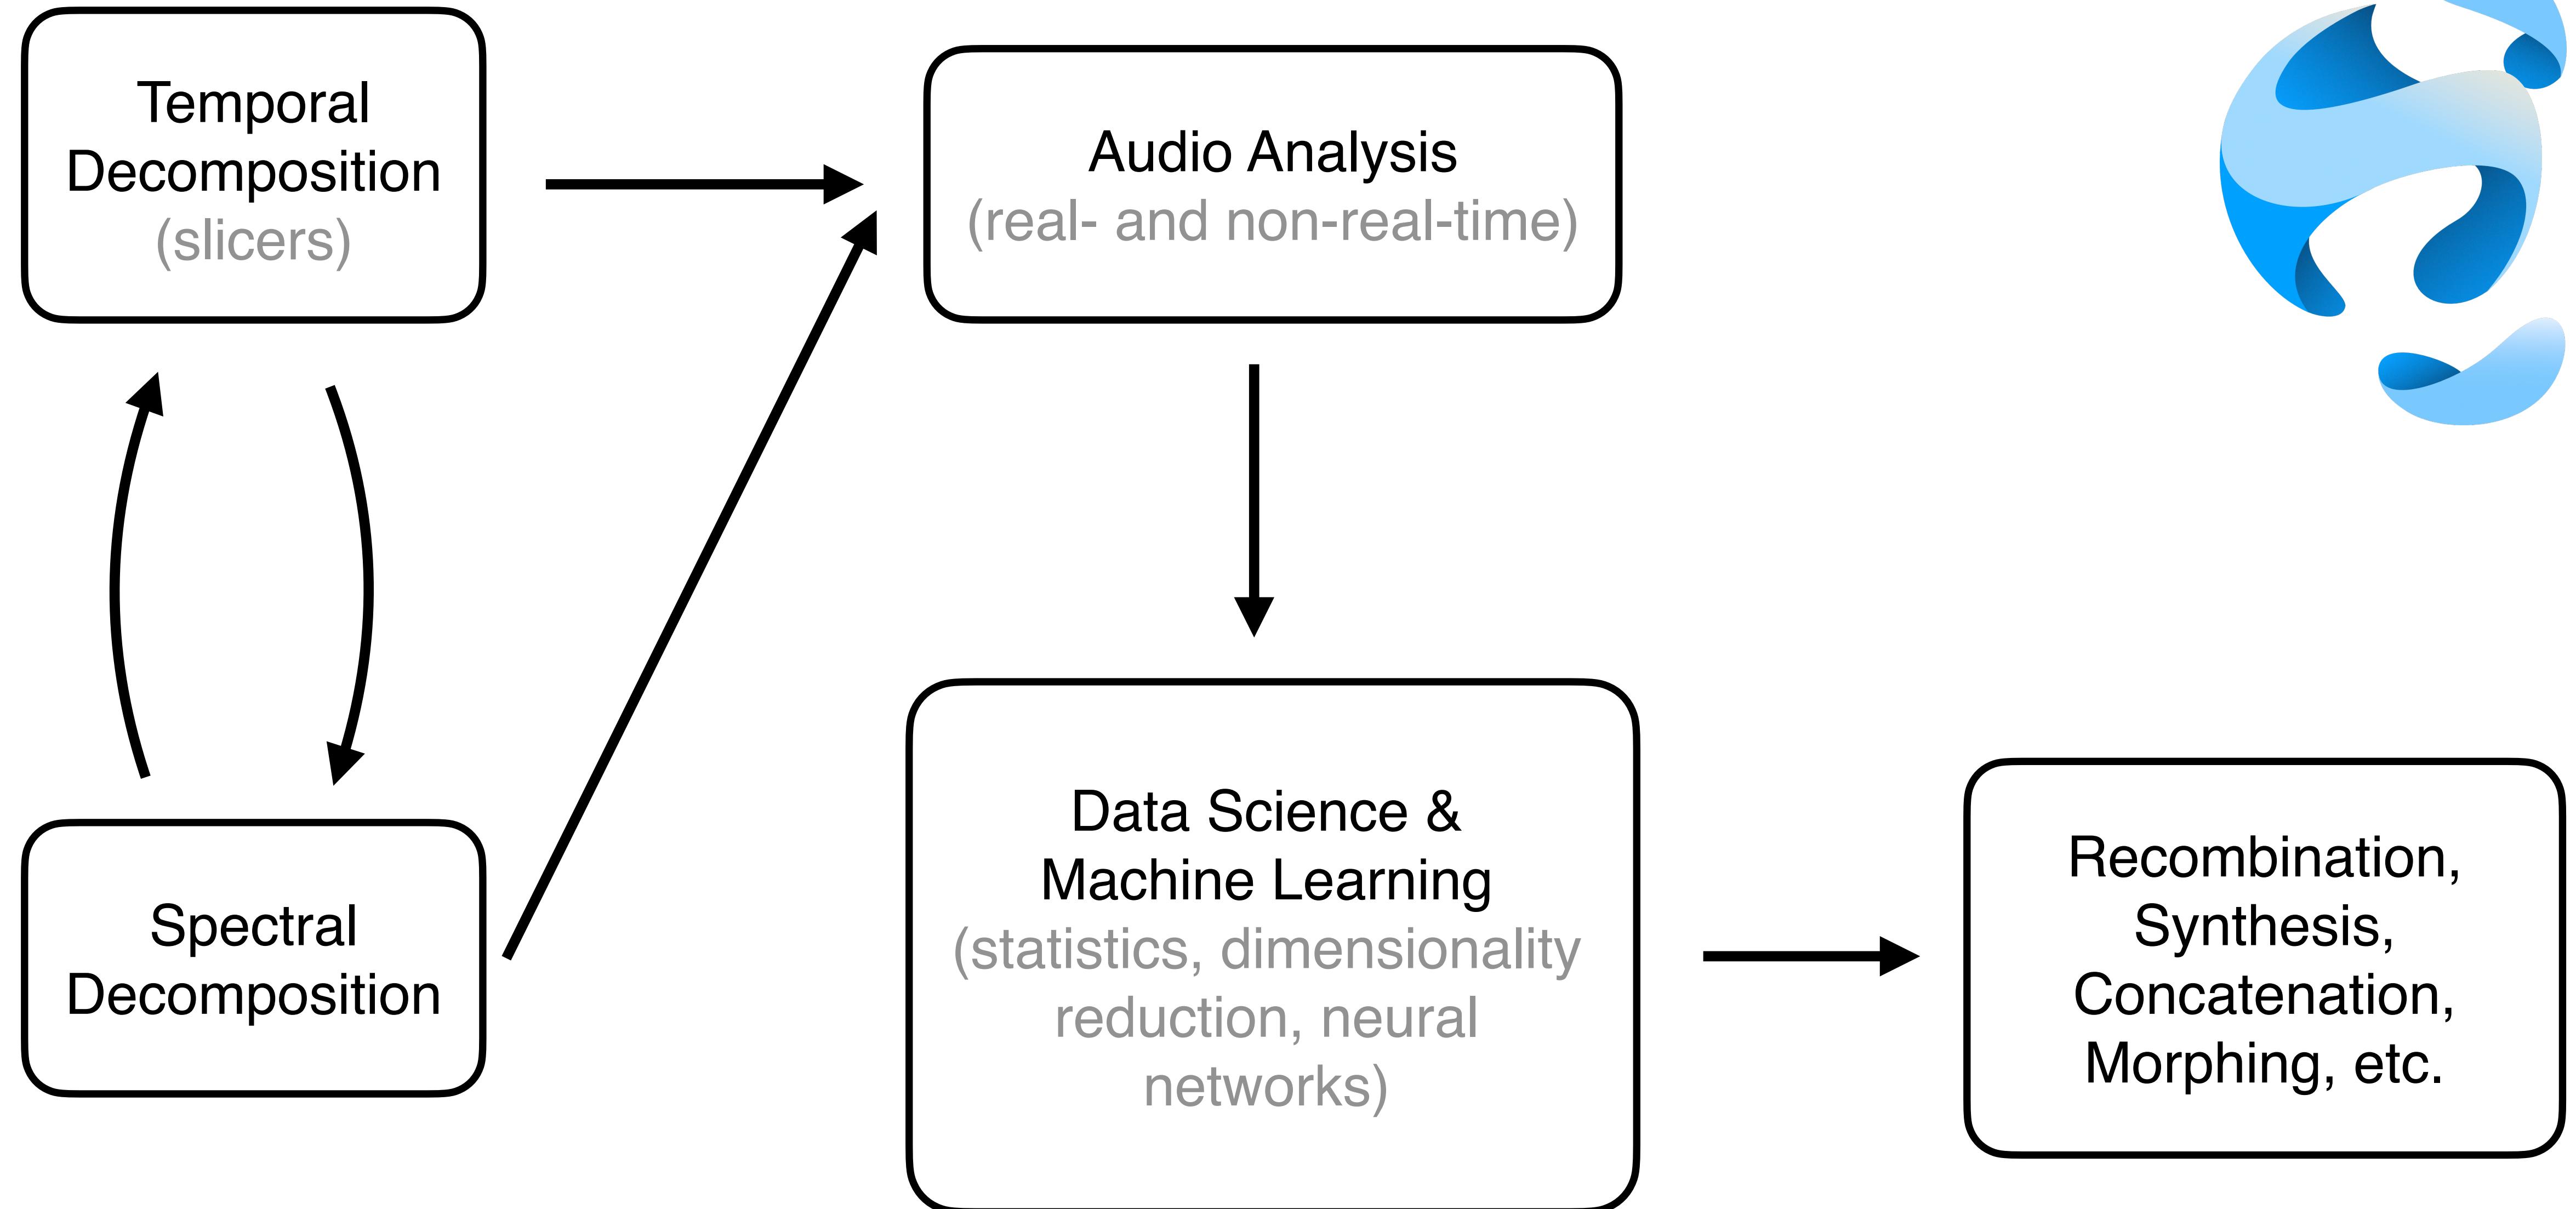
- audio transport
- non-negative matrix factorisation filters & morphing

Analyse Data

- datasets
- labelsets
- statistical analyses
- normalization
- standardization
- robust scaler
- principal component analysis

- MDS
- KDTree
- K Nearest Neighbours
- neural networks
- SQL-type query
- KMeans
- UMAP
- grid

... and more



MLPRegressor

control *many* synthesizer
parameters from a smaller
control space



fluid.mlpregressor~



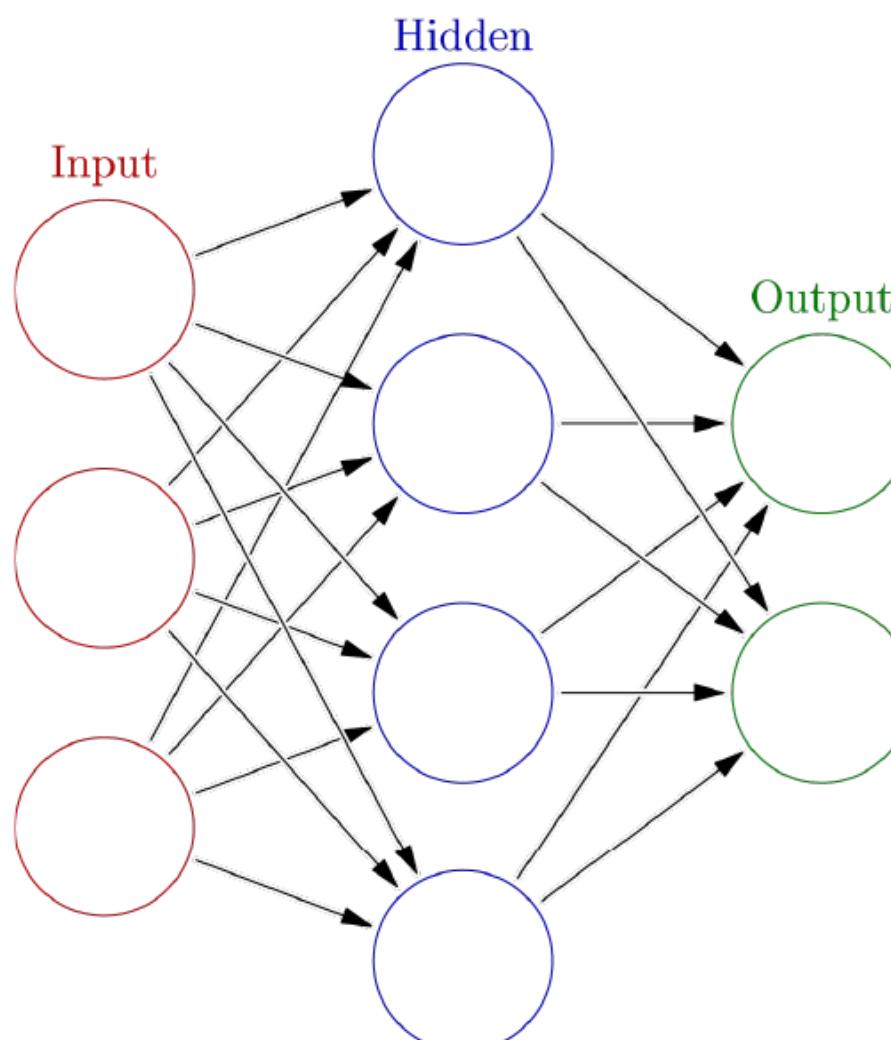
Neural Network Training a Regressor

identifier

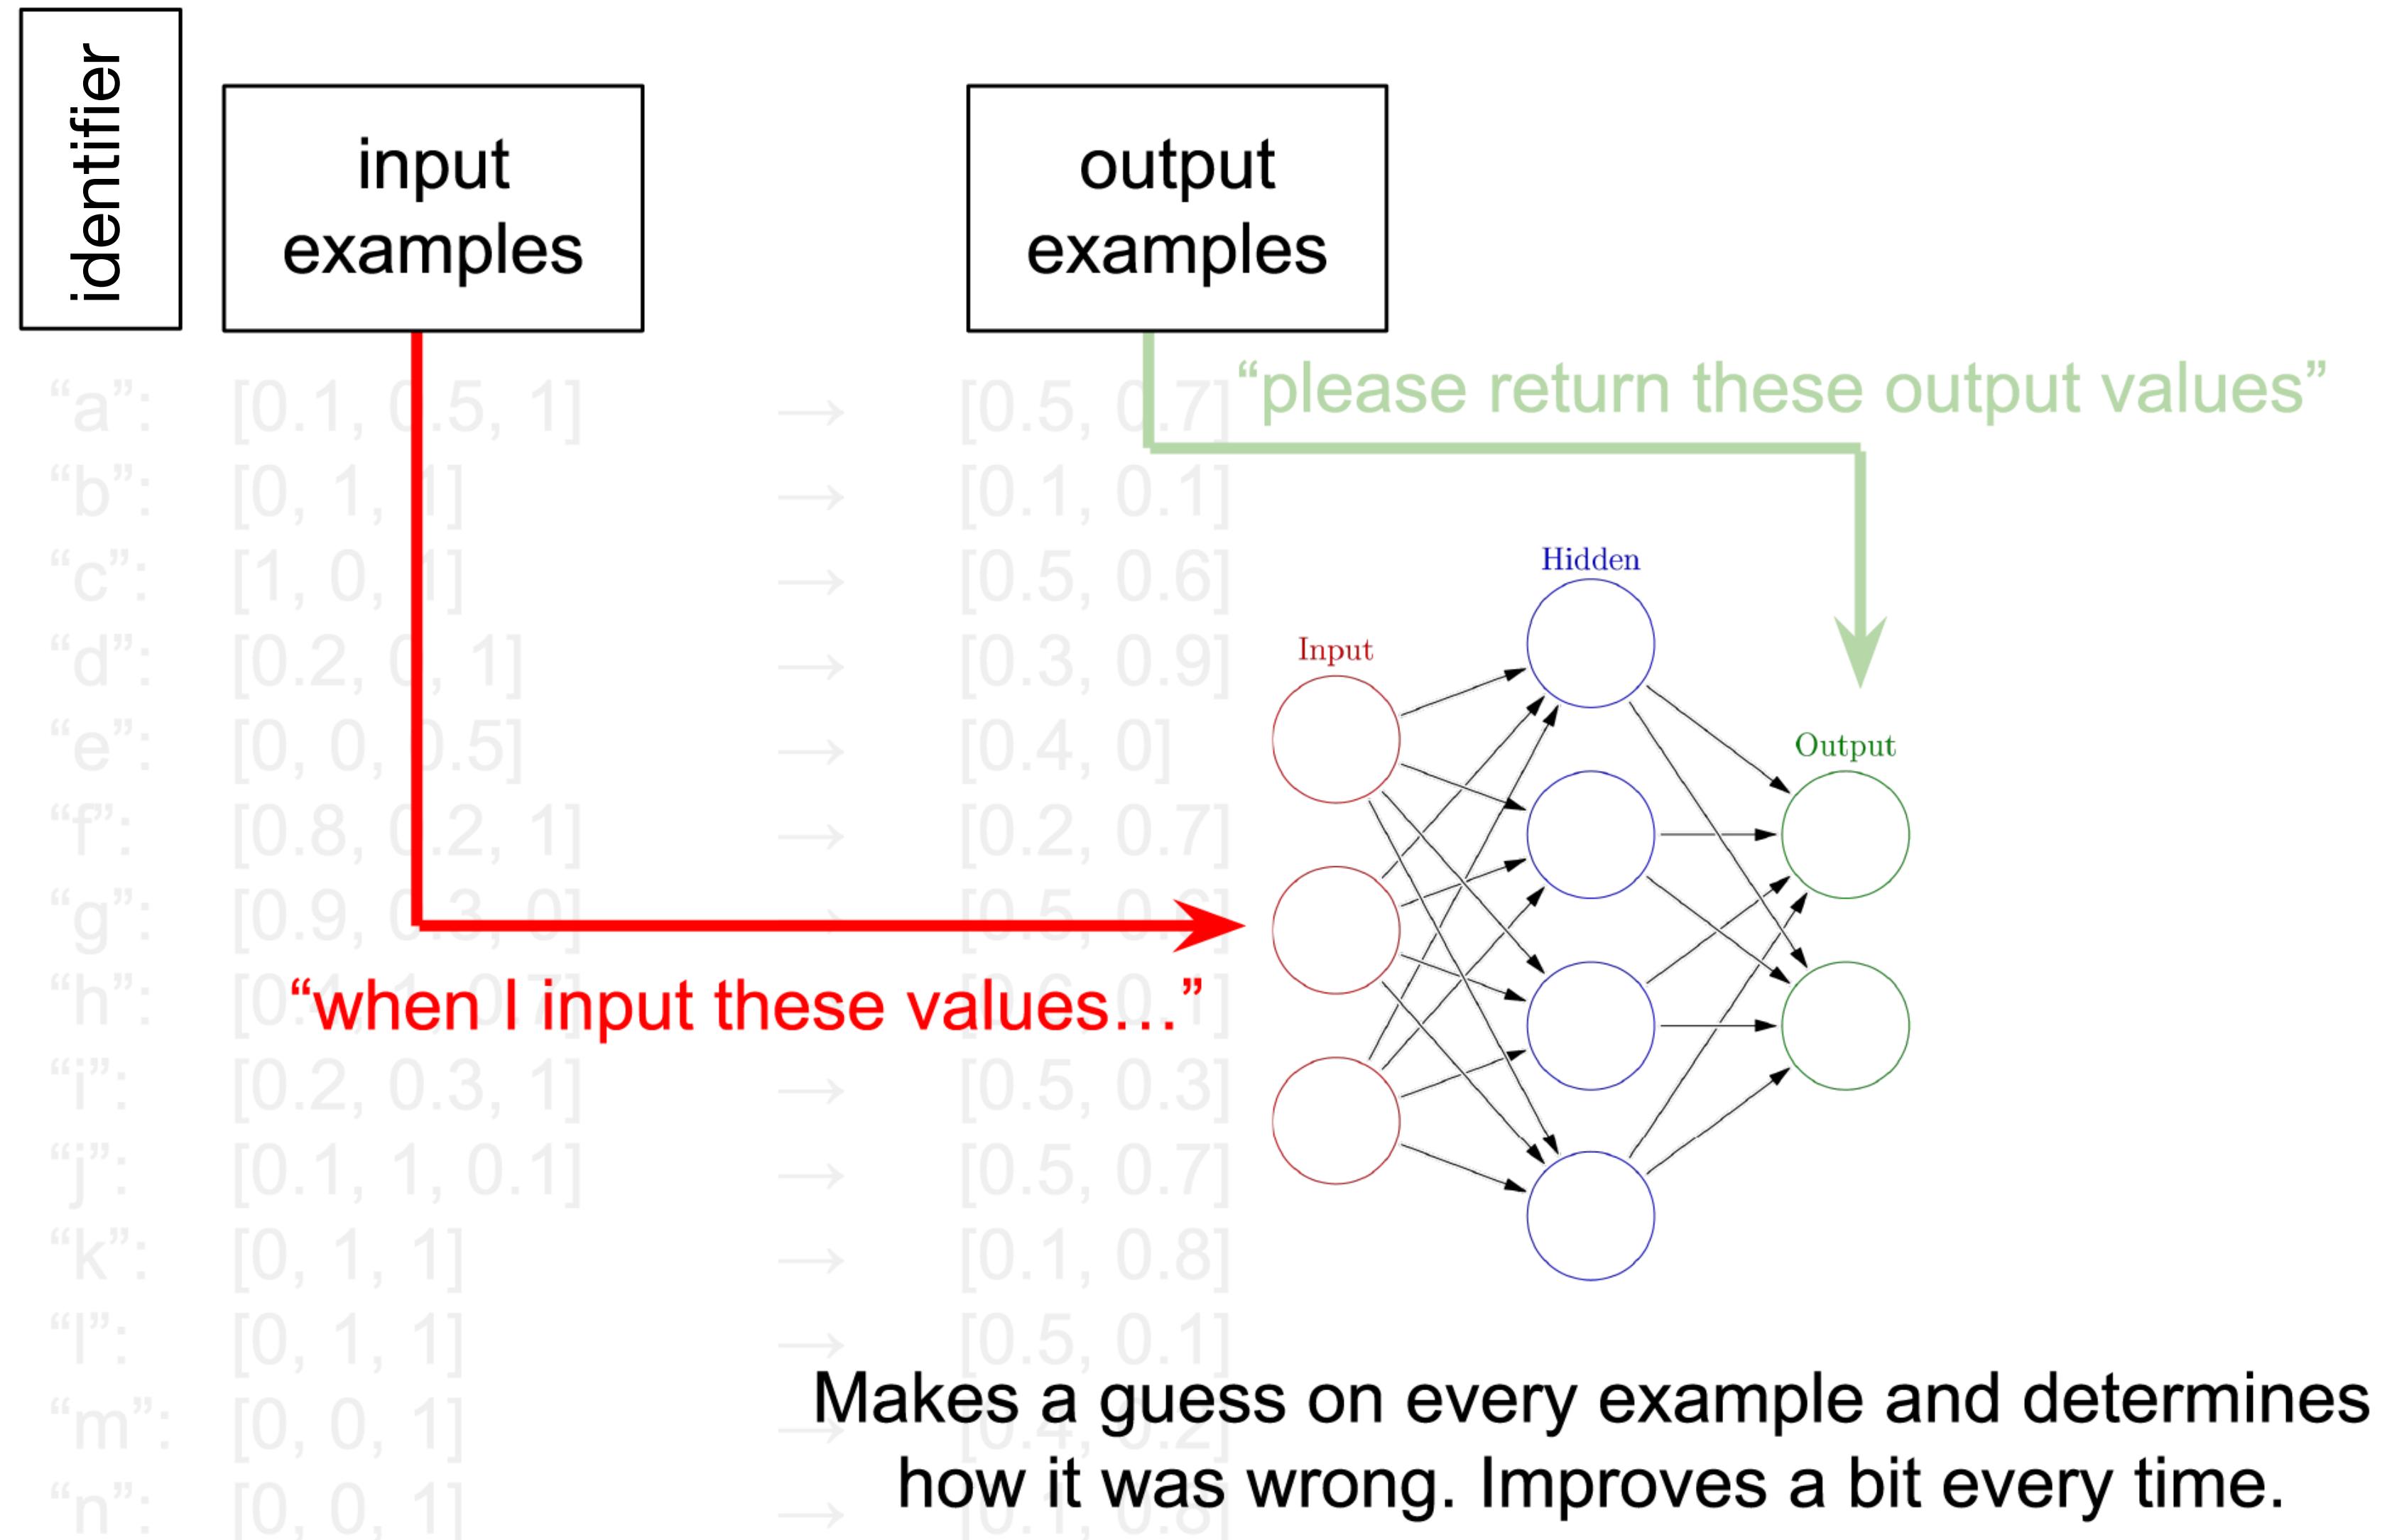
input
examples

“a”:	[0.1, 0.5, 1]	→	[0.5, 0.7]
“b”:	[0, 1, 1]	→	[0.1, 0.1]
“c”:	[1, 0, 1]	→	[0.5, 0.6]
“d”:	[0.2, 0, 1]	→	[0.3, 0.9]
“e”:	[0, 0, 0.5]	→	[0.4, 0]
“f”:	[0.8, 0.2, 1]	→	[0.2, 0.7]
“g”:	[0.9, 0.3, 0]	→	[0.5, 0.6]
“h”:	[0.4, 1, 0.7]	→	[0.6, 0.1]
“i”:	[0.2, 0.3, 1]	→	[0.5, 0.3]
“j”:	[0.1, 1, 0.1]	→	[0.5, 0.7]
“k”:	[0, 1, 1]	→	[0.1, 0.8]
“l”:	[0, 1, 1]	→	[0.5, 0.1]
“m”:	[0, 0, 1]	→	[0.4, 0.2]
“n”:	[0, 0, 1]	→	[0.1, 0.8]

output
examples



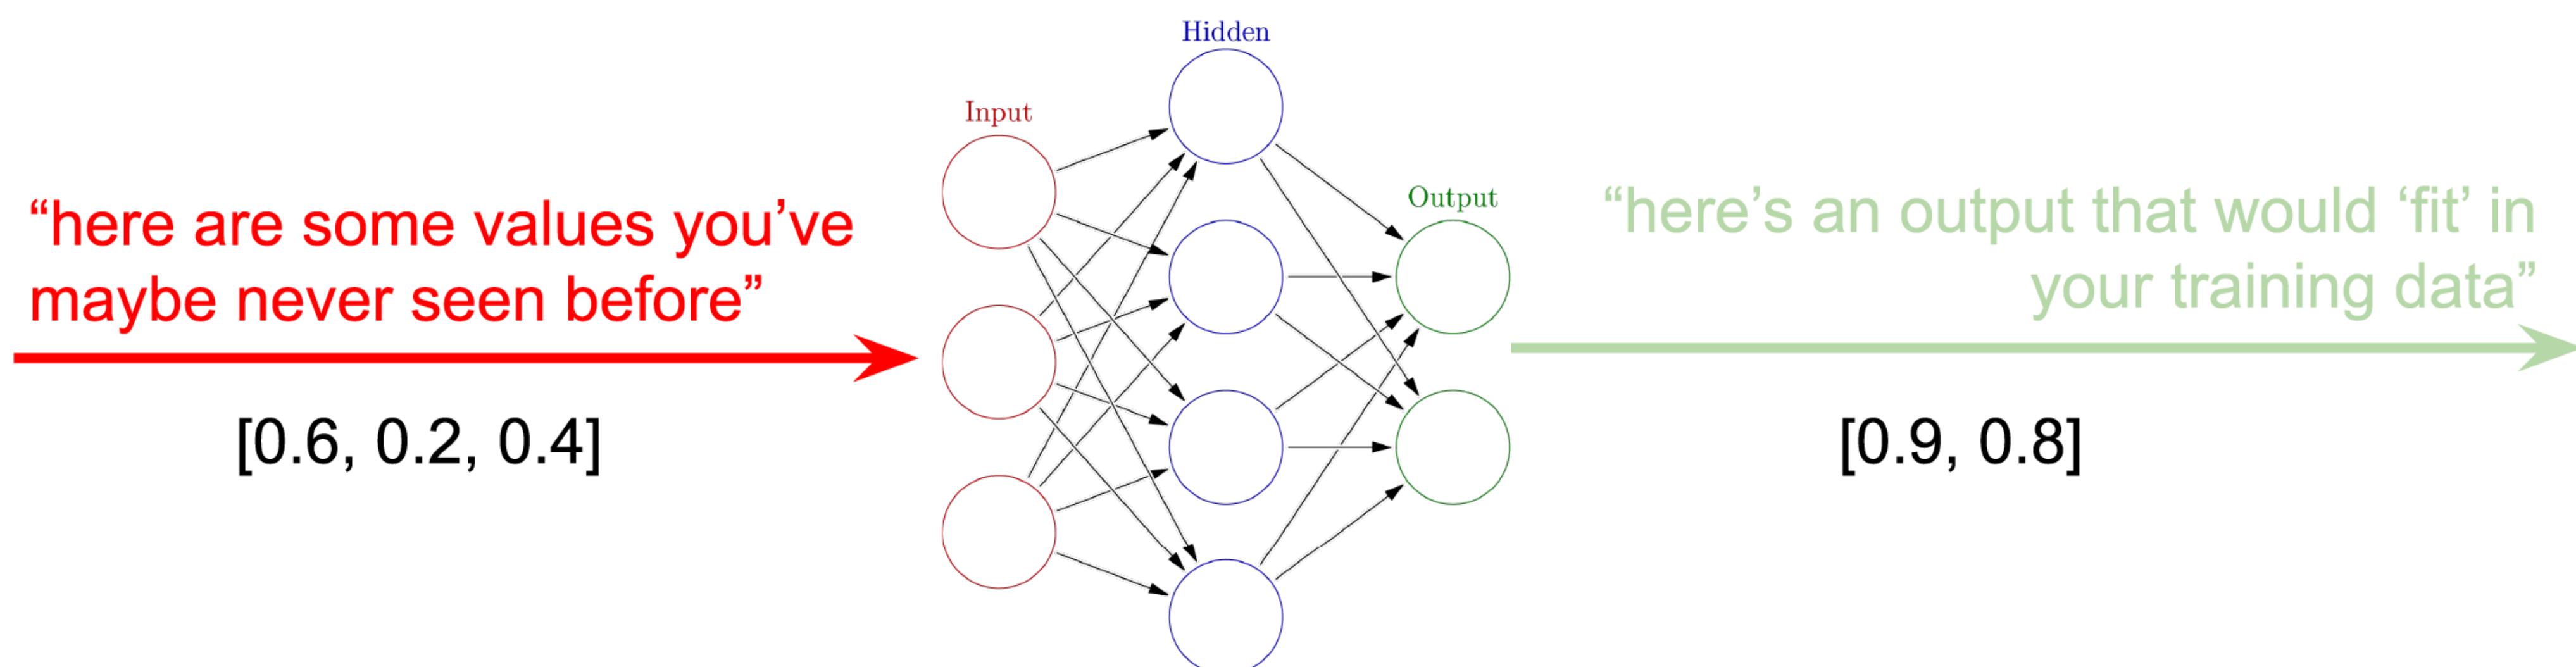
Neural Network *Training* a Regressor



Neural Network *Predicting with Regression*

input
examples

output
examples



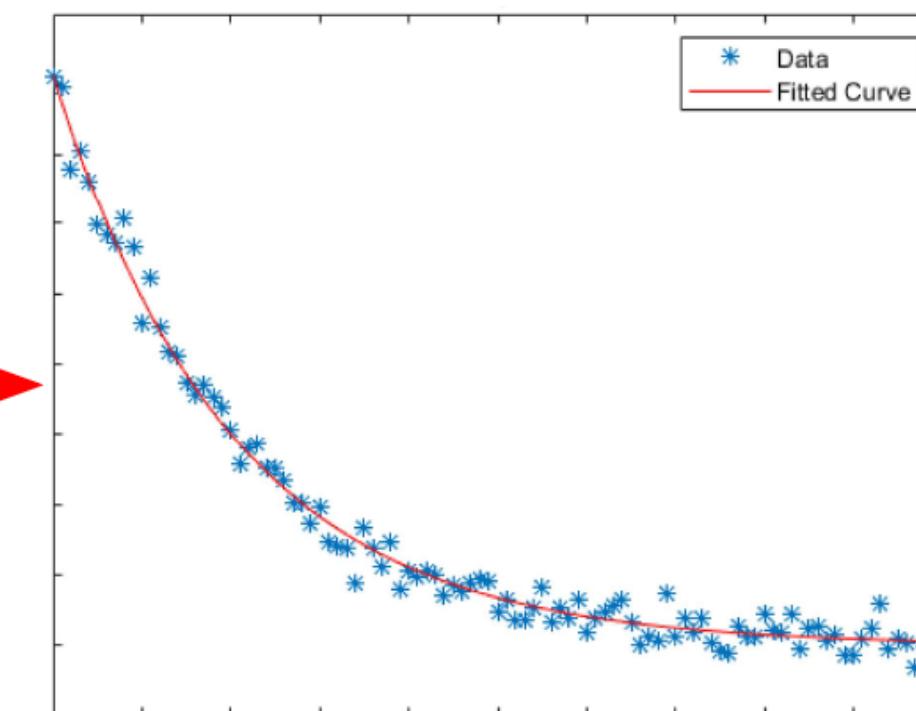
Neural Network *Predicting with Regression*

input
examples

output
examples

“here are some values you’ve
maybe never seen before”

[0.6, 0.2, 0.4]



“here’s an output that would ‘fit’ in
your training data”

[0.9, 0.8]

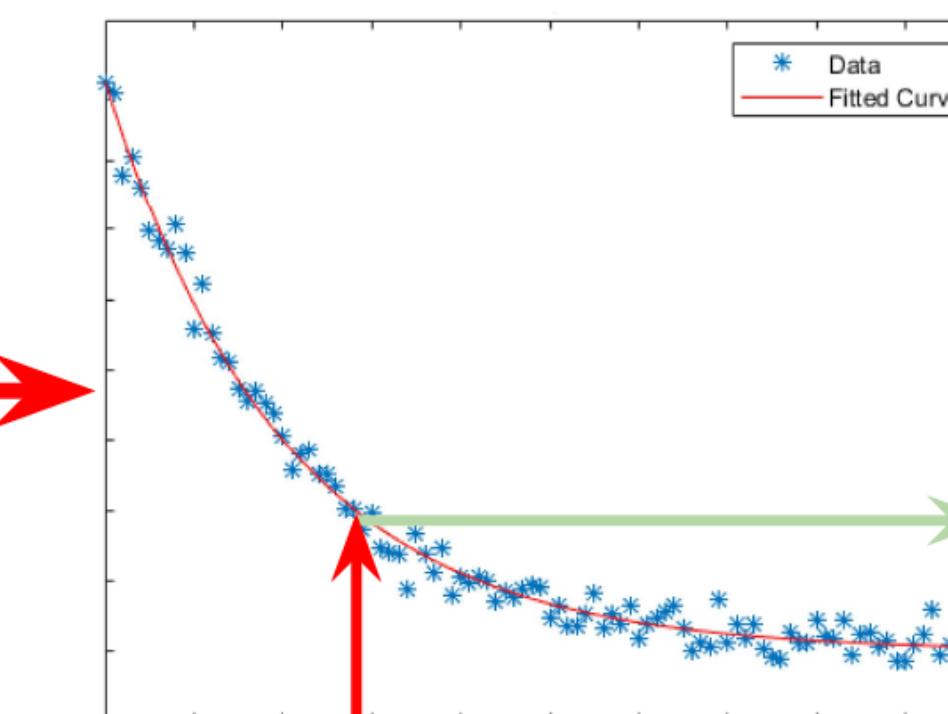
Neural Network *Predicting with Regression*

input
examples

output
examples

“here are some values you’ve
maybe never seen before”

[0.6, 0.2, 0.4]

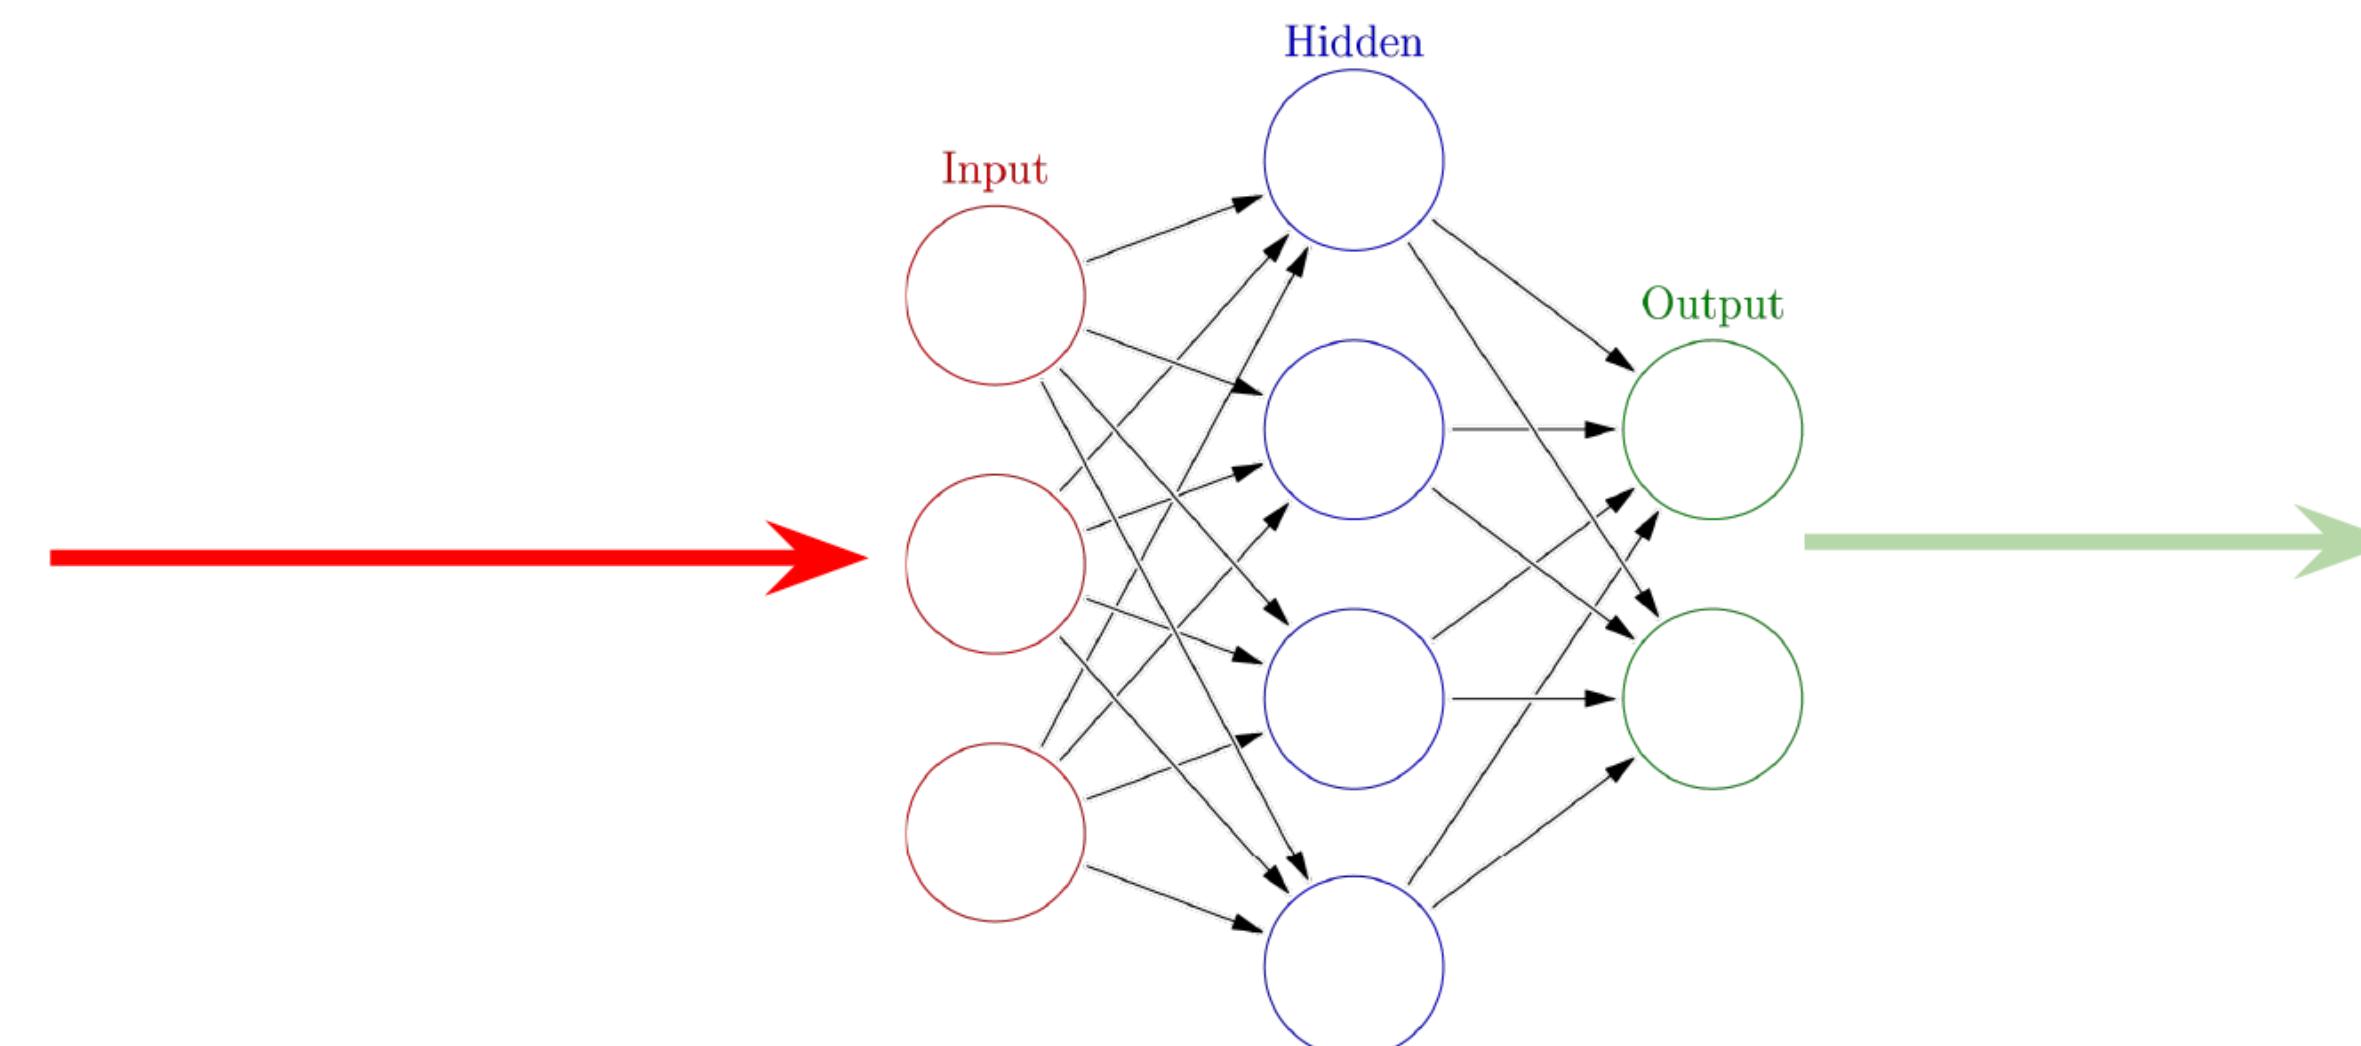


“here’s an output that would ‘fit’ in
your training data”

[0.9, 0.8]

Neural Network Predicting with Regression

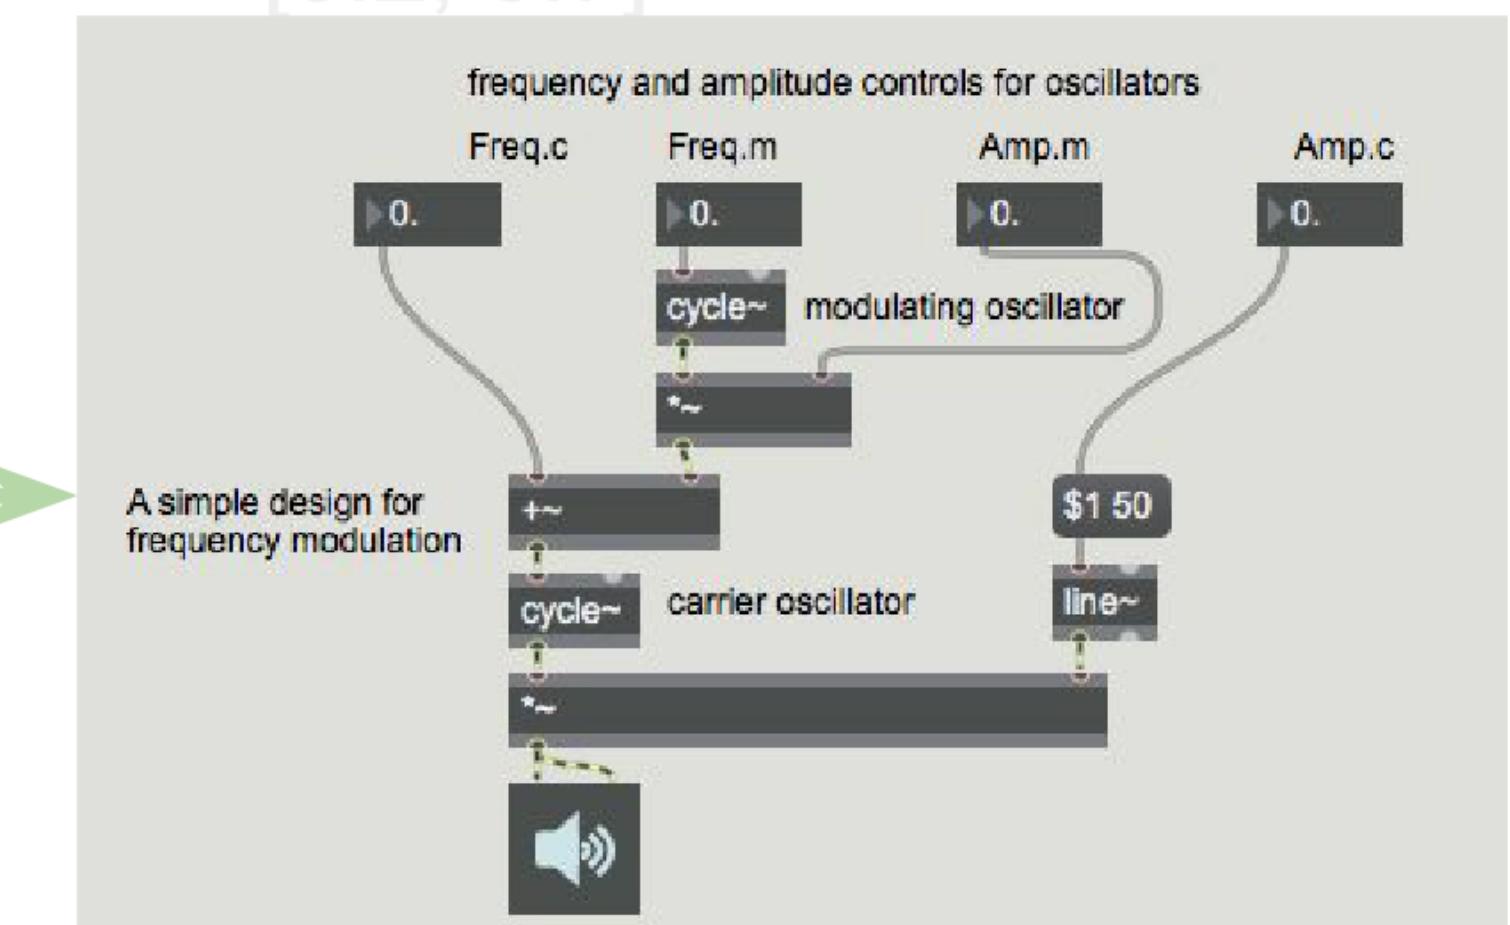
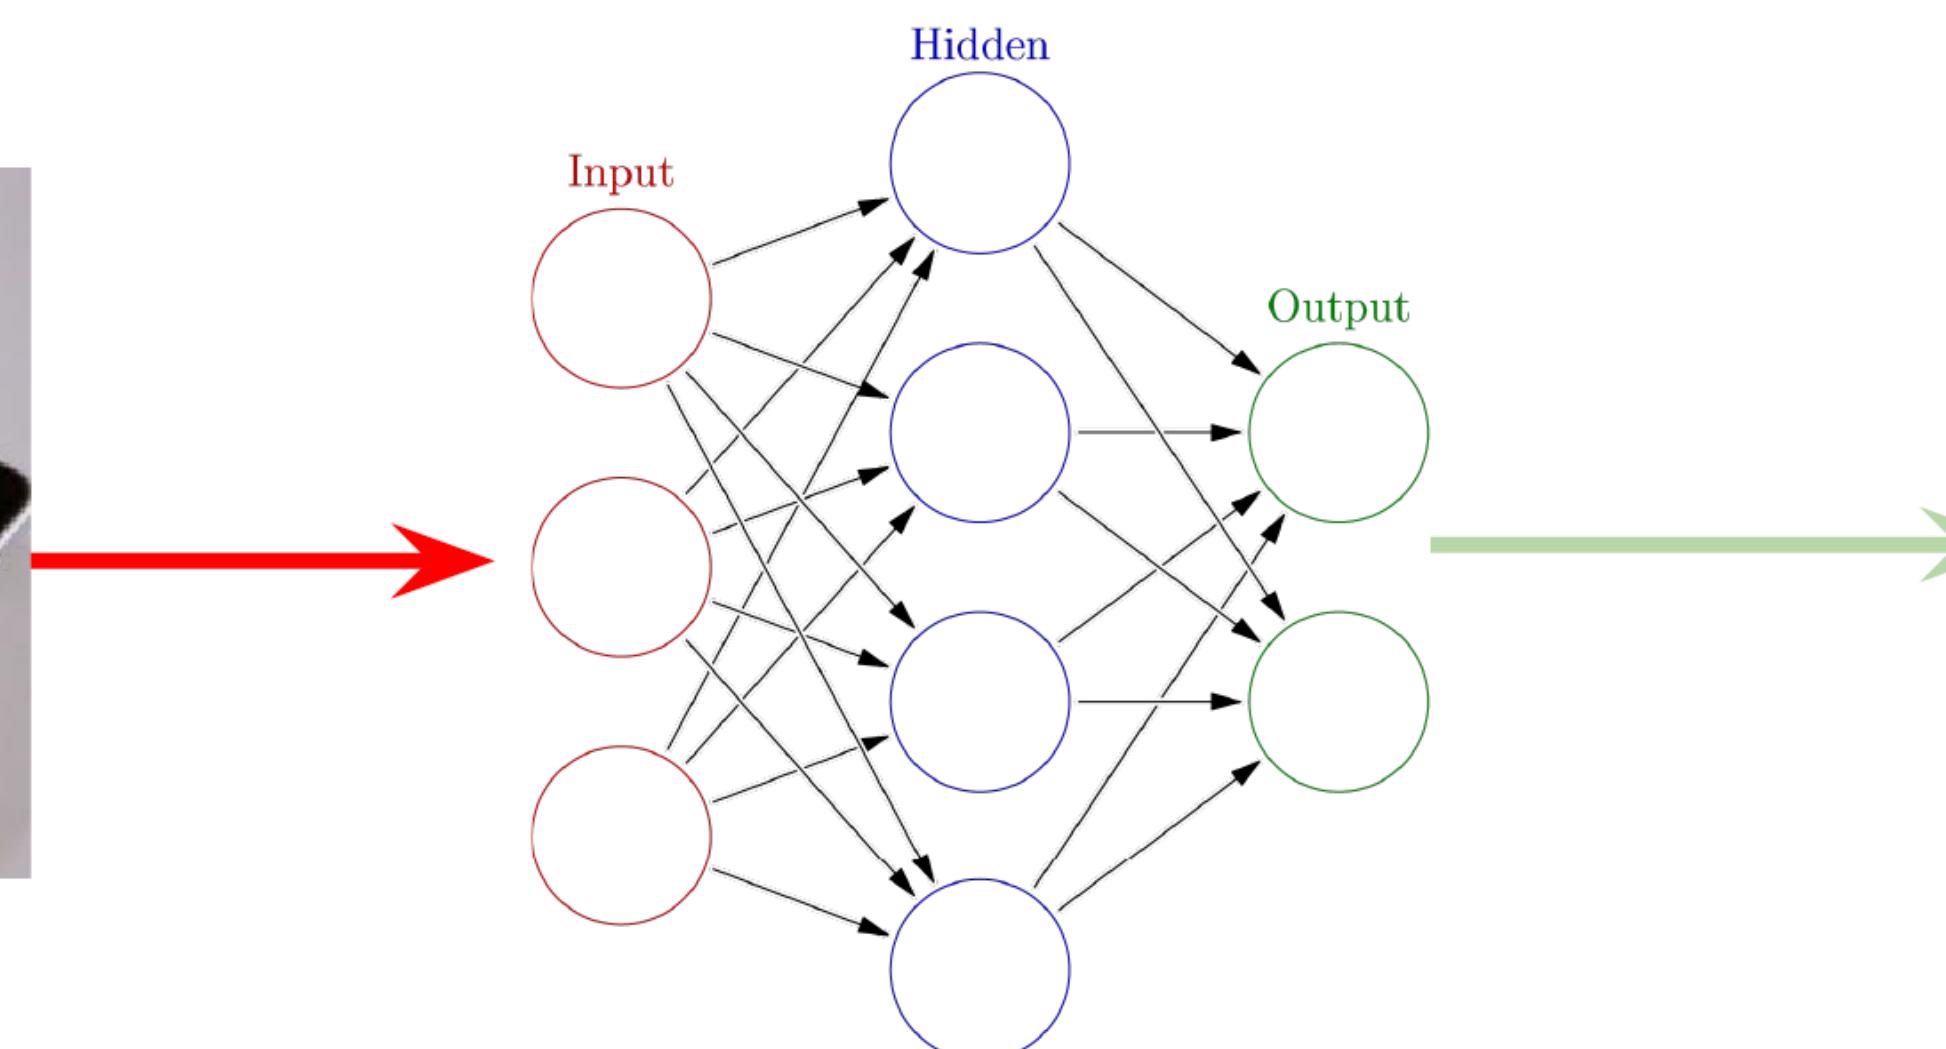
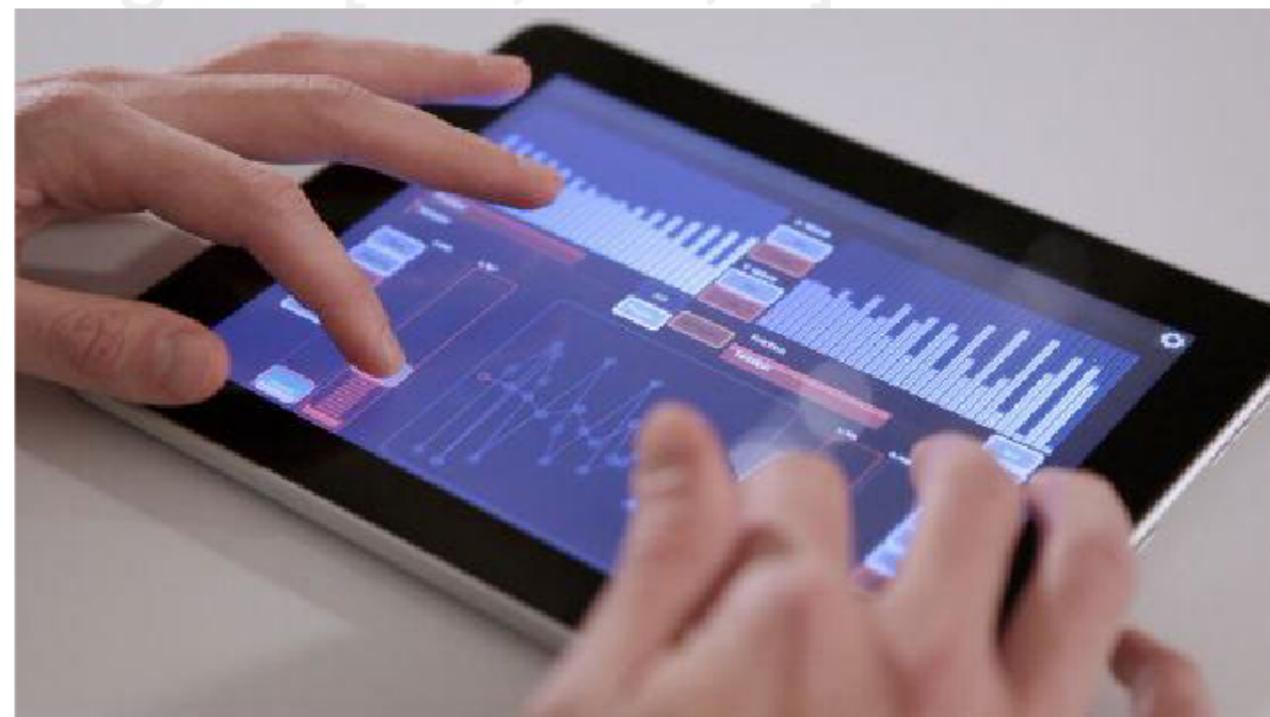
“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“g”: [0.9, 0.3, 0]
“h”: [0.4, 1, 0.7]
“i”: [0.2, 0.3, 1]
“j”: [0.1, 1, 0.1]
“k”: [0, 1, 1]
“l”: [0, 1, 1]
“m”: [0, 0, 1]
“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]



[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0]
[0.2, 0.7]
[0.5, 0.6]
[0.6, 0.1]
[0.5, 0.3]
[0.5, 0.7]
[0.1, 0.8]
[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0.3]
[0.2, 0.7]

Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“g”: [0.9, 0.3, 0]

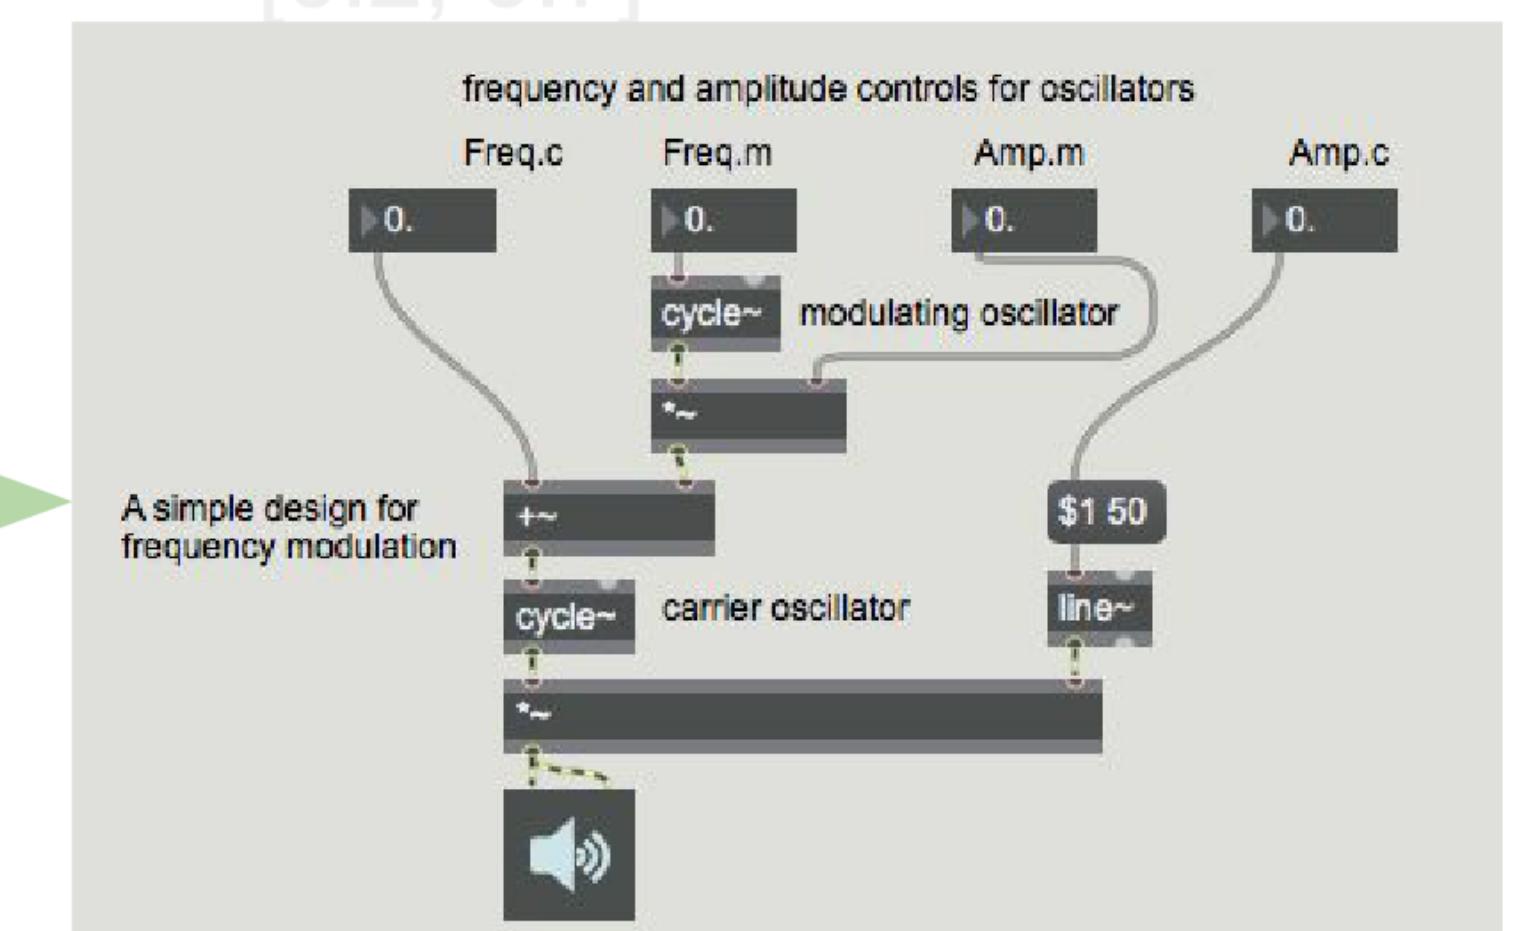
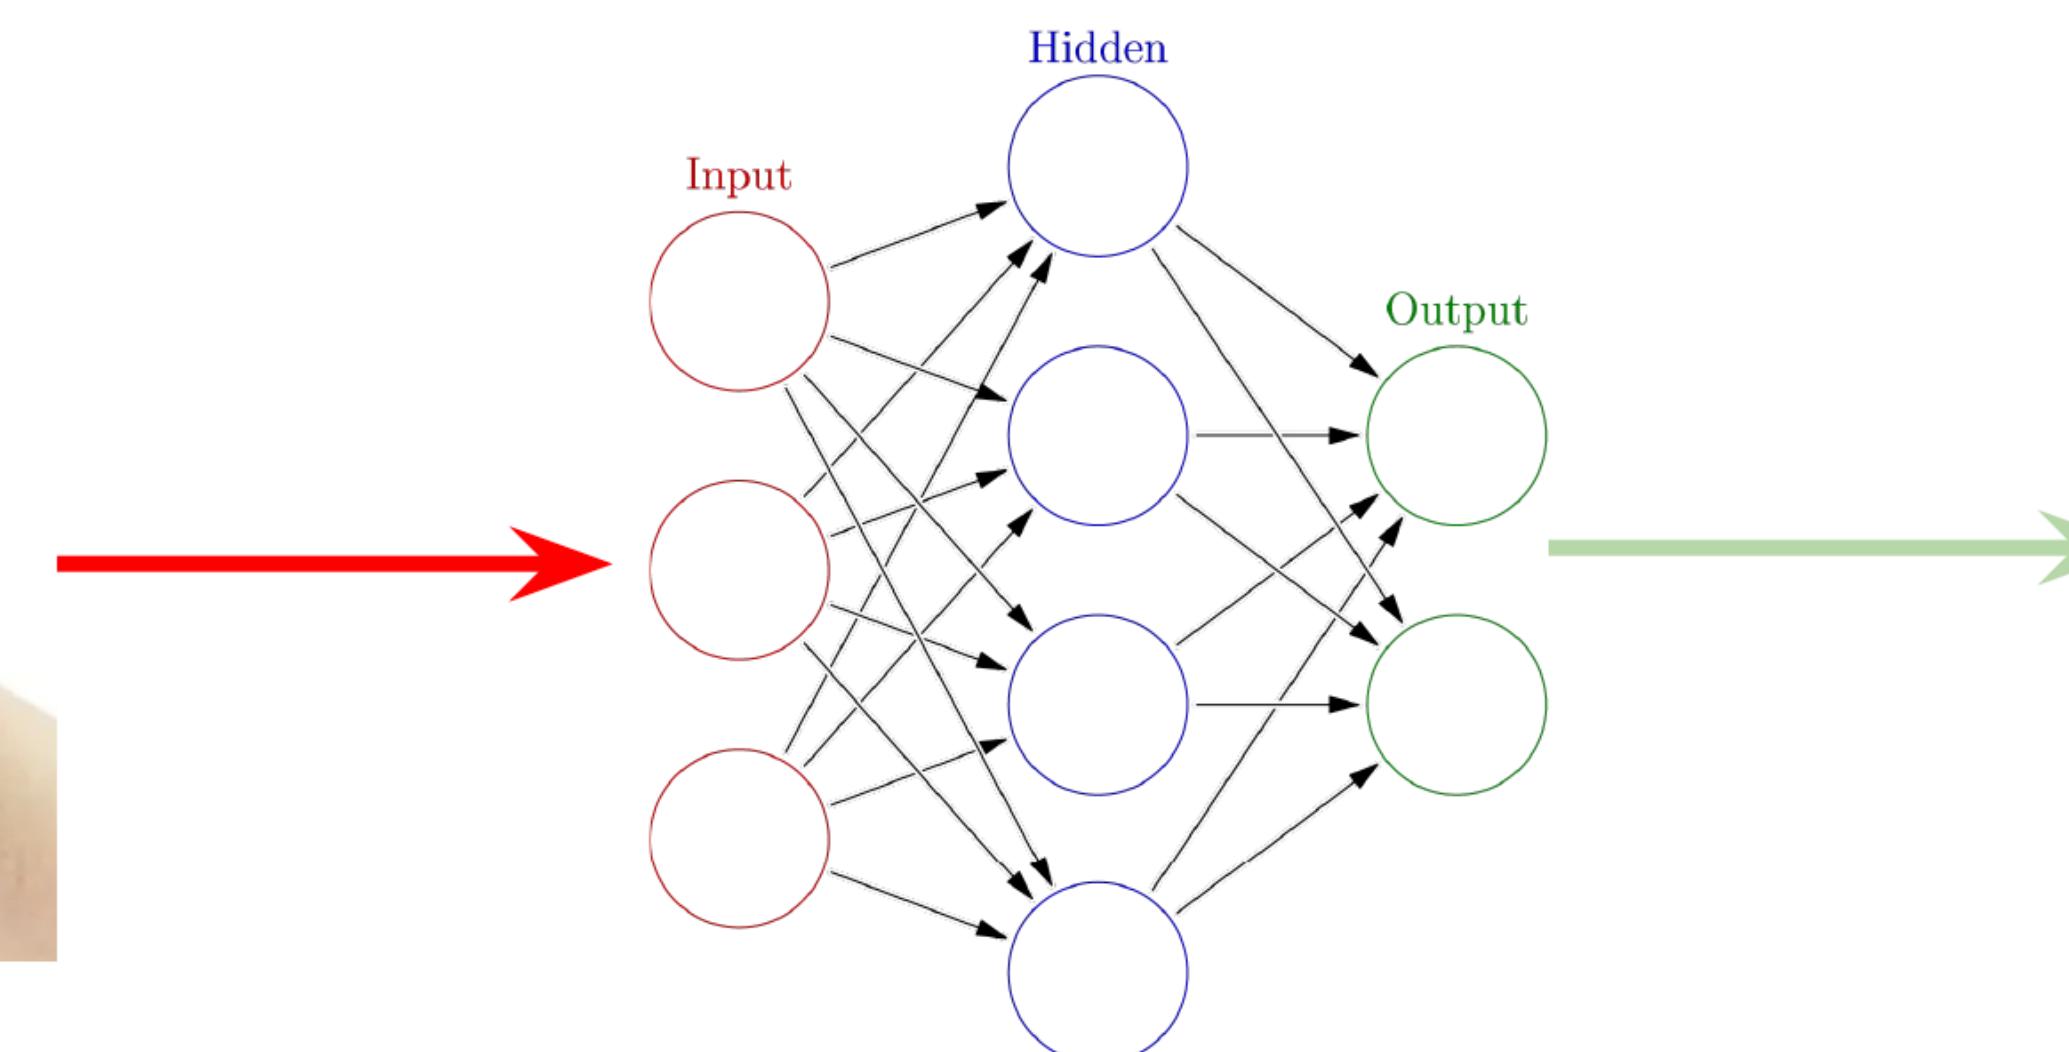
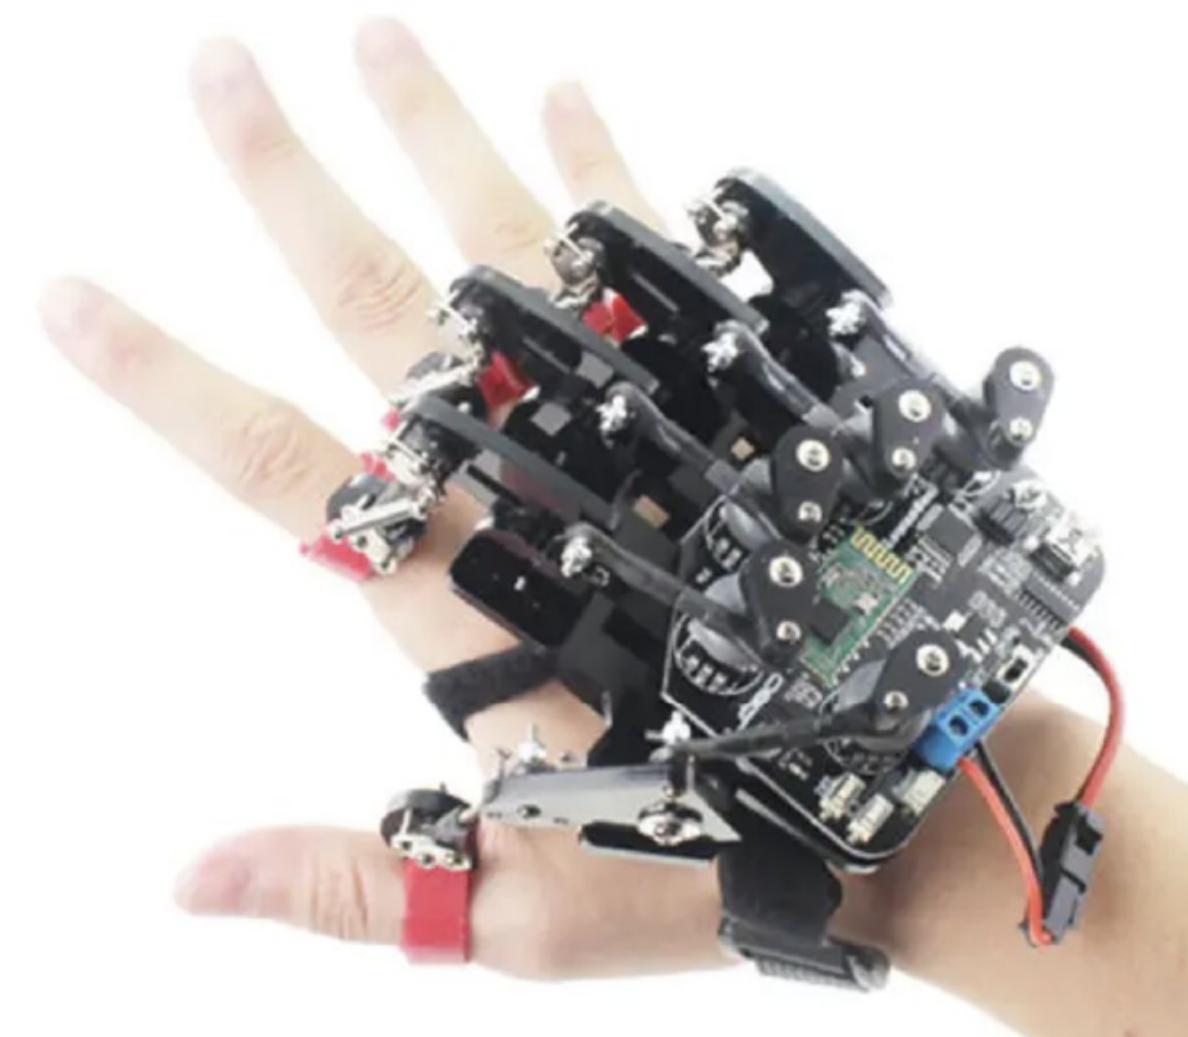


“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]

[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0]
[0.2, 0.7]

[0.5, 0.7]
[0.1, 0.8]
[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]

Neural Network Predicting with Regression



Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]

“b”: [0, 1, 1]

“c”: [1, 0, 1]

“d”: [0.2, 0, 1]

“e”: [0, 0, 0.5]

“f”: [0.8, 0.2, 1]



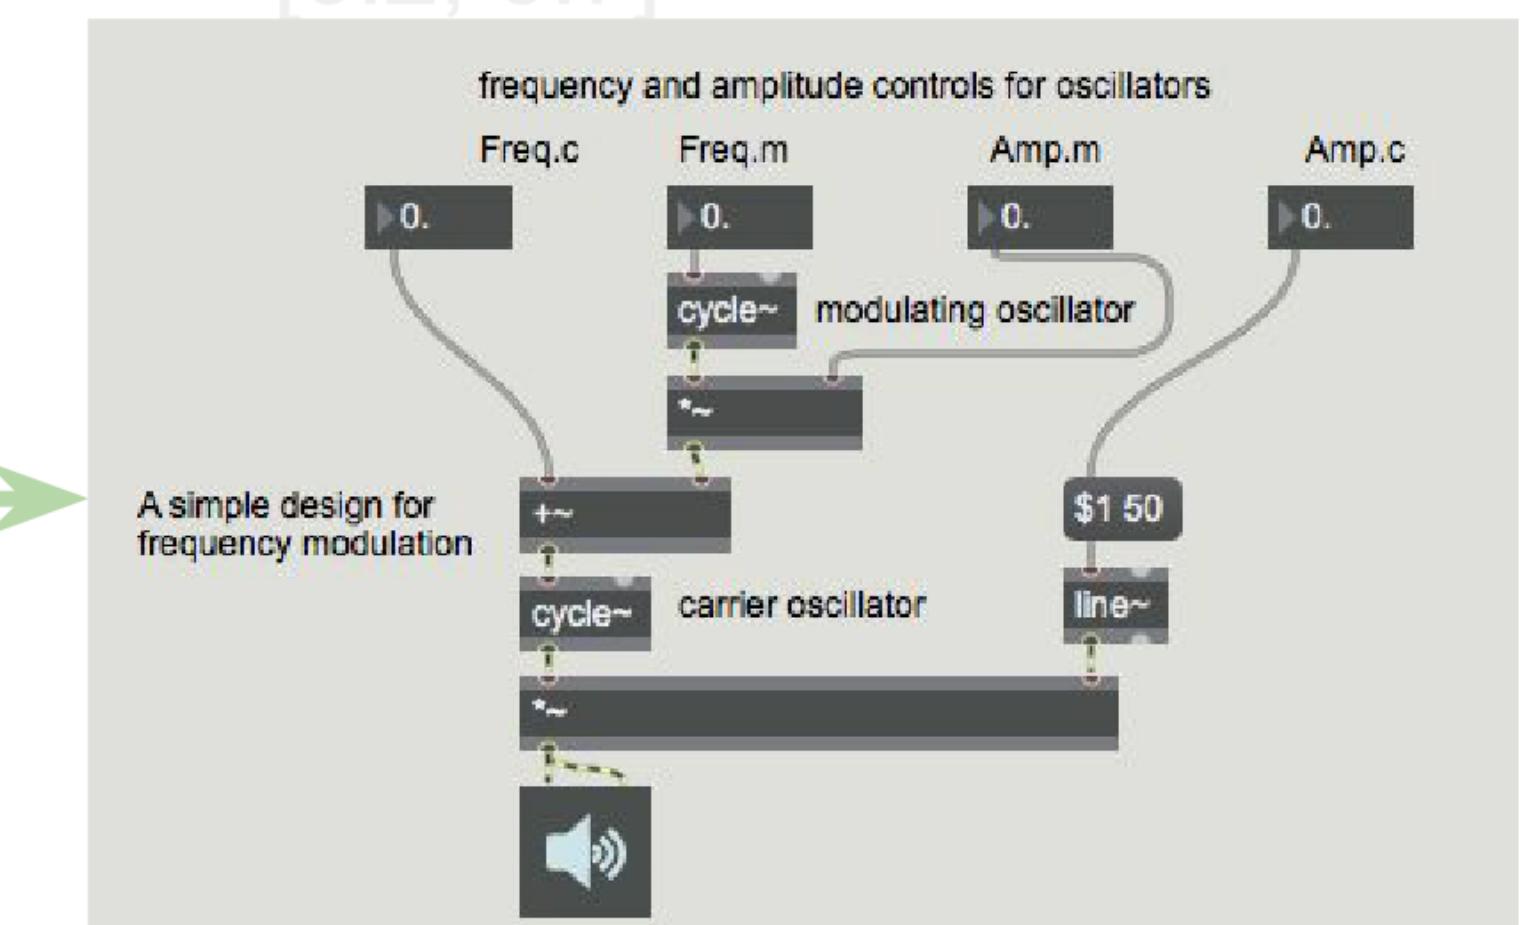
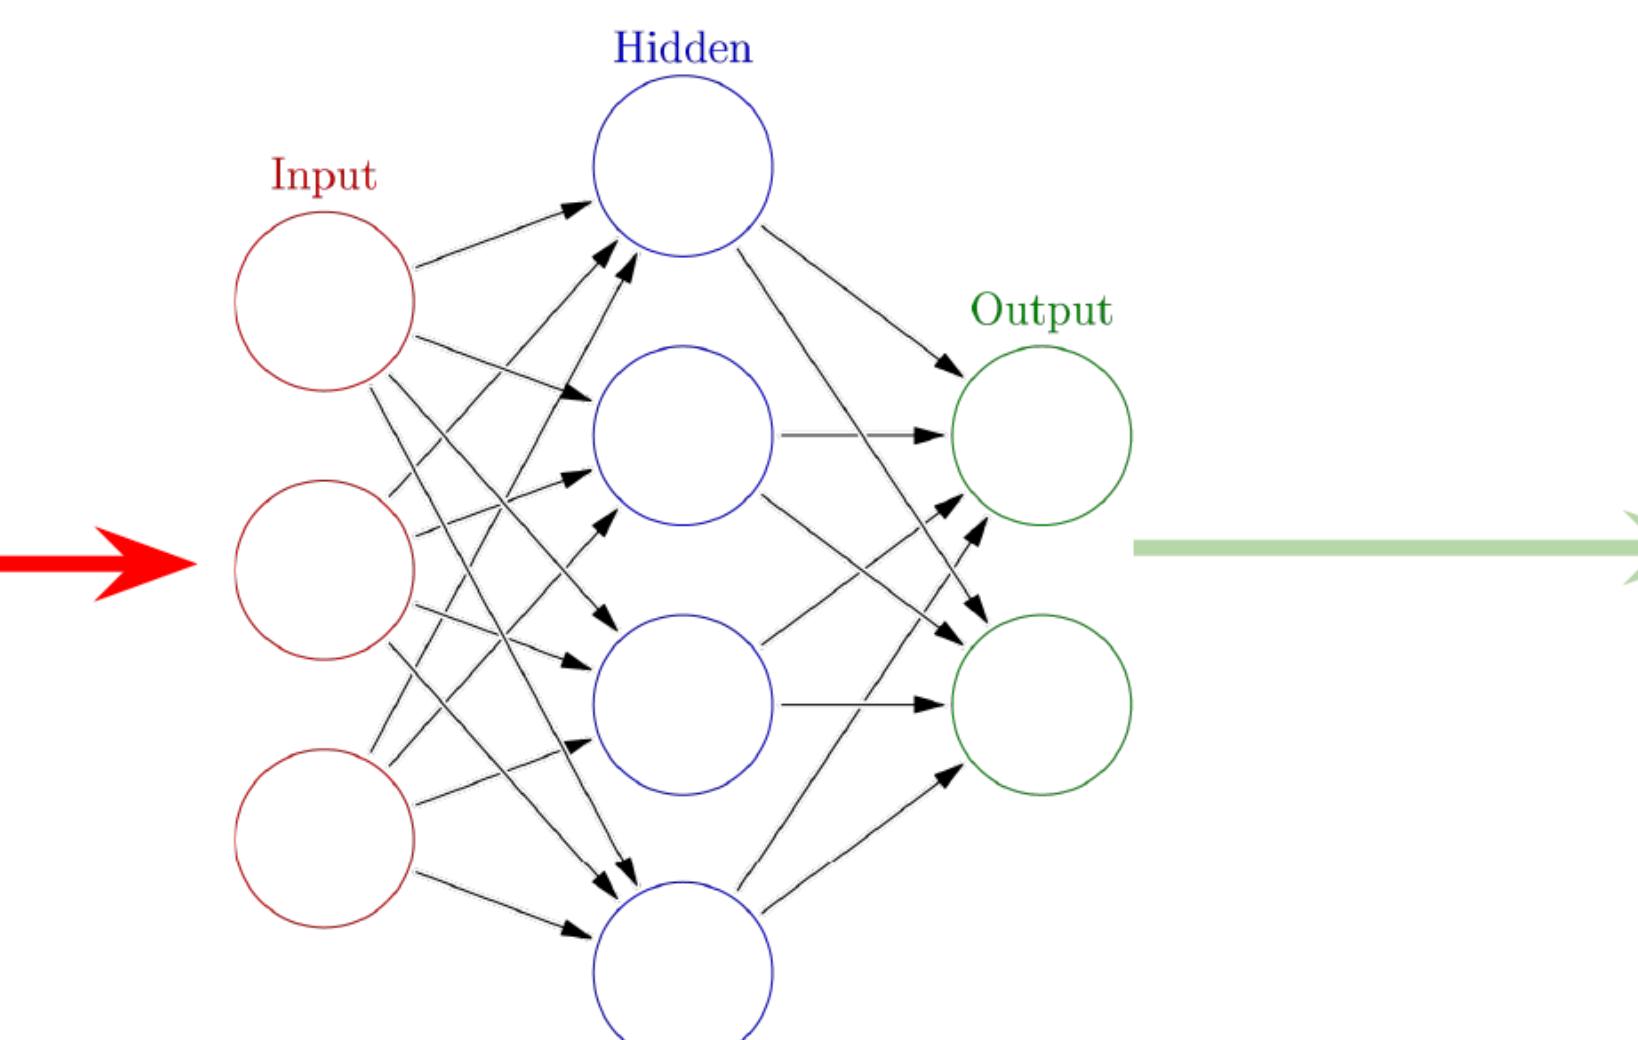
e.g.:

pitch

loudness

spectral centroid

spectral flatness



“g”: [0.1, 0.5, 1]

“h”: [0, 1, 1]

“i”: [1, 0, 1]

“j”: [0.2, 0, 1]

“k”: [0, 0, 0.5]

“l”: [0.8, 0.2, 1]

[0.5, 0.7]

[0.1, 0.1]

[0.5, 0.6]

[0.3, 0.9]

[0.4, 0]

[0.2, 0.7]

[0.5, 0.7]

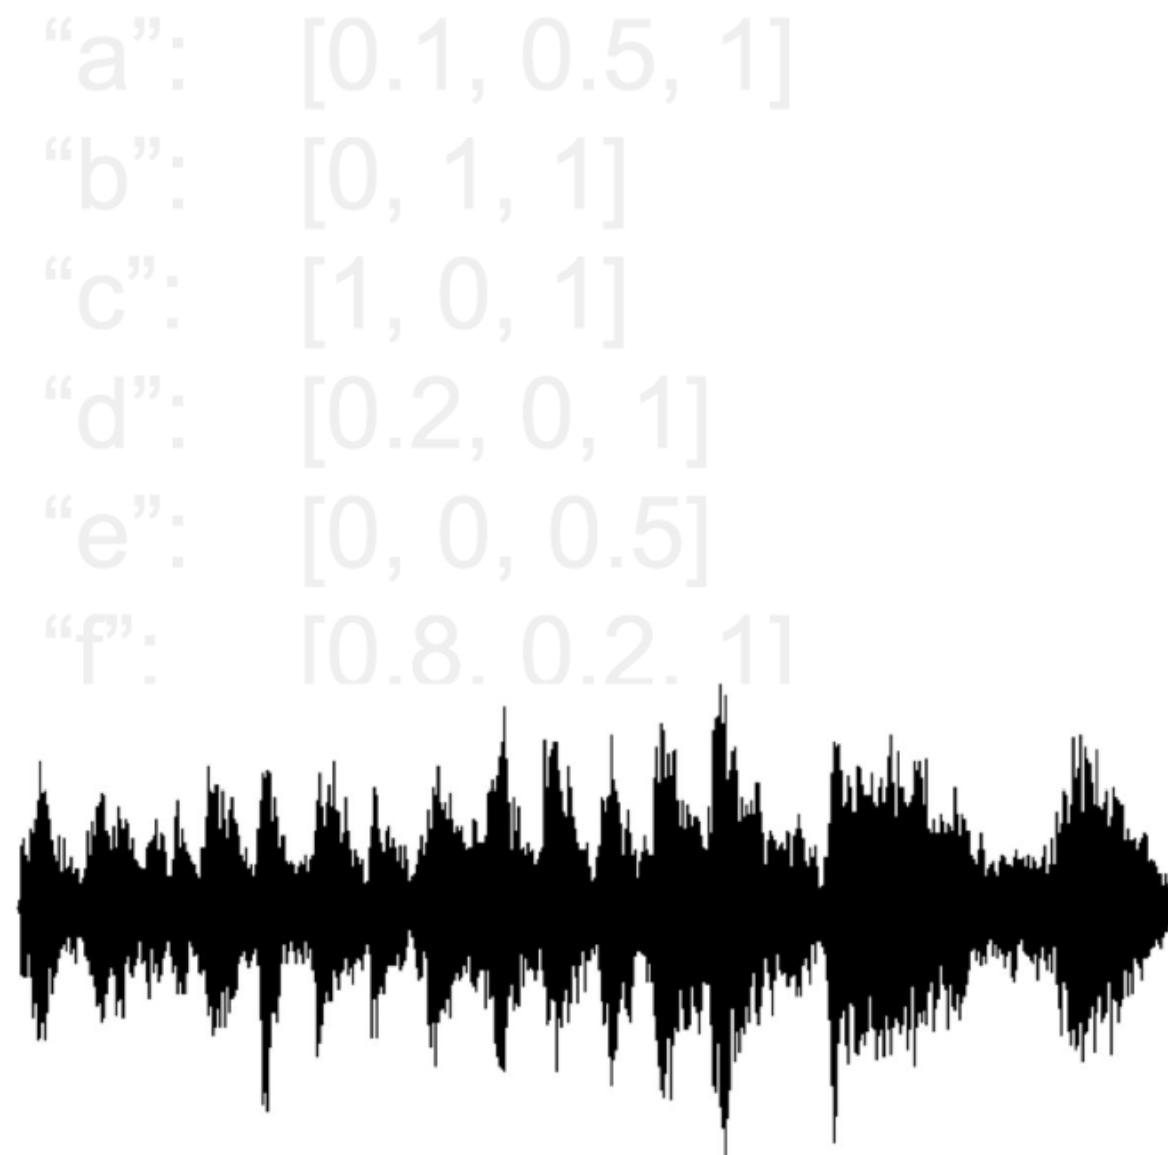
[0.1, 0.8]

[0.5, 0.1]

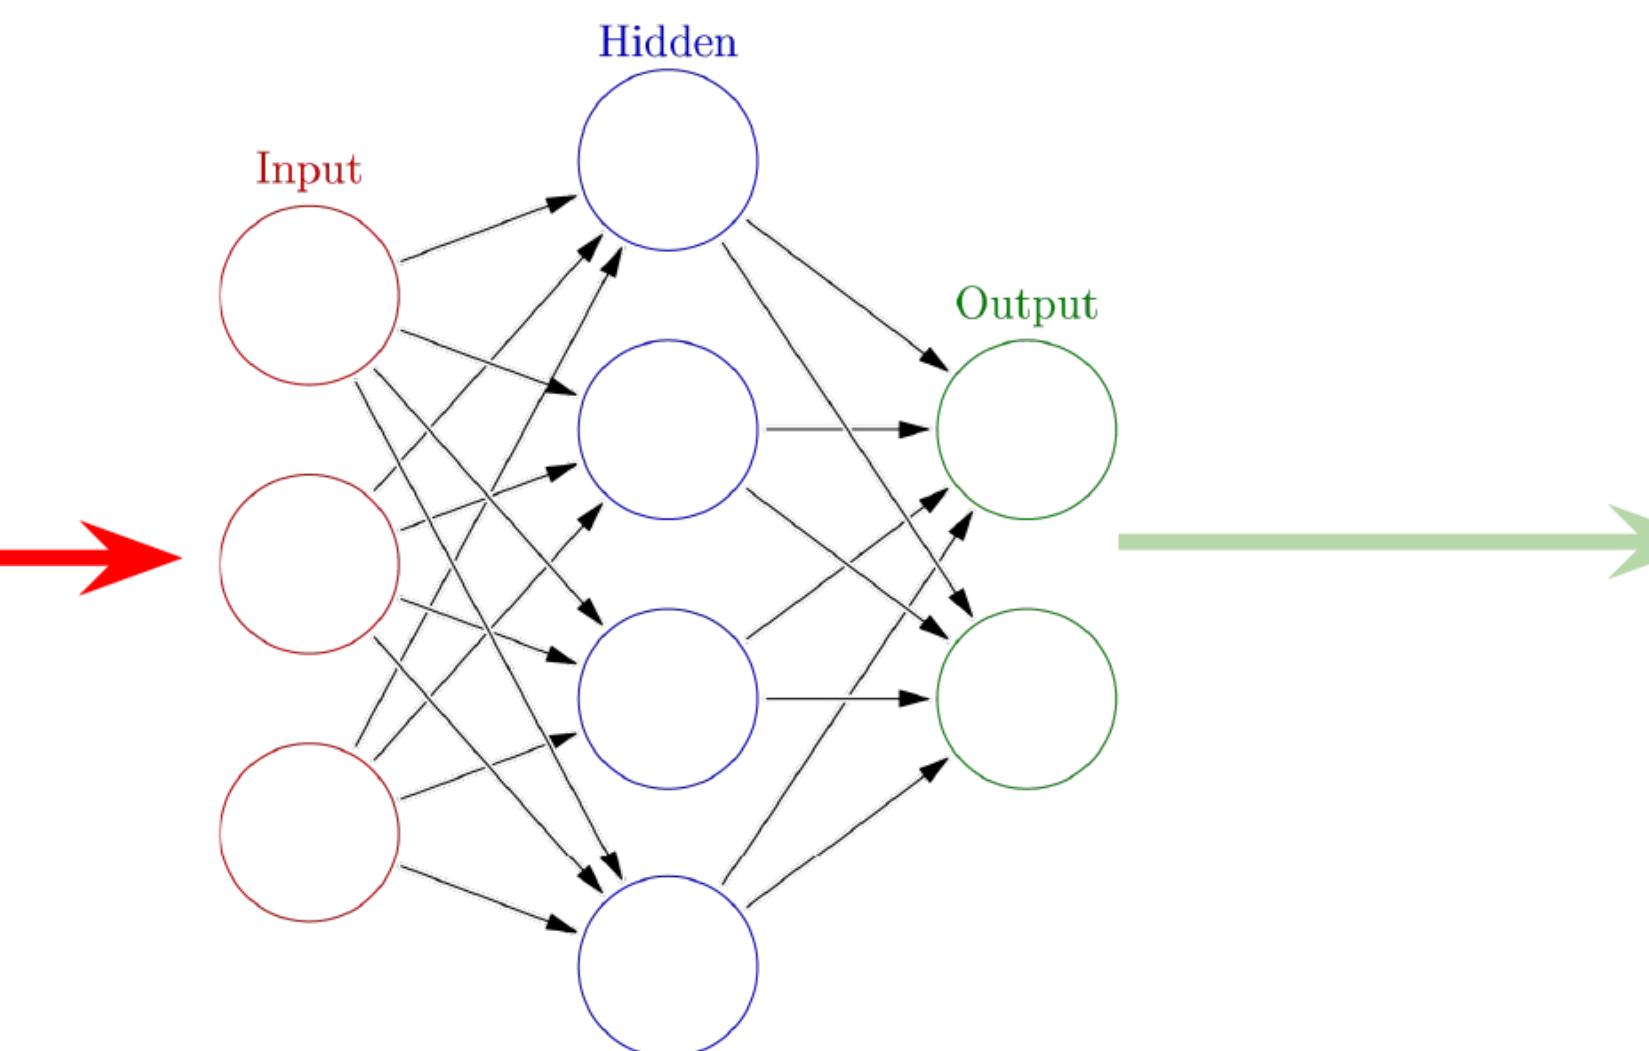
[0.4, 0.2]

[0.1, 0.8]

Neural Network Predicting with Regression



“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“j”: [0, 0, 0.5]
“k”: [0, 1, 1]
“l”: [0, 1, 1]
“m”: [0, 0, 1]
“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]

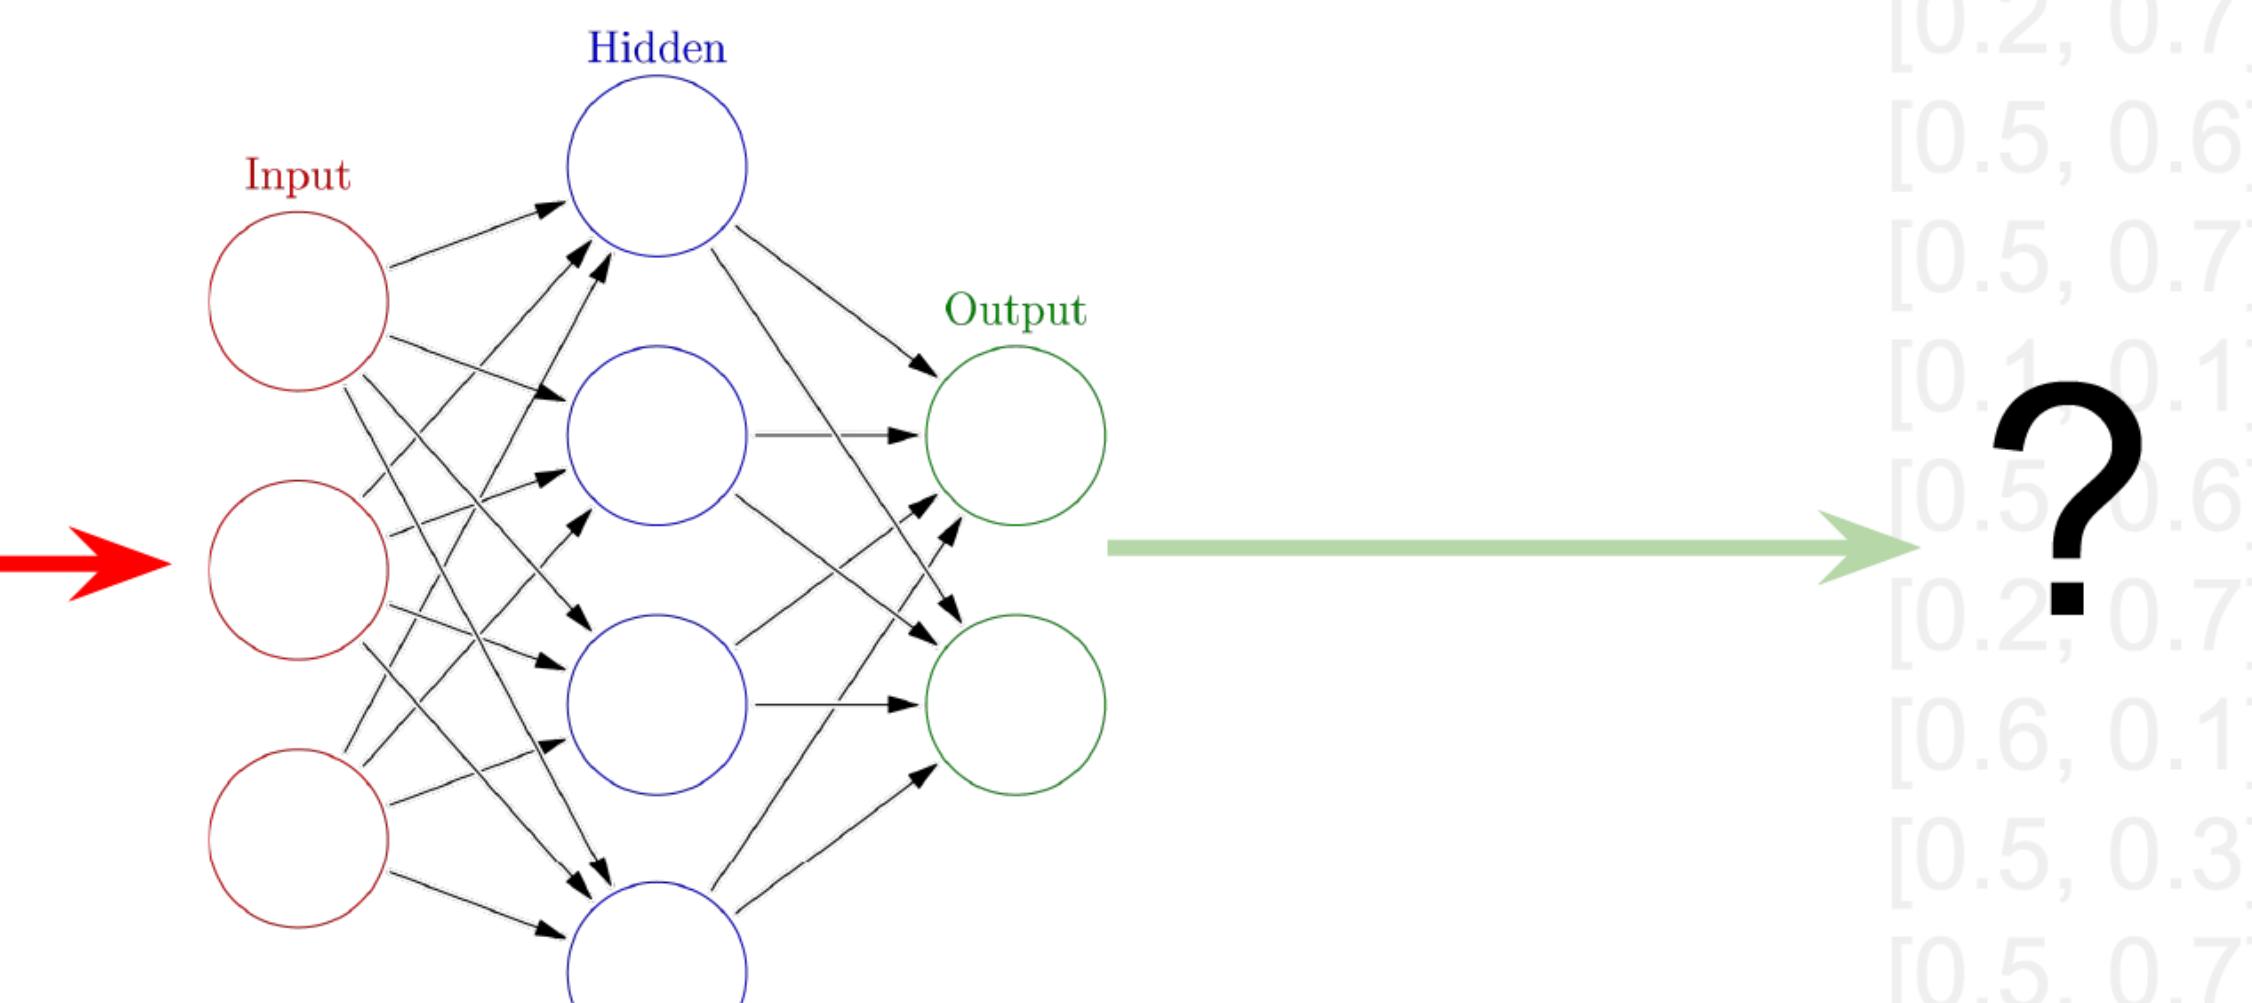


[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]

[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]

Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]
“b”: [0, 1, 1]
“c”: [1, 0, 1]
“d”: [0.2, 0, 1]
“e”: [0, 0, 0.5]
“f”: [0.8, 0.2, 1]
“g”: [0.9, 0.3, 0]
“h”: [0.4, 1, 0.7]
“i”: [0.2, 0.3, 1]
“j”: [0.1, 0.1, 0.1]
“k”: [0, 1, 1]
“l”: [0, 1, 1]
“m”: [0, 0, 1]
“n”: [0.4, 1, 0.7]
“o”: [0.2, 0.3, 1]
“p”: [0.1, 1, 0.1]
“q”: [0, 1, 1]
“r”: [0, 1, 1]



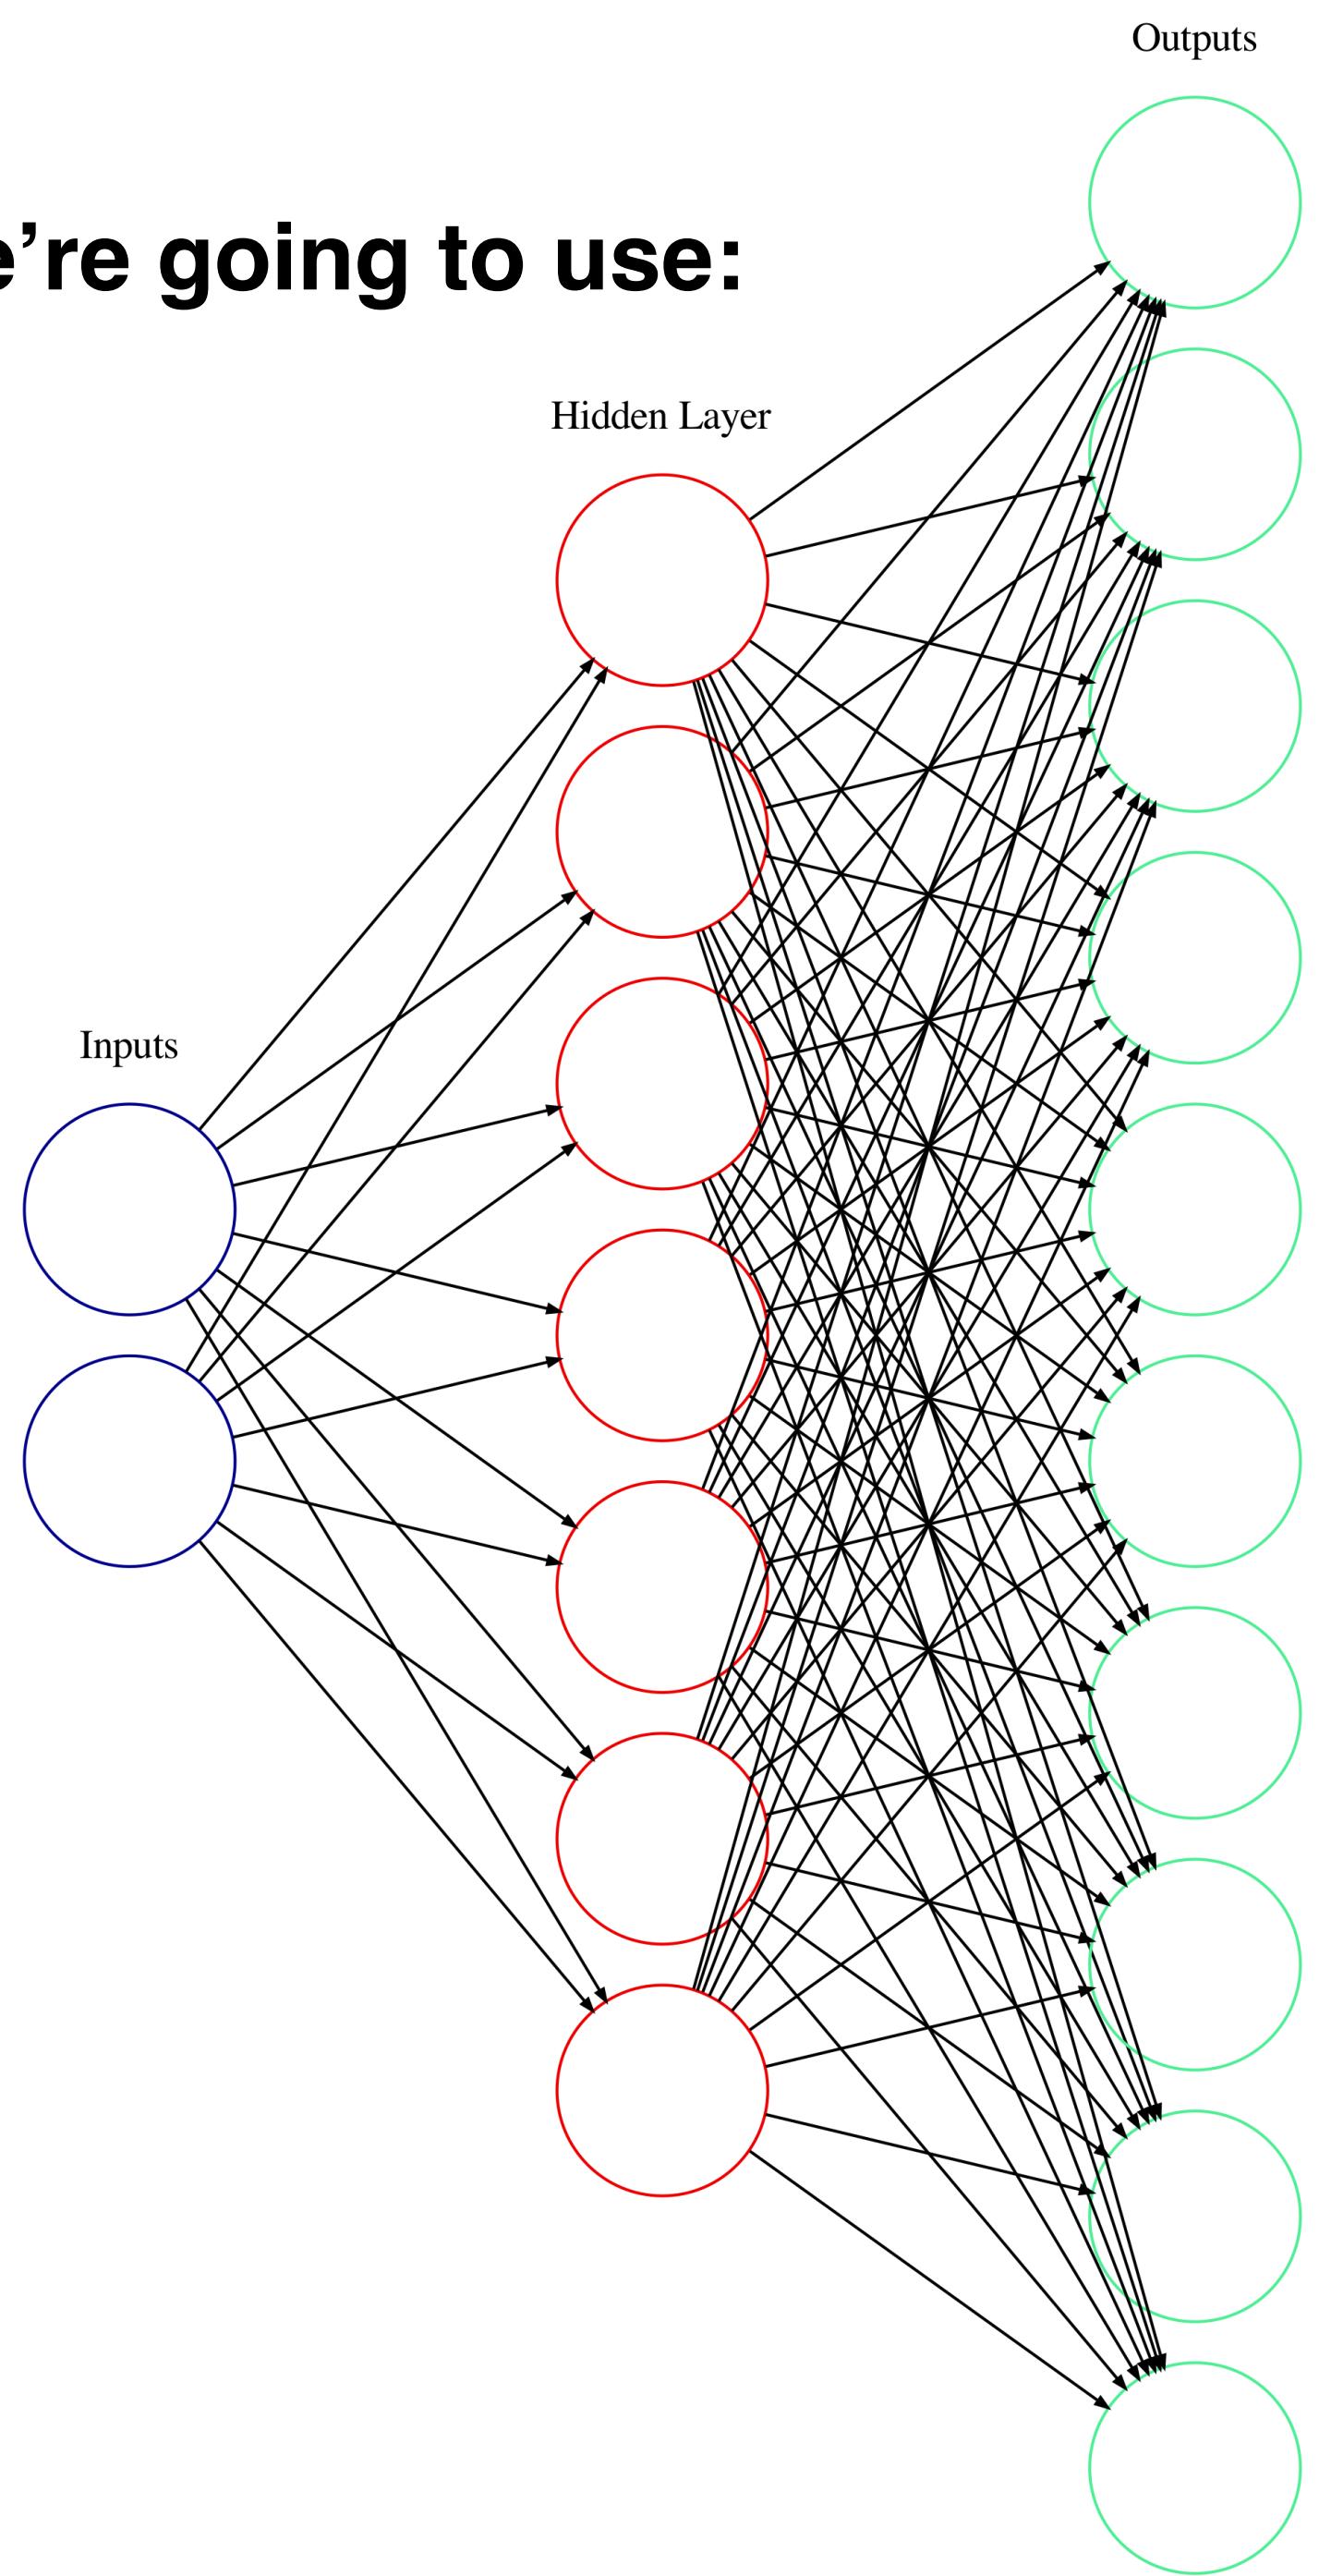
[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.3, 0.9]
[0.4, 0]
[0.2, 0.7]
[0.5, 0.6]
[0.5, 0.7]
[0.1, 0.1]
[0.5, 0.6]
[0.2, 0.7]
[0.6, 0.1]
[0.5, 0.3]
[0.5, 0.7]
[0.1, 0.8]
[0.5, 0.1]
[0.4, 0.2]
[0.1, 0.8]

Structure of the neural network we're going to use:

2 inputs

1 hidden layer of 7 nodes

10 outputs



fluid.mlprogression



There are certain things that we care about, as musicians for example, that are really hard to articulate in code. It's hard for me to talk about what kind of quality of sound I want and then translate that into a set of filter coefficients. It's hard for me to talk about how I want a performer to move on stage and then translate that into some sort of mathematical equation for their trajectory. But it's a lot easier for me to either find examples of sounds that have a particular quality or to give examples of movements or if I'm using other types of modalities, often curating or creating examples are just way easier for us as people. And this relates to the types of tacit knowledge and embodied knowledge we bring to creative practices.

-Rebecca Fiebrink

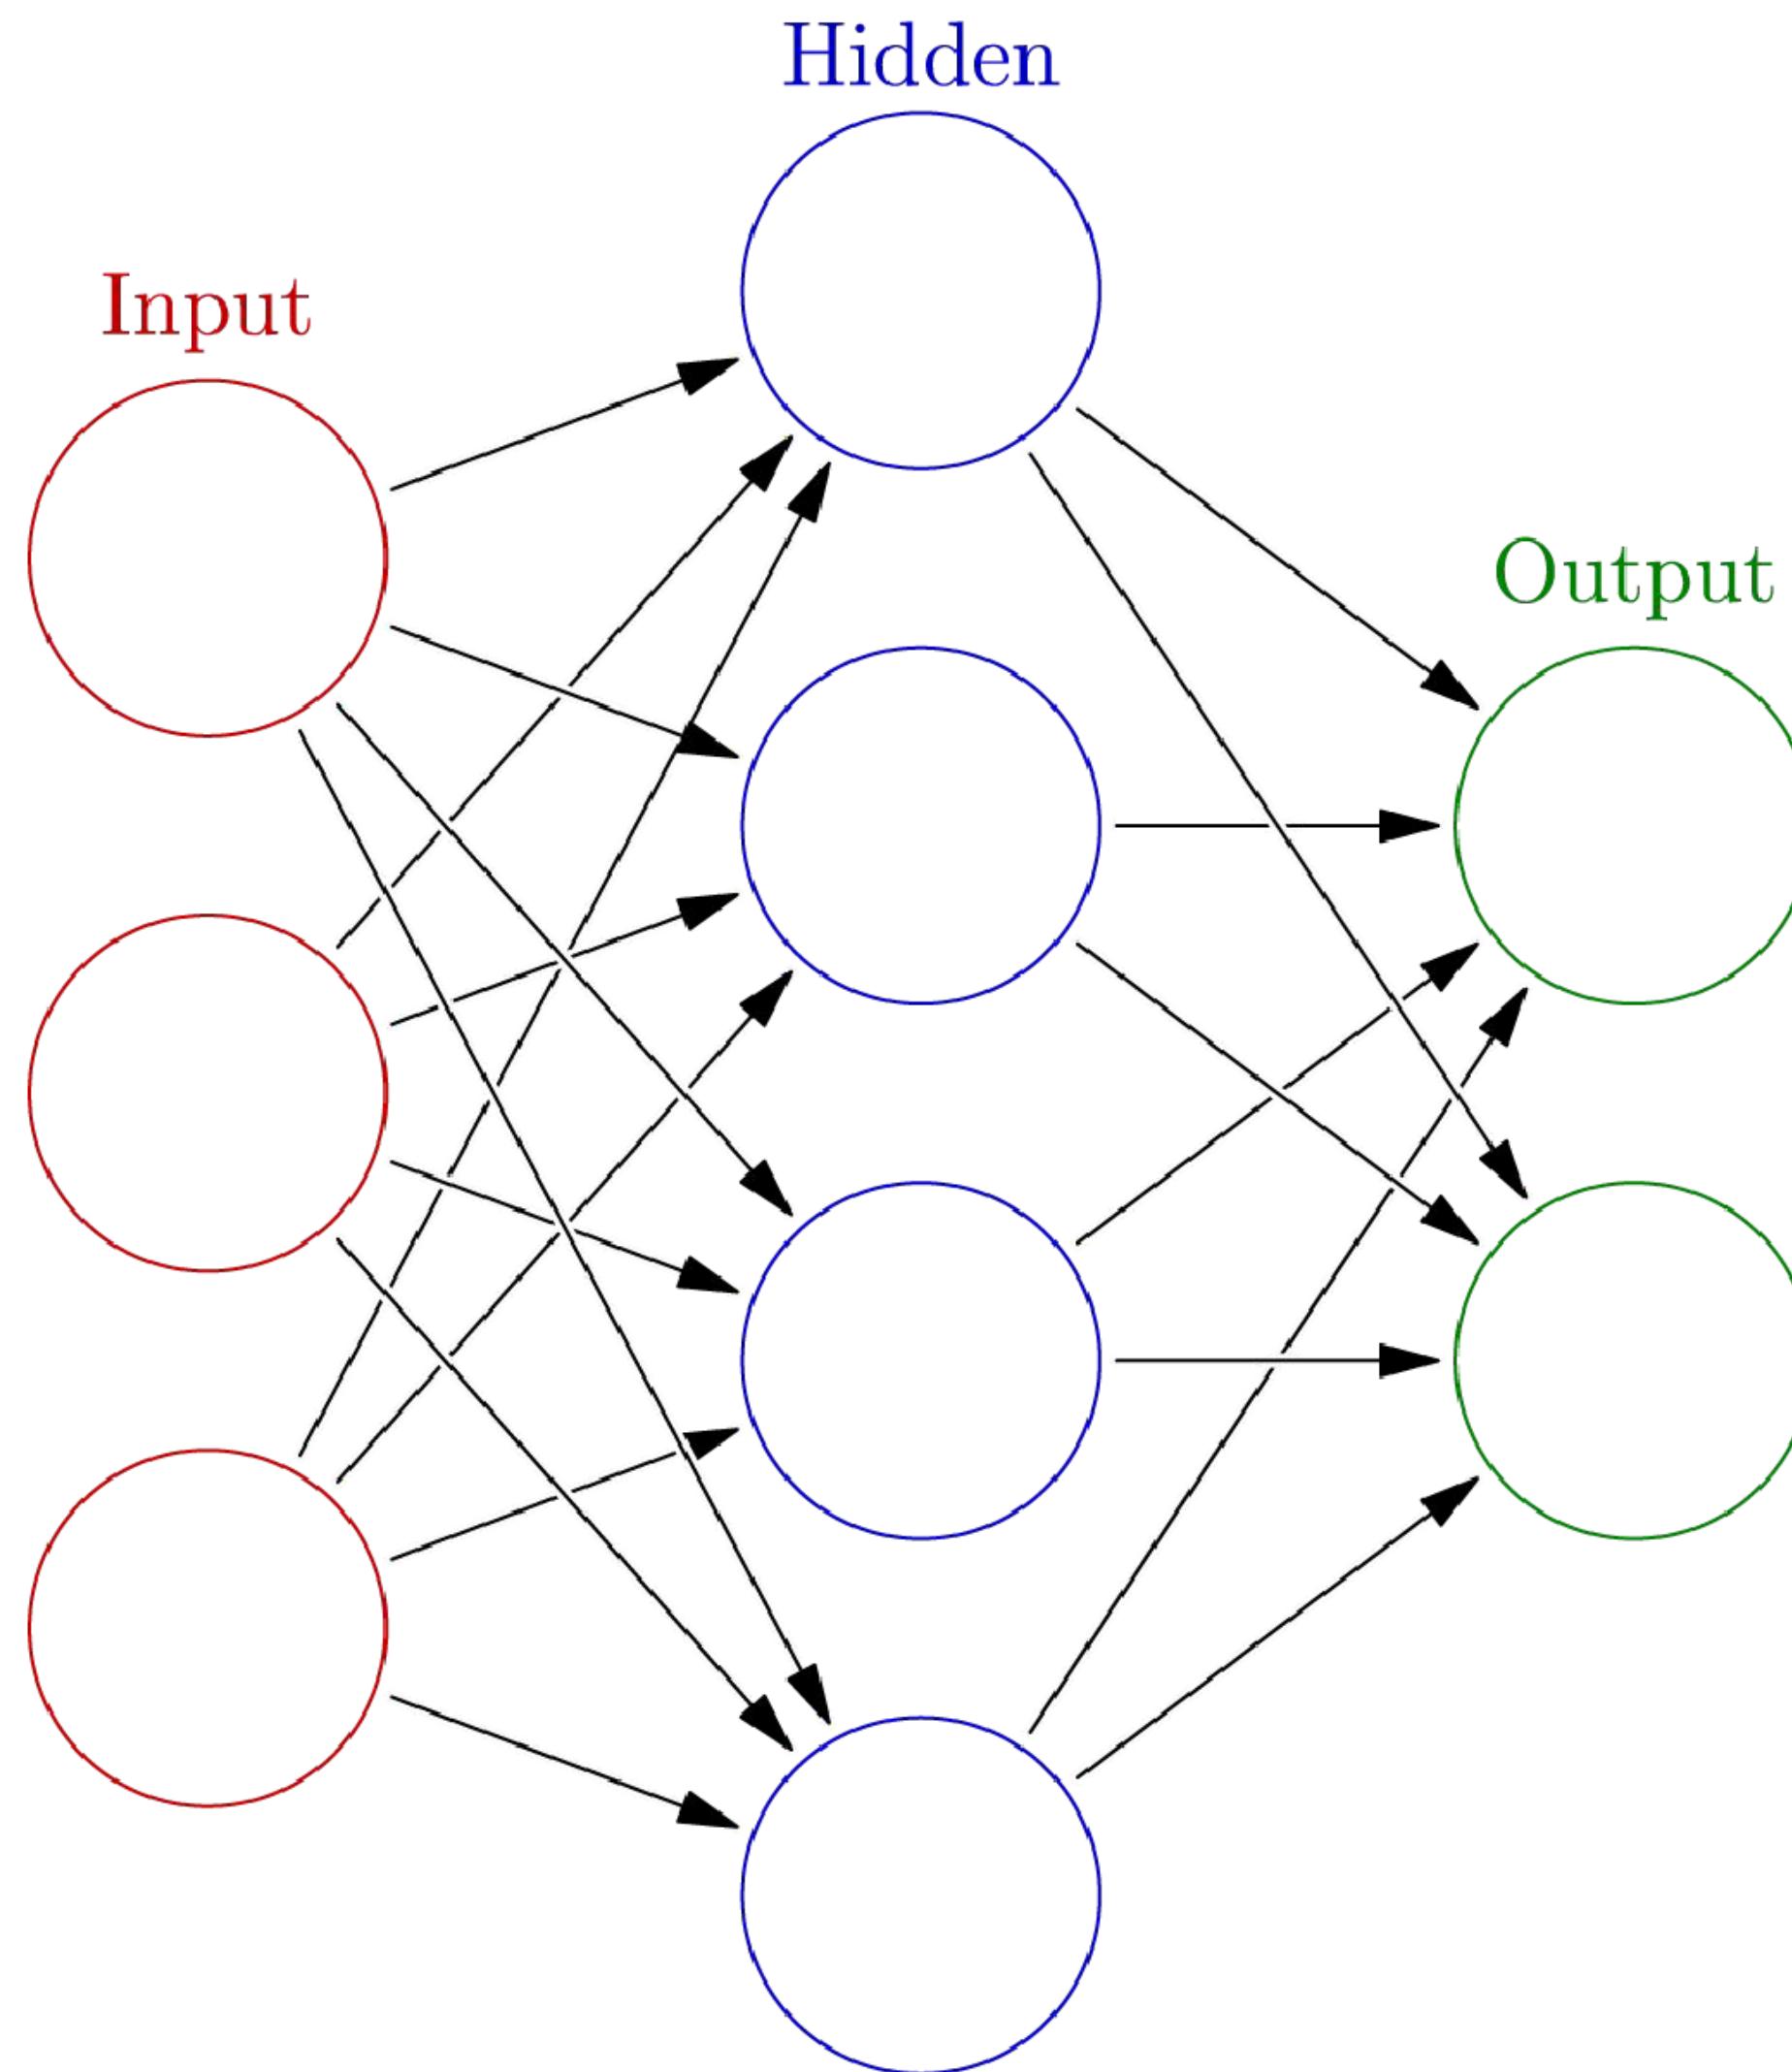
CeReNeM. 2019. Rebecca Fiebrink: Machine Learning as Creative Design Tool. Centre for Research in New Music. YouTube. Visited on 11/23/2020. <https://www.youtube.com/watch?v=Qo6n8MuEgdQ>.

FluidMLPClassifier

classify sounds by
timbre in real-time



Neural Network (Multi-Layer Perceptron)

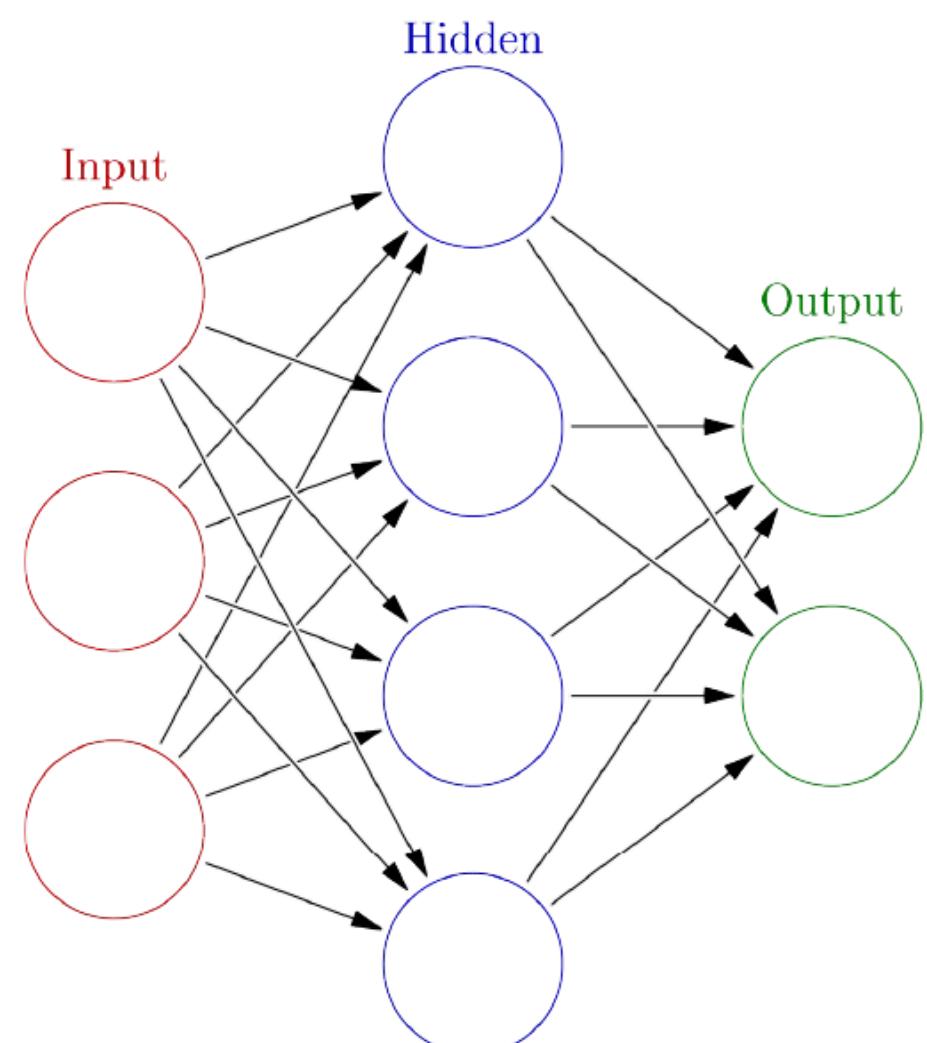


Classification:
a neural network predicts
**which category (or “class”)
an input belongs to**

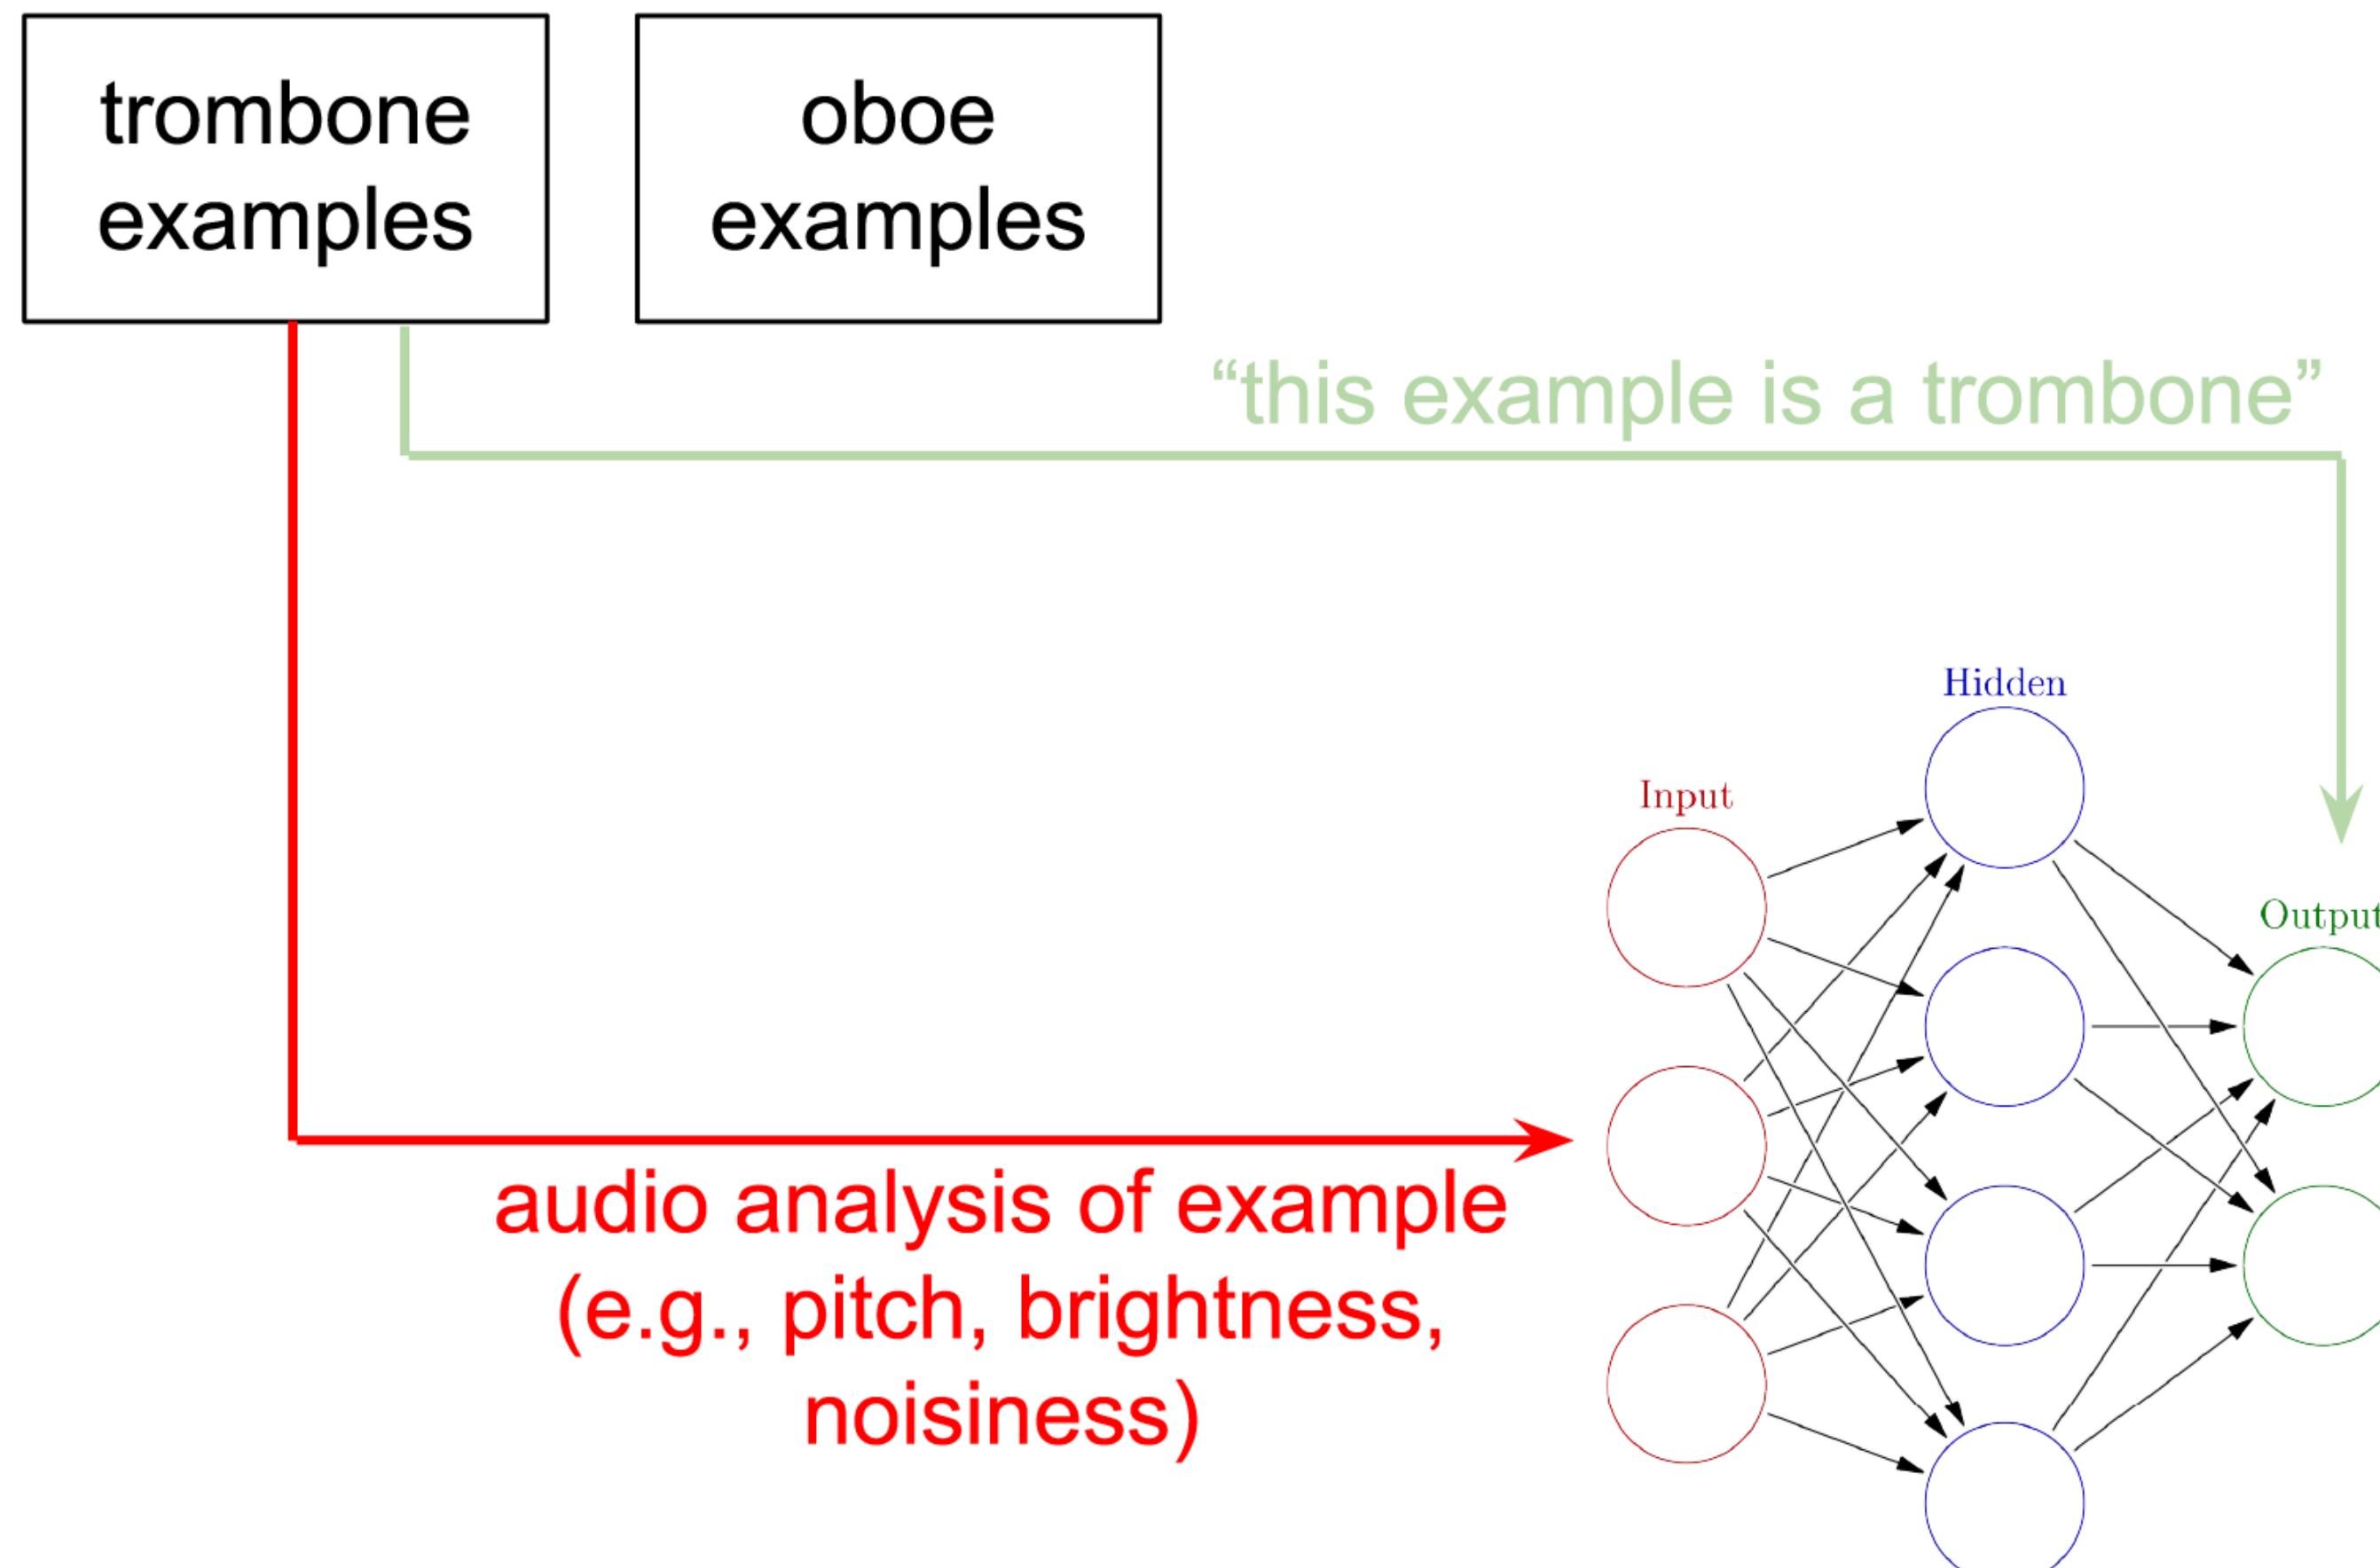
Neural Network ***Training a Classifier***

trombone
examples

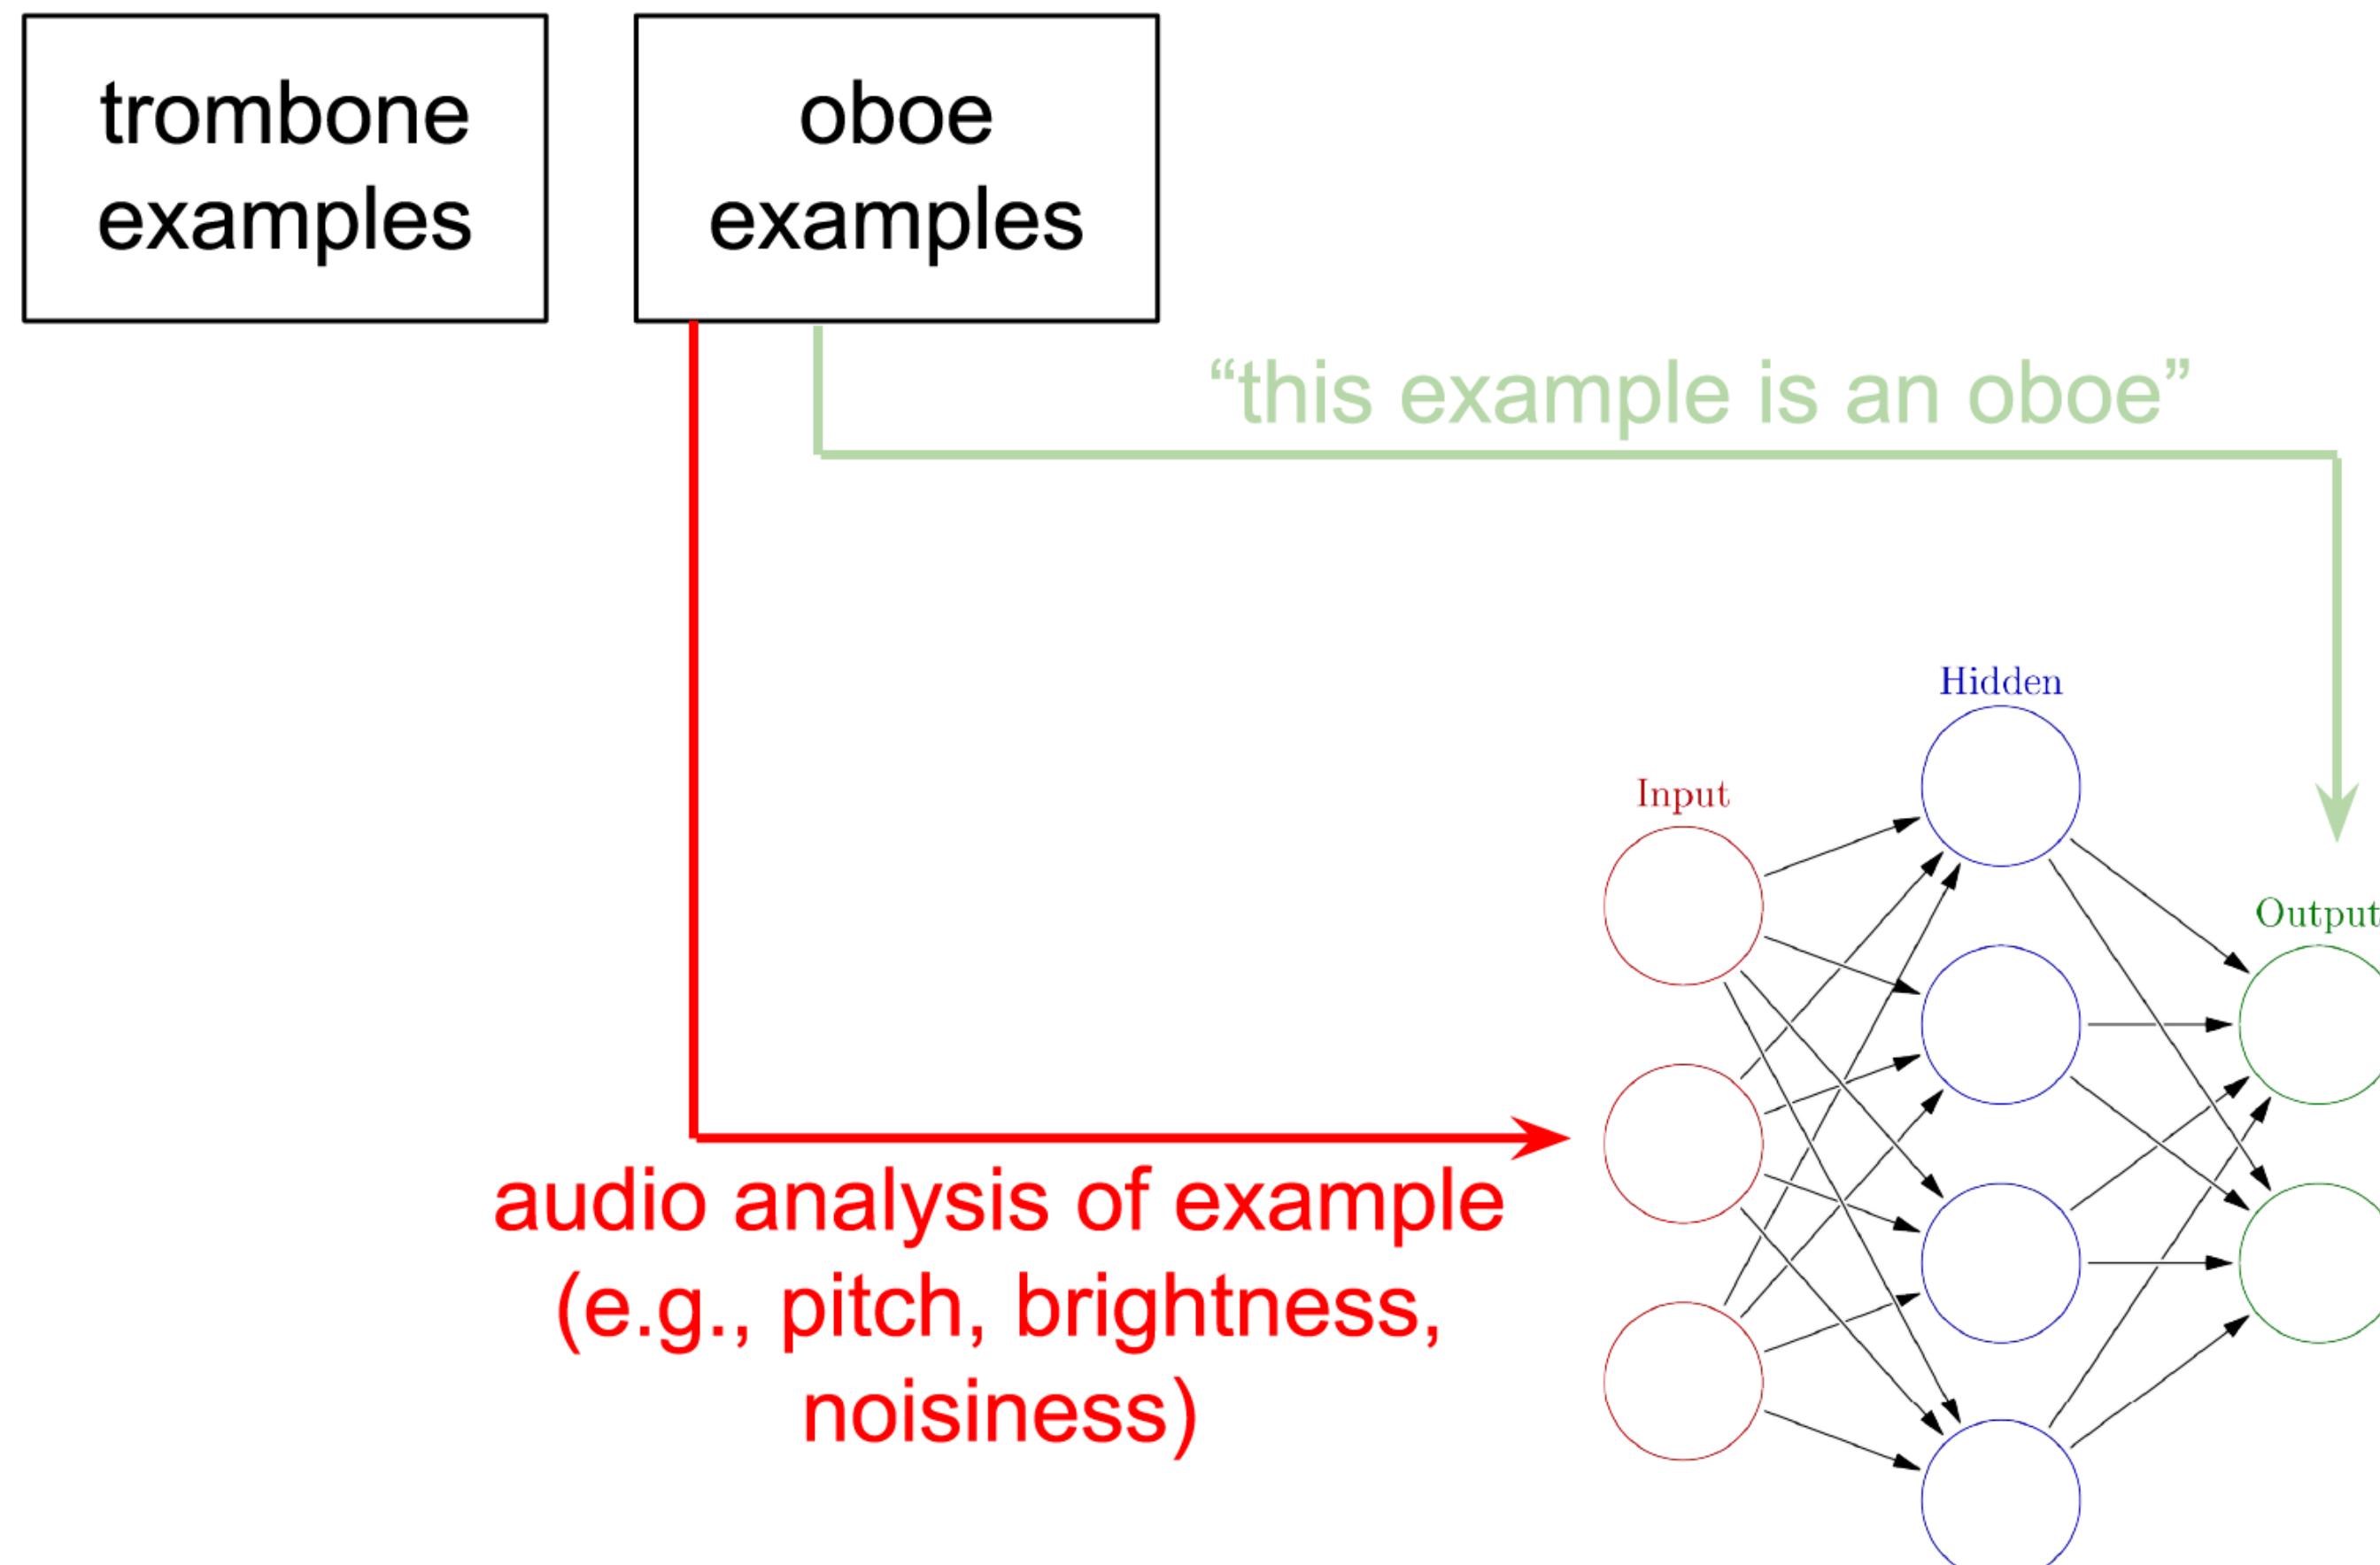
oboe
examples



Neural Network **Training a Classifier**



Neural Network **Training a Classifier**



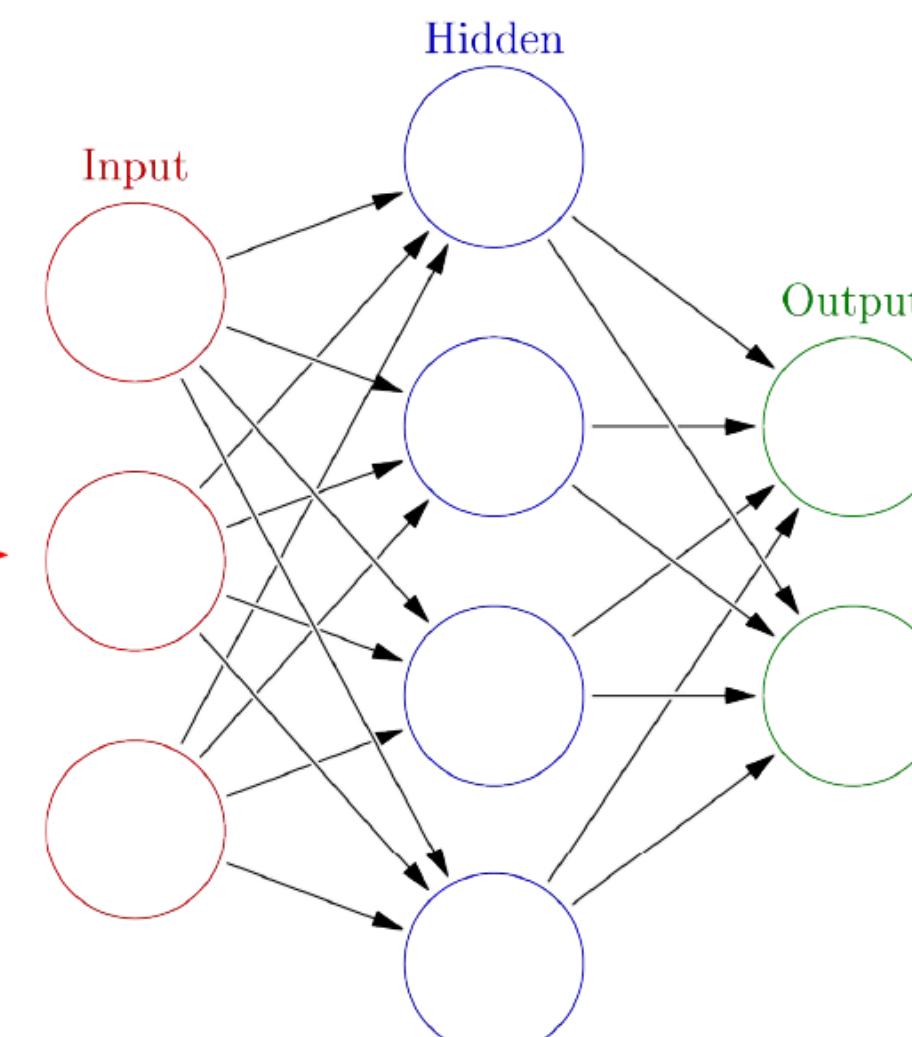
Neural Network Predicting a Classification

trombone
examples

oboe
examples

new example it has
never seen before

audio analysis of example
(e.g., pitch, brightness,
noisiness)



“this new example is most
like the oboe examples you
showed me before”
(or trombone...)

fluid.mlpclassifier

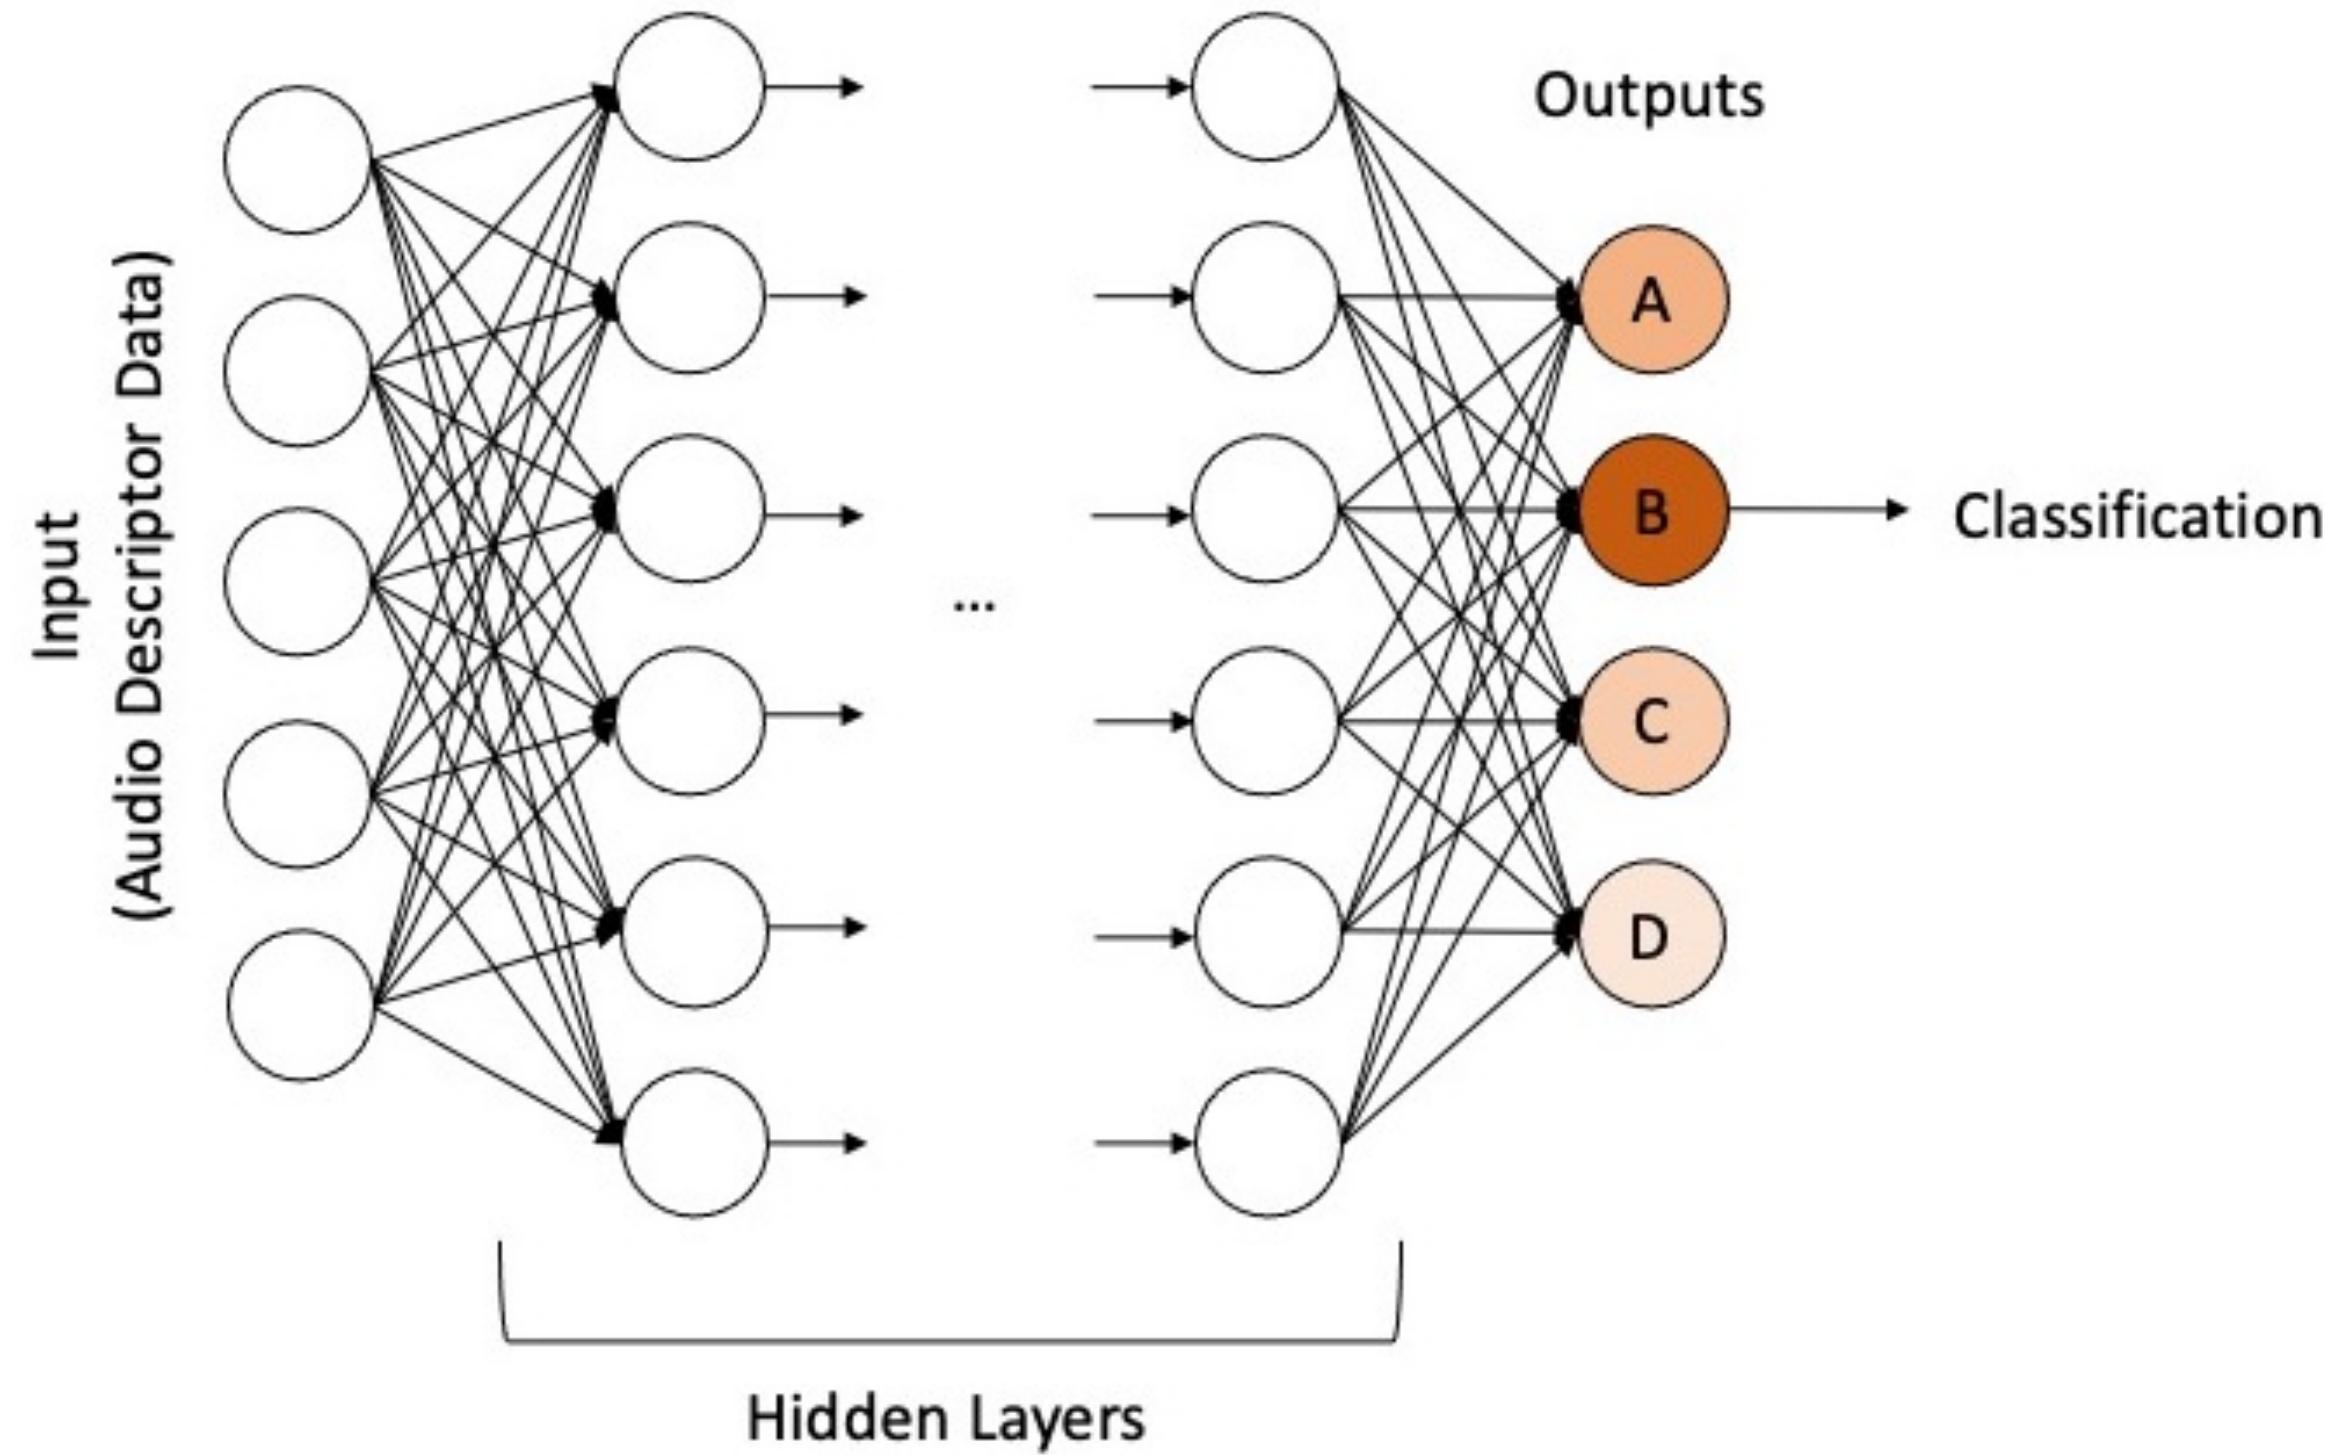


fluid.mlpclassifier

validation



- **overfitting**
 - excellent on training data, poor on anything else
- **underfitting**
 - poor on everything...keep training!



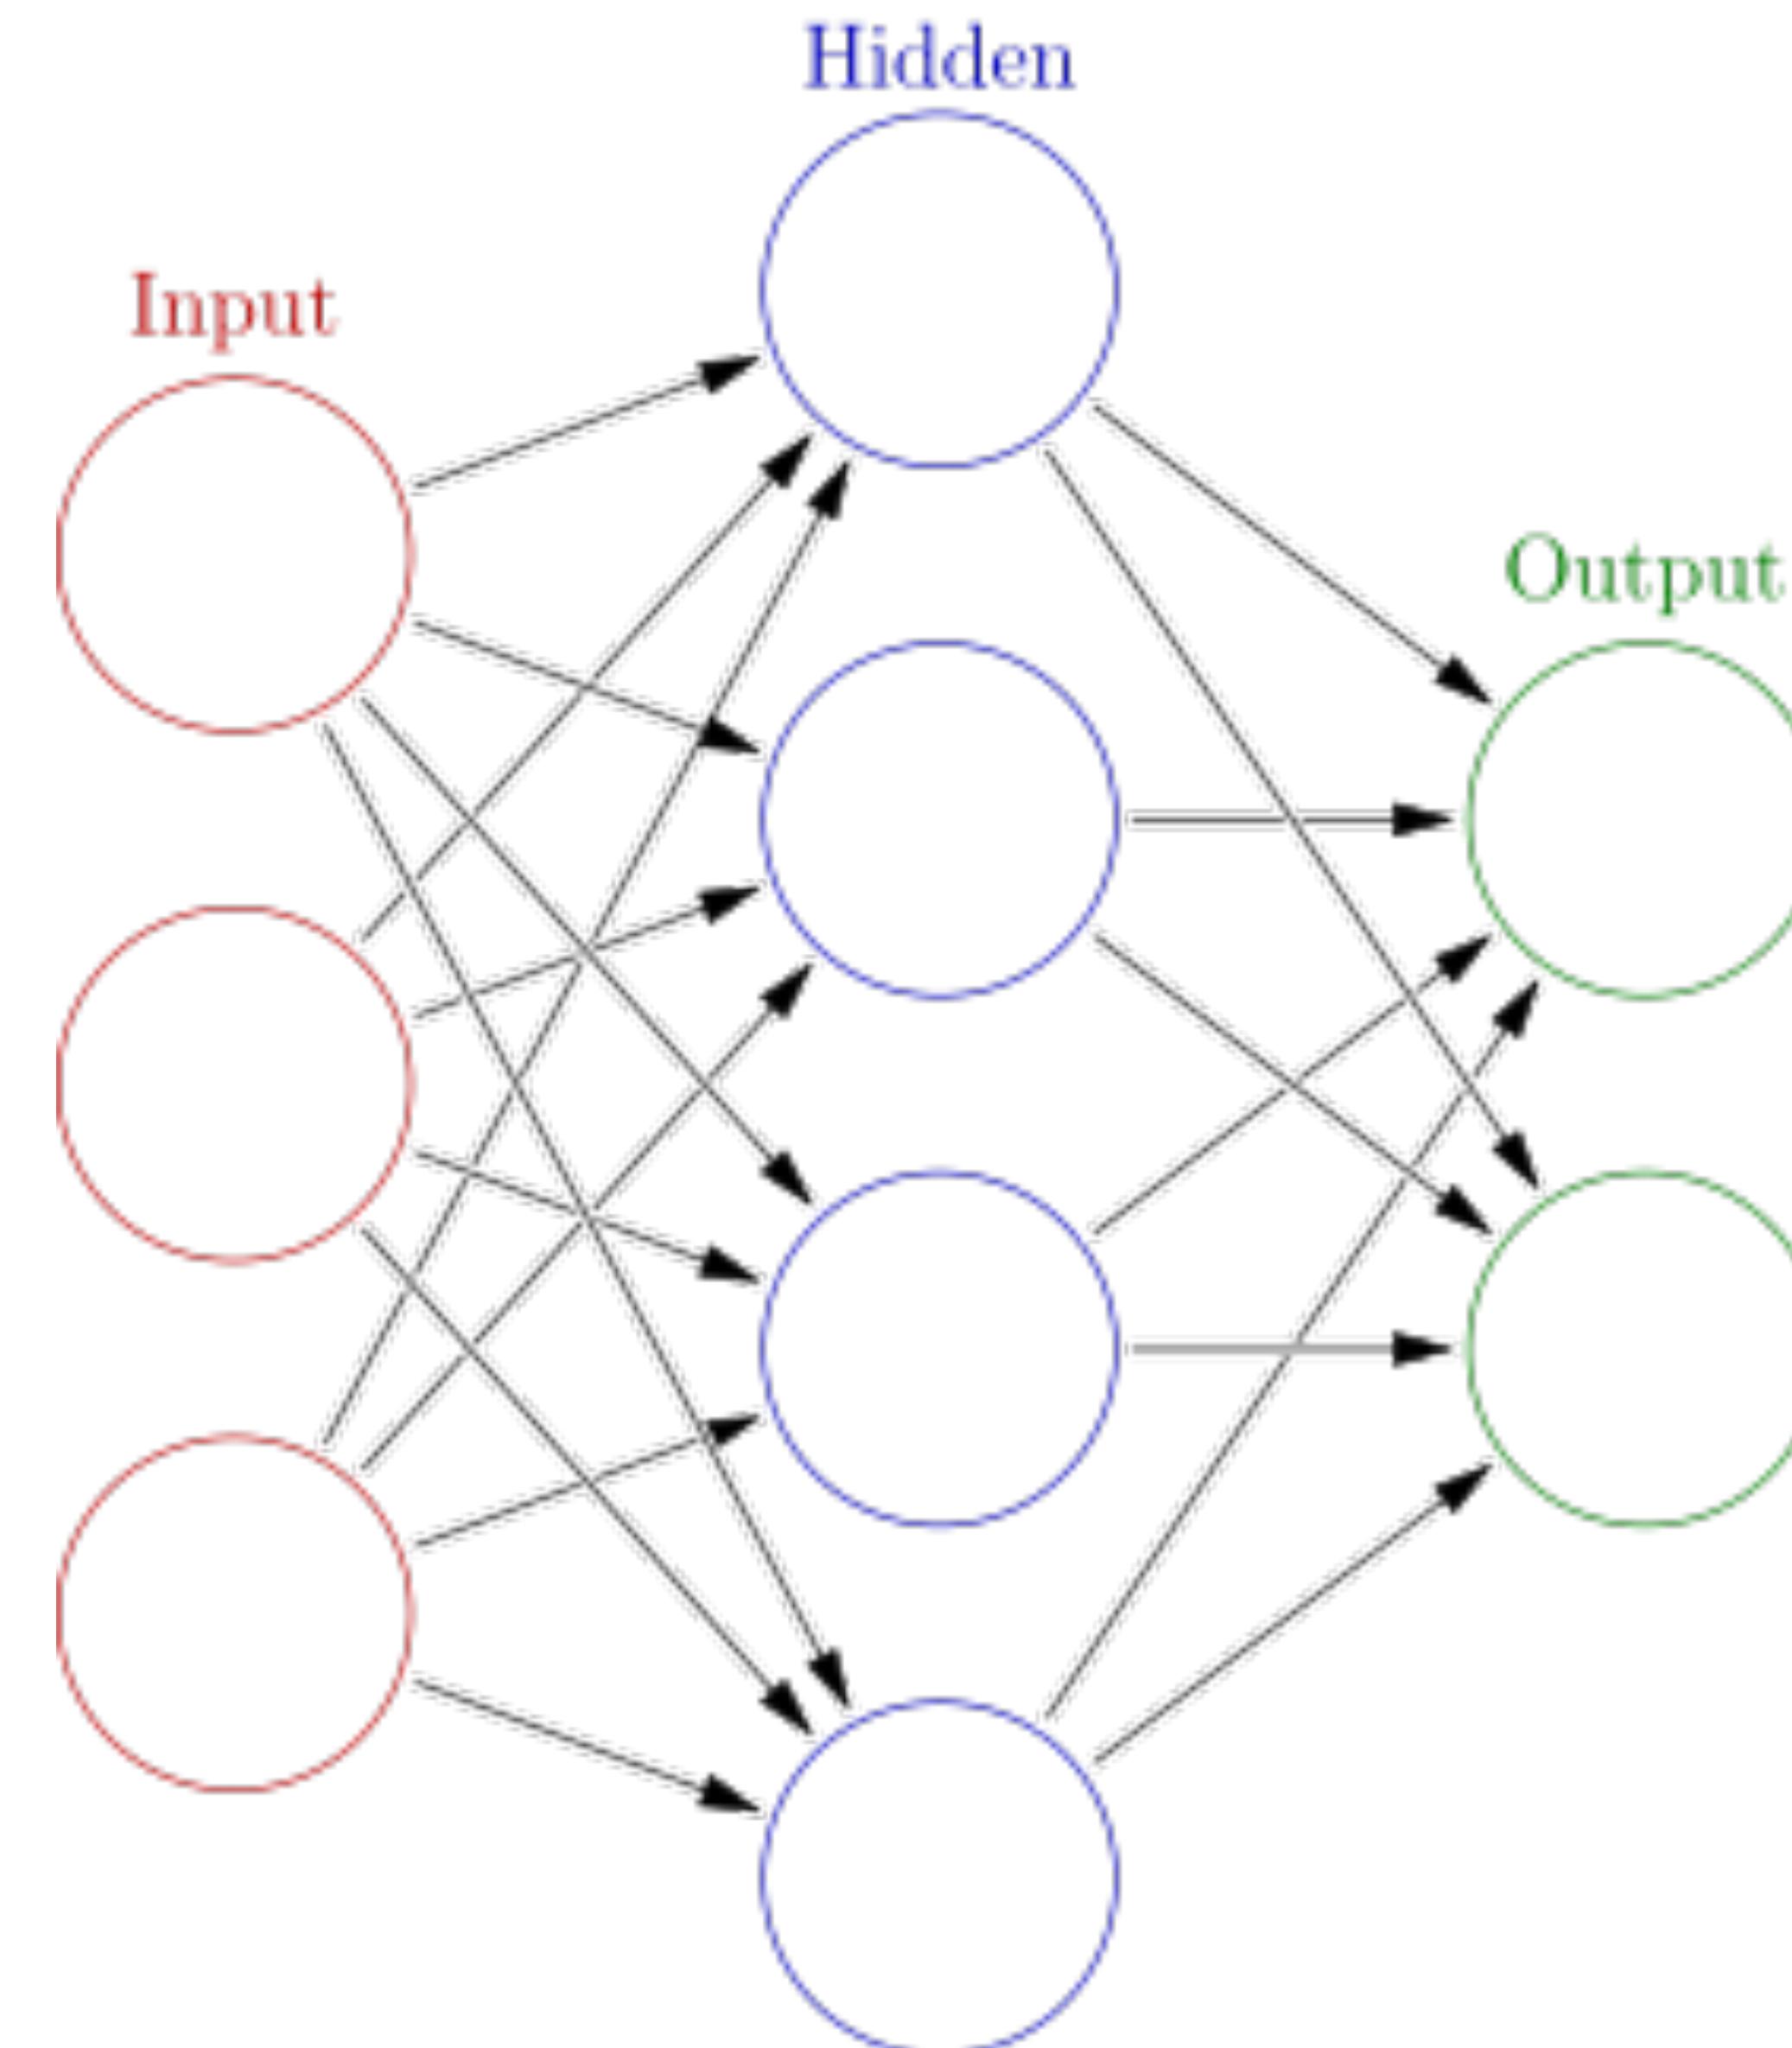
Drift Shadow by Alex Harker

learn.flucoma.org/explore/harker

MLP Parameters

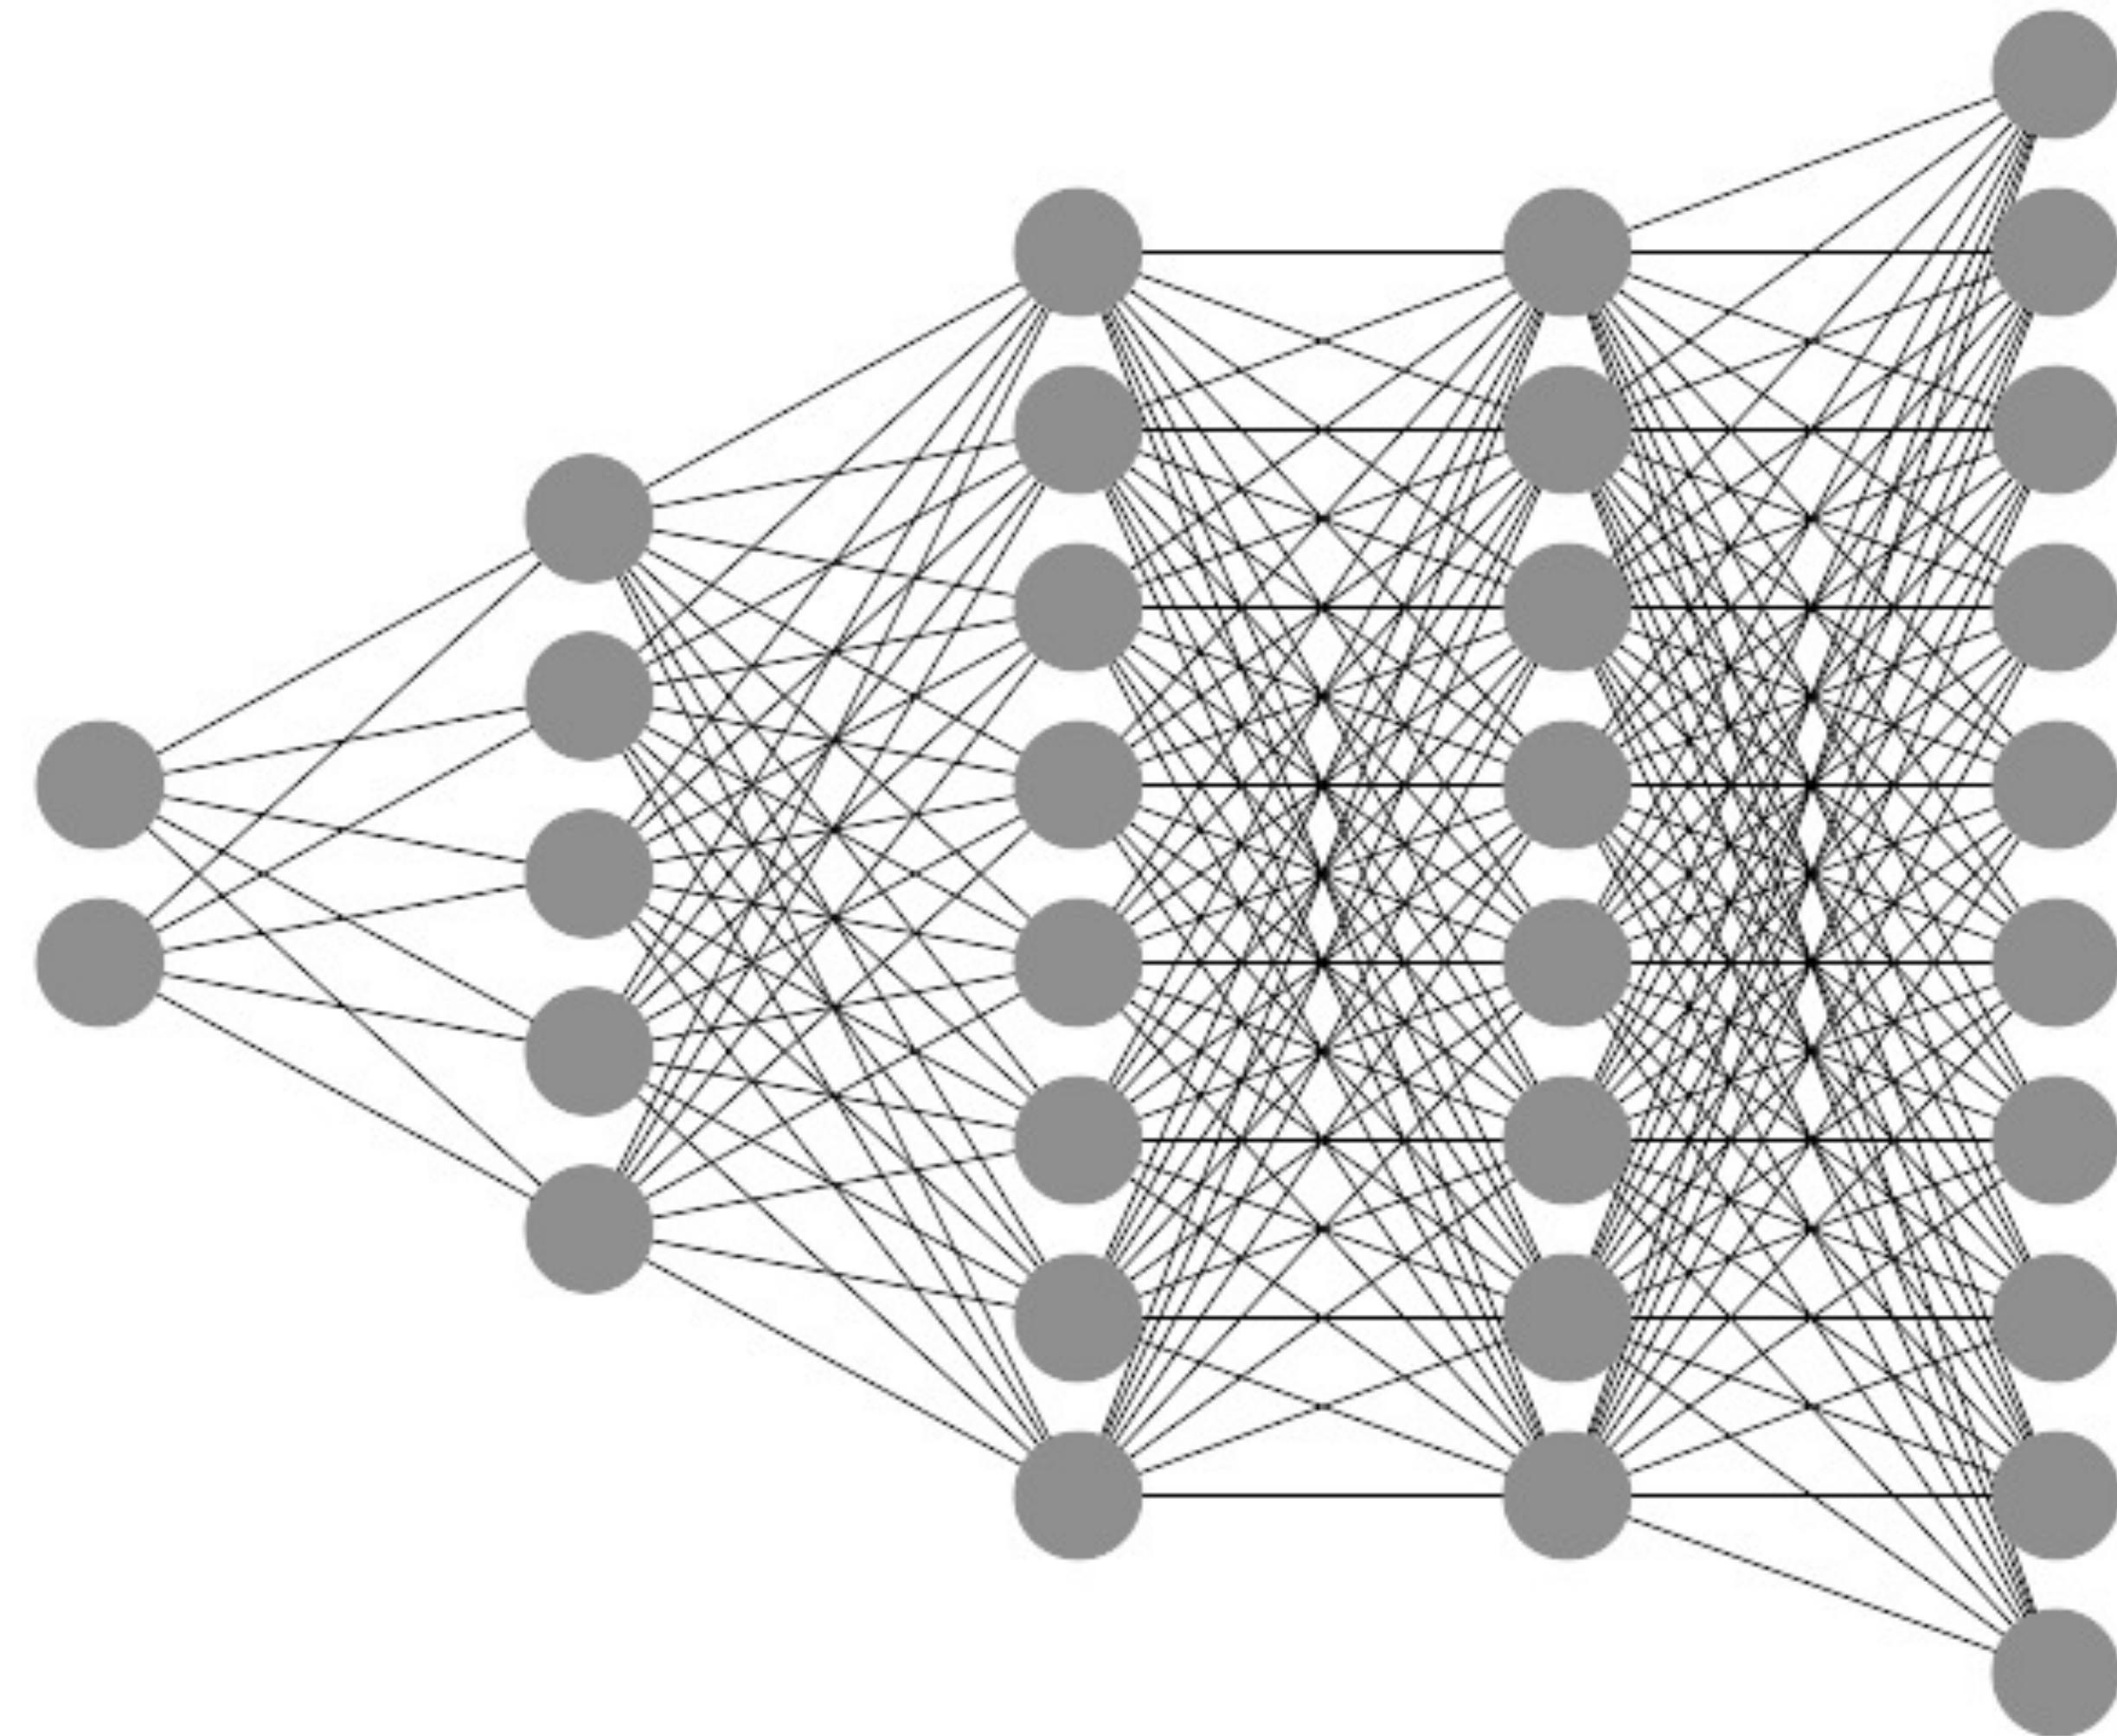


@hiddenlayers (list)
n elements = n layers
integers = n neurons

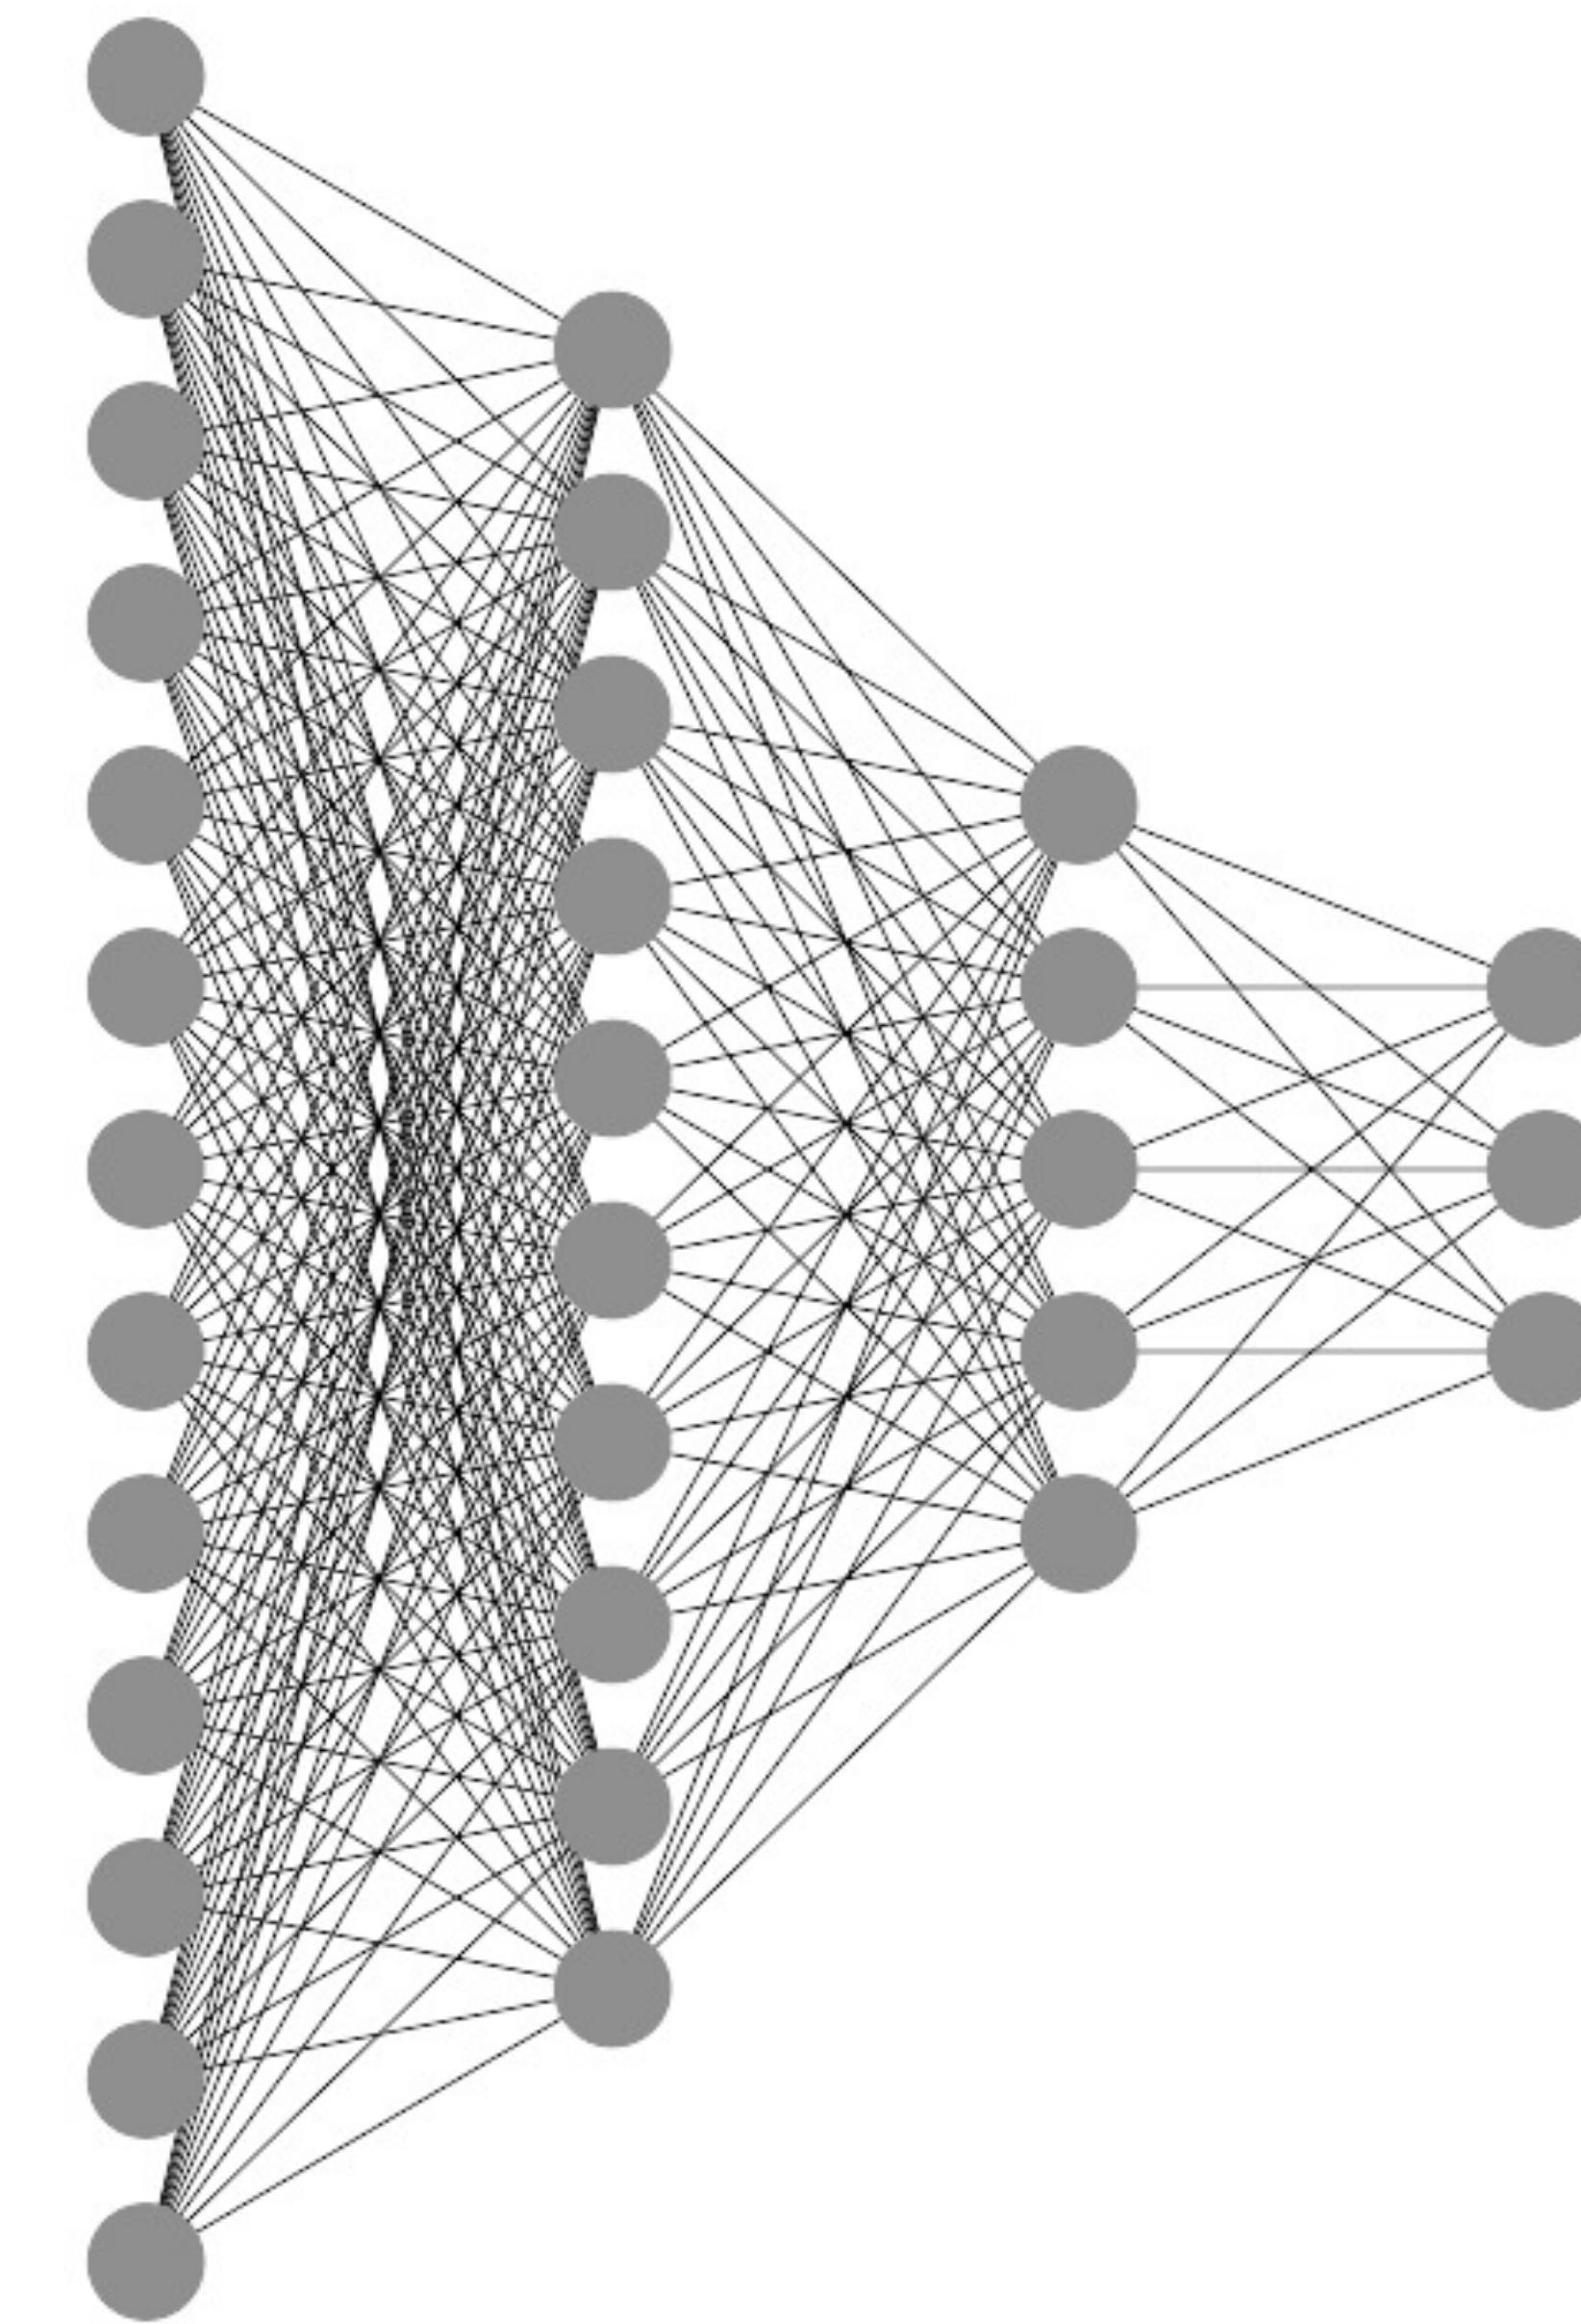


(@hiddenlayers 4)

will reset network



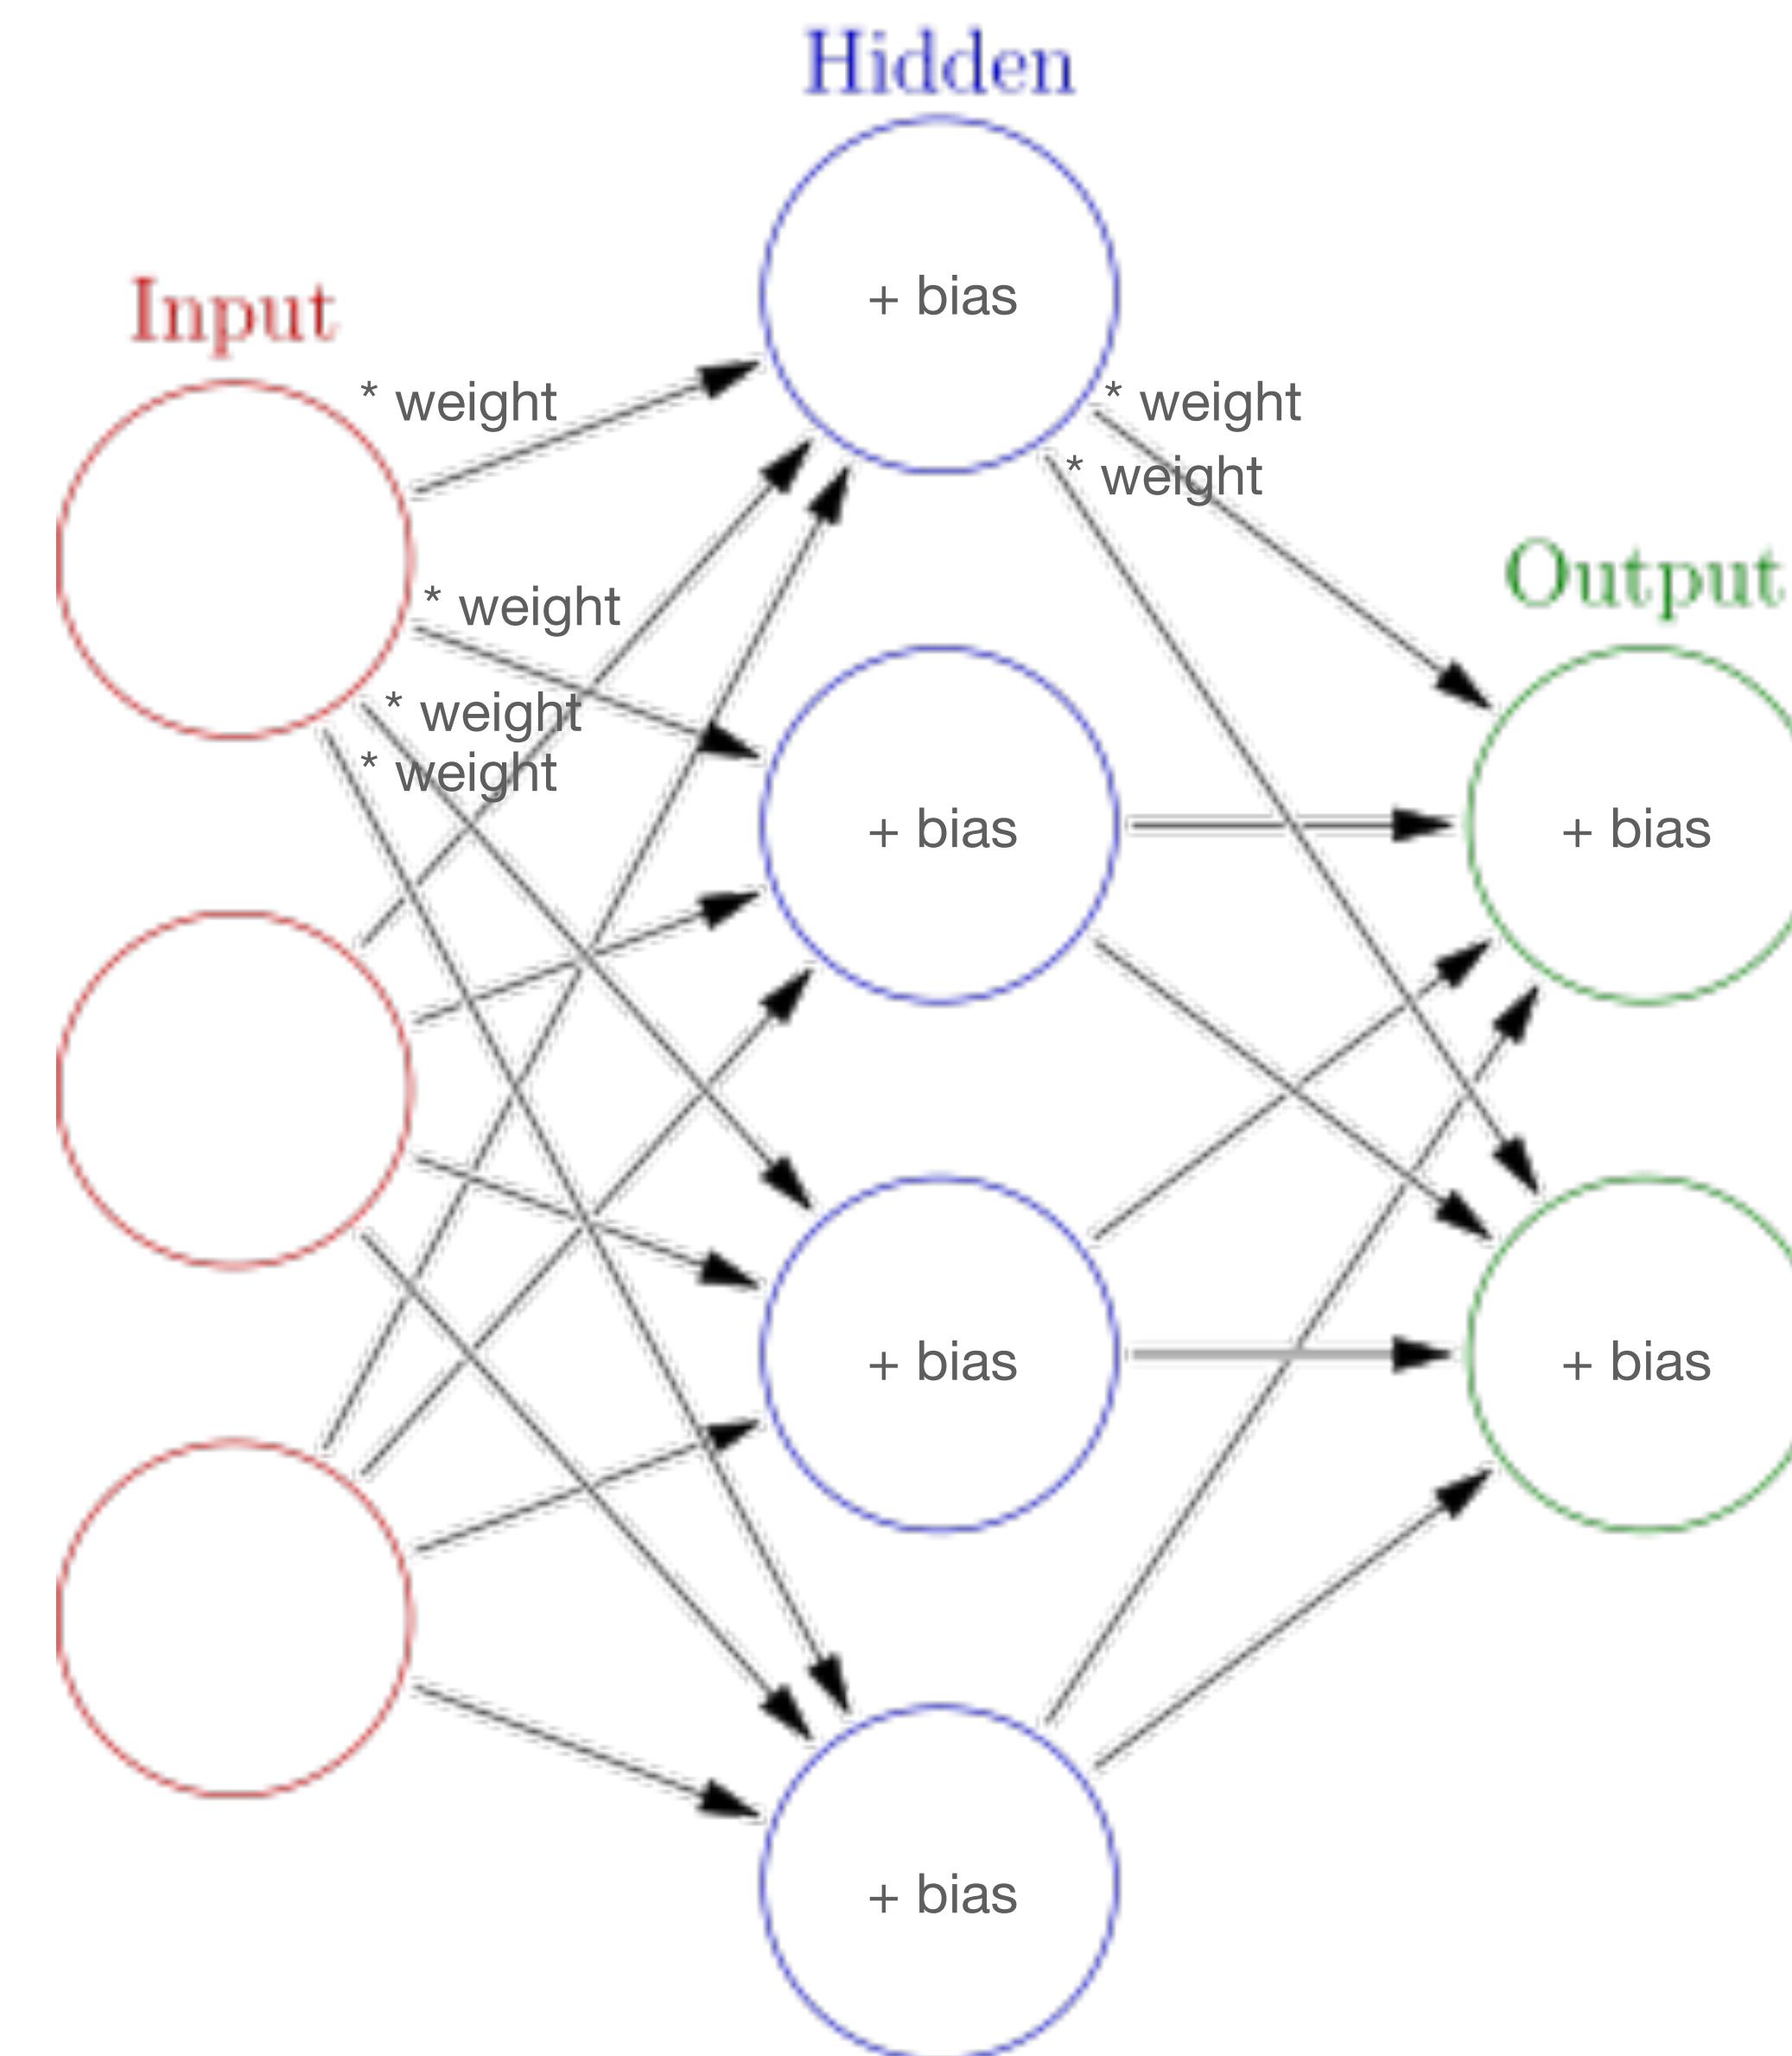
(@hiddenlayers 5 8 8)



(@hiddenlayers 10 5)

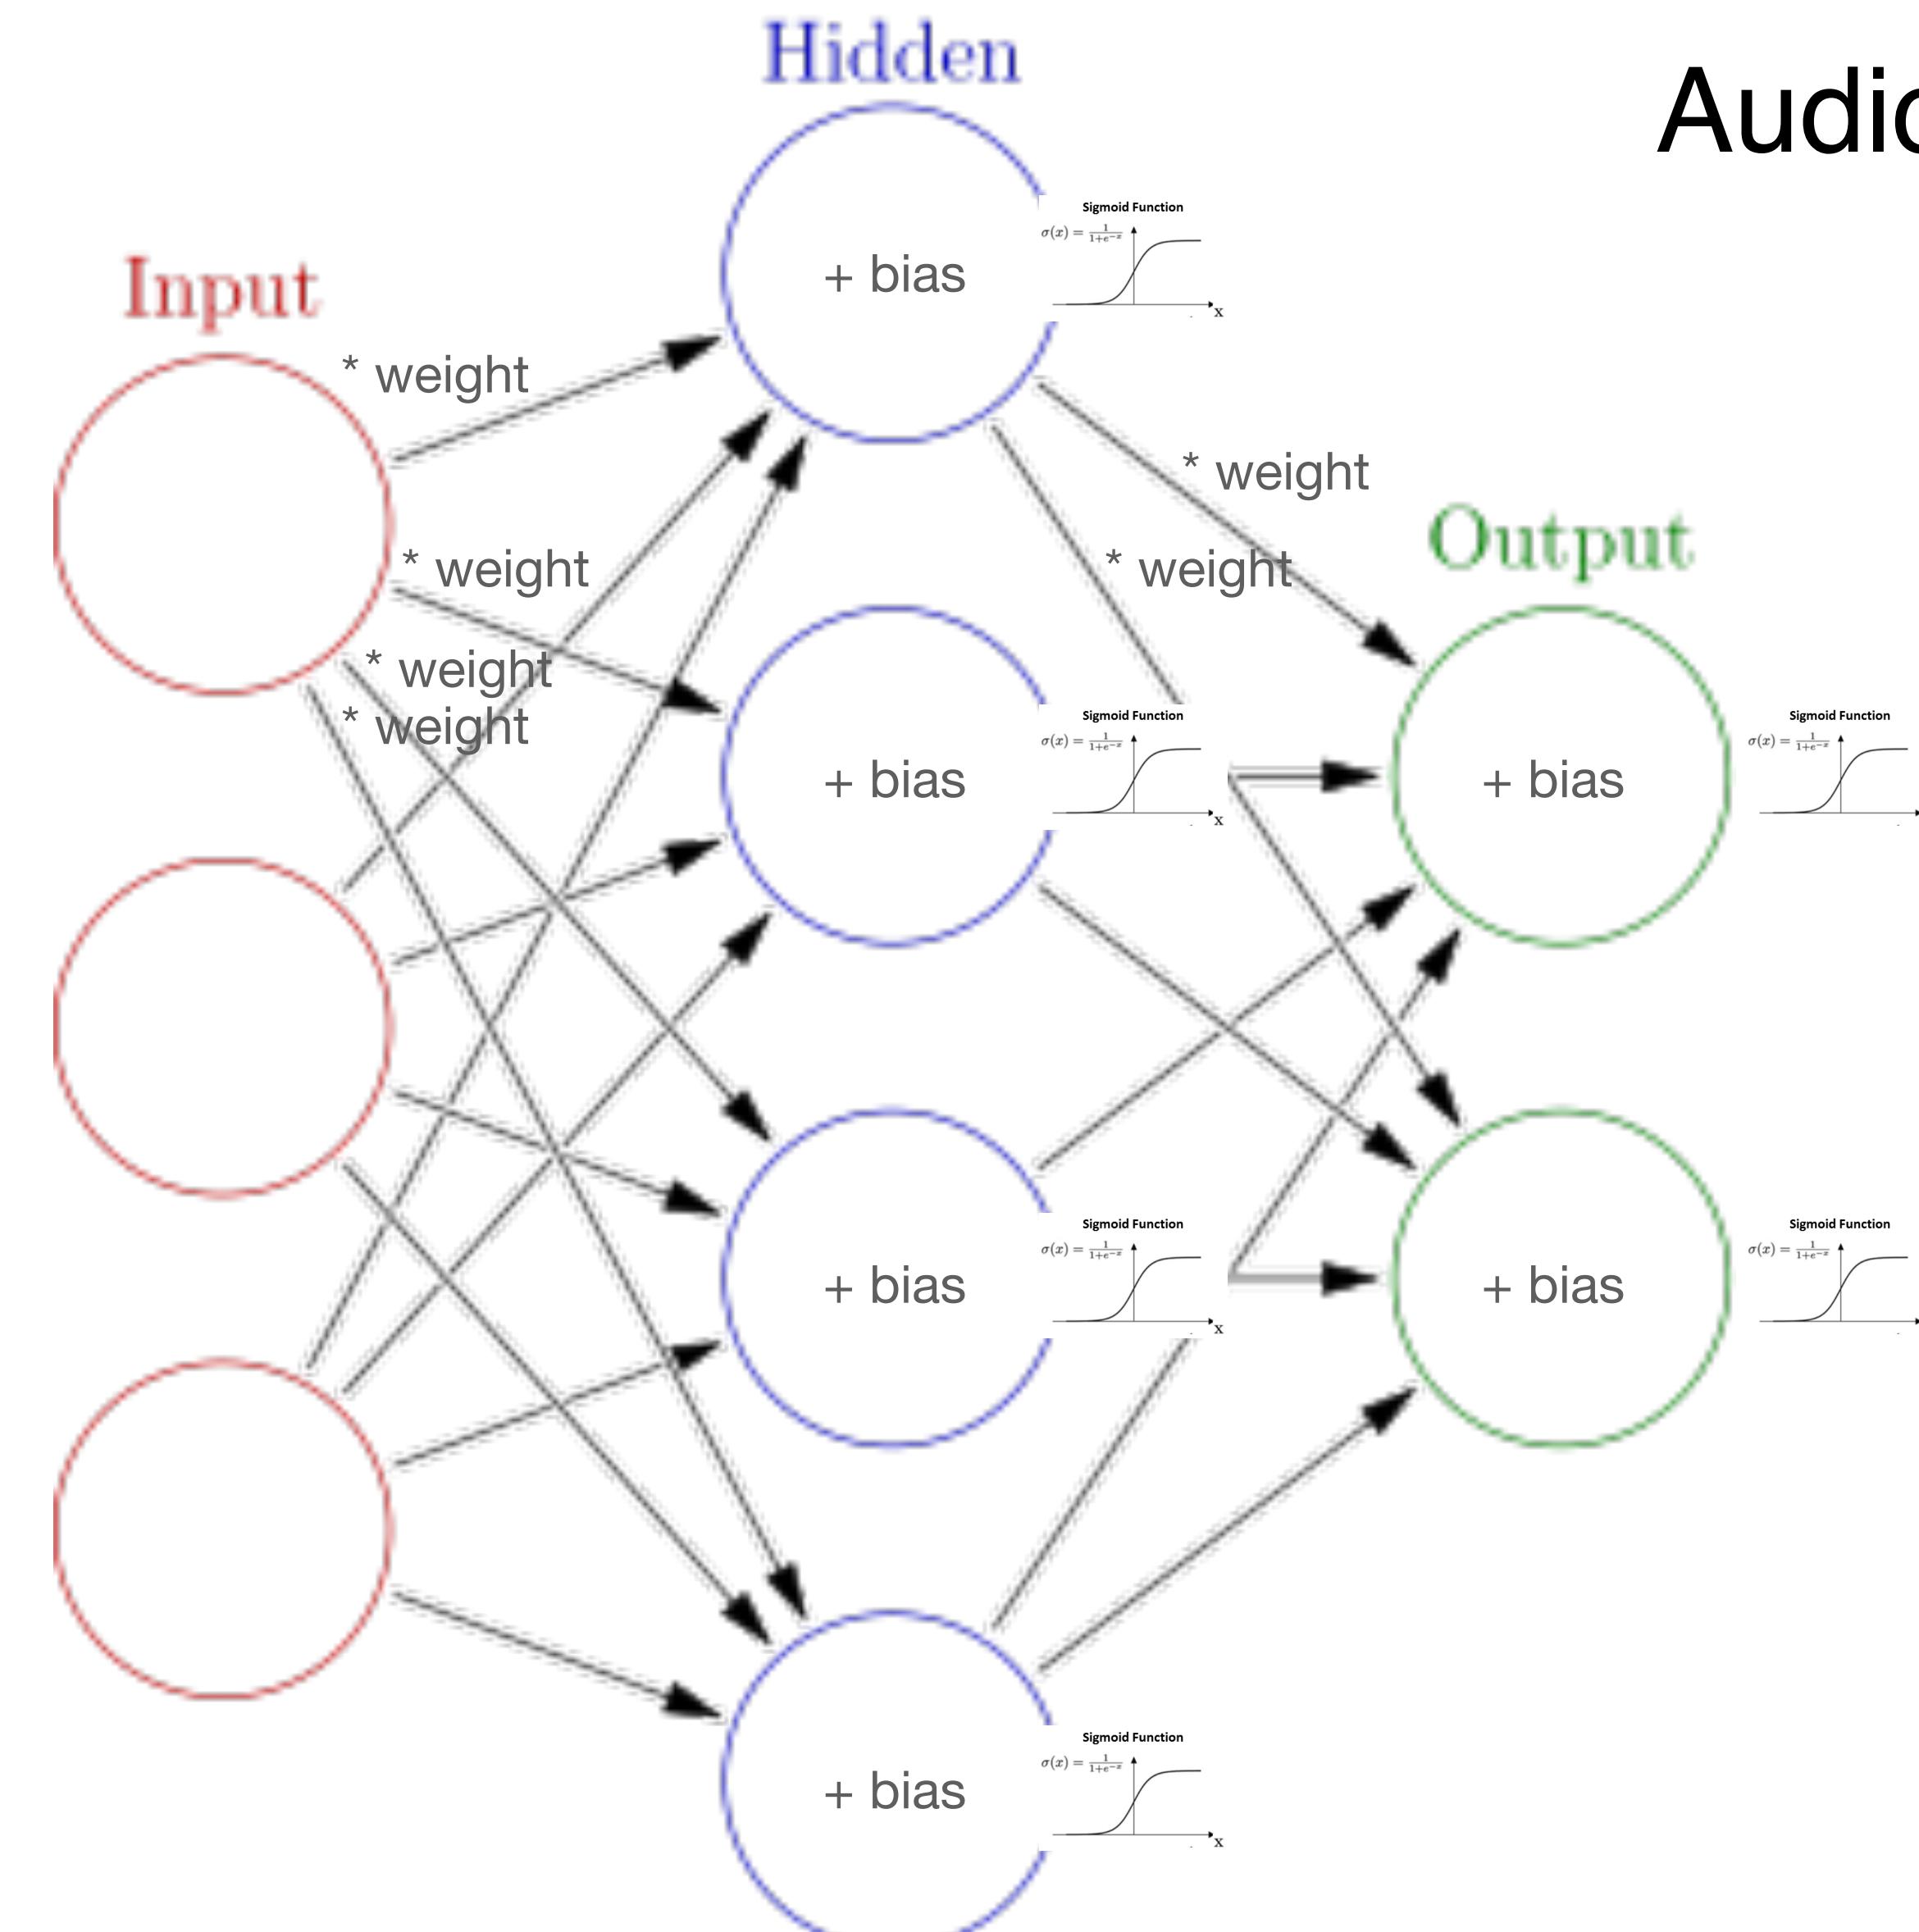
Internal Parameters

weights & biases



@activation
@outputactivation

Audio Mixer Metaphor

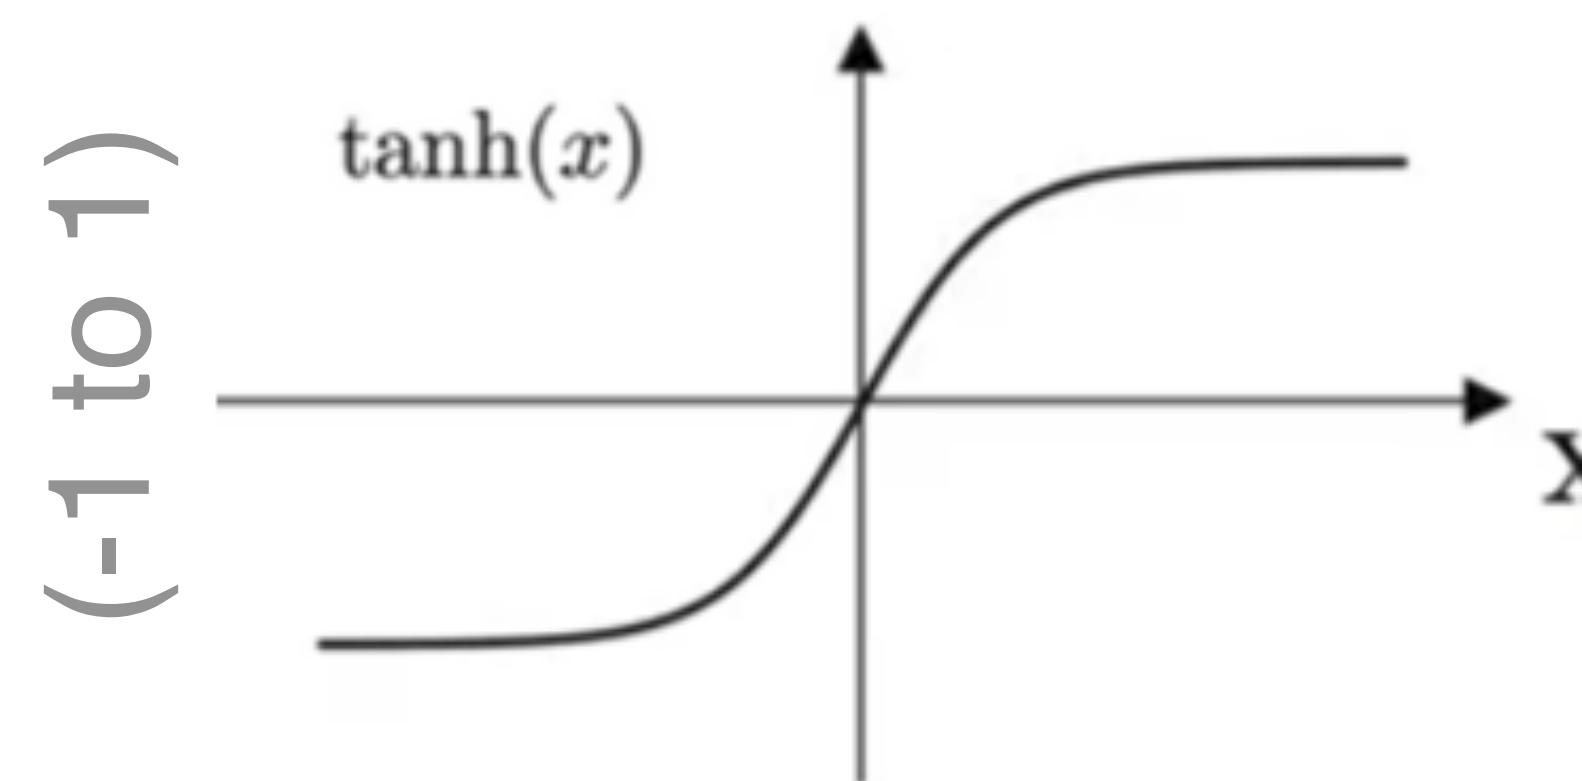


will reset network

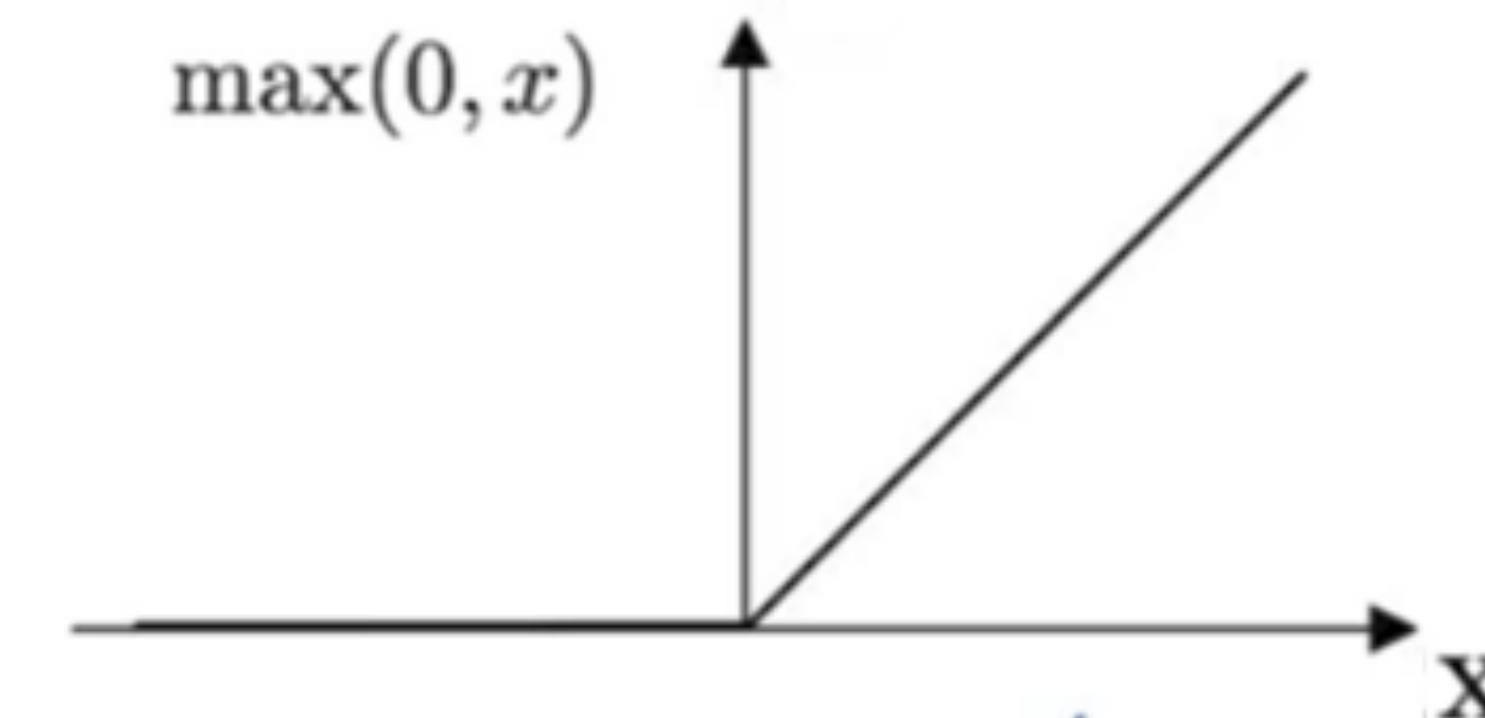
@activation

@outputactivation

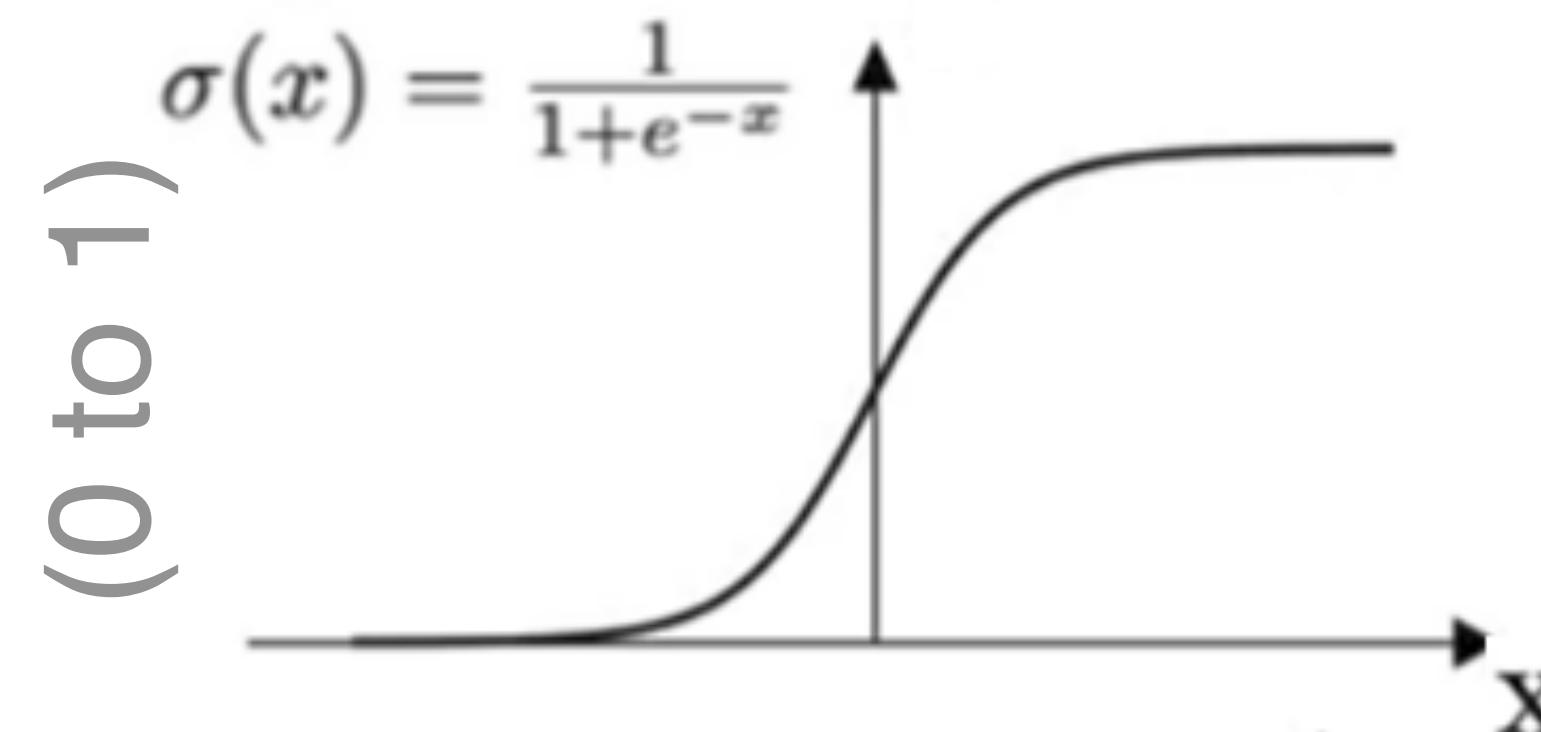
(3) Hyper Tangent Function



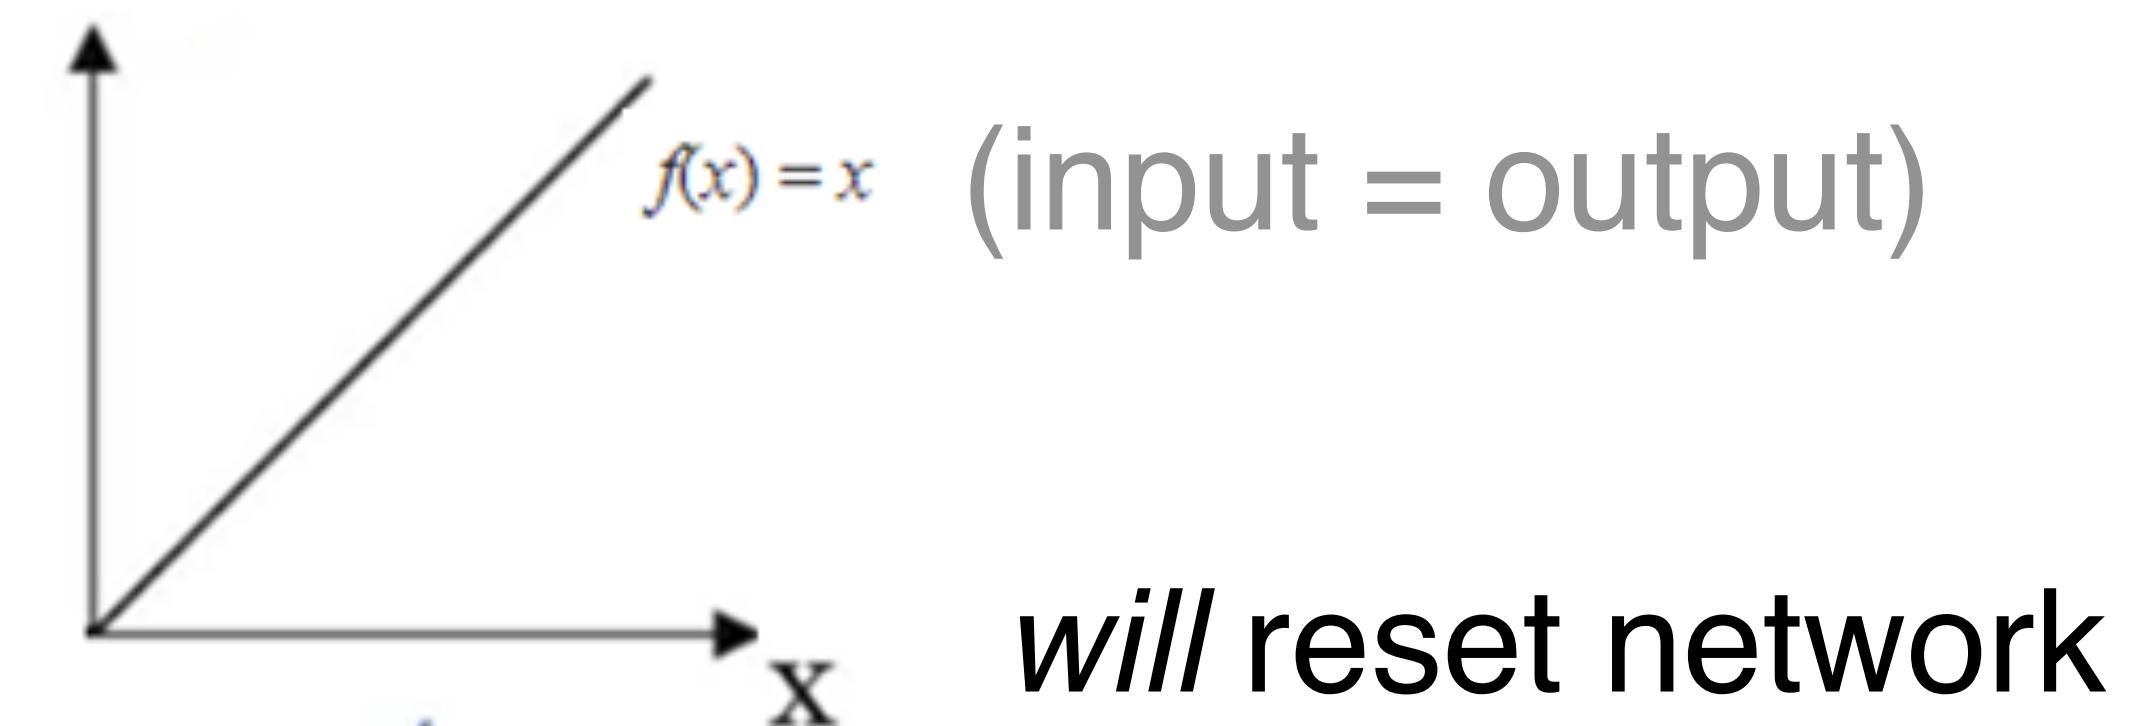
(2) ReLU Function



(1) Sigmoid Function



(0) Identity Function



`@learnrate`

a scalar for indicating how much the neural network should adjust its internal parameters during training (during back-propagation)—starting around 0.1 and decreasing over time is usually a good start

`@maxiter`

number of times to use all the examples on each “fit” message—also called an “epoch”

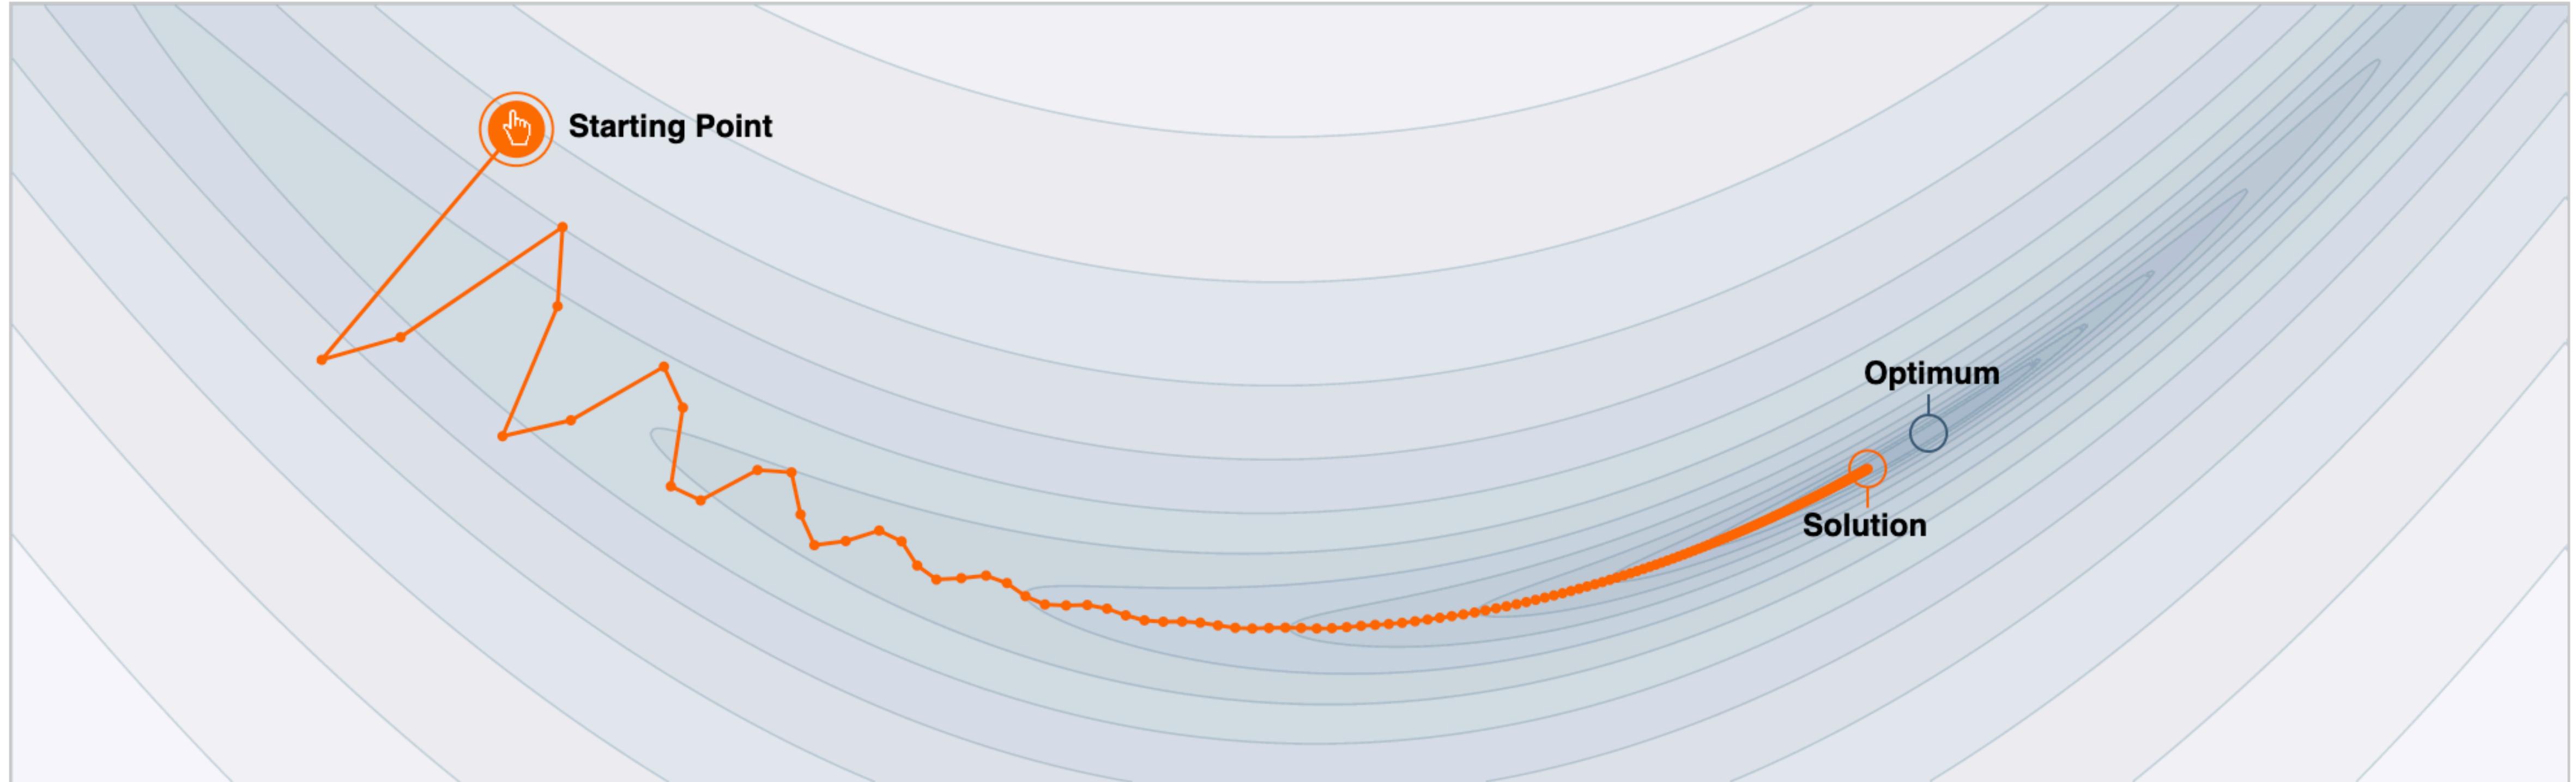
`@batchsize`

number of input examples feed-forward before checking their output pairs and adjusting

`@validation`

percentage (as a decimal) of training data to hold back (not train on) then test with—when there’s *lots* of data, may be helpful to prevent overfitting

@momentum apply smoothing adjustments



Regression with Audio Analyses

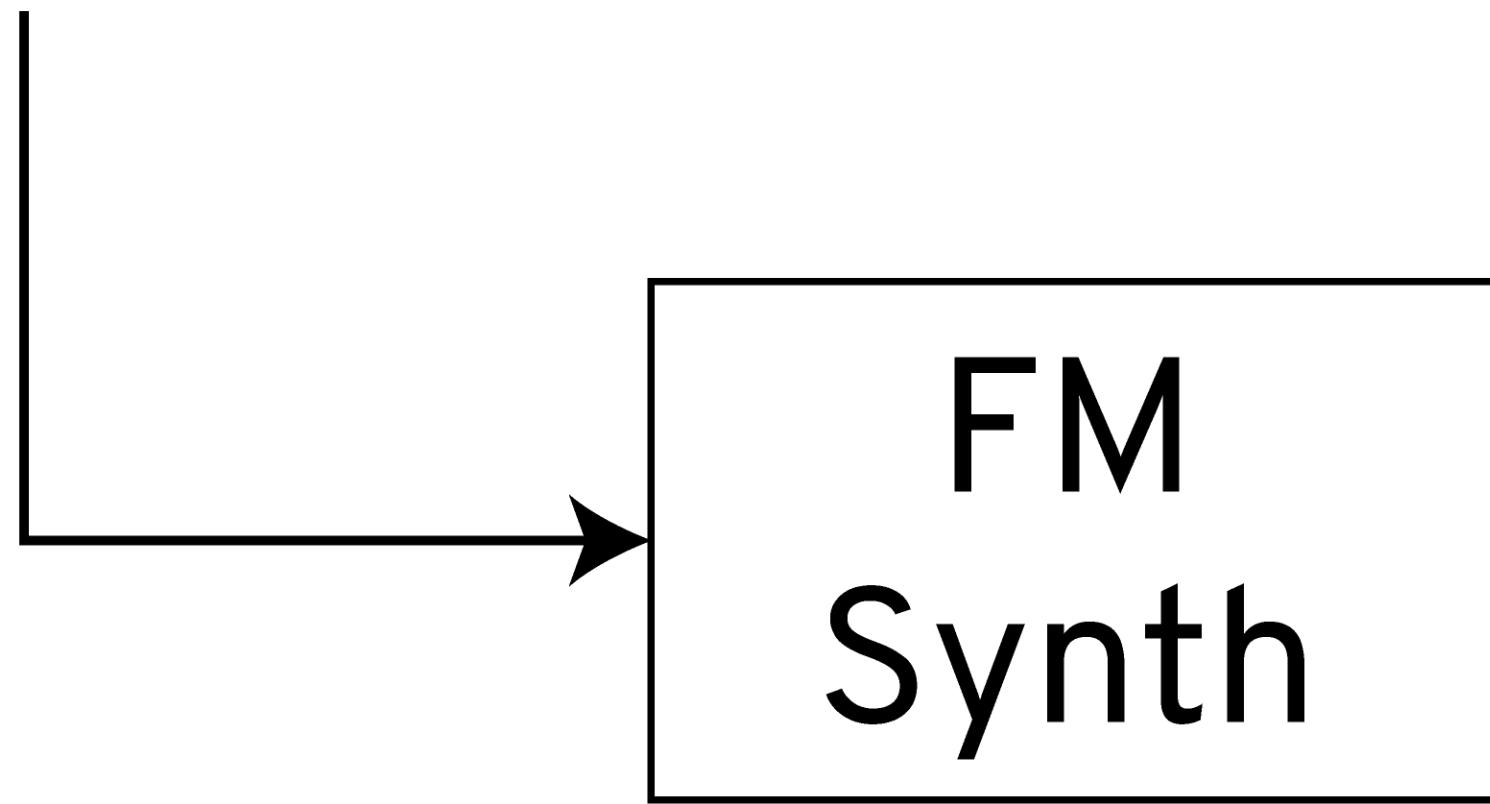


Making the Datasets

Freq. Mod. Params

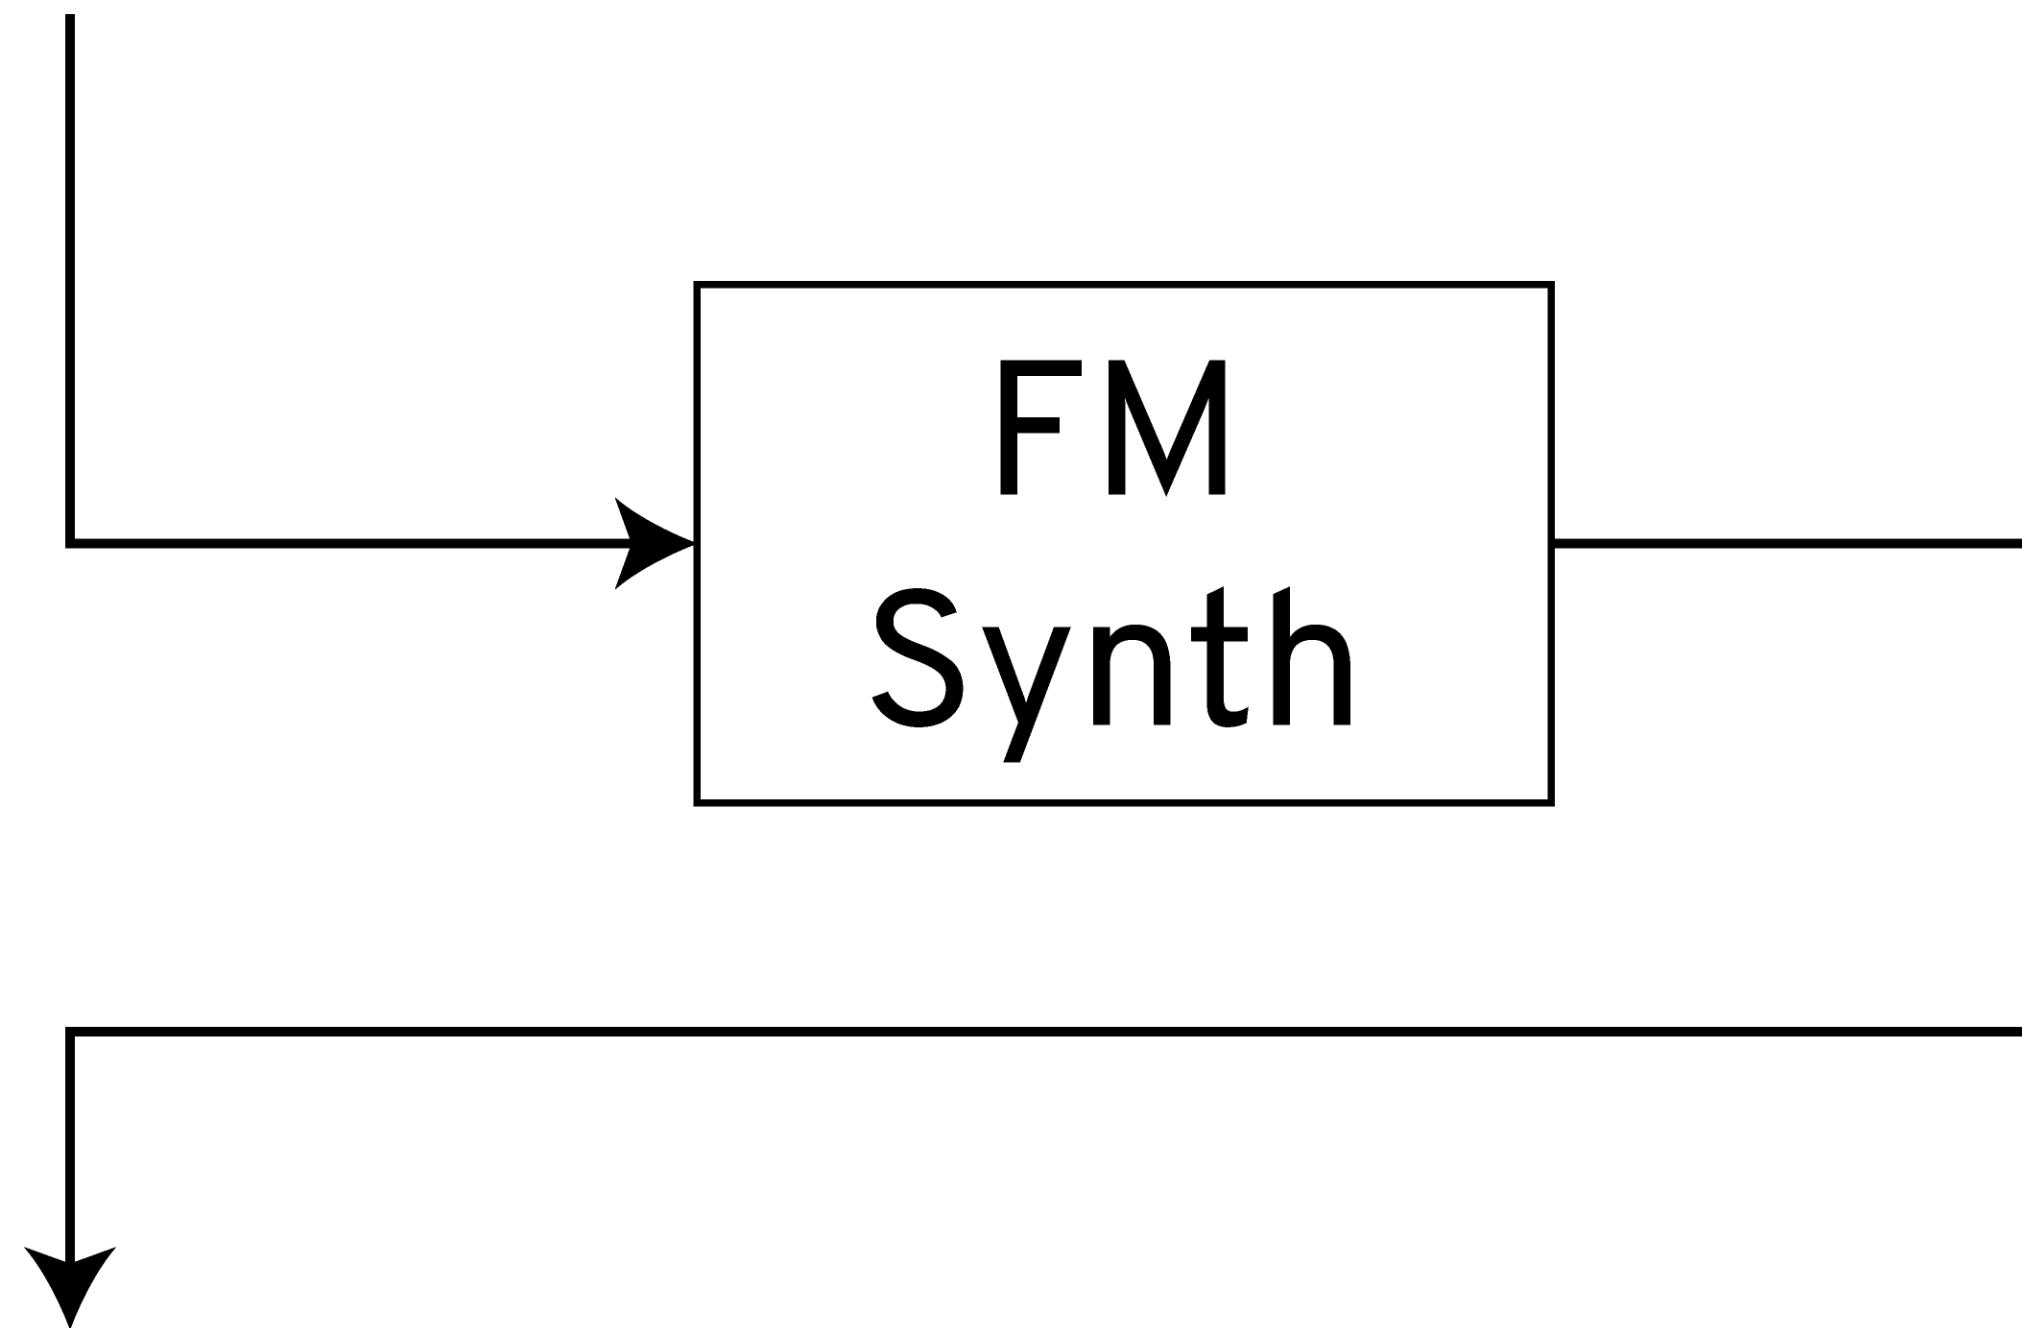
Making the Datasets

Freq. Mod. Params



Making the Datasets

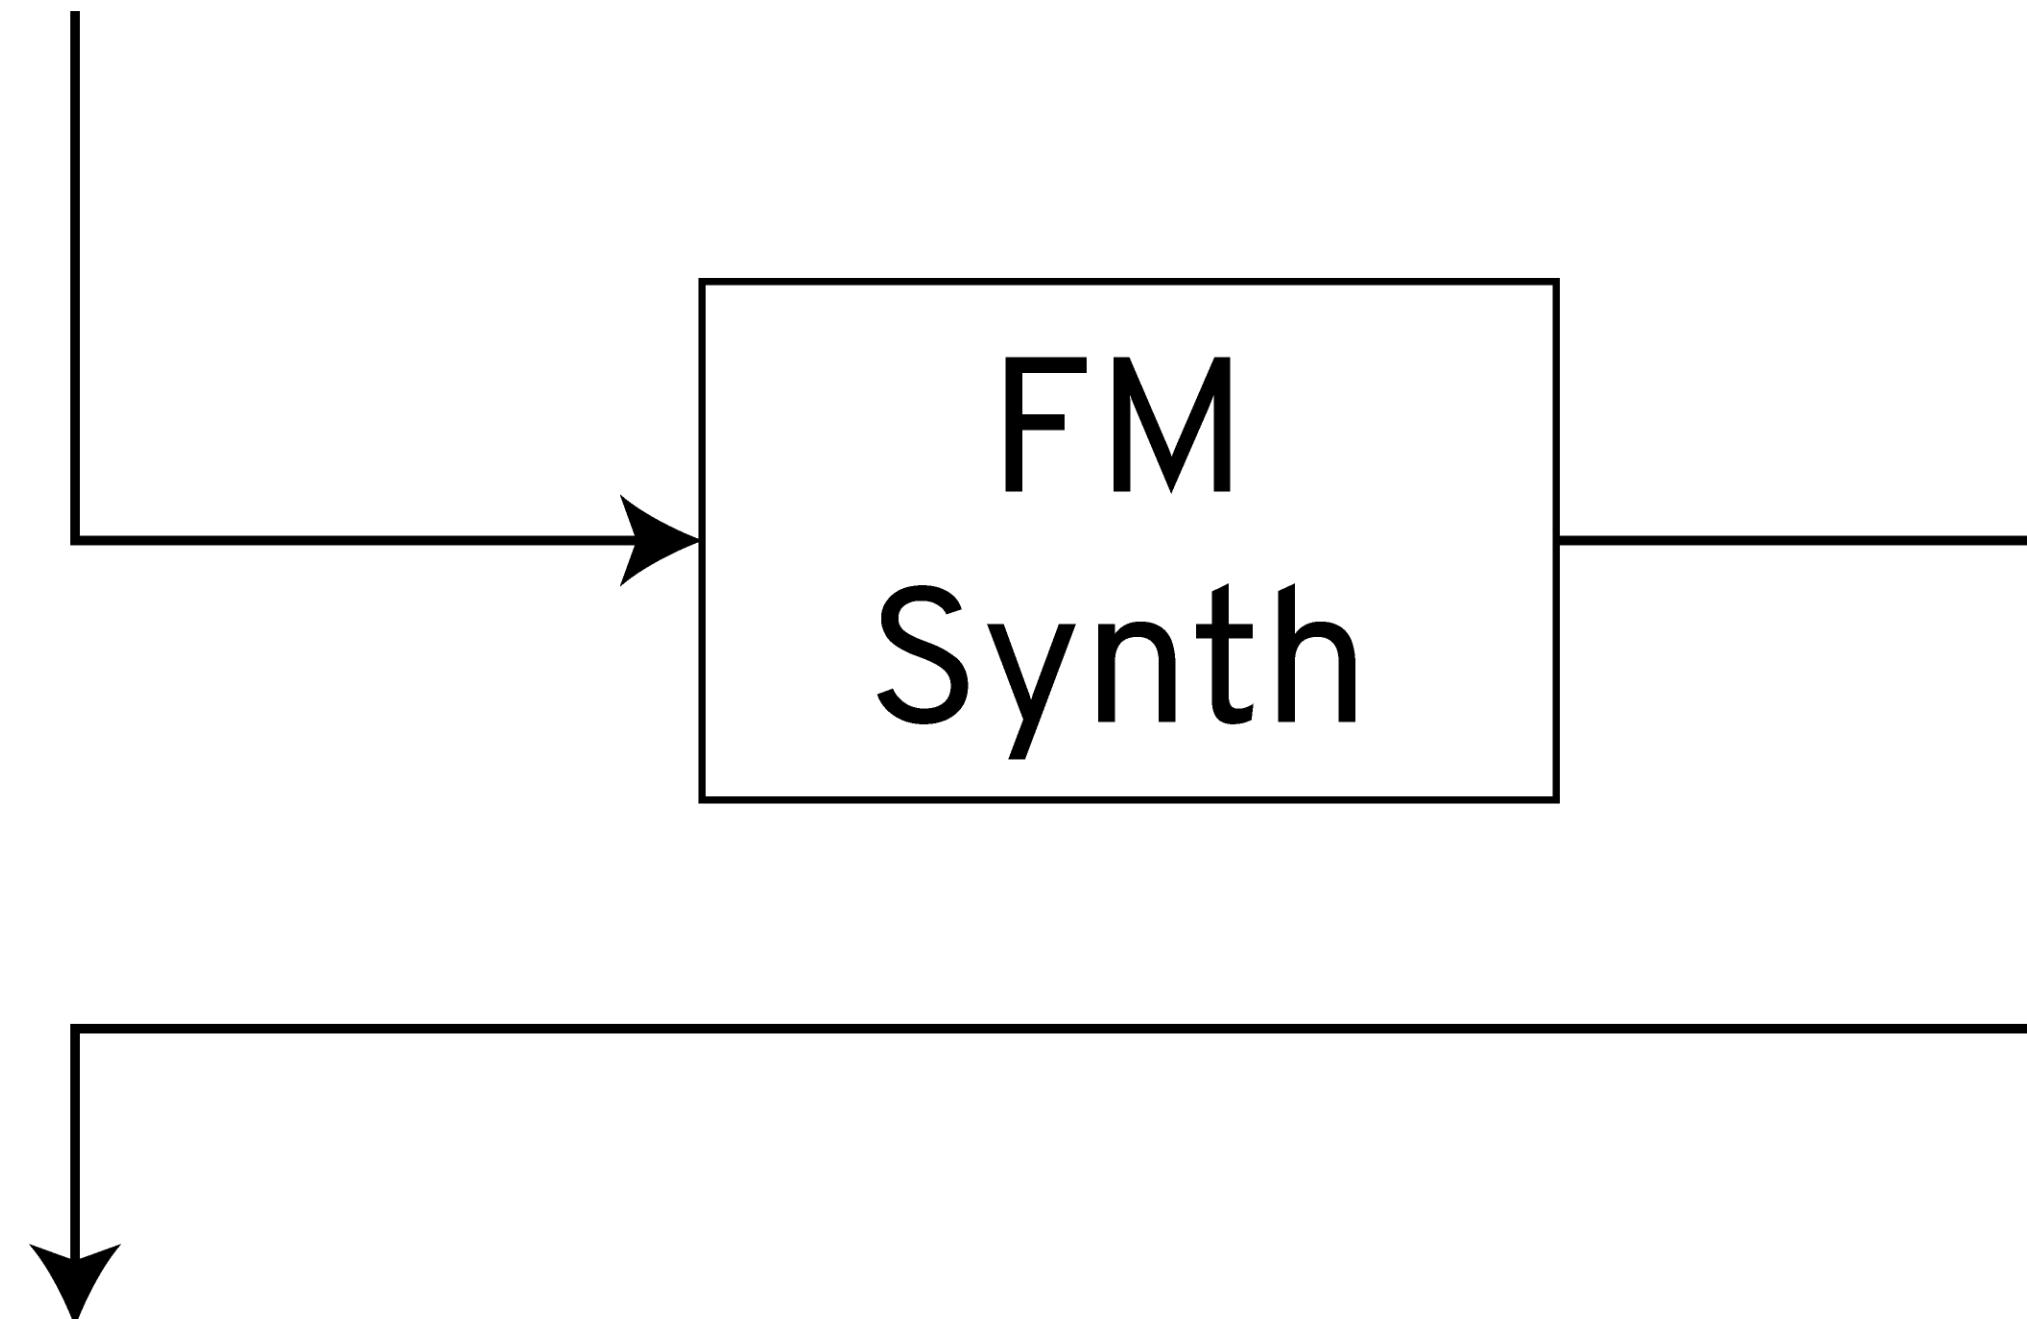
Freq. Mod. Params



Audio Signal

Making the Datasets

Freq. Mod. Params



Audio Signal —————→ 13 MFCCs

Making the Datasets

Freq. Mod. Params

13 MFCCs

Making the Datasets

Freq. Mod. Params

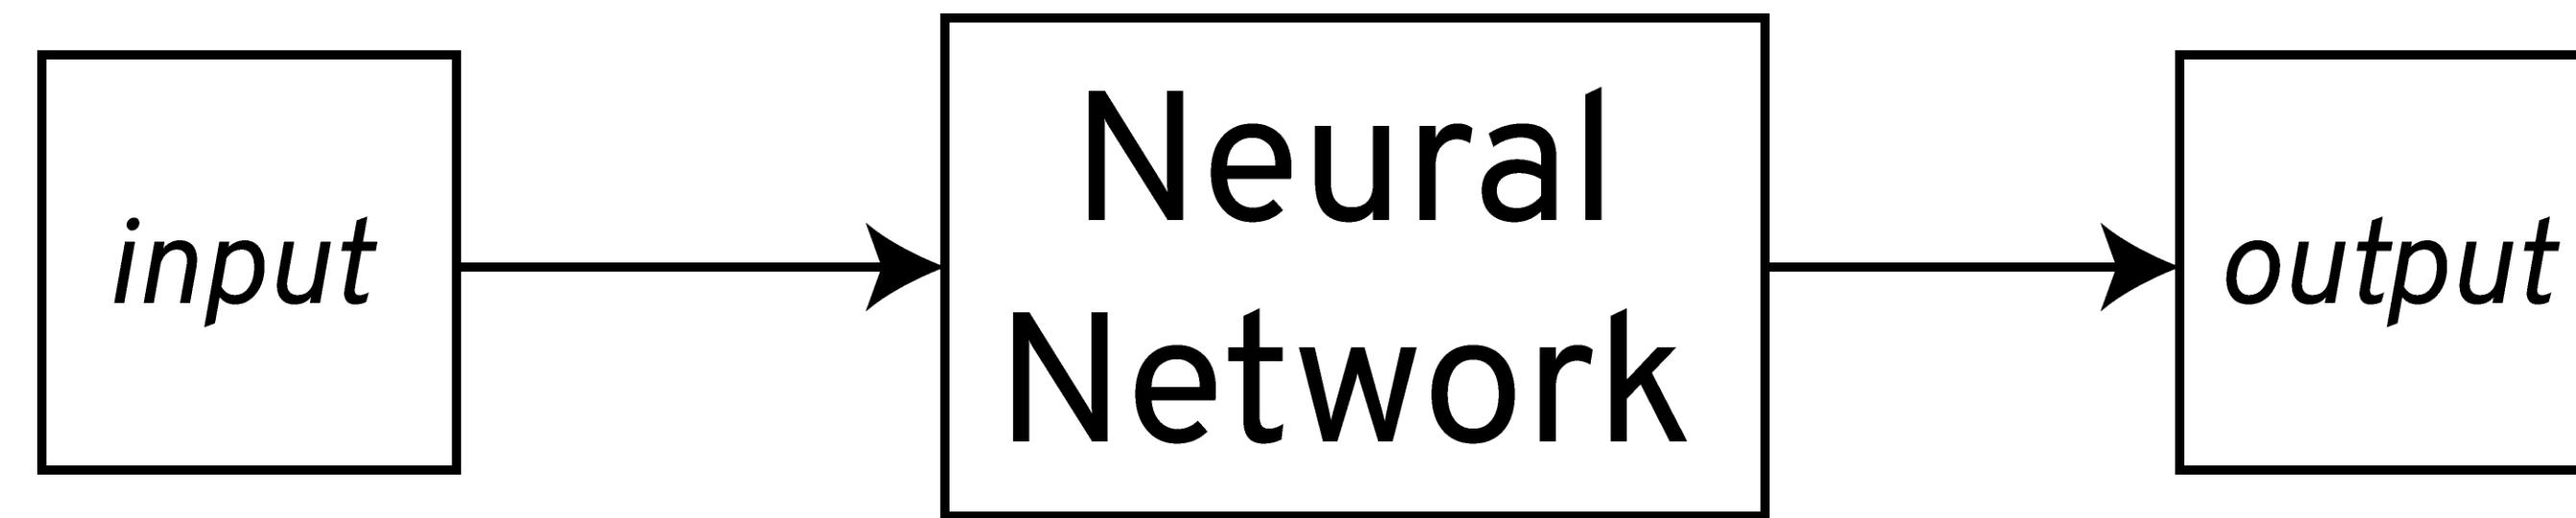
```
DataSet 64:  
rows: 100 cols: 3  
point-0      2411.1    206.48   3.9116  
point-1      69.214     65.334   3.862  
point-2      214.7     67.041   6.1554  
...  
point-97     2114.2    82.574   7.8388  
point-98     307.08    307.08   15.905  
point-99     3797.7    60.559   15.826
```

```
DataSet 63:  
rows: 100 cols: 13  
point-0      183.98    -204.99   -97.85    ...    17.116    6.7544   -2.6597  
point-1      238.64    92.367    67.025    ...    -2.1895   0.26965   3.3148  
point-2      279.43    96.95     44.27     ...    9.4531    2.0989   -1.2892  
...  
point-97     162.65    -185.08   -142.91    ...    -19.588   9.1036   14.504  
point-98     171.71    -210.44   64.047    ...    0.6732    -15.17   -8.3311  
point-99     4.409     -235.47   69.209    ...    -12.299   -6.0117   28.718
```

13 MFCCs

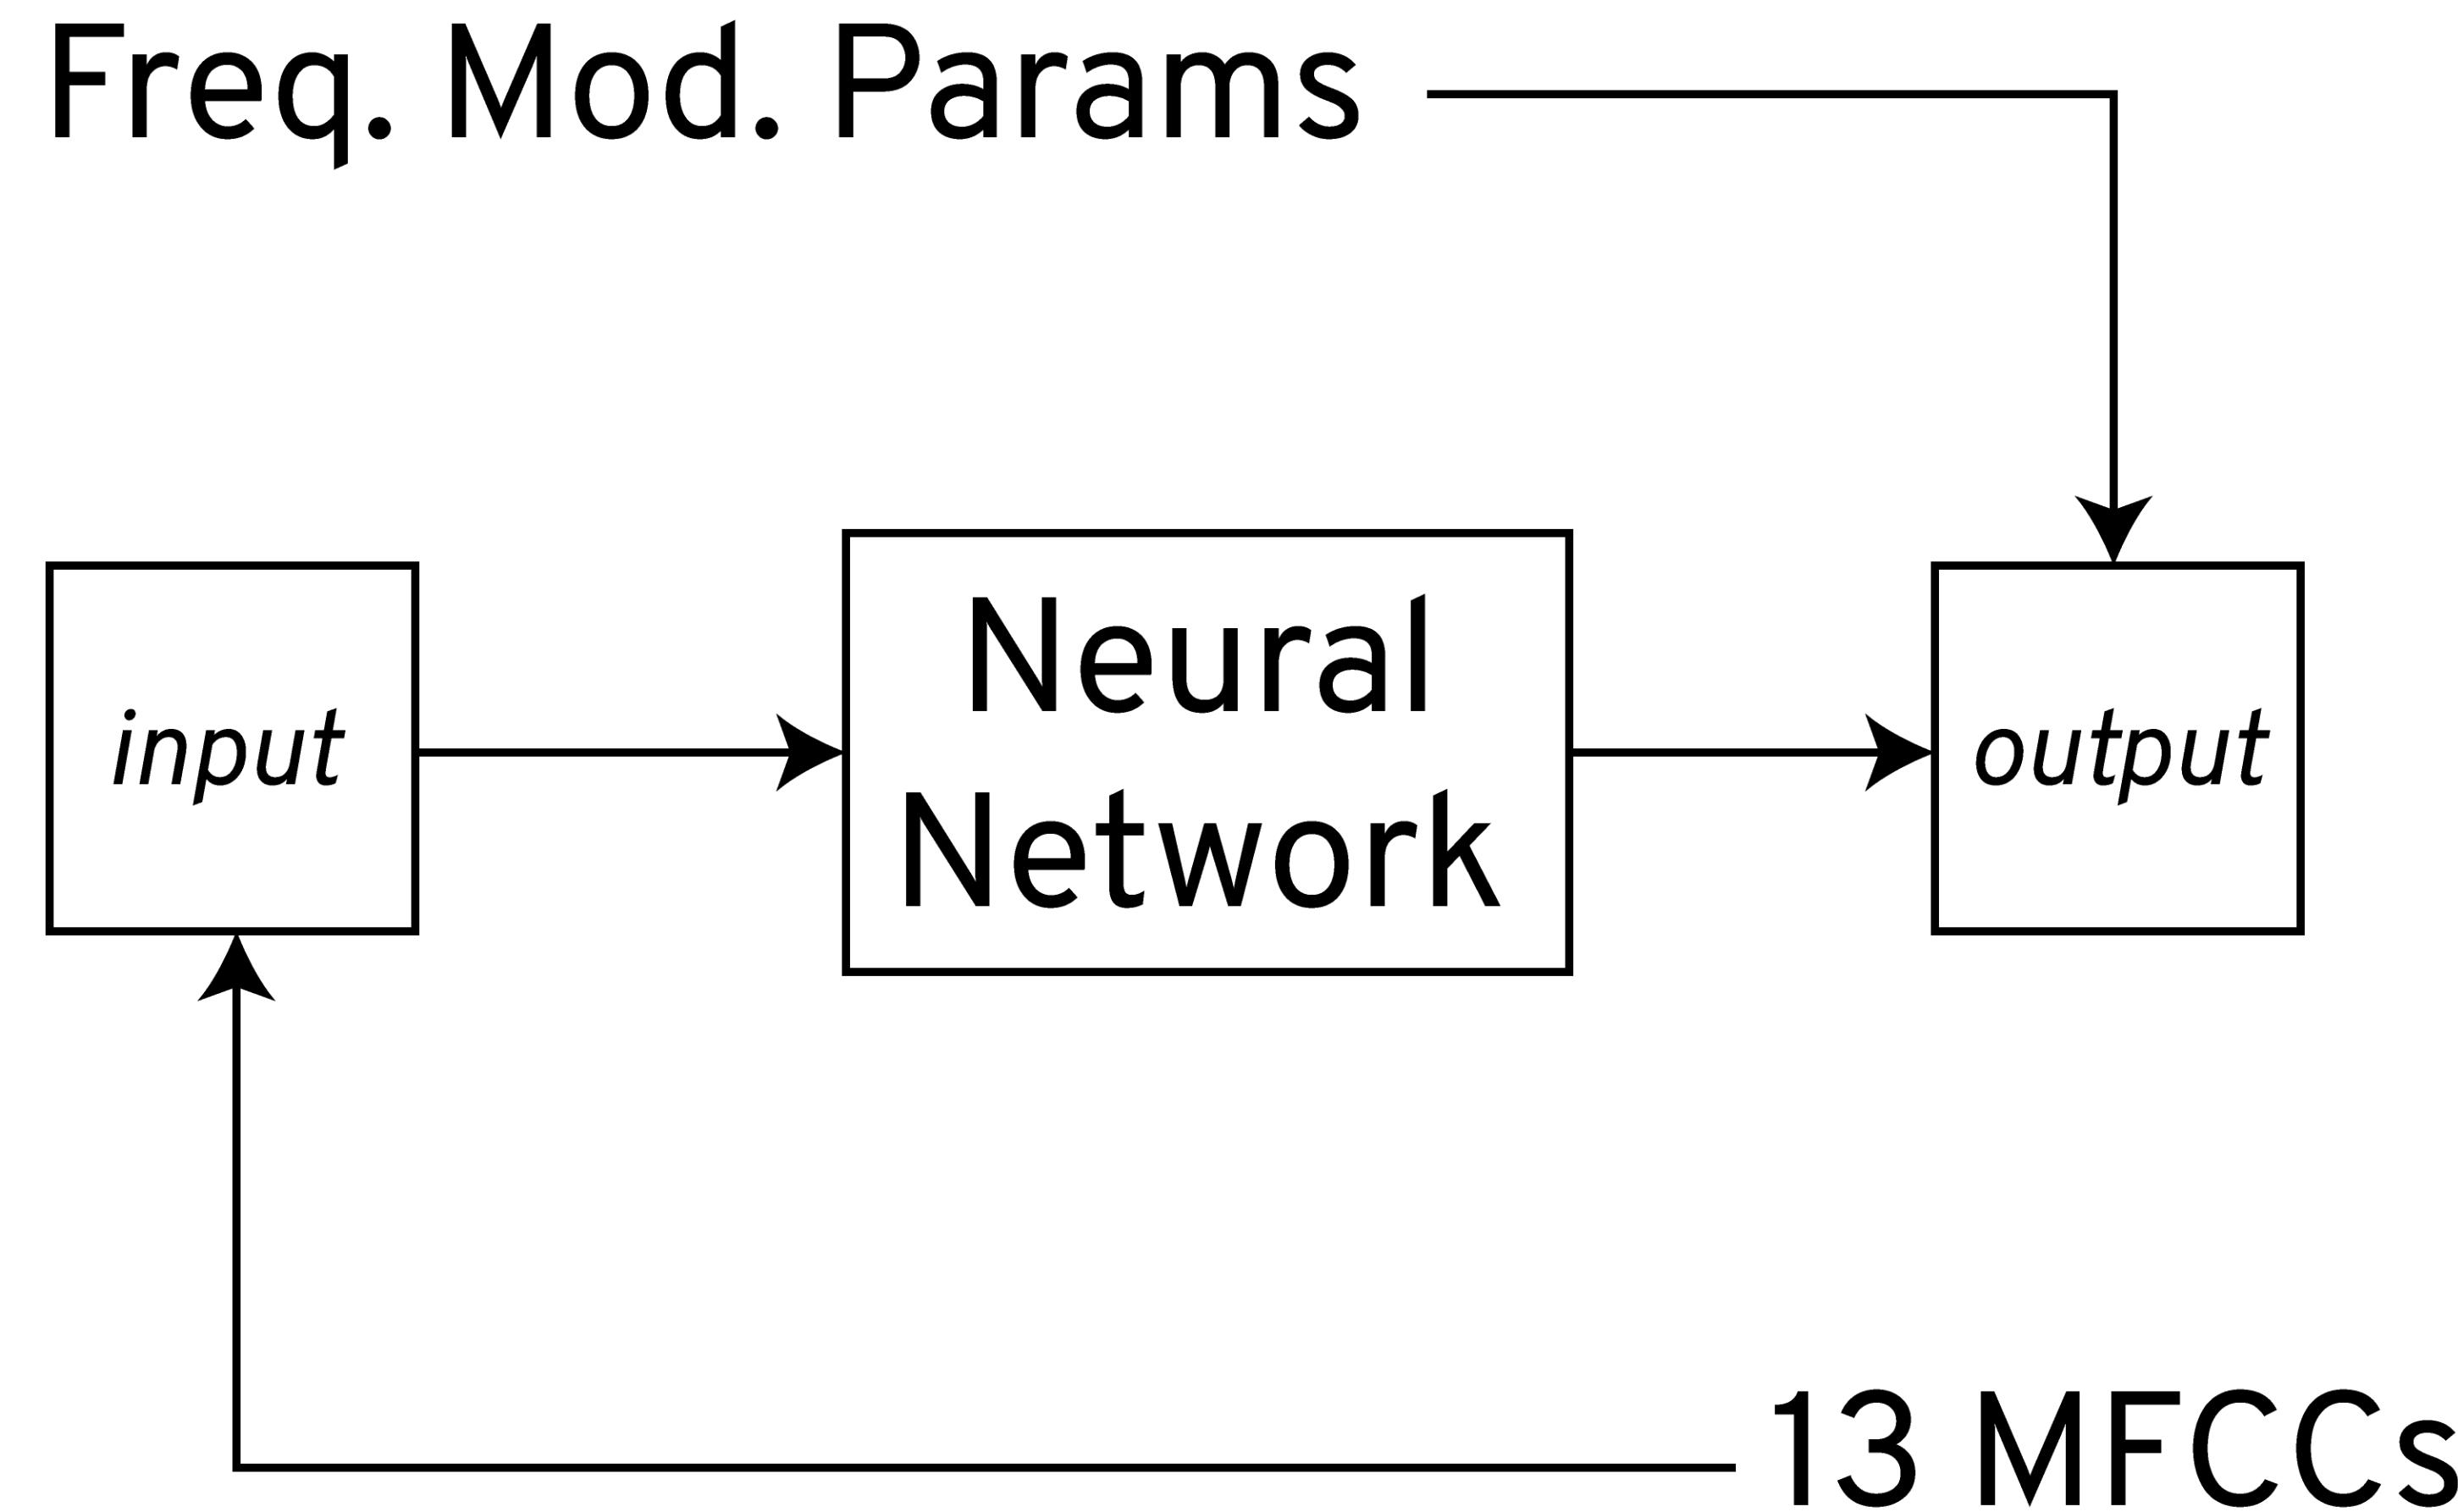
Training

Freq. Mod. Params



13 MFCCs

Training



Performing

